

TRYST

Product and Technical Specification, and Terms of Service

Table of Contents

Document control

Consolidated edition v1.0 — 4 October 2026 — Confidential

Owner: Office of the Group Chief Executive

Delivery vehicle: Ring-fenced SPV, not branded to the parent group

Status: Pre-incorporation. Legal opinion gating build start.

Discretion is the product. Everything else is plumbing.

How to read this document

This consolidated edition brings together the two working documents for the TRYST venture.

Part A — Product & Technical Specification (sections 1–21) is the developer-ready build document: market thesis, agentic architecture, behaviour-learning models, matching, discretion architecture, data model, API surface, stack, compliance envelope, business model, roadmap, screening and exclusion design, liability architecture, and the Guardian classifier annex.

Part B — Terms of Service (clauses 1–29) is a drafting-stage document prepared by the product team for review by qualified UK counsel. It is not legally reviewed and must not be published in this form. Clause cross-references inside Part B refer to Part B clause numbers; section cross-references inside Part A refer to Part A section numbers. The two numbering schemes are independent.

Gating note. No engineering work should begin before a written legal opinion and a completed Data Protection Impact Assessment are in place. Section 14 of Part A sets out what counsel must be instructed on; clause 29 of Part B lists the ten open questions.

All regulatory, payment and jurisdictional statements require confirmation by qualified UK counsel before build commences. Unit economics are illustrative and assumption-driven, and should be replaced with measured cohorts after closed beta. Third-party market figures are as reported by the cited sources and have not been independently audited. Statutory positions, particularly commencement dates, must be verified as at the date of use.

Part A — Product & Technical Specification

1. Thesis & Market Gap

The discreet-connection market is large, proven, and structurally badly served. The incumbents monetise frustration. TRYST monetises resolution.

What the market actually looks like

Two incumbents define the space, and both have just handed a challenger an opening.

Feeld owns ethical non-monogamy and couple-seeking-third. Revenue grew 26% to £48.9M in 2024 per Companies House filings, and it reached 2M+ active users by 2026 while ranking in the Google Play top ten for its category. Its weakness is well documented in its own reviews: the free tier is too limited to evaluate whether the app works for you, discovery is high-noise, and it is a *browsing* product, not a *resolution* product.

Ashley Madison owns the affair position — and is walking away from it. On 24 February 2026 it announced a global rebrand as a discreet dating app, marking a shift away from married dating, with the tagline "Where Desire Meets Discretion" and an emphasis on "Ethical Discretion". Its economics are the pay-to-initiate credit model: for a conversation between two members, one member — always the man — must pay eight credits to initiate. That model is extractive and it shows: public user discussion repeatedly complains about cost and bots, especially from male users.

THE STRATEGIC READ

The category leader in affairs is vacating the affair position under regulatory, payments and app-store pressure. That is simultaneously **the opening** (an 80M-member brand abandoning its own core intent) and **the warning** (they are not doing it for fun). TRYST must take the position *and* survive the pressure that pushed them off it. Sections 9 and 14 are therefore not compliance chores — they are the moat.

The four failures TRYST is built to fix

Failure	Why it persists	TRYST answer
Stated ≠ real intent. People misrepresent what they want and what their partner knows. Every mismatch ends in a dead thread.	Profiles are self-authored marketing. No incumbent models revealed preference against stated preference.	Behaviour-learned RevealedPreferenceVector (\$6) matched against IntentVector; divergence is a first-class feature, not noise.
The funnel is 200 messages to zero meets. Logistics, discretion windows and boundaries are never resolved before people invest.	Engagement-maximising ranking is rewarded for long, unresolved threads.	Envoy agent-to-agent handshake (\$5.4) pre-qualifies intent, logistics and limits before humans meet. Reward function penalises time-in-app (\$6.4).
Attention collapse. A small	Unconstrained ranking is a	Exposure-fairness

Failure	Why it persists	TRYST answer
number of profiles absorb nearly all inbound. Everyone else churns.	rich-get-richer loop; couple broadcast makes it worse.	constraint + inbound caps + asymmetric Key pricing on couple→solo outreach (\$7.4).
Discretion is a marketing word, not an architecture. The 2015 Ashley Madison breach is the category's defining event and nobody re-architected for it.	Retention-by-default and server-side plaintext are cheaper to build.	Breach-assumed design: E2EE messaging, cryptographic erasure, token-only verification, PSI exclusion ring (\$9).

Operating stance (write this into the constitution)

TRYST does not moralise and does not police whether a user's partner knows. It is not the platform's business, it is unverifiable, and pretending otherwise produces bad product. What TRYST *does* enforce, without exception, is the line that keeps it licensable, bankable and distributable:

- Verified adults only, enforced by highly effective age assurance.
- Consent between every party *on the platform*. Couples must be dual-verified; no one is listed by a partner without co-signature.
- Zero tolerance: image abuse, sextortion, coercion, commercial sexual services, minors.

That is the defensible line. "We are indifferent to your marriage; we are absolutist about consent and safety on our surface."

2. The Name

TRYST

tryst (*n.*) — a private, prearranged meeting between lovers.

One word. One syllable. Five letters. It already means precisely what the product does, which means zero semantic education spend. Hard *T* open, hard *ST* close — it reads sharp rather than seedy, which is exactly the register this category never achieves. It verbs naturally. It is deniable on a lock screen in a way that "Ashley Madison" is catastrophically not. And it is old — 14th-century Middle English — which buys it an elegance that no coined tech name can.

Brand codes

Oxblood and bone. Serif display, tight tracking. No pink, no hearts, no silhouettes of anonymous couples, no lipstick. Reference set: a private members' club, not a nightclub. The competitive set looks like adult entertainment; TRYST should look like a members' bar in Marylebone.

Fallbacks if `tryst.com`/trademark clears poorly

- **ASIDE** — theatrical term for words spoken to be heard by some and not others. Conceptually perfect, commercially colder.
- **DUPLEX** — two dwellings, one structure. Reads as the double life; strong for the couples/third segment. Weaker for solo affair intent.

Domain strategy: secure `tryst.app` / `tryst.co` / `trystapp.com` plus defensive registrations before any UKIPO Class 45 / Class 9 filing becomes public.

3. Segments & Jobs

Six segments. Every one is first-class in the data model — no segment is an edge case, because edge-case treatment is exactly why couples get harassed and solos get flooded on existing apps.

Code	Segment	Job to be done	Primary anxiety
S1	Attached, undisclosed	Meet one person, discreetly, without altering my primary life	Exposure. Everything else is secondary.
S2	Attached, disclosed / open	Find quality connections without re-explaining my structure every time	Judgement and time-waste
S3	Couple seeking a third	Find someone both of us want, who wants both of us	Being ghosted; internal partner friction
S4	Solo open to couples	Be treated as a person, not a utility, with a veto that is respected	Harassment volume; pressure and bait-and-switch
S5	Couple seeking a couple	Four-way compatibility with low coordination cost	Coordination collapse (the hardest match in the market)
S6	Curious / never acted	Understand whether this is for me, at low risk	Committing before they are ready; being discovered mid-exploration

DESIGN CONSEQUENCE: S4 IS THE SCARCE ASSET

S4 supply determines whether S3 and S5 have a product at all. Treat S4 inbound volume as a rationed resource, priced and capped (\$7.4). Any ranking change that raises S4 inbound above the cap is a regression, regardless of what it does to session metrics.

4. Product Loops

LOOP A – INTAKE (once, ~7 min conversational)

Age assurance → Cartographer interview → IntentVector v0
→ ExclusionRing enrolment → DiscretionProfile → first slate

LOOP B – DISCOVERY (daily, target < 6 min/session)

Curated slate (12-18 profiles, not infinite)
→ Intent sent (costed) / passed / saved
→ Mirror updates RevealedPreferenceVector
→ slate recomputed

LOOP C – HANDSHAKE (the differentiator)

Mutual intent → Envoy_A ≠ Envoy_B negotiate
→ structured Brief delivered to BOTH humans
→ human accepts / declines / amends
→ E2EE thread opens ONLY on double-accept

LOOP D – RESOLUTION

Thread → logistics agreed → Safe Meet checklist
→ meet confirmed → Aftercare prompt (48h)
→ reward signal closes the training loop


LOOP E – DISCRETION (continuous, background)

Curtain scores session risk → geofence, collision,
notification masking, retention TTL enforcement

Loop C is the product. Everything else is table stakes that a competent team ships in four months. If Loop C does not measurably raise meets-per-intent-sent over a control cohort, TRYST is a me-too and should not raise a Series A.

5. Agentic Architecture

Seven agents. Each is an independently deployable service with its own tool surface, memory scope, policy file and eval suite. Agents communicate over a typed internal bus — never by passing free text.





5.1 Cartographer — intake

A conversational interview, not a form. Forms produce socially performed answers; interviews produce usable ones. Elicits, over ~7 minutes:

- **Structure:** relationship status, disclosure state, partner-awareness (self-declared, never verified, never displayed to others)
- **Intent shape:** one-off / recurring / ongoing-parallel / exploratory-only / couple-third / couple-couple
- **Desire taxonomy:** ~140 tags across dynamic, activity, context. Multi-select with a "curious" tier distinct from "yes" tier.
- **Hard limits:** absolute exclusions. Never overridden by ranking, ever.
- **Discretion geometry:** home exclusion radius, work exclusion radius, availability windows, device-risk level, notification tolerance
- **Logistics reality:** travel radius, hosting capability, typical notice needed, budget band for venues

Output: a signed `IntentVector` (§10). Re-run quarterly or on detected drift.

5.2 Mirror — behaviour learning

Detailed in §6. Runs continuously, owns `RevealedPreferenceVector`, and is the only agent permitted to write to it.

5.3 Broker — matching

Detailed in §7. Two-stage retrieve-then-rank, with hard-limit filtering and exposure-fairness constraints applied as hard gates before scoring, never as a post-hoc re-sort.

5.4 Envoy — the agent-to-agent handshake

FLAGSHIP FEATURE

When two users express mutual intent, their Envoys negotiate before either human types a word. The negotiation runs over a **constrained schema, not free-form chat** — this keeps it auditable, cheap, and immune to prompt injection.

Negotiation covers exactly six dimensions:

1. **Intent alignment** — is one side seeking one-off and the other ongoing? Surface it in 3 seconds, not 3 weeks.
2. **Structure alignment** — couple ↔ solo expectations, veto rights, whether both partners are participating
3. **Boundary intersection** — computed intersection of hard limits and desire tags; flags any collision
4. **Logistics feasibility** — availability window overlap, travel radius intersection, notice requirements, hosting

5. **Discretion compatibility** — exclusion zones intersected, photo-disclosure pace, device constraints
6. **Verification parity** — does each side's verification tier meet the other's stated minimum?

Envoy emits a `Brief` to each side. Each `Brief` is asymmetric: it contains what *that* user needs, never what the other user's agent disclosed in confidence. Fields flagged `private:true` in an `IntentVector` are usable in computation but never in output text.

```
// Envoy negotiation turn — strict schema, no free text
{
  "turn": 3,
  "from": "envoy:usr_7f3a",
  "to": "envoy:usr_91bc",
  "assert": {
    "intent_shape": "one_off",
    "structure": "attached_undisclosed",
    "notice_required_hours": 48,
    "availability": [{"dow": [2,3], "start": "18:00", "end": "23:00", "tz": "Europe/London"}],
    "travel_radius_km": 25,
    "hard_limits": ["no_photos_exchanged", "no_home_visit", "no_overnight"],
    "verification_tier": "V2",
    "min_counterparty_tier": "V2"
  },
  "query": ["intent_shape", "notice_required_hours", "hosting_capability"],
  "private_fields": ["exclusion_zones"] // computable, never emitted
}

// Resulting Brief (user-facing, generated, human-reviewable)
{
  "compatibility": 0.81,
  "aligned": ["intent_shape", "availability", "travel_radius", "verification"],
  "friction": [{"field": "hosting", "detail": "Neither side can host. Venue required.",
    "suggested_resolution": "hotel_or_venue"}],
  "blocking": [],
  "recommended_next": "open_thread",
  "estimated_meet_probability": 0.34
}
```

NON-NEGOTIABLE ENVOY CONSTRAINTS

- Envoy **never impersonates the user in conversation with another human**. It negotiates agent-to-agent only, and every `Brief` is labelled as agent-generated. Undisclosed AI impersonation destroys trust in this category faster than any other and invites a bot scandal of exactly the kind that already damages the incumbent.
- Envoy cannot open a thread. Only a double human accept opens a thread.
- Envoy output is never the last word — the human sees the friction and decides.
- All counterparty-agent input is untrusted data. Schema-validated, never executed as instruction.

5.5 Curtain — discretion

Scores every session and every action for exposure risk, and has authority to degrade the experience to protect the user. Controls: geofence suppression, contact-collision blocking (§9.3), notification masking, decoy skin activation, media-reveal gating, session-risk nudges ("you are inside your declared work exclusion zone — hiding profile for 2h").

5.6 Guardian — safety

Runs inference on every message, image, profile edit and report. Has a **veto over every other agent** — Guardian can halt a handshake, freeze a thread, or suspend an account without Broker or Envoy consent. Maps directly to the risks the regulator names for this category: Ofcom identified romance fraud, intimate image abuse, sextortion, harassment, stalking and grooming as key risks commonly associated with dating services. Each gets a dedicated classifier and a documented escalation path (§8.4).

5.7 Aftercare — outcome capture

48 hours after a scheduled meet, a two-tap private prompt: did it happen, was it as represented, would you repeat. This is the only honest source of ground truth in the category and it is the single highest-value dataset TRYST will own. Responses are never shown to the counterparty and never aggregated into a public rating — ratings invite retaliation and coercion. They feed the reward function only.

6. Behaviour Learning

Core premise: in this category, **the gap between what a user says and what a user does is larger than in any other consumer vertical**, and that gap is the signal.

6.1 Event taxonomy

Emitted to Kafka, schema-registry enforced, pseudonymous at rest.

Class	Events	What it reveals
Attention	card_impression, dwell_ms, photo_index_viewed, bio_scroll_depth, profile_revisit, expand_tag	Real attraction and real curiosity, uncoloured by self-presentation
Decision	intent_sent, pass, save, undo, hesitation_ms	Revealed preference; hesitation is a strong ambivalence signal
Conversation	msg_sent, reply_latency_s, msg_len, typing_abandoned, question_ratio, thread_decay	Genuine interest vs politeness; effort asymmetry predicts ghosting
Handshake	brief_viewed, brief_accepted, friction_dismissed, amend_requested	Which frictions actually matter to this user vs which they wave through
Resolution	logistics_agreed, meet_scheduled, meet_confirmed, no_show, aftercare_score	Ground truth. The only events in the reward function with positive weight above 0.3.

Class	Events	What it reveals
Discretion	panic_invoked, session_aborted, geofence_hit, screenshot_detected, decoy_opened	Real-world risk posture; drives Curtain aggressiveness
Negative	report, block, unmatched, ghost_detected, refund_request	Harm and dissatisfaction signals; heavily weighted negative

6.2 The two vectors

IntentVector — stated. Dim 128. Source: Cartographer. Updates: quarterly/manual.
 RevealedPreferenceVector — revealed. Dim 256. Source: Mirror. Updates: nightly batch + hourly delta.

DivergenceScore = 1 - cosine(project(IntentVector), RevealedPreferenceVector)

Divergence is **used, not corrected**. A high-divergence user is not lying — they are exploring, or they wrote their profile for an audience. Broker blends the two with a learned per-user weight α :

EffectiveVector = $\alpha \cdot \text{RevealedPreferenceVector} + (1 - \alpha) \cdot \text{project(IntentVector)}$
 $\alpha = \sigma(w \cdot [\text{interaction_count}, \text{days_active}, \text{divergence_stability}, \text{aftercare_count}])$
 $\alpha \in [0.15, 0.85]$ // stated preference NEVER drops below 15% weight

The floor on stated preference is a product decision, not a modelling one. Fully surrendering to revealed behaviour produces a system that decides what a user wants better than they do, which in this category is both creepy and wrong — people's stated boundaries are often the thing keeping their life intact.

HARD-LIMIT INVIOLABILITY

Hard limits, exclusion zones and structure-incompatibility are applied as **pre-retrieval filters**. No learned signal, no matter how strong, can surface a candidate that violates one. Any model change that could route around a hard limit is a P0 incident. Write this as a test, not a policy.

6.3 Model inventory

ID	Model	Type	Serving	Target metric
M1	Intent classifier	Multi-label transformer, 9 classes	Batch + on-edit	Macro F1 > 0.82
M2	Two-tower retrieval	Dual encoder, 256-d, ANN index	< 40 ms p95	Recall@200 > 0.91
M3	Cross-encoder ranker	Gradient-boosted + neural blend	< 120 ms p95 for 200 cands	Mutual-reply AUC > 0.78
M4	Meet propensity	Survival model (time-to-meet)	Nightly	C-index > 0.70

ID	Model	Type	Serving	Target metric
M5	No-show / flake risk	GBDT on conversation features	Pre-meet	Precision@20% > 0.55
M6	Harassment & coercion	Fine-tuned encoder + rules	Inline, < 80 ms	Recall > 0.95 @ FPR < 0.02
M7	Authenticity (bot/catfish)	Ensemble: image, text, device, graph	On signup + continuous	Catch rate > 0.93
M8	Discretion risk	Contextual bandit	Per session	Panic-invocation reduction
M9	Exposure fairness monitor	Gini + constraint solver	Hourly	Inbound Gini < 0.55

6.4 The reward function — and what it deliberately does not reward

```
R = 0.10·mutual_reply
+ 0.15·sustained_thread(≥6 turns, both sides)
+ 0.20·logistics_agreed
+ 0.35·meet_confirmed
+ 0.25·aftercare_positive
- 0.40·report
- 0.20·block
- 0.15·no_show
- 0.10·ghost_after_intent
- 0.05·session_time_beyond_8min      // ← time in app is a COST
```

Read the last term again. Every competitor in this market optimises for session length and message volume, because that is what monetises under a pay-per-message or ad-adjacent model. TRYST's pricing (\$15) is deliberately structured so that resolving a user quickly is *more* profitable than detaining them — subscription plus success-weighted retention, never pay-per-message. The economic model and the loss function have to agree or the loss function loses.

6.5 Cold start & exploration

- Turns 1-20: content-based from IntentVector only; $\alpha = 0.15$.
- Exploration: Thompson sampling, **15% of every slate is exploratory** and permanently so. Without a persistent exploration floor, the system converges on a narrow type and the user's world shrinks — the classic failure of every mature dating recommender.
- New-profile boost: 72-hour decaying multiplier, capped by the inbound limiter so a boost cannot itself cause attention collapse.
- Couples (S3/S5): the couple entity carries a blended vector from both partners plus a learned `joint_veto_model` — historically the hardest cold start in the category, since a couple's true preference is the *intersection* of two, not the union.

7. Matching & Ranking

STAGE 0 HARD GATES (non-scored, non-bypassable)

- └ age assurance tier ≥ required
- └ hard-limit intersection = ∅ collisions
- └ structure compatibility matrix
- └ ExclusionRing check (PSI) – §9.3
- └ geofence exclusion (both sides)
- └ block / report history
- └ jurisdiction gate

STAGE 1 RETRIEVAL (M2) ~50k → 200 candidates, < 40 ms

STAGE 2 RANKING (M3+M4) 200 → 40, < 120 ms

STAGE 3 CONSTRAINTS exposure fairness, inbound caps, diversity

STAGE 4 SLATE 18 cards + 15% exploration slots

7.4 Exposure fairness — the mechanism that makes or breaks S4 retention

Three coupled controls:

1. **Inbound cap.** Hard ceiling on intents a profile can receive per 24h (default 12, user-adjustable 4–40). Overflow is queued and re-scored, not delivered. Prevents the 400-message inbox that drives S4 churn.
2. **Gini constraint.** M9 monitors inbound-intent distribution per city per segment. If Gini > 0.55, the ranker's fairness penalty term increases until it recovers. Exposure equality is a *constraint*, not an objective — the ranker still ranks on quality, but within a bounded distribution.
3. **Asymmetric outreach cost.** Couple→solo (S3→S4) costs 3× the Keys of solo→solo. Broadcast-style couple outreach is the dominant harassment vector in this market, and the cheapest fix is price, not moderation. Outreach is *fully refunded in Keys* on a reply — so genuine, targeted outreach is free in practice and spray-and-pray is expensive.

8. Trust, Verification, Consent

8.1 Verification tiers

Tier	Method	Retained	Unlocks
V0	Email/phone, device fingerprint	Hashed identifier	Browse only. No outbound intent.
V1	Liveness selfie matched to profile photos	Boolean + match score. No biometric template.	Outbound intent, threads
V2	Highly effective age assurance via third-party provider	Boolean is_adult + provider token only. No ID image, no DOB.	Media exchange, Envoy, meets
V3	Optional: bank/open-banking or digital ID	Boolean + token	Trust badge; required to set high min_counterparty_t

Tier	Method	Retained	Unlocks
VC	Couple co-signature — both partners independently V2, each confirms the link from their own device	Link record, both user IDs	ier Couple profile. Mandatory — no exceptions.

Ofcom's assessment framework is explicit about what counts: age assurance methods are evaluated against technical accuracy, robustness against circumvention, reliability and fairness, and must meet all four to be considered highly effective; the published non-exhaustive list includes Open Banking, photo ID matching, facial age estimation and digital identity services. Build to that list, with a documented provider due-diligence file.

VC IS THE SINGLE MOST IMPORTANT TRUST FEATURE IN THE PRODUCT

Every couples platform is plagued by one person creating a "couple" profile their partner has not consented to, or a solo man posing as a couple. Requiring independent V2 verification of *both* partners, each confirming from their own device, eliminates the entire class. It will cost 20–30% of couple signups in year one. Pay it. It is also the cleanest answer you have when a regulator asks how you handle consent.

8.2 Consent ledger

Every consent is an immutable, timestamped, hash-chained record: Art. 9 explicit processing consent, media-view consent, meet confirmation, couple link, data-retention election. Exportable on SAR request in under 30 days. Revocation writes a new record; it never mutates the old one.

8.3 Anti-fraud

Ofcom's own research on the scale of the problem: 29% of adult internet users have experienced romance or dating scams, and 6% reported encountering fraud on a dating website or app. Controls: M7 ensemble; reverse-image and AI-generation detection on every uploaded photo; device-graph clustering; velocity limits; payment-instrument fingerprinting; refundable verification deposit (\$15) as economic friction against burner accounts.

8.4 Escalation matrix

Signal	Automatic	Human SLA	External
Suspected minor	Immediate full freeze, both accounts, content preserved under legal hold	15 min	Mandatory report; law-enforcement liaison
Intimate image abuse / NCII	Media quarantine, hash to StopNCII, sender frozen	1 h	Takedown log; victim support routing

Signal	Automatic	Human SLA	External
Sextortion / coercion	Thread freeze, evidence preserved, victim safety banner	1 h	Action Fraud referral path
Threat of violence / stalking	Account suspend, mutual geo-suppression, device ban	1 h	Police referral on request
Romance fraud pattern	Payment-link stripping, shadow-limit, cohort sweep	4 h	Bank/PSP intel share
Commercial sexual services	Immediate removal	4 h	—
Harassment (M6 > 0.9)	Send-block, warning, strike	24 h	—

9. Discretion Architecture

Design from the assumption that you will be breached, subpoenaed, and that a suspicious partner already has the user's phone in their hand.

9.1 Threat model

ID	Adversary	Likelihood	Primary control
A1	Partner with physical device access	Very high — the daily case	Decoy skin, duress PIN, biometric gate, masked notifications, no app in recents, neutral icon/name
A2	Partner reading a shared bank statement	Very high	Neutral billing descriptor, voucher/gift-code payment path, no auto-renew without 7-day notice
A3	Recognition by shared social graph	High	ExclusionRing PSI (9.3), progressive photo reveal, face-blur default, geofence
A4	Mass data breach	Medium / existential	E2EE messaging, client-side media

ID	Adversary	Likelihood	Primary control
			encryption, no plaintext identity ↔ profile join, cryptographic erasure, 30-day TTL
A5	Targeted extortion by a counterparty	Medium	Watermarked media, screenshot detection, no-download viewer, Guardian sextortion classifier
A6	Civil discovery in divorce proceedings	Medium	Minimal retention, cryptographic erasure on deletion, transparency report, published legal-process policy
A7	Insider threat	Medium	Zero standing access, break-glass with dual authorisation, full audit log, staff cannot read message content by design
A8	Face-search de-anonymisation of profile photos	Medium and rising	Perturbation/adversarial-noise layer on public-tier photos, private-tier photos never served unencrypted

9.2 Cryptographic design

MESSAGING	MLS (RFC 9420) or libsignal. Server stores ciphertext only. Server CANNOT read message content. Not "does not" – cannot.
MEDIA	Client generates per-media AES-256-GCM key. Ciphertext → object store. Key wrapped per-recipient, delivered in-band. Revoke = destroy wrapped key. Ciphertext becomes noise instantly.
IDENTITY SPLIT	Two stores, separate credentials, separate encryption domains: IDENTITY-DB : phone/email hash, verification tokens, billing ref PROFILE-DB : profile, vectors, behaviour, media refs Join key is an HSM-held pepper. Neither dump alone is identifying.


```

    assurance_at      TIMESTAMPTZ,
    billing_ref       TEXT,                                -- PSP token. No PAN, ever.
    jurisdiction       CHAR(2) NOT NULL,
    root_key_ref       TEXT NOT NULL,                      -- KMS handle; destroy = erasure
    created_at         TIMESTAMPTZ DEFAULT now(),
    deleted_at         TIMESTAMPTZ
);

-- ===== PROFILE-DB (join to identity only via HSM-held pepper) =====
CREATE TABLE profile (
    profile_id         UUID PRIMARY KEY,
    segment            TEXT CHECK (segment IN ('S1','S2','S3','S4','S5','S6')),
    entity_type        TEXT CHECK (entity_type IN ('solo','couple')),
    display_name       TEXT,                                -- pseudonym; real name never required
    pronouns           TEXT,
    gender_identity    TEXT,                                -- free text, not enum
    seeking            TEXT[],
    age_band           TEXT,                                -- band only; exact DOB never stored
    city_geohash5      TEXT,                                -- ~5km precision. Never exact coords.
    bio_ct             BYTEA,                                -- encrypted at rest
    visibility         TEXT CHECK (visibility IN ('public','matched','incognito')),
    status             TEXT DEFAULT 'active',
    created_at         TIMESTAMPTZ DEFAULT now()
);

CREATE TABLE couple_link (
    couple_id          UUID PRIMARY KEY,
    profile_a          UUID NOT NULL REFERENCES profile,
    profile_b          UUID NOT NULL REFERENCES profile,
    a_cosigned_at      TIMESTAMPTZ NOT NULL,                -- both NOT NULL or link is invalid
    b_cosigned_at      TIMESTAMPTZ NOT NULL,
    a_device_hash      BYTEA NOT NULL,                      -- must differ from b_device_hash
    b_device_hash      BYTEA NOT NULL,
    veto_mode          TEXT DEFAULT 'either',               -- 'either' | 'both'
    dissolved_at       TIMESTAMPTZ,
    CONSTRAINT distinct_devices CHECK (a_device_hash <> b_device_hash)
);

CREATE TABLE intent_vector (
    profile_id         UUID PRIMARY KEY REFERENCES profile,
    intent_shape       TEXT[],                              -- one_off|recurring|ongoing|exploratory|third|
couple
    structure          TEXT,
    desire_tags        JSONB,                                -- {tag: yes|curious|no}
    hard_limits        TEXT[] NOT NULL,                    -- NEVER overridable
    discretion         JSONB,                                -- exclusion zones, windows, device risk
    logistics          JSONB,                                -- radius, hosting, notice_hours, venue budget
    min_counterparty_tier TEXT DEFAULT 'V1',
    embedding          VECTOR(128),
    private_fields     TEXT[],                              -- computable, never emitted by Envoy
    updated_at         TIMESTAMPTZ
);

CREATE TABLE revealed_preference (
    profile_id         UUID PRIMARY KEY REFERENCES profile,
    embedding          VECTOR(256),
    alpha              REAL CHECK (alpha BETWEEN 0.15 AND 0.85),
    divergence         REAL,
    confidence         REAL,
    n_signals          INT,
    computed_at        TIMESTAMPTZ
);

```



```

CREATE TABLE handshake (
  handshake_id      UUID PRIMARY KEY,
  a_profile          UUID, b_profile UUID,
  transcript_ct      BYTEA,                -- encrypted agent turns, 30d TTL
  brief_a_ct         BYTEA, brief_b_ct BYTEA,
  compatibility       REAL,
  meet_probability   REAL,
  a_decision          TEXT, b_decision TEXT, -- accept|decline|amend|pending
  opened_thread       UUID,                -- non-null ONLY on double accept
  created_at          TIMESTAMPTZ
);

CREATE TABLE exclusion_ring (
  profile_id          UUID, excluded_profile_id UUID,
  source              TEXT,                -- psi_contact|device_graph|manual|payment
  confidence           REAL,
  PRIMARY KEY (profile_id, excluded_profile_id)
); -- suppression is ALWAYS bidirectional and ALWAYS silent

CREATE TABLE consent_ledger (
  consent_id          UUID PRIMARY KEY,
  profile_id           UUID NOT NULL,
  kind                 TEXT NOT NULL,      -- art9_processing|media_view|meet|couple_link|retention
  granted              BOOLEAN NOT NULL,
  scope                JSONB,
  prev_hash            BYTEA, entry_hash BYTEA NOT NULL, -- hash chain; append-only
  created_at           TIMESTAMPTZ DEFAULT now()
);

CREATE TABLE aftercare (
  meet_id             UUID PRIMARY KEY,
  reporter             UUID, counterparty UUID,
  occurred             BOOLEAN,
  as_represented       SMALLINT,          -- 1..5
  would_repeat         BOOLEAN,
  safety_flag          BOOLEAN,
  created_at           TIMESTAMPTZ
); -- NEVER surfaced to counterparty. NEVER a public rating.

```

11. API Surface

REST/JSON at the edge, gRPC internally. Every endpoint: OAuth2 + DPoP-bound tokens, per-endpoint rate limits, jurisdiction gate middleware, Art. 9 consent check middleware.

AUTH & IDENTITY

POST	/v1/auth/start	phone email → OTP
POST	/v1/auth/verify	→ session + device key registration
POST	/v1/assurance/session	→ third-party HEAA handoff URL
POST	/v1/assurance/callback	provider token → is_adult boolean only
DELETE	/v1/account	→ cryptographic erasure, < 60s

PROFILE & INTENT

POST	/v1/intake/turn	Cartographer conversational turn
GET	/v1/intake/vector	→ IntentVector
PATCH	/v1/profile	
POST	/v1/profile/media	client-encrypted blob + wrapped keys
POST	/v1/couple/invite	→ co-sign token (expires 24h)
POST	/v1/couple/cosign	MUST originate from a distinct device

DISCOVERY

GET	/v1/slate	→ 18 cards + exploration slots
POST	/v1/intent	{target, keys_spent} → costed outbound
POST	/v1/pass POST /v1/save	
POST	/v1/signals:batch	behaviour events (batched, ≤ 50)

HANDSHAKE

POST	/v1/handshake	on mutual intent → Envoy negotiation
GET	/v1/handshake/{id}/brief	→ asymmetric Brief for caller only
POST	/v1/handshake/{id}/decision	accept decline amend
↳ thread opens ONLY when both decisions == accept		

THREADS (E2EE – server relays ciphertext and nothing else)

POST	/v1/threads/{id}/messages	ciphertext envelope
GET	/v1/threads/{id}/messages	
POST	/v1/threads/{id}/meet	propose confirm cancel
POST	/v1/threads/{id}/burn	revoke keys both sides

DISCRETION

POST	/v1/exclusion/psi/blind	OPRF round 1
POST	/v1/exclusion/psi/commit	OPRF round 2 → suppression pairs
PUT	/v1/discretion/zones	home/work exclusion geofences
POST	/v1/panic	wipe local, revoke device, sign out all

SAFETY

POST	/v1/report	→ Guardian queue with SLA routing
POST	/v1/block	

COMMERCE (web-first; see §14.3)

GET	/v1/entitlements	
POST	/v1/keys/purchase	
POST	/v1/subscription	neutral descriptor enforced at PSP
POST	/v1/subscription/cancel	one-click, no retention gauntlet

12. Stack & Infrastructure

Layer	Choice	Why
iOS / Android	Native — Swift + Kotlin	Secure Enclave / StrongBox, FLAG_SECURE, screenshot detection, decoy skin and backup exclusion are all unreliable through a cross-platform shell. Security requirements decide this, not velocity.
Web	Next.js PWA, installable	Primary revenue surface (§14.3) and the only surface permitted to carry explicit media
Edge API	Go, gRPC-gateway	Concurrency, small attack surface, fast cold start

Layer	Choice	Why
Agents	Python, FastAPI, LangGraph-style typed state machines	Deterministic orchestration; every agent turn is replayable for eval
LLM — profile-touching	Self-hosted open-weights, UK region, vLLM	Art. 9 special-category data must not leave the controller's region or enter a third-party API without explicit consent and a DPA. Non-negotiable.
LLM — generic copy	Hosted frontier model, PII-redacted at boundary	Quality where no personal data is involved
Primary store	PostgreSQL 16 + pgvector (two isolated clusters)	Identity/profile split (\$9.2) enforced at infrastructure level
Vector search	Qdrant	Filtered ANN at the scale of hard-gate pre-filtering
Events	Redpanda + Schema Registry	Kafka API, lower ops burden
Analytics	ClickHouse	Behavioural aggregation; raw events purged at 180d
Feature store	Feast	Train/serve parity — critical for M3/M4
Objects	S3-compatible, client-side encrypted	Server never holds plaintext media
Secrets / keys	AWS KMS + CloudHSM	Root-key destruction is the erasure primitive
E2EE	MLS (RFC 9420) via OpenMLS, or libsignal	Standardised, audited, supports group threads for couples
Age assurance	Two providers (Yoti-class + ID-match), abstracted behind one interface	Ofcom expects layered protections and regular due diligence on external age assurance providers, with responsibility remaining with the service even when outsourced. A single provider is a single point of regulatory failure.
Region	AWS eu-west-2 (London), UK/EU residency only	Art. 9 data residency posture
IaC / CI	Terraform, GitHub Actions,	Auditable change control

Layer	Choice	Why
	Argo CD	

Engineering non-negotiables

- No third-party analytics SDK in the client. None. Not one. Every SDK is an Art. 9 leak and a breach vector.
- No advertising SDK, no attribution SDK that transmits behavioural data off-platform.
- Zero standing production data access. Break-glass requires dual authorisation and writes an immutable audit entry.
- Staff cannot read message content — not a policy, an architectural property.
- Quarterly external pen test; annual red-team exercise scoped explicitly to A4 and A7.

13. MLOps

- **Offline eval before every model ship:** replay on held-out cohorts, measure ΔR (§6.4), not AUC alone. AUC improvements that reduce meet rate are rejected.
- **Agent evals:** 400-case golden set per agent — Envoy negotiation fidelity, Guardian recall on adversarial phrasing, Cartographer intent extraction accuracy. Runs on every PR touching a prompt or policy file.
- **Prompt-injection suite:** counterparty-authored text (bios, messages, agent turns) is hostile input by default. Red-team corpus of injection attempts run on every Envoy and Guardian release.
- **Shadow deploy** → 5% → 25% → 100%, with automatic rollback on report-rate or S4-churn regression.
- **Fairness audit:** quarterly, per protected characteristic and per segment. A ranker that systematically buries a gender identity or body type is both a legal exposure and a product failure.
- **Feedback-loop guard:** monthly counterfactual audit for popularity collapse. Persistent 15% exploration floor (§6.5) is the structural defence.
- **Model cards + DPIA annex** maintained per model. The regulator will ask.

14. Compliance

This section gates the build. No engineering start before a written legal opinion and a completed DPIA. The cost of getting this wrong is not a fine — it is the whole company. Everything below requires sign-off from UK counsel with OSA and Art. 9 experience; it is a specification of what to instruct them on, not a substitute for the instruction.

14.1 Online Safety Act 2023

TRYST is unambiguously in scope. The OSA covers user-to-user services with links to the UK, including dating and social discovery apps where users create profiles, exchange messages, share media and connect with strangers through matching systems, and Ofcom has classified dating apps as services where minors face a risk of significant harm, which triggers mandatory age assurance requirements. In May 2026 Ofcom published guidance written specifically for dating platforms.

Required, before public launch:

- **Illegal-content risk assessment** and **children's access assessment**, documented and dated.
- **Highly effective age assurance** to the four-criteria standard (§8.1). Critically, this applies to the message inbox as much as the public feed — platforms that assume age assurance is only about the sign-up screen are wrong. Gate media exchange at V2, not just registration.
- Verify the user is an adult before displaying explicit material, and prevent that material appearing during the verification process itself.
- **Clear terms, accessible in-app reporting, rapid review and removal** — the baseline Ofcom requires of in-scope services.
- Controls mapped 1:1 to Ofcom's named dating risks (§8.4).
- Named senior manager accountable. Penalties reach 10% of global annual turnover or £18 million, whichever is higher, with criminal liability for senior managers in cases of wilful non-compliance.

Enforcement is live, not theoretical: as of February 2026 Ofcom had opened investigations into more than 90 platforms and issued six fines, including £800,000 against Kick Online Entertainment and £1 million against an adult website operator.

14.2 UK GDPR

- **Article 9.** Data concerning sex life and sexual orientation is special category. Lawful basis: **explicit consent, Art. 9(2)(a)** — granular, separate from the ToS accept, freely withdrawable, logged in the consent ledger.
- **DPIA is mandatory** (Art. 35): large-scale special-category processing plus systematic behavioural evaluation. Two triggers, not one.
- **Data minimisation as architecture.** No legal name. No exact DOB. No ID images. Geohash-5, never coordinates. Age band, never age.
- **Automated decision-making (Art. 22).** Ranking is not a legal-effect decision, but account suspension is. Human review path required and documented.
- **DSR machinery:** SAR export, rectification, erasure (cryptographic, < 60s), portability. Build these in sprint one, not year two.
- **Residency:** UK/EU only. No US sub-processor for Art. 9 data without a completed TIA.
- **72-hour breach playbook**, rehearsed twice yearly.

14.3 Distribution — the constraint that shapes the whole product

Apple's rule is explicit: Guideline 1.1.4 — "Hookup" apps that may include pornography or be used to facilitate prostitution will be rejected. Google has also tightened: the Age-Restricted Content and Functionality policy requires apps with matchmaking or dating features to use Play Console features to block minors, and new requirements and restrictions apply to anonymous and random chat apps.

Therefore, the two-surface strategy is a requirement, not an optimisation:

- **Native apps (store builds):** sanitised. No nudity, no explicit imagery, no explicit copy in listing or screenshots. Matched-only messaging — never random or anonymous chat, which is now a named Play risk category. 18+ rating, Play Console minor-blocking enabled, Apple's updated age

questionnaire completed. Positioned as "discreet dating", following the vocabulary the incumbent itself moved to in February 2026.

- **Web PWA:** the full product. All explicit media lives here behind V2. All billing lives here.
- Ship the boring compliance basics that cause most rejections anyway: in-app account deletion, a web deletion link, privacy manifests for every SDK, a working reviewer demo account, and a Data Safety form that exactly matches the privacy policy.

14.4 Payments

Dating is structurally high-risk and TRYST will be underwritten as such. Dating services process under MCC 7273, which acquirers treat as high-risk regardless of how cleanly an individual platform runs, and Visa's Integrity Risk Program places adult services and dating in Tier 1 — activity that could harm individuals — requiring acquirer registration before processing. Stripe, Square and PayPal all prohibit adult content; a dedicated high-risk merchant account is required, with typical rates of 3–7%.

The category-specific risk is the one that matters most here: spousal chargebacks — a partner discovers the transaction on a shared bank statement and disputes it out of anger or embarrassment. The discreet nature of the service works against the merchant. Mitigations are product features, not finance admin:

- **Neutral billing descriptor** with a customer service URL. A discreet descriptor reduces privacy-driven disputes by 30–50%.
- **Alternative payment paths:** prepaid vouchers / gift codes purchasable offline. This is the single highest-impact discretion feature in the commerce stack.
- Pre-renewal email notification, genuinely one-click cancellation, and Ethoca/Verifi chargeback alerts to refund before a dispute becomes a formal chargeback.
- **Backup MID** from day one. A single processor pause must not stop revenue.
- Correct MCC coding — mis-coding violates scheme rules and incurs penalties.

14.5 Jurisdiction gating

Hard geo-block, enforced at the edge and at signup, for any market where adultery is criminalised, same-sex conduct is criminalised, or the platform category is prohibited. This protects users from prosecution and the company from liability — it is a safety feature, not a legal formality. Each market requires a written opinion before enabling.

Launch set: UK, IE. **Wave 2:** NL, BE, DE, SE, DK, ES, PT. **Wave 3:** CA, AU, NZ.

Explicit recommendation: no Africa or MENA expansion, at any stage, under any brand. And this entity must sit in a ring-fenced SPV with its own controller registration, its own brand, and no public linkage to the parent group's other ventures. The category's reputational and regulatory risk should not be able to touch fintech, edtech or government-facing work. Separate cap table, separate counsel, separate PR.

15. Business Model

15.1 Principle

The incumbent model charges men to send messages. It is profitable and it is a trap: it rewards the platform for unresolved conversation, it drives the bot and fake-profile accusations that dog the category, and it makes the loss function (\$6.4) fight the P&L. TRYST charges for **agency and discretion**, and makes resolution the profitable outcome.

15.2 Tiers

Tier	Price (UK)	Includes
Verified (free)	£0 + refundable deposit	Full profile, be discovered, reply to anyone free , 3 outbound intents/week. Replying is always free — that is what keeps the scarce side present.
TRYST+	£24.99/mo · £14.99/mo annual	Unlimited intents, filters, incognito, travel mode, extended media, read-receipt control
ENVOY	£49.99/mo	The agent tier. Envoy runs outreach and pre-qualification, produces Briefs, screens risk, negotiates logistics. The thing no competitor sells.
DUO	£39.99/mo (couple)	Linked VC profile, dual consent controls, joint-veto model, third-screening Envoy
Keys (credits)	£9.99 / 60	For non-subscribers: outbound intent = 5 Keys (solo→solo), 15 (couple→solo). Fully refunded on reply . Prices spam, not conversation.
Verification deposit	£9.99 → 100% Keys credit	Economic friction against burners and bots. Costs the honest user nothing.
GHOST add-on	£7.99/mo	Decoy skin, duress PIN, voucher billing, aggressive TTLs, extended exclusion

Tier	Price (UK)	Includes
		ring

Never: advertising, data sale, affiliate placement, engagement-driven paywall interruption. Art. 9 data plus an ad stack is a regulatory suicide pact.

15.3 Illustrative unit economics

Assumption-driven and to be replaced with real cohorts after beta. Shown so the model can be argued with.

Blended CAC (paid + referral, UK)	£38	
Signup → V2 verified	52%	
Verified → payer (m3)	11%	
Blended ARPPU	£31.50 /mo	
Payer gross churn	19% /mo	(avg life 5.3 mo)
Gross revenue per payer (life)	£167	
less PSP (high-risk, 5.0%)	-£8	
less verification (2 checks @ £1.10)	-£2	
less infra + LLM inference	-£14	
Contribution per payer	£143	
CAC per payer (£38 / 0.11 / 0.52)	£664	← paid acquisition alone DOES NOT WORK

∴ Viability requires blended CAC < £12 via referral, invite-loop and organic. Model paid acquisition as a seeding cost, never a growth engine. Any plan that assumes paid CAC scales is wrong.

The one number that decides this business. At an 11% payer rate, paid acquisition is fatal — and this is true for every player in the category, which is why they are invite-driven and word-of-mouth-driven. Two levers only: raise verified→payer toward 18% (ENVOY tier is the instrument), and drive blended CAC below £12 through invite mechanics. Track both weekly from beta day one. Everything else in this document is secondary to these two numbers.

16. Liquidity & GTM

Marketplace liquidity is the only thing that kills this product, and it kills it in month four. Rules, in priority order:

1. **Density gate.** No city opens below 2,000 verified profiles. A thin city produces an empty slate, which produces churn you never recover. Waitlist is a better experience than an empty grid.
2. **Balance gate.** Hard throttle on the oversupplied side to hold the ratio at or below 2.2:1 per city. Excess signups go to a genuine, moving waitlist. This is unpopular internally and correct. The incumbent's roughly 5:1 problem is the direct cause of its bot accusations and its pricing model.
3. **Seed the scarce side first.** S4 and women-identifying users get invite waves, free ENVOY for 6 months, and founding-member status. Open the other side only once they are present.
4. **Invite-only for 9 months.** Each member gets 3 invites, refreshed monthly on good standing. Scarcity is on-brand here in a way it is not in mainstream dating, and invite-driven CAC is the only CAC that works (\$15.3).
5. **Launch cities:** London first (density, anonymity, ENM-literate), then Manchester, Birmingham, Brighton, Dublin. Not nationwide.

6. **Channel:** paid social is largely closed to this category — assume Meta and Google will reject the ad account. Plan for it. Real channels: podcast host-reads in adjacent verticals, ENM/relationship creators, press on the *technology* story (agent handshake, breach-proof architecture) rather than the sex story, and a serious editorial presence.
7. **Positioning line:** "Discretion is the product." Never "cheat better." The first is defensible to a journalist, an acquirer and a regulator. The second is a headline you lose.

17. Roadmap, Team, Budget

Phase	Window	Gate to exit
P0 Legal & foundation	M0–M2	SPV incorporated · UK counsel opinion · DPIA complete · OSA risk assessments drafted · high-risk PSP + backup MID approved · two age-assurance providers contracted · threat model signed off
P1 Core MVP	M2–M6	Closed beta, London, 2,000 verified. Loops A/B/D. E2EE live. V2 gating live. VC couple co-sign live. No Envoy yet — establish the control cohort.
P2 The differentiator	M6–M10	Envoy shipped. Mirror v1. ExclusionRing PSI. Exit gate: meets-per-intent-sent up ≥ 40% vs P1 control. If not met, stop and rethink.
P3 Public UK + IE	M10–M15	Invite-only public. DUO + GHOST. Native apps approved on both stores. 25,000 verified. Payer rate ≥ 11%.
P4 EU wave 2	M15–M24	NL/BE/DE/SE/DK/ES/PT. Per-market legal opinions. Localised assurance. 150,000 verified. Contribution-positive.

Team to P3 (17 people)

PRODUCT & ENG (12)

1 Head of Product · 1 Principal Eng (security-led — hire first) · 2 Backend (Go) · 1 iOS · 1 Android · 1 Web · 2 ML/Agents · 1 Data Eng · 1 Security Eng · 1 Design

TRUST, LEGAL & GROWTH (5)

1 Head of Trust & Safety (hire in P0, not P2) · 2 T&S ops (24/7 rota from P3) · 1 DPO/Counsel (fractional to P2) · 1 Growth/Community

INDICATIVE 24-MONTH BUDGET — People £3.4M · Legal, DPIA, counsel £260k · Age assurance (≈£1.10/check) £180k · Infra + inference £420k · Security audits & pen tests £150k · Seeding & community £480k · PSP reserves and fees £220k · Contingency 15% £780k · **TOTAL ≈ £5.9M**

18. KPIs & Risk Register

North star

QCAM

Quality Connections per Active Month — confirmed meets with a positive aftercare score, per active user, per month. Not DAU. Not messages. Not session length. If a proposed change raises engagement and lowers QCAM, it does not ship.

Metric	P1	P3	Why it matters
QCAM	0.18	0.45	North star
Meets per intent sent	1.8%	4.5%	Envoy's proof of value
Mutual reply rate	22%	34%	Match quality
Inbound Gini (per city, per segment)	< 0.60	< 0.55	Attention collapse guard
Gender ratio (per city)	≤ 2.2:1	≤ 2.0:1	Liquidity health
S4 30-day retention	45%	62%	The scarce side. Leading indicator of everything.
Report rate / 1k threads	< 12	< 6	Safety and regulatory posture
V2 completion	52%	68%	Compliance + funnel
Chargeback ratio	< 0.65% always		Above 0.9% you lose the MID and the business stops
Erasure latency	< 60s p99		The trust promise, measured

Risk register

Risk	Sev	Mitigation
Breach — existential, category-defining precedent	Critical	E2EE, identity split, cryptographic erasure, 30d TTL, zero standing access, quarterly pen test, annual red team scoped to A4/A7, rehearsed 72h playbook
App store rejection or removal under Apple 1.1.4	High	Sanitised store build, explicit media web-only, matched-only messaging, web-first revenue so removal is survivable rather than fatal
Loss of merchant account	High	Backup MID from day one, chargeback ratio alerting at 0.5%, Ethoca/Verifi, one-click cancel, pre-renewal notice, voucher path
Ofcom enforcement	High	Two age-assurance providers, inbox-level gating, documented risk assessments, named accountable manager, quarterly compliance review
Liquidity failure in month 4	High	Density gate, balance gate, scarce-side-first seeding, invite-only — accept slow growth over dead cities
Envoy fails to move meets-per-intent	High	P1 control cohort exists precisely to measure this. Explicit P2 exit gate at +40%. Kill the thesis honestly rather than ship a me-too.
Agent perceived as bot / impersonation scandal	High	Envoy never converses with a human as the user. Every Brief labelled agent-generated. Publish the design.
Harm incident attributed to the platform	High	T&S hired in P0, 24/7 rota from P3, published

Risk	Sev	Mitigation
		escalation SLAs, Safe Meet checklist, transparency reporting
Reputational contagion to the parent group	Medium	Ring-fenced SPV, separate brand, separate counsel, no public linkage (\$14.5)
Feed or Match Group ships an agent layer	Medium	Speed to P2; the defensible moat is the discretion architecture and the aftercare dataset, not the agent itself

19. Admission Control: Screening & Exclusion

The instruction is to keep sex offenders, domestic abusers, fake profiles and declared-unsafe people off the platform. Here is the honest engineering answer: in the UK **you cannot buy that screen**. It has to be built, and it has to be built differently from how the Americans do it.

19.1 What is actually available — the hard constraint

Source	Available to TRYST?	Reality
UK sex offender register (ViSOR)	No	Not publicly accessible. No legal API, no commercial feed, no lawful route for a private company. Unlike the US, where state registries are public, the UK has no equivalent. Any vendor offering you "UK sex offender screening" is selling scraped, unofficial, error-prone data — do not buy it.
DBS checks (Standard/Enhanced)	No	Only available to employers for eligible roles in regulated activity. A dating platform cannot lawfully request one on a user. Tinder explored criminal-record checking and shelved it — this is why.

Source	Available to TRYST?	Reality
Basic DBS	Indirectly — user-initiated only	An individual can request their own Basic DBS (unspent convictions and conditional cautions only). TRYST can invite and fund voluntary self-submission. It cannot compel it, and it covers only unspent convictions.
Clare's Law (Domestic Violence Disclosure Scheme)	Facilitate only — never receive	A Right to Ask held by the individual at risk, made to police, disclosed confidentially to the person best placed to protect. Not redistributable, not transferable to a platform. TRYST can surface it, explain it, and pay for the user's time — and must never ask for or store the answer.
Sarah's Law (Child Sex Offender Disclosure Scheme)	Facilitate only	Same structure, child-safeguarding scoped.
US-style commercial screening (Checkr-class)	Not for UK users	Depends on SSN trace and public registries that do not exist here. Relevant only if wave 3 opens North America, and even then see 19.2.
Criminal-conviction data generally	Restricted	Art. 10 UK GDPR: processing personal data on convictions and offences requires official authority or a lawful condition. DPA 2018 Sch. 1 conditions must be identified in writing by counsel before a single record is touched.

19.2 The cautionary tale — and the rule it produces

Match Group partnered with Garbo to bring sex-offender background checks into Tinder and Hinge. The partnership collapsed after two years and Garbo wound down, citing a lack of commitment from its

partners — and specifically that the dating apps wanted to reduce screening to a **badge on a profile indicating a "clean" check**.

Rule, absolute: TRYST never issues a "clean", "safe" or "background-checked" badge.

The reason is not squeamishness, it is arithmetic. The overwhelming majority of people who will harm another user on this platform have no criminal record at all. A green tick converts *absence of a record* into *presence of safety*, and a user who trusts that tick takes risks they would not otherwise take. A badge system does not just fail to protect — it actively raises the harm rate while transferring liability onto us for a representation we cannot support. It is the worst possible design and it is the one every competitor drifts toward.

19.3 What we build instead: five-layer admission control

L1 — Identity binding: one human, one account, permanently

The highest-yield safety control in this category is not a registry lookup. It is making **ban evasion economically and technically infeasible**. A removed predator who re-registers in four minutes under a new email has not been removed. Binding, at V2:

- Liveness capture with randomised challenge-response (pose/phrase), defeating replay and pre-recorded video
- ID-document match, processed by the assurance provider — TRYST retains the boolean and a one-way anchor, never the document
- Device attestation: Play Integrity / App Attest, emulator and root/jailbreak detection
- Payment-instrument fingerprint (token, never the PAN)
- The £9.99 refundable verification deposit (\$15) — a bot farm needs a distinct real card per account, which is the single cheapest anti-fraud control in the product

L2 — The Ban Anchor: exclusion that survives deletion

```
On permanent removal, write to BAN_ANCHOR (separate store, separate key domain):
anchor_id_doc      = HMAC(doc_number, HSM_pepper_ban)      -- irreversible
anchor_biometric   = HMAC(provider_face_template_ref, ...) -- if lawful basis holds
anchor_device      = HMAC(attestation_id, ...)
anchor_payment     = HMAC(psp_instrument_token, ...)
reason_code        = enum      -- NOT free text, NOT the evidence
tier               = permanent | law_enforcement_referred
reviewed_by        = staff_id   -- human decision, always
expires_at         = NULL for tier=permanent
```

Signup checks every anchor BEFORE account creation. Match → silent refusal.
Never tell a rejected applicant which anchor matched (it teaches evasion).

READ THIS CAREFULLY — IT IS A DECLARED EXCEPTION TO §9

§9 promises cryptographic erasure on deletion. **Ban anchors are the one carve-out**. They must survive account deletion or the entire exclusion system is theatre. The lawful basis is the substantial-public-interest condition for safeguarding individuals at risk (DPA 2018 Sch. 1 Pt. 2) — counsel must confirm the specific condition and it must be written into the DPIA, the privacy notice and the ToS in plain language. Do not bury it. A promise of total erasure that has an undisclosed exception is a worse problem than the exception itself.

L3 — Voluntary attestation, honestly labelled

- **Self-submitted Basic DBS.** User-initiated, TRYST-funded, verified through the issuing route rather than an uploaded PDF (uploads are forgeable). Displayed as a dated factual statement — "*Submitted a Basic DBS dated 14 Mar 2026*" — never as a safety rating.
- **Clare's Law facilitation.** Before any first meet, a one-tap explainer: what the Domestic Violence Disclosure Scheme is, that it is their right, how to make the request to their local force. TRYST never asks for the outcome and has no field to store it.
- **Conduct attestation.** At V2, an explicit, separately-logged declaration: no unspent convictions for sexual or violent offences; not subject to a Sexual Harm Prevention Order, restraining order or non-molestation order. Legally this is weak on its own — but a knowingly false declaration converts a later incident into a clean, evidenced ground for permanent removal and a referral, and it materially strengthens our position in any claim.

L4 — Behavioural risk detection (where the actual yield is)

Most dangerous users have no record. They have a *pattern*, and the pattern is observable in-product long before an incident. This layer is worth more than every registry lookup combined, and it is the one thing a well-built agentic platform can do that a screening vendor cannot.

Model	Detects	Action
M6a Coercive control	Isolation language, monitoring demands, guilt-leverage, jealousy escalation, "don't tell anyone" framing	Score → Guardian review → thread freeze at high confidence
M6b Boundary-violation pattern	Escalation <i>after</i> a stated refusal; repeat requests for declined content; pressure against a hard limit	Send-block + strike. This is the highest-signal predictor of offline harm.
M6c Off-platform pressure	Early, insistent moves to unmonitored channels; refusal to use in-product safety features	Friction + user-facing caution banner
M6d Location / identity pressure	Demands for home address, workplace, real name, or a home-only meet against stated limits	Redaction + warning + strike
M6e Image coercion	Pressure for intimate images; threat language; requests to disable screenshot protection	Immediate freeze, sextortion escalation path (§8.4)
M7b Ban-evasion graph	Writing-style similarity, device/network adjacency, behavioural fingerprint	Shadow-limit → human review → anchor extension

Model	Detects	Action
	against banned cohort	

Targets: M6b recall > 0.90 at FPR < 0.03. Every model in this family is human-reviewed before a permanent action — Art. 22 requires it and good practice demands it.

L5 — Cross-platform consortium

A predator removed from TRYST moves to the next app and vice versa. Join or convene an industry hash-sharing scheme (Online Dating Association as the natural home) exchanging **only** one-way anchors plus a reason code — never evidence, never identity, never message content. Gated on a competition-law and data-protection opinion; a shared exclusion list between competitors is a joint controllership arrangement and needs a proper legal frame. Model it on the NCII hash-sharing schemes that already work.

19.4 Fake-profile prevention stack

SIGNUP	liveness challenge-response · ID-doc match · GAN/deepfake detector on selfie + video · device attestation · emulator and root detection · refundable deposit (real card required)
MEDIA	C2PA provenance where present · reverse-image search · EXIF anomaly detection · face-consistency across the photo set · AI-generation classifier on EVERY upload, not just the first
NETWORK	datacentre-IP and VPN scoring — SCORE, never block. Our own users have legitimate reasons to use a VPN; blocking VPNs on a discretion product is self-harm.
GRAPH	signup-burst clustering · shared device/payment across accounts · referral-chain anomaly
BEHAVIOUR	typing cadence · message templating · copy-paste detection · response-latency uniformity · slate-interaction entropy
ARMS RACE	deepfake detection degrades. Budget a quarterly model refresh and an adversarial red-team retainer. Treat a static detector as a known future failure, not a shipped feature.

19.5 "Declared unsafe" — operational definitions

Moderators cannot act on a vibe. Removal grounds, defined so they are auditable:

1. **Self-declared intent** to coerce, deceive into sexual contact, or harm
2. **Corroborated reports** — two or more independent, unconnected reporters describing consistent conduct
3. **Court order** served on the platform: SHPO, restraining order, non-molestation order
4. **Law-enforcement notification**
5. **Model signal** above threshold, confirmed by human review
6. **Knowingly false conduct attestation** (L3), discovered

Ladder: warn → shadow-limit → suspend → permanent + Ban Anchor → law-enforcement referral. Grounds 1, 3 and 4 skip straight to permanent.

19.6 Due process — because the abuse vector here is specific

On this platform more than any other, a **suspicious or vengeful spouse is a predictable false-reporting vector**, and so is a counterparty who was turned down. Design for it or the exclusion system becomes a weapon.

- Permanent bans require a named human decision-maker and a written reason code. Never fully automated.
- Appeal route with a different reviewer. 10 working day SLA.
- Reporter-reputation scoring: coordinated or retaliatory reporting is itself a removable offence.
- Evidence preserved under legal hold, outside the 30-day TTL, on a documented basis.
- Published transparency report: removals by ground, appeals, overturn rate. Publishing the overturn rate keeps the system honest.

19.7 The safety floor — active protection, not just exclusion

- **Safe Meet:** public-venue-first for any first meet; a check-in timer with silent escalation if missed.
- **Share My Plan:** encrypted, out-of-band share of venue and time with a trusted contact — *designed for someone who cannot tell their partner where they are*. Neutral message wording, no TRYST branding, auto-deleting. This is the single most product-specific safety feature we can build, and no competitor has a reason to build it.
- **Panic:** in-app escalation to 999 with location, distinct from the discretion Burn control (§9.4) so neither is ever hit by mistake.
- **Emergency Curtain override:** safety beats discretion, always. If a user triggers Panic, location and evidence are preserved and released to emergency services regardless of TTL settings. Write this in the ToS in plain words.
- **Post-incident:** victim-support routing, evidence preservation, guaranteed no-retaliation (aftercare is never visible to the counterparty).
- **24/7 Trust & Safety rota** from P3, with the SLAs in §8.4 published, not internal.

19.8 Claims discipline

Never say	Say
"Background checked" / "Safe" / "Vetted"	"Identity verified" / "Age assured"
"No sex offenders on TRYST"	"We permanently remove and block re-registration where we find harmful conduct"
"Screened for domestic violence"	"We detect and act on coercive and abusive behaviour, and we'll show you how to use Clare's Law"
"Guaranteed real people"	"Every member completes liveness and ID verification"

Never say	Say
Green "verified safe" tick	Factual, dated attestations only

Overclaiming here is not just a safety failure — under the CPUT Regulations and the DMCC Act consumer regime it is a misleading commercial practice, and it is the first thing a claimant's solicitor will point at after an incident.

20. Liability & Disclaimer Architecture

THE UNCOMFORTABLE PART, STATED PLAINLY

A blanket "we are not responsible for what happens between users" is **not enforceable in the UK against consumers**, and drafting one is actively counterproductive.

Consumer Rights Act 2015, s.65(1): a trader cannot, by a contract term or a consumer notice, exclude or restrict liability for death or personal injury resulting from negligence. And s.65(2) closes the back door: a person is *not* taken to have voluntarily accepted a risk merely because they agreed to or knew about the term. "Personal injury" expressly includes impairment of mental condition.

s.62: any term causing a significant imbalance to the consumer's detriment, contrary to good faith, is unfair and unenforceable — and a court assesses the clause you actually wrote. An overbroad exclusion does not get read down to something reasonable; it gets struck out entirely, and it reads as evidence of bad faith while it is being struck.

And: OSA duties are statutory. No term can contract out of them.

20.1 What a disclaimer can and cannot do

Objective	Achievable?	Mechanism
Not liable for users' conduct toward each other	Largely yes	Correct characterisation of the service + no assumption of a duty we haven't taken on + user indemnity for economic loss
Not liable for the truth of what users say about themselves	Yes	Explicit no-verification clause, backed by never claiming otherwise (§19.8)
Not liable for consequences to a user's relationship or marriage	Yes	Express clause; no third party has a contract with us and there is no free-standing duty
Cap financial exposure	Yes, if reasonable	Cap at fees paid in the preceding 12 months, floor of £100

Objective	Achievable?	Mechanism
Exclude liability for death/personal injury from <i>our</i> negligence	No. Never.	s.65. Barred outright. The answer is not to be negligent, plus insurance.
Exclude OSA / data-protection obligations	No	Statutory

The strategic insight: your real protection is \$19, not \$20. Liability for a harm on this platform will turn on whether TRYST took reasonable care — whether the screening, detection, response and removal systems were adequate. A well-drafted disclaimer wins you the economic-loss argument. Doing the trust-and-safety work properly is what wins you the negligence argument. Budget accordingly: the T&S hire in P0 is worth more than any clause counsel can draft.

20.2 Draft clause architecture

Drafting instructions for counsel, not final terms. Every clause requires UK consumer-law review before publication.

CLAUSE 1 – NATURE OF THE SERVICE

TRYST is an introduction and discovery platform. We provide the means for adults to find and communicate with one another. We are not an agency, a matchmaker, an escort service, or a party to any arrangement you make. We do not select, endorse, recommend or approve any member for you.

CLAUSE 2 – WHAT WE DO AND DO NOT VERIFY

WE DO verify: that you are an adult; that you are a real, living person; that your photographs are of you; that a couple profile has been independently confirmed by both partners.

WE DO NOT verify: your relationship or marital status; whether a partner knows about or consents to your use of TRYST; your employment, income, or personal circumstances; your criminal record; the truth of anything else you tell another member.

Membership of TRYST is not a statement that a member is safe, honest, available, or suitable for you.

CLAUSE 3 – YOUR CONDUCT AND YOUR DECISIONS

You decide who to speak to, whether to meet, where, and what happens. Those are your decisions and your responsibility. You agree not to hold us responsible for another member's conduct, statements or omissions, and you indemnify us against economic loss arising from your own breach of these terms or of the law.

CLAUSE 4 – CONSEQUENCES OUTSIDE TRYST

We are not responsible for the effect of your use of TRYST on your relationships, marriage, family, employment or reputation. Those consequences are yours. Nobody outside this agreement acquires rights under it.

CLAUSE 5 – FINANCIAL LIMIT

Where we are liable to you for financial loss, our total liability is limited to the greater of (a) the fees you paid us in the 12 months before the event, or (b) £100. We are not liable for loss of profit, opportunity, reputation, or for indirect or consequential loss.

CLAUSE 6 – WHAT WE DO NOT EXCLUDE ← INCLUDE THIS. IT IS LOAD-BEARING.

Nothing in these terms limits or excludes our liability for:

- (a) death or personal injury caused by our negligence;
- (b) fraud or fraudulent misrepresentation;
- (c) our obligations under the Online Safety Act 2023 and data protection law;
- (d) anything else that cannot lawfully be limited or excluded.

CLAUSE 7 – REMOVAL AND PERMANENT EXCLUSION

We may suspend or permanently remove you. Where we remove you for conduct that endangers others, we retain irreversible identifiers so that you cannot rejoin. These are kept after your account is deleted, to protect other people. You may appeal (clause 8).

CLAUSE 8 – APPEAL

Named human review, different reviewer, 10 working days.

CLAUSE 9 – EMERGENCIES OVERRIDE PRIVACY

Our privacy and deletion commitments are strong, but if there is a risk to life or of serious harm, we will preserve and disclose what is necessary to emergency services or the police. Safety comes first.

WHY CLAUSE 6 MAKES CLAUSES 3–5 STRONGER, NOT WEAKER

A limitation clause that carves out what the law says cannot be excluded reads as drafted in good faith and is far more likely to survive the s.62 fairness test. A clause that tries to exclude everything reads as an attempt to overreach and takes the enforceable parts down with it. Founders instinctively want Clause 6 deleted. It is the clause that makes the rest work.

20.3 The in-product safety notice

Separate from the ToS, shown at intake and again before a first meet. Plain language — required by s.62's transparency limb and, more importantly, it is the version anyone actually reads.

Before you meet anyone.

We verify that members are real adults. We do not — and cannot — check anyone's criminal record, their honesty, or whether they are who they say they are once you are off this app. No UK platform can. Meet in public the first time. Tell someone where you are going; Share My Plan does it privately. Trust your instincts and leave if something feels wrong. If anyone pressures you, threatens you, or pushes past a boundary you set, report them — we act on it, and we do it fast.

20.4 Terms to keep out

- Any exclusion of liability for death or personal injury. Barred by s.65 and it taints the whole document.
- "Use at your own risk" as an operative clause. s.65(2) means it buys nothing, and it reads as an abdication.
- Mandatory arbitration with a class-action waiver. Vulnerable to s.62 and likely to attract CMA attention.
- Unilateral, no-notice amendment of material terms. Named in Sch. 2 as indicatively unfair.
- Any term suggesting members are screened, vetted, or checked (§19.8).
- Forfeiture of prepaid Keys on suspension before an appeal concludes.

20.5 The real risk transfer: insurance and structure

The mechanism that actually protects the company is not a clause. It is:

- **Cyber and data liability** — sized against an Art. 9 breach, not a generic SaaS breach. This is the policy that matters most.
- **Technology E&O / professional indemnity**
- **Public and product liability** — the cover that responds where s.65 leaves you exposed
- **D&O** — non-optional given named-senior-manager liability under the OSA
- **Media liability** — defamation and privacy exposure from user content
- **The ring-fenced SPV (§14.5)** — structural containment of the whole risk surface

INDICATIVE ANNUAL PREMIUM AT P3 SCALE: £85k-£140k. Add to the \$17 budget. Brokers will need the DPIA, the OSA risk assessments and this document's §9 and §19 to quote sensibly — good architecture materially reduces the premium, which is one more reason to build it first.

21. Guardian Classifier Spec (M6a-M6e)

Guardian is platform functionality, not an add-on: §5.6 gives it veto over every other agent. This annex specifies the five behavioural models and — the harder problem — how they are trained without causing harm in the process.

21.1 The central technical problem, stated first

CONSENSUAL POWER EXCHANGE LOOKS IDENTICAL TO COERCION AT THE TOKEN LEVEL

TRYST serves people who negotiate dominance, submission, degradation and control *as an agreed activity*. A naive coercion classifier trained on domestic-abuse corpora will flag a perfectly healthy, well-negotiated exchange between two enthusiastic adults, and it will flag it constantly. That failure mode is not a nuisance — it destroys the product for exactly the users the platform exists for, and it teaches moderators to dismiss the model's output, which then misses the real cases.

Therefore Guardian does not classify content. It classifies trajectory. The signal is never *what was asked for*. The signal is **what happened after a boundary was stated**. A request for control following an explicit yes is a negotiated activity. The identical request following an explicit no is a boundary violation. Every model below is built on that distinction, and any model that can be shown to separate the two only by topic vocabulary is rejected at eval.

CORE FEATURE, computed per thread, not per message:

```
consent_state(t) = f(explicit_affirmations, explicit_refusals,
                     limit_declarations, enthusiasm_markers,
                     reciprocity, initiative_balance)
```

RISK fires on: refusal(t) $\wedge$ escalation(t+k) $\leftarrow$ the primary signal
 limit_stated(t) $\wedge$ limit_approached(t+k)
 disengagement(t) $\wedge$ persistence(t+k)
 asymmetric initiative sustained over window

RISK does NOT fire on: intensity, vocabulary, kink taxonomy, power-exchange framing, explicitness

21.2 Model specifications

ID	Target	Positive-class definition	Threshold action
M6a	Coercive control	Isolation pressure, monitoring demands, guilt-leverage, jealousy escalation, secrecy framing directed at the counterparty ("don't tell anyone about us" where the counterparty has not framed it that way), identity erosion	>0.75 score → Guardian queue; >0.92 → thread freeze pending review
M6b	Boundary violation	Escalation after a stated refusal; repeat of a declined request; approach to a declared hard limit; reframing a no as a maybe. Highest-value model in the system.	>0.70 → send-block + warning; second instance → strike; third → suspend
M6c	Off-platform pressure	Early, insistent migration to unmonitored channels combined with refusal to use safety features or resistance to a public first meet	>0.65 → caution banner to counterparty; feeds M4 no-show and M6a
M6d	Identity / location pressure	Demands for home address, workplace, legal name, or insistence on a private-residence first meet against a stated preference	Auto-redact in transit + caution banner + strike at repeat
M6e	Image coercion & sextortion	Pressure for intimate images; threat to share;	Any hit >0.60 → immediate freeze, no gradient.

ID	Target	Positive-class definition	Threshold action
		requests to disable screenshot protection; "I already have" claims; payment demands	Sextortion escalates fast and a warning ladder costs lives.

M6e is the exception to everything. Every other model warns, scores and escalates. M6e freezes on first plausible hit and routes to a human within the hour. The asymmetry is deliberate: a false positive on M6e costs one user one frozen thread and an apology. A false negative has produced fatalities in this category, repeatedly, in adjacent platforms. Tune M6e for recall and accept the precision cost.

21.3 Training data strategy

WHAT WE WILL NOT DO

We will not train on real user conversations scraped from the platform without an explicit lawful basis and consent. We will not purchase abuse-conversation corpora of uncertain provenance. We will not train on survivors' disclosures collected for another purpose. Beyond the data-protection exposure — this is Art. 9 material about identifiable people at their most vulnerable — a model built that way cannot be defended to a regulator, a journalist, or the DV sector whose cooperation we need.

Five permitted sources, in build order

#	Source	Volume target	Notes
D1	Expert-authored synthetic corpus. Written by DV specialists and kink-community educators under paid contract, to a taxonomy we co-design. Crucially includes matched pairs: the same content in a consensual negotiated frame and in a coercive frame.	12,000 threads	The matched pairs are the whole point. They are what forces the model onto trajectory rather than vocabulary. Partner candidates: SafeLives, Women's Aid, Refuge, Galop (LGBT+ abuse), plus paid consultants from the ENM/kink education world.
D2	Published research taxonomies , used as labelling schema rather than data:	—	Gives the label set academic defensibility. Adapt carefully: these

#	Source	Volume target	Notes
	Stark's coercive-control framework, Johnson's typology of intimate partner violence, Biderman's Chart of Coercion, the Duluth power-and-control wheel.		describe cohabiting relationships, not early-stage strangers, and the translation is non-trivial.
D3	Transfer learning from public harassment and toxicity datasets, plus open LLM pretraining.	—	Provides general language competence. Explicitly <i>not</i> a source of the coercion label — public toxicity datasets encode exactly the vocabulary bias we are trying to avoid.
D4	Reported content, with consent. Where a user reports a message and separately opts in to its use for safety-model improvement, that content becomes training data. Consent is granular, revocable, and never a condition of the report being acted on.	Grows from P3	The highest-quality source and the slowest to accumulate. Lawful basis needs its own counsel sign-off. Reporter identity is stripped before it reaches the labelling queue.
D5	Active learning on the moderator queue. Uncertainty sampling surfaces the cases the model is least sure about for human label, which is where marginal accuracy	Continuous	Prioritise near-threshold cases and appeal overturns. An overturned ban is the most informative label in the whole system.

#	Source	Volume target	Notes
	actually comes from.		

Bootstrapping sequence

P1	rules + D3 transfer baseline (intentionally crude; exists to generate the moderator queue)	→ recall ~0.55, precision ~0.30
P2	+ D1 expert corpus, matched pairs + D2 taxonomy-aligned labels	→ recall ~0.82, precision ~0.61 consent-state feature introduced
P3	+ D4 consented reports, + D5 active	→ recall ~0.91, precision ~0.74
P4	+ multilingual (EU wave 2)	→ per-language eval, no shared threshold

21.4 Annotation

- **Trained annotators only**, briefed by DV specialists and by kink educators. An annotator who reads negotiated power exchange as abuse will poison the label set, and this is the most common failure in practice.
- **Triple annotation** on the coercion classes, with Krippendorff's $\alpha > 0.70$ required before a batch is accepted. Below that, the taxonomy is wrong, not the annotators.
- **Adjudication** by a senior annotator on disagreement; disagreements are themselves logged as a signal that the boundary is genuinely ambiguous, which is useful.
- **Annotator welfare is a budget line, not a nicety.** Rotation caps, mandatory breaks, clinical supervision available on demand, and a hard limit on continuous exposure to M6e material. Content-moderation trauma is well documented and it is both an ethical obligation and a retention problem.

21.5 Fairness — the failure modes to test for explicitly

Bias	Failure	Test
Kink vocabulary	Negotiated D/s flagged as coercion	Matched-pair eval set (D1). Model must score the consensual variant below threshold and the coercive variant above, on identical vocabulary. Ship gate.
Dialect and register	AAVE, MLE, regional and working-class register over-flagged — a documented, repeatable failure of toxicity models	Stratified eval; FPR parity within 1.5× across dialect groups
Non-native English	Directness read as aggression; missing hedges read as pressure	Eval slice by stated language background; parity check
Gender	Identical message scored differently by sender gender; female-perpetrated coercion under-detected	Counterfactual swap test on the full eval set

Bias	Failure	Test
Orientation / structure	Queer and polyamorous relationship language misread; "my partner and I" patterns flagged	Segment-stratified eval across S1-S6
Age	Older users' register over- or under-flagged	Age-band eval slice

Quarterly fairness audit, published internally, reviewed by the Head of Trust & Safety. A model that passes accuracy and fails the matched-pair gate does not ship.

21.6 Serving, thresholds and human-in-the-loop

LATENCY M6a-M6e inline, < 80 ms p95, batched per thread turn
RUNS ON plaintext ON DEVICE for E2EE threads – the server cannot read messages (§9.2), so Guardian's inline classifiers run client-side and emit only a SCORE and a REASON CODE to the server. Message content never leaves the device for scoring. Full content is transmitted only when the USER REPORTS, which is an explicit act of disclosure.

ESCALATION score → server → Guardian queue → HUMAN
 NO permanent action is ever taken by a model alone (Art. 22)

THRESHOLDS per-model, per-locale, per-segment. Never a single global number. Reviewed monthly against overturn rate.

CALIBRATION isotonic regression on a held-out set; scores must be interpretable as probabilities or moderators cannot triage

The on-device constraint is a feature, not a compromise. Because §9 makes server-side message reading impossible, Guardian has to run client-side. That means the platform can say, truthfully: we detect abuse without reading your conversations. It also means an attacker who breaches the server gets scores and reason codes, not the content behind them. Build it this way from P1 — retrofitting client-side inference onto a server-side pipeline is a rewrite.

21.7 Adversarial resistance

- **Red-team retainer**, quarterly, with a brief to evade each model: euphemism substitution, deliberate misspelling, image-embedded text, code-switching, splitting a coercive sequence across sessions or across multiple accounts.
- **Cross-message and cross-session windows.** A threat split across six messages evades any per-message classifier. Score the trajectory over a rolling window, not the turn.
- **Cross-account correlation.** The same actor running the same script against twenty people is the strongest possible signal and is invisible to any single-thread model. This lives server-side on reason codes alone.
- **Drift monitoring.** Weekly distribution checks on score, report rate and overturn rate. A sudden drop in positives usually means evasion has been discovered and shared, not that behaviour improved.

21.8 Evaluation gates

Gate	Requirement
Matched-pair separation (D1)	> 0.85 AUC separating consensual from coercive on identical vocabulary. Hard ship gate.
M6b recall	> 0.90 at FPR < 0.03
M6e recall	> 0.95, precision floor 0.35 accepted
Dialect FPR parity	within 1.5× across groups
Counterfactual gender swap	mean score delta < 0.05
Appeal overturn rate	< 8% of permanent removals, trending down
Moderator agreement with model	> 75% on queued items — below that the model is wasting human time

21.9 What Guardian is never allowed to do

- Take a permanent action without a named human decision (§19.6, ToS cl. 17.2).
- Score a user on protected characteristics, or on inferred ones.
- Produce a public or counterparty-visible risk score. A visible score is a retaliation and coercion tool, and it would convert our detection into the "safe badge" failure of §19.2.
- Read message content server-side. Ever.
- Flag on kink vocabulary, explicitness or intensity absent a trajectory signal.

ONE-PARAGRAPH SUMMARY FOR THE BOARD

TRYST is a discretion-first connection platform for attached adults and couples, built on an agent-to-agent handshake that resolves intent, boundaries and logistics before two people ever speak — and on a breach-assumed architecture that makes the category's defining disaster technically impossible to repeat. The category leader in affairs has just vacated the position under regulatory and payments pressure; the leader in non-monogamy is a browsing product with a resolution problem. The opportunity is real. The difficulty is not the software — it is the compliance envelope, the marketplace liquidity, and a CAC model that only works through invitation. Build the envelope first, seed the scarce side first, and let the agent layer earn its place against a measured control.

TRYST — PRODUCT & TECHNICAL SPECIFICATION v1.0

17 SEPTEMBER 2026 · CONFIDENTIAL

All regulatory, payment and jurisdictional statements require confirmation by qualified UK counsel before build commences. Unit economics are illustrative and assumption-driven; replace with measured cohorts after closed beta. Third-party market figures are as reported by the sources cited and have not been independently audited.

Part B — Terms of Service

Draft 0.9 — for counsel review. Not for publication.

Prepared by: Product. For review by: UK consumer, Online Safety Act and data-protection counsel.

Governing law assumed: England & Wales.

Status: Not legally reviewed. Do not publish or rely on in this form.

NOTE TO COUNSEL

This draft has been written by the product team to be reviewed, not to be approved. Clauses carrying the highest legal risk are marked with a **REVIEW** — box and a specific question. Three areas need particular attention: **clause 15** (retention of exclusion identifiers after account deletion — the lawful-basis question), **clause 16** (the subscription regime, whose commencement date is unsettled), and **clauses 21–23** (the liability structure, drafted deliberately narrow so it survives s.62 rather than broad so it gets struck).

1. These terms, and who we are

1.1 TRYST is operated by [SPV NAME] LIMITED, a company registered in England and Wales (No. [—]), registered office [—]. In these terms, "we", "us" and "our" mean that company. "You" means the person using TRYST.

1.2 These terms are a contract between you and us. They apply whenever you use the TRYST website, apps or any related service.

1.3 They are written to be readable. Where a clause matters to you, we have added a plain-English box explaining it. Those boxes are part of these terms and can be relied on.

1.4 Our Privacy Notice, Community Rules and Safety Policy form part of these terms.

2. Definitions

2.1 Member — a person with a TRYST account. **Couple Profile** — a profile representing two Members who have each independently confirmed the link. **Keys** — the prepaid credits described in clause 16. **Intent** — an outbound expression of interest. **Brief** — the agent-generated summary described in clause 7. **Exclusion Identifier** — the irreversible identifier described in clause 15.

3. Who can use TRYST

3.1 You must be 18 or over. We verify this using age-assurance methods that meet the "highly effective" standard under the Online Safety Act 2023. If we cannot verify that you are an adult, you cannot use TRYST, and there is no appeal against that.

3.2 You must have the legal capacity to enter this contract, and you must not be prohibited from using a service of this kind by any court order, licence condition or notification requirement.

3.3 TRYST is available only in the countries listed at [URL]. We block access elsewhere, including in countries where using a service of this kind could expose you to criminal liability. You must not attempt to circumvent that block.

3.4 You may hold one account. If you previously held an account that we permanently closed under clause 17, you may not open another.

4. What TRYST is

4.1 TRYST is an introduction and discovery platform. We provide software that lets adults find one another, communicate, and decide for themselves whether to meet.

4.2 We are not a matchmaking agency, an introduction agency acting on your behalf, an escort or companionship service, or a party to anything you arrange with another Member. We do not act as your agent and we owe you no duty to find you anyone.

4.3 We do not select, endorse, recommend, approve or vouch for any Member. Our software ranks and suggests profiles; that is a computational output, not a recommendation from us that a person is suitable, honest or safe.

IN PLAIN ENGLISH — We introduce people. What happens after that is between you and them. Nothing we show you is us telling you that someone is a good idea.

5. What we verify — and what we do not

5.1 We do verify, before a Member can send an Intent or exchange media: that they are an adult; that they are a real, living person, by liveness check; that their photographs are of them; and, for a Couple Profile, that both partners have independently confirmed the link from their own separate devices.

5.2 We do not verify, and cannot verify:

- your relationship or marital status;
- whether a partner of yours knows about, or consents to, your use of TRYST;
- your criminal record, or whether you are subject to any order or notification requirement;
- your employment, income, health or personal circumstances;
- the truth of anything you or any other Member says.

5.3 There is no lawful route by which a UK platform can check a Member against the sex offenders register or obtain a criminal-record check on them. We do not claim to do so, and you should treat any platform that claims otherwise with caution.

5.4 Membership of TRYST is not, and must never be read as, a statement by us that a Member is safe, honest, available, unattached, or suitable for you.

IN PLAIN ENGLISH — We check you are a real adult who looks like your photos. We do not check your marriage, your record, or your honesty. No UK app can check a criminal record on you — and we will never pretend we have.

REVIEW — Clause 5 is drafted to be the anchor for the whole liability structure and to keep us clear of the CPUT Regulations and the DMCC Act consumer regime on misleading practices. Please confirm the wording is strong enough to defeat an argument that our verification badges, ranking, or marketing create an implied representation of safety.

6. Your account

6.1 Everything you tell us about yourself must be accurate, and your photographs must be current and of you. Profiles that misrepresent who you are will be removed.

6.2 Keep your access credentials, PIN and device secure. Tell us immediately if you think someone else has access to your account.

6.3 You are responsible for everything done through your account, unless someone gained access through our failure to take reasonable care.

6.4 You may not sell, transfer, rent or share your account.

7. Couple Profiles, and our agent features

7.1 A Couple Profile requires both partners to hold their own verified accounts and each to confirm the link from their own device. We will not create a Couple Profile on one person's say-so. Either partner may dissolve the link at any time, and we will act on that immediately without telling the other partner why.

7.2 Some features use automated agents. Before you and another Member speak, our agents may exchange structured information about intentions, boundaries, availability and verification, and produce a Brief for each of you.

7.3 Our agents never pose as you, and never speak to another Member as if they were you. Every Brief is labelled as automatically generated. A Brief is information to help you decide; it is not advice, and it is not a judgement about whether someone is safe.

7.4 A conversation only opens when both of you have seen your Brief and accepted.

IN PLAIN ENGLISH — Our software does the awkward pre-screening so you don't waste weeks. It never pretends to be you, and it never tells you someone is safe.

8. Community Rules — what you must not do

You must not, on TRYST or in any contact that begins on TRYST:

- be under 18, or attempt to contact anyone you believe or suspect is under 18;
- share, request, or threaten to share intimate images of anyone without their consent;
- pressure anyone for images, contact, money, or sexual activity, or continue after someone has said no;
- threaten, harass, stalk, intimidate, or attempt to coerce or control another person;
- misrepresent your identity, age, or appearance, or use another person's images;

- offer or seek sexual services in exchange for money or anything of value;
- attempt to identify, locate, contact, or expose a Member outside TRYST without their agreement, or disclose that someone is a Member to anyone;
- solicit money, investment, or financial information;
- record, screenshot or copy another Member's private content or messages;
- use TRYST for any commercial, promotional or research purpose;
- scrape, reverse-engineer, or automate access to TRYST;
- evade a suspension or removal.

8.1 Exposing a Member's use of TRYST to a third party is one of the most serious things you can do on this platform, and we treat it as grounds for immediate permanent removal.

9. Your conduct and your decisions

9.1 You decide who to speak to, what to tell them, whether to meet, where, when, and what happens when you do. Those are your decisions.

9.2 You are responsible for your own conduct. Another Member is responsible for theirs. We are not responsible for what any Member says, does, or fails to do — subject always to clause 22.

9.3 You accept that other Members may misrepresent themselves despite our checks, and you agree to take reasonable care of your own safety, including using the safety features in clause 11.

10. Consequences outside TRYST

10.1 We are not responsible for the effect of your use of TRYST on your relationships, marriage, family, employment, finances or reputation. Those consequences are yours.

10.2 No person who is not a party to this contract acquires any right under it. The Contracts (Rights of Third Parties) Act 1999 does not apply.

REVIEW — Clause 10.2 is intended to head off a claim by a Member's spouse or partner. Please advise whether it is sufficient and whether any free-standing duty could arise in tort notwithstanding it.

11. Safety features — and their limits

11.1 We provide safety features including Safe Meet guidance, Share My Plan (which privately tells a contact of your choosing where you are going), a check-in timer, in-app reporting, blocking, and a Panic control that contacts emergency services.

11.2 These features are aids. They depend on your device, your network, and third parties we do not control. **We cannot guarantee that any of them will work at the moment you need it, and you must not rely on any of them as your only safeguard.** If you are in danger, call 999.

11.3 Using a safety feature does not make a meeting safe, and choosing not to use one does not make what happens to you your fault.

REVIEW — Clause 11.2 is a statement about technical limits, not an exclusion of liability for negligent design or operation of the features themselves — which we accept we cannot exclude under CRA 2015 s.65 where personal injury results. Please confirm the drafting does not stray across that line, since a term that did would be void and would jeopardise clause 21.

12. Your content

12.1 You keep ownership of everything you upload. You give us a limited, non-exclusive licence to host, store, encrypt, transmit and display your content solely to operate TRYST for you. That licence ends when your content is deleted, save as set out in clause 15.

12.2 Your private messages are end-to-end encrypted. We relay them; we cannot read them. This means we cannot recover them for you, and it also means we cannot moderate their content directly — we rely on your reports and on signals available to us outside the message content.

12.3 Media you send is encrypted on your device. You can revoke access at any time, which makes it permanently unreadable to the recipient. We cannot stop someone photographing their own screen with another device, and you should assume anything you send may be kept.

IN PLAIN ENGLISH — We genuinely cannot read your messages. That protects you, and it also means we can only act on what you report. Assume any photo you send could be kept, whatever our software does.

13. Intimate images

13.1 Sharing or threatening to share an intimate image of someone without their consent is a criminal offence. On TRYST it results in immediate permanent removal, preservation of evidence, and referral to the police.

13.2 If an intimate image of you has been shared without your consent, report it to us. We will act within the time stated in our Safety Policy, support you in submitting a hash to StopNCII where you wish, and preserve evidence for any investigation.

14. Your personal data

14.1 Information about your sex life and sexual orientation is special category data under the UK GDPR. We process it only on your explicit consent, which you give separately from accepting these terms, and which you can withdraw at any time.

14.2 Our Privacy Notice explains what we collect, why, how long we keep it, and your rights. It is not a summary — it is the operative document, and it prevails over anything in these terms about personal data.

14.3 We do not sell your data, we do not show advertising, and we do not embed third-party tracking in our apps.

15. Keeping and deleting your information

15.1 You can delete your account at any time. When you do, we destroy the encryption key that protects your data, which makes it permanently unreadable. This happens in under a minute and cannot be reversed. We cannot restore a deleted account.

15.2 Messages are deleted after 30 days by default, or after 7 or 90 days if you choose. Media is deleted after 7 days by default.

15.3 There is one exception, and we want you to read it. If we permanently remove you because your conduct endangered another person, we keep a small set of **Exclusion Identifiers** — irreversible one-way codes derived from your identity document, device and payment method. We cannot reverse them into your personal details. We keep them so that you cannot open another account, and we keep them **after your account is deleted**. We do this to protect other people. We keep no evidence, no messages and no profile alongside them — only the codes, a category of reason, and the date.

15.4 If you are removed on that basis you may appeal under clause 18, and if your appeal succeeds we destroy the Exclusion Identifiers.

IN PLAIN ENGLISH — Delete your account and it is genuinely gone. The one thing that survives is a code that stops a removed predator coming straight back under a new name. If that is you and we got it wrong, appeal and we delete it.

REVIEW — This is the highest-risk clause in the document. We need: (a) the specific DPA 2018 Sch. 1 condition relied on — we have assumed Pt. 2 para. 18 (safeguarding of individuals at risk) — confirmed in writing; (b) a view on whether retention post-erasure is defensible against an Art. 17 request; (c) confirmation that the disclosure here is sufficiently prominent, given that the rest of clause 15 promises complete erasure. We have deliberately refused to bury this. Please tell us if it needs to be more prominent still.

16. Paying for TRYST

16.1 Basic membership is free. Replying to any Member is always free. Paid tiers and Keys are described at [URL] with prices shown before you pay.

16.2 Billing discretion. Your card or bank statement will show a neutral descriptor that does not name TRYST. You can also pay using prepaid vouchers bought elsewhere. We do this because your privacy matters, not to obscure what you are paying for — the descriptor includes a customer service address.

16.3 Cancelling. You can cancel a subscription at any time, from inside the app or website, in one step. We will not ask you to call us, email us, or answer questions to cancel. Cancellation takes effect at the end of your current paid period.

16.4 Your 14-day right to cancel. Under the Consumer Contracts Regulations 2013 you normally have 14 days to cancel a distance contract. If you ask us to start the service immediately, you may lose that right for digital content, or owe us a proportionate amount for services already supplied. We tell you this clearly at checkout and ask you to confirm it separately.

16.5 Reminders. We email you before any renewal and before any trial converts to a paid subscription. The reminder says what you will be charged, when, and how to stop it.

16.6 Keys. Keys spent on an Intent are refunded in full if the Member replies. Unused Keys do not expire while your account is open. If we permanently remove you, unused Keys are refunded unless removal was on a ground in clause 17.3.

16.7 Refunds. Where we owe you a refund we pay it to the original payment method within 14 days.

REVIEW — The DMCC Act 2024 subscription regime is not yet in force and its commencement date has moved repeatedly — announced for spring 2026, then autumn 2026, then spring 2027, with more recent reporting suggesting it may be brought forward again. We have drafted to comply with the full regime now regardless, on the view that pre-contract information, reminder notices, cooling-off notices, end-of-contract notices and one-step cancellation are all things we want anyway and that retrofitting a payments flow is expensive. Please confirm the current commencement position and whether clause 16 already meets it, particularly on the form and prominence of the notices. Note CMA fining powers of up to 10% of global turnover.

17. Suspension and removal

17.1 We may warn you, limit your account, suspend it, or remove you permanently if you break clause 8, these terms, or the law, or if we reasonably believe you present a risk to another person.

17.2 We will tell you what ground we acted on, unless telling you would put someone at risk or prejudice an investigation. A permanent removal is always decided by a named person at TRYST, never by software alone.

17.3 Immediate permanent removal, with Exclusion Identifiers retained under clause 15.3, applies where: you are under 18 or sought contact with someone you believed to be; you shared or threatened to share intimate images without consent; you coerced, threatened or attempted to control another person; you exposed a Member's use of TRYST; you are subject to a relevant court order; or you knowingly gave a false declaration under clause 3.2.

17.4 We may preserve evidence relating to a removal beyond the periods in clause 15, where we need it for an investigation, an appeal, or a legal claim.

17.5 Reporting another Member falsely, or organising others to report someone, is itself grounds for removal.

18. Appeals

18.1 You can appeal any suspension or removal by [route]. A different person from the one who made the decision will review it, and we will respond within 10 working days.

18.2 If we were wrong, we restore your account, refund anything you lost, and destroy any Exclusion Identifiers.

18.3 We publish, twice a year, how many accounts we removed, on what grounds, how many appealed, and how many appeals succeeded.

19. Reporting, and what we do about it

19.1 Report anything from inside any conversation or profile. Reporting takes one tap and you never have to explain yourself to the person you are reporting.

19.2 Our response times for each category of report are published in our Safety Policy. We do not hide them.

19.3 We will never tell the person you reported that it was you.

20. Emergencies come before privacy

20.1 Our privacy and deletion commitments are strong. But if there is a risk to someone's life or of serious harm, or where the law requires it, we will preserve and disclose whatever is necessary to the emergency services, the police, or a court.

20.2 We publish what we disclose, in aggregate, in our transparency report.

IN PLAIN ENGLISH — If someone's life is at risk, safety wins over discretion. Every time.

21. Our liability to you

21.1 We will supply TRYST with reasonable care and skill, as the Consumer Rights Act 2015 requires. Nothing in these terms limits that obligation.

21.2 We are not liable to you for the acts or omissions of another Member, or for anything that happens between you and another person, except where our own failure to take reasonable care caused it.

21.3 Where we are liable to you for financial loss, our total liability is limited to the greater of (a) the amount you paid us in the 12 months before the event, or (b) £100.

21.4 We are not liable for loss of profit, loss of opportunity, loss of reputation, or any indirect or consequential loss.

21.5 We are not liable for failures caused by events outside our reasonable control, or by your device, network or operating system.

22. What we do not exclude

22.1 Nothing in these terms limits or excludes our liability for:

- death or personal injury caused by our negligence, including any impairment of mental health;
- fraud or fraudulent misrepresentation;
- any breach of the obligations the Consumer Rights Act 2015 says we cannot exclude;
- our duties under the Online Safety Act 2023 and data protection law;
- anything else that cannot lawfully be limited or excluded.

22.2 Your statutory rights as a consumer are not affected by these terms. Agreeing to these terms does not mean you have accepted any risk.

REVIEW — Clause 22 is deliberately included and deliberately broad. We understand it removes any argument that we tried to exclude what CRA 2015 s.65 bars, and that s.65(2) means a "use at your own risk" framing would buy us nothing anyway. Our working assumption is that a narrow, honestly-carved limitation survives the s.62 fairness test where an all-encompassing one would be struck out in full, taking clause 21 with it. Please confirm — and please push back if you think clause 21.3's cap is too low to be reasonable given what Members pay.

23. Your liability to us

23.1 If you break these terms or the law and that causes us financial loss — including a claim by another Member, a fine, or the cost of responding to an investigation — you agree to cover that loss.

23.2 This covers financial loss only. It does not apply where the loss was caused by our own failure to take reasonable care, and it is limited to what is reasonably foreseeable.

24. Changing these terms

24.1 We may change these terms. If a change materially affects you, we will tell you at least 30 days before it takes effect and explain what has changed and why.

24.2 If you do not accept a material change, you can close your account before it takes effect and we will refund any unused prepaid amount in full.

24.3 We will not change these terms retrospectively.

25. Ending this agreement

25.1 You can close your account at any time, from inside the app, in one step.

25.2 Clauses 9, 10, 12, 13, 15, 17.4, 21, 22 and 23 survive the end of this agreement.

26. Complaints

26.1 Complain to [email]. We acknowledge within 2 working days and respond substantively within 10.

26.2 If you are not satisfied, you may refer the matter to [ADR provider], or contact Citizens Advice. Nothing here stops you going to court.

26.3 Our Online Safety Act complaints route, and the name of the senior manager accountable for it, are at [URL].

27. Law and courts

27.1 These terms are governed by the law of England and Wales.

27.2 You can bring proceedings in the courts of England and Wales. If you live in Scotland or Northern Ireland, you can also bring proceedings there. If you live elsewhere in a country where consumer law gives you the right to sue locally, that right is not affected.

27.3 There is no arbitration clause and no class-action waiver in these terms.

REVIEW — We have deliberately not included mandatory arbitration or a class-action waiver. Our view is that both are vulnerable under CRA 2015 s.62 and Sch. 2 in a consumer contract, and that including them invites CMA attention for limited benefit. Please confirm or overrule.

28. Contacting us

28.1 General: [email]. Safety and urgent reports: [route], monitored 24 hours.

28.2 Data protection: our Data Protection Officer at [email]. You can also complain to the Information Commissioner's Office.

28.3 Online Safety Act accountable senior manager: [name/role], [email].

29. Open questions for counsel

#	Question	Clause
Q1	Confirm the DPA 2018 Sch. 1 condition for retaining Exclusion Identifiers post-erasure, and its resilience against an Art. 17 request.	15.3
Q2	Does any part of clause 11 stray into excluding liability for negligent design of a safety feature, which would be void under s.65 and could taint clause 21?	11.2
Q3	Is the £100 / 12-months cap reasonable, or does the low free-tier ARPU make it vulnerable under s.62?	21.3
Q4	Current commencement position for the DMCC subscription regime, and whether clause 16 meets it in full.	16
Q5	Could a Member's spouse establish a free-standing duty of care in tort	10.2

#	Question	Clause
	notwithstanding clause 10.2?	
Q6	Do our verification badges, ranking outputs or marketing create an implied representation of safety that clause 5 fails to displace?	5
Q7	Does clause 12.2 (we cannot read E2EE messages) weaken our OSA position, and how should the risk assessment address content moderation without content access?	12.2
Q8	Is a cross-platform exclusion consortium (spec §19.5) viable under competition law and joint controllership?	—
Q9	Confirm the geo-block list and whether operating in any listed market creates exposure for directors.	3.3
Q10	Review the Community Rules and Safety Policy as incorporated documents for consistency with these terms.	1.4

TRYST — TERMS OF SERVICE · DRAFT 0.9 · 17 SEPTEMBER 2026

PREPARED BY PRODUCT FOR LEGAL REVIEW · NOT LEGAL ADVICE · NOT FOR PUBLICATION

This draft reflects the product team's understanding of the Consumer Rights Act 2015, the Consumer Contracts Regulations 2013, the Digital Markets, Competition and Consumers Act 2024, the Online Safety Act 2023, the UK GDPR and the Data Protection Act 2018. It has not been reviewed by a qualified lawyer. Statutory positions, particularly commencement dates, must be verified as at the date of publication.