



Product & Technical Specification

Version 1.2 — Baseline · 4 October 2026 · Confidential

Private chemistry. Intelligent discretion.

TRYST Specification v1.3 — Baseline

Private chemistry. Intelligent discretion. · *Discretion is the product.*

Version	1.3 — Baseline (stable). v1.2 added the brand identity and account-holder-only login; v1.3 adds the internal operating system (09; 06 §4)
Date	4 October 2026
Status	Stable for engineering planning and estimation. Build start remains gated on the written UK legal opinion and completed DPIA (Phase P0; see 01 §14).
Owner	Office of the Group Chief Executive → TRYST SPV (pre-incorporation)
Classification	Confidential

1. Document set

The specification is split by audience. Each document owns its topics; the others link to it rather than restating it.

#	Document	Audience	Owns
00	README (this file)	Everyone	Index, source precedence, change control
01	Product	Leadership, product, investors, counsel	Thesis, segments, modes, principles, loops, agent roles, business model, unit economics, GTM, roadmap, team, budget, KPIs, risks, compliance envelope, admission policy, liability architecture
02	Shared Contracts	Backend and frontend	Glossary, enums, verification tiers and gates, data model, API contract, event taxonomy, agent message schemas, retention, error model, requirement catalogue (FR/NFR) with owners
03	Backend	Backend, ML, security, SRE, T&S engineering	Services, stack, agents, learning models, matching pipeline, cryptography, identity split, erasure, PSI server, Guardian server side, admission control, MLOps, commerce, backend backlog
04	Frontend	iOS, Android, web, design	Two-surface strategy, client stacks, screens and flows, client discretion, client cryptography, on-device Guardian, PSI client, safety UX, copy rules, store compliance, frontend backlog
05	Terms of Service — draft 0.9.3	Counsel	Consumer terms for counsel review. Not for publication.
06	Decisions, errata and change log	Product owner, tech leads	How v1.1 reconciled its sources, resolved and open decisions, errata fixed
07	Source B — founder brief (verbatim) and traceability	Everyone	The founder's 17 Sep 2026 specification word for word, plus where each part is implemented, the agent mapping and the FR-ID crosswalk
08	Source A — traceability	Everyone	Where each section of the consolidated Spec + ToS lives in the spec and the code
09	Operating System (internal)	Leadership, engineering, security, payments, operations	Agent workforce and governance (autonomy levels, always-human actions, kill switch, audit chain), command centres for every user type, self-managing platform layer, BitriPay API door, connector register, ops and partner schemas and APIs, B2B licensing and ACU, Admin Super Control Centre, roadmap

A combined PDF of 00–09 is at [TRYST_Spec_v1.3.pdf](#). Brand assets (logo exactly as supplied, colour tokens) are in [assets/brand/](#).

2. Sources and precedence

v1.1 merges three inputs. Where they conflict, the higher-ranked source wins unless 06 records a decision otherwise.

Rank	Source	File
A	Product & Technical Specification and Terms of Service — consolidated edition (4 Oct 2026; Part A §1–21, Part B ToS draft 0.9)	source/TRYST_Spec_and_ToS_Consolidated_2026-10-04.pdf
B	AI Discreet Relationship Platform — Developer Specification (17 Sep 2026)	source/TRYST_Developer_Specification_2026-09-17.docx

Rank	Source	File
C	PRD & Business Plan draft v0.1 (4 Oct 2026)	archive/TRYST_PRD_v0.1.md

Both A and B call themselves "v1.0", so v1.1 refers to them as **Source A** and **Source B**. In v1.1, references such as [A §9.3] point to the source section a requirement came from.

3. Identifier conventions (stable across documents)

Prefix	Meaning	Defined in
S1-S6	Member segments	02 §2.2
V0-V3, VC	Verification tiers	02 §3
A1-A8	Threat-model adversaries	03 §6.1
M1-M9, M6a-M6e, M7b	Models	03 §4.3
L1-L5	Admission-control layers	03 §8
FR-nnn, NFR-nn	Functional and non-functional requirements	02 §10
D-nn	Decisions	06 §2
L0-L3 (agents)	Agent autonomy levels	09 §5.2
E-nn	Errata fixed in v1.1	06 §3
G-xx	Phase exit gates and no-go conditions	01 §12-13

IDs are never reused or renumbered. A retired requirement keeps its ID with status Retired.

4. Change control (stabilisation rules)

- 1. **v1.1 is a baseline.** Changes go through a pull request that edits the owning document and adds a row to 06 §4.
- 2. **Contract changes come first.** Any change to an API, entity, enum, event or schema is made in 02 Shared Contracts before 03 or 04 depend on it.
- 3. **Open decisions** (D-nn, status *Open*) carry a stated default. Engineering may build against the default, and the owner must close the decision before the phase gate shown.
- 4. **Consistency check.** python3 scripts/check_spec.py checks that every FR/NFR/D/E ID referenced anywhere is defined exactly once, every FR has an owner, and relative links resolve. It must pass before merge.
- 5. **Version bumps:** editorial → 1.1.x; new or changed requirement → 1.2; change to a principle in 01 §4 → 2.0 (requires sign-off from the Group CEO).

5. Source register (revalidate before launch)

Policies, prices and legal requirements must be revalidated before launch, and again at each phase gate.

- **Source B's register** (accessed 17 Sep 2026) is reproduced verbatim in 07 §3: Ashley Madison, Gleeden, Victoria Milan, the FTC Ashley Madison settlement, ICO special-category guidance, UK GDPR Art. 9, Apple App Review Guidelines, Google Play inappropriate-content policy, NIST AI RMF and OWASP MAS.
- **Source A's citations** (Ofcom OSA and dating guidance, CRA 2015, CCR 2013, DMCC Act 2024, DPA 2018, card-scheme rules) are in the consolidated PDF.

Added for v1.2 — not yet checked, verify in P0:

Topic	Source to check	Why
Stripe acceptance	https://stripe.com/legal/restricted-businesses and a written confirmation from Stripe	D-22: Stripe as primary processor (could not be fetched from the build environment)
Search-engine spam and AI-content policies	https://developers.google.com/search/docs/essentials/spam-policies	D-21: Herald link and content rules
Passkeys / WebAuthn	https://www.w3.org/TR/webauthn-3/ and https://fidoalliance.org/passkeys/	D-16: login factors
Liveness PAD certification	ISO/IEC 30107-3; the provider's iBeta (or equivalent) Level 2 report	NFR-17
Biometric data	ICO biometric data guidance	D-17, D-18, ToS Q14

Topic	Source to check	Why
BitriPay API and webhook formats	The BitriPay API reference and OpenAPI document (supplied developer docs, 4 Oct 2026)	D-26: refund field names and the BitriPay-Signature layout are marked ASSUMED in the adapter
EU AI Act obligations for P4 markets	Regulation (EU) 2024/1689 and Commission guidance	09 §13.4

01 — TRYST Product Specification v1.1

Part of the TRYST v1.1 baseline. Contracts live in 02; implementation in 03 and 04.

1. Executive summary

PRODUCT SCOPE

TRYST is designed specifically for adults who are already in a relationship and want a discreet one-off affair, a continuing extramarital connection, another couple, or a threesome. Existing-partner permission is not a condition of membership. Every person using TRYST must nevertheless be an independently consenting adult; the platform prohibits coercion, stalking, blackmail, impersonation, minors, non-consensual imagery and commercial sexual services.

Founder scope statement, verbatim (Source B; re-confirmed 4 Oct 2026). Applied through §4, §5 and D-01.

TRYST is a discretion-first connection platform for attached adults, solo adults open to couples, and couples. It covers one-off encounters, ongoing extramarital connections, couples seeking a third, and couple-to-couple connections, for all genders and orientations.

Two things make it different:

- 1. **An agent-to-agent handshake (Envoy).** It resolves intent, boundaries, logistics, discretion needs and verification parity *before* two people speak.
- 2. **A breach-assumed architecture.** End-to-end encrypted messaging, client-side encrypted media, split identity and profile stores, cryptographic erasure in under 60 seconds, and partner-collision avoidance through private set intersection. Together these make the category's defining disaster — the 2015 Ashley Madison breach — technically very hard to repeat.

Board summary [A §21]. The category leader in affairs has vacated that position under regulatory and payments pressure. The leader in non-monogamy is a browsing product with a resolution problem. The opportunity is real. The hard part is not the software but the compliance envelope, marketplace liquidity, and an acquisition-cost model that only works through invitation. Build the envelope first, seed the scarce side first, and let the agent layer earn its place against a measured control.

Gating. No engineering build starts before a written UK legal opinion and a completed DPIA (§14).

1.1 Executive product decision

Founder text, verbatim (Source B; re-confirmed 4 Oct 2026).

Build TRYST as a privacy-first, affair-specific connection platform rather than a generic dating app. Its strategic advantage is an AI relationship-intelligence layer that learns attraction, discretion requirements, availability, risk tolerance and interaction quality while deliberately collecting less identifying data than incumbents.

Decision	Specification	Why it wins
Name	TRYST	One sharp word meaning a private romantic meeting; instantly communicates the category.
Audience	Partnered adults of any gender/orientation; couples seeking a third or another couple	Clear, high-intent niche rather than mass-market dating.
Promise	Private chemistry. Intelligent discretion.	Combines emotional benefit with the functional moat.
Core modes	One-off, Ongoing, Couple+, Threesome, Couple-to-Couple	Intent is explicit before matching.
Moat	Discretion Graph + Behavioural Compatibility Engine	Learns fit without exposing identity or partner data.
Revenue	Subscription-led with credits and premium privacy	Predictable recurring revenue plus high-margin intent transactions.
Launch	Web/PWA first, then native apps subject to store review	Reduces store dependency and permits rapid privacy iteration.

ONE-SENTENCE PITCH

TRYST is the intelligent private network where attached adults and couples find precisely matched affairs, extra-relationship connections and threesome partners without exposing more of their identity than necessary.

How this decision maps onto the rest of v1.2:

Decision row	Implemented as
Core modes	One-off → SPARK; Ongoing → EMBER; Couple+ and Threesome → THIRD (Couple+ = a couple seeking a third; Threesome = a solo joining a couple); Couple-to-Couple → QUAD; plus OPEN for exploration (§5.2)
Moat	Discretion Graph = Discretion Policy + ExclusionRing + zones enforced on every read (Curtain); Behavioural Compatibility Engine = Mirror + Broker + Envoy (§7)
Revenue	"Credits" = Keys; "premium privacy" = the GHOST add-on, subject to D-05 (§10)
Launch	Web PWA is the full product and the revenue surface; native apps follow store review, approved by G-P3 (§12, §14.3)
Audience	See D-01 on unattached solos

2. Thesis and market gap

The discreet-connection market is large, proven and badly served. The incumbents monetise frustration; TRYST monetises resolution. [A §1]

Founder market thesis, verbatim (Source B §1; re-confirmed 4 Oct 2026):

The category is already validated. Ashley Madison markets private, anonymous connections and claims more than 91 million members. Gleeden uses a credit economy for extramarital connections; Victoria Milan markets a rapid-exit "panic button." The opportunity is not to prove demand but to rebuild the category around modern privacy engineering, verified authenticity, gender-neutral pricing, transparent AI and higher-quality matching.

Incumbent pattern	Observed weakness	TRYST response
Large anonymous inventory	Fake profiles, low trust and expensive dead-end messaging	Liveness, authenticity confidence and conversation-quality scoring.
Credit-heavy messaging	Revenue can conflict with real connection outcomes	Subscription provides core value; credits enhance but do not manufacture access.
Basic privacy theatre	A panic button cannot protect server-side data	Pseudonymity, data minimisation, isolated vaults, short retention and discreet billing.
Pairwise matching	Poor fit for couples and threesomes	Native multi-person matching and room consent.
Swipe engagement	High noise and fatigue	Small daily set of high-confidence, intent-aligned introductions.

Membership figures differ between sources: Source B cites Ashley Madison's claim of 91M+ members, Source A cites 80M+. Both are the company's own unaudited claims; use the current figure from Ashley Madison's site, with an access date, in any external material (E-32). "Gender-neutral pricing" supports the default in D-06.

2.1 Market

Player	Position	Observed weakness	TRYST response
Ashley Madison	Affair category leader (claims 80M+ members). Rebranded globally as a "discreet dating app" on 24 Feb 2026, moving away from married dating.	Pay-to-initiate credits (men pay to start conversations); recurring bot and cost complaints; 2015 breach (FTC: about 36M users exposed, deletion failures).	Take the affair position it is vacating, and survive the pressure through §9 and §14. Subscription-led pricing; replying is always free.
Feeld	Owns ethical non-monogamy and couple-seeks-third. Revenue £48.9M in 2024 (+26%, Companies House); 2M+ active users by 2026.	Thin free tier; high-noise discovery; a browsing product, not a resolution product.	Envoy handshake; daily curated slate; meets as the north star.
Gleeden, Victoria Milan, Illicit Encounters	Regional affair sites (credits; panic button).	Fake profiles; "privacy theatre" that does not protect server-side data.	Liveness plus age assurance for everyone; real cryptography.

Player	Position	Observed weakness	TRYST response
3Fun / #Open / swinger communities	Couples and thirds niche.	Fake "couples", moderation gaps, dated UX.	Mandatory dual-verified couple profiles (VC).
Mainstream apps (Tinder, Hinge, Bumble, Grindr)	Mass market.	Not discreet; exposure through social graphs.	ExclusionRing; decoy skin; no social login.

Third-party figures are as reported by the sources cited in Source A and have not been independently audited.

2.2 Four failures TRYST fixes

Source: A §1

Failure	Why it persists	TRYST answer
Stated $\neq$ real intent	Profiles are self-authored marketing	Behaviour-learned RevealedPreferenceVector matched against IntentVector; divergence is a feature
200 messages, zero meets	Engagement-maximising ranking rewards long unresolved threads	Envoy handshake pre-qualifies; the reward function penalises time in app
Attention collapse	Rich-get-richer ranking; couples broadcast to everyone	Inbound caps, Gini constraint, asymmetric Key pricing
Discretion as marketing	Retention by default and server plaintext are cheaper	E2EE, cryptographic erasure, token-only verification, PSI ExclusionRing

3. Name and brand

Source: A §2

TRYST — *a private, prearranged meeting between lovers*. One syllable, five letters, no explanation needed, and deniable on a lock screen. Tagline: **Private chemistry. Intelligent discretion.** Internal line: **Discretion is the product.**

- **Brand codes:** oxblood and gold on night; serif display, tight tracking. No pink, hearts, silhouettes or lipstick. Reference: a private members' club, not a nightclub.
- **Fallback names** if trademark clearance fails: ASIDE, DUPLEX.
- **Domains:** tryst.app / tryst.co / trystapp.com plus defensive registrations before any UKIPO Class 9/45 filing becomes public.

3.1 Logo



The master logo is `assets/brand/tryst-logo.png` (1254 × 1254 px, transparent background), supplied by the founder (D-19).

- **Use it exactly as supplied.** Do not redraw, recolour, crop, stretch, re-shadow or separate the mark from the wordmark. The repository keeps the original file byte-for-byte (SHA-256 5f2ceb55...eae2761).

- **Background:** the logo is designed for the dark night ground (§3.2), and is only placed on night, oxblood-deep or oxblood.
- **Never shown where it would expose a member (P3, P9):**
 - on a device in decoy mode;
 - on the lock screen or in notifications;
 - in the app-switcher snapshot;
 - in billing descriptors;
 - in the home-screen icon of a member who has chosen a decoy skin;
 - in Share My Plan messages.
- **Where it appears:**
 - the web PWA (landing page, sign-in, header);
 - inside the native apps after unlock;
 - store listings, subject to the sanitised-store rules (§14.3);
 - investor and press materials.

3.2 Colour palette

The palette is sampled from the logo. Tokens live in `assets/brand/tokens.json` and `assets/brand/tokens.css`.

Token	Hex	Role	Contrast on night
night	#0B0607	Primary background	—
oxblood-deep	#3F0000	Deep surfaces, pressed states	—
oxblood	#560001	Brand surface, primary buttons	1.34 (surfaces only)
oxblood-light	#702A2E	Hover, borders, dividers	1.97 (never text)
bronze	#836042	Secondary lines, disabled states	3.56 (large text only)
gold	#CB944F	Accent, links, icons, headings	7.56 AA/AAA
champagne	#F5D3AC	Highlight text, active states	14.18
cream	#FCF2DD	Body text; light surface (replaces "bone")	18.1

Text pairs that pass WCAG 2.2 AA (NFR-07):

- cream, champagne or gold on night;
- cream (13.5), champagne (10.6) or gold (5.65) on oxblood;
- oxblood on cream (13.5).

Oxblood and oxblood-light are never used for text on dark grounds.

3.3 No "AI" in the name or brand (founder, v1.3)

TRYST is never named, branded or marketed as AI (D-33).

- The words "AI", "A.I." and "artificial intelligence" never appear in the name, tagline, logo, app-store listings, marketing site, Herald content, ads, press materials, product UI or notifications.
- The product describes what members get, not the technology: "your agent compares intentions and limits", "we settle the details before you speak".
- Internal engineering documents (02, 03, 09) may describe models and agents in technical terms. That vocabulary never reaches a member or the public.
- Enforced by `web/test/brand.test.ts`, which fails the build if any user-facing source contains these words.

4. Operating stance and principles

4.1 Stance

Source: A §1

We are indifferent to your marriage; we are absolutist about consent and safety on our surface.

TRYST does not police whether a member's partner knows; that is unverifiable and not the platform's business. It does enforce, without exception:

- Verified adults only, through highly effective age assurance.

- Consent from every party *on the platform*. Couples are dual-verified; no one is listed by a partner without co-signature. No permission is required from anyone not on TRYST.
- Zero tolerance for image abuse, sextortion, coercion, commercial sexual services, and minors.

4.2 Product principles (changing one requires v2.0)

#	Principle	Source
P1	Resolution over engagement. Optimise for confirmed, positively rated meets (QCAM), never time in app or message volume.	A \$6.4, \$18
P2	Hard limits are inviolable. No learned signal can surface a candidate who violates a hard limit, exclusion zone or structure incompatibility. Enforced by tests, not policy.	A \$6.2
P3	Breach-assumed by design. Assume breach, subpoena, and a suspicious partner holding the phone.	A \$9
P4	Agents never impersonate. Agents negotiate only with other agents; every Brief is labelled as machine-generated; only a double human accept opens a thread.	A \$5.4
P5	Stated preference keeps a floor. Revealed behaviour never drops stated preference below 15% weight; members can view, correct and reset learning.	A \$6.2, B FR-010
P6	No safety theatre. No "safe", "vetted" or "background-checked" badges, ever. Only factual, dated statements.	A \$19.2, \$19.8
P7	Safety beats discretion in an emergency. Panic preserves and releases what emergency services need, regardless of retention settings.	A \$19.7
P8	Revenue ethic. No advertising, data sale, affiliates, fake profiles, unreadable credit mechanics, engagement paywalls, or charging for deletion or data access. Replying is always free.	A \$15, B \$11
P9	Truthful privacy claims. Never promise an affair is undetectable. Disclose the limits (device owner, bank, network, recipient screenshots, legal process) and the declared exceptions (ban anchors, legal holds).	B \$4, A \$19.3

4.3 Never

- Never the "cheat better" message; the positioning is "Discretion is the product." [A \$16]
- Never a "catch a cheater" feature, partner search or exposure tool.
- Never random or anonymous chat; messaging is between matched members only. [A \$14.3]
- Never operate in markets that criminalise adultery or same-sex conduct, and never in Africa or MENA under any brand. [A \$14.5]

5. Segments, modes and intent shapes

5.1 Segments

Source: A §3

Every segment is first-class in the data model.

Code	Segment	Job to be done	Primary anxiety
S1	Attached, undisclosed	Meet one person discreetly without changing my primary life	Exposure
S2	Attached, disclosed / open	Quality connections without re-explaining my structure	Judgement, wasted time
S3	Couple seeking a third	Someone both of us want, who wants both of us	Ghosting; friction between partners
S4	Solo open to couples	Be treated as a person, not a utility; a veto that is respected	Harassment volume; bait-and-switch
S5	Couple seeking a couple	Four-way compatibility with low coordination cost	Coordination collapse
S6	Curious / never acted	Find out whether this is for me, at low risk	Committing too early; discovery

S4 is the scarce asset. S4 supply decides whether S3 and S5 have a product. S4 inbound is rationed, priced and capped (03 \$5.4). Any ranking change that pushes S4 inbound above the cap is a regression, whatever it does to session metrics.

People in open relationships are welcome (founder, v1.3): consensually non-monogamous and polyamorous members, solo or as a couple, join as S2 (solo) or S3/S5 (couples), mainly through the OPEN, SPARK and EMBER

modes. Nothing in the product assumes secrecy from a partner. Solo members may be attached or unattached. Partner awareness is self-declared, never verified, and never shown to others. (D-01)

5.2 Modes (member-facing) ↔ intent shapes (data)

Source B's mode names are kept as the **member-facing labels**; Source A's intent shapes are the **canonical data values** (02 §2.3).

Mode (UI)	Intent shape(s)	Typical segments	Notes
SPARK	one_off	S1, S2, S4, S6	Matched on availability, distance and chemistry
EMBER	recurring, ongoing	S1, S2	Matched on cadence, emotional depth, communication windows
THIRD	third	S3 ↔ S4	Couple profile must be VC; the third is an equal member with a respected veto
QUAD	couple	S5 ↔ S5	Four independent verified accounts; match only after each couple's veto mode is satisfied
OPEN	exploratory	S2, S6	Member picks allowed configurations and visibility

A member may enable several modes, each with its own visibility and filters.

6. Product loops

Source: A §4

Loop	Cadence	Flow
A — Intake	Once (~7 min, conversational)	Age assurance → Cartographer interview → IntentVector v0 → ExclusionRing enrolment → Discretion Policy → first slate
B — Discovery	Daily (< 6 min per session target)	Curated slate (up to 18 cards, 15% exploration) → Intent sent (costed) / pass / save → Mirror updates RevealedPreferenceVector → slate recomputed
C — Handshake (the product)	Per mutual intent	Envoy A ⇄ Envoy B negotiate over a fixed schema → asymmetric Brief to each human → accept / decline / amend → E2EE thread opens only on double accept
D — Resolution	Per connection	Thread → logistics agreed → Safe Meet checklist → meet confirmed → Aftercare prompt at 48 h → reward signal
E — Discretion	Continuous	Curtain scores session risk → geofence, collision suppression, notification masking, retention enforcement

Loop C is the product. If it does not measurably raise meets per intent sent over a control cohort (gate **G-P2**), TRYST is a me-too and should not raise a Series A.

7. Agents (roles)

The AI agents facilitate matching; members make every decision. Agents learn, rank, pre-qualify and draft. They never send an intent, open a conversation, reveal a member or act as a member without that member's explicit action (P4, FR-055, FR-066).

Seven agents, each an independently deployable service with its own tool surface, memory scope, policy file and eval suite. They talk over a typed internal bus, never by passing free text. [A §5] Implementation: 03 §3.

Agent	Role	Autonomy boundary
Cartographer	Conversational intake → signed IntentVector; re-run quarterly or on drift	Member confirms every structured field [B]
Mirror	Behaviour learning; sole writer of RevealedPreferenceVector	Approved features only; message content never used
Broker	Retrieval and ranking under hard gates and fairness constraints	Cannot bypass Stage 0 gates; paid status cannot override eligibility
Envoy	Agent-to-agent handshake; asymmetric Briefs; logistics and timing negotiation	Never talks to a human as the member; cannot open a thread; output is advisory
Curtain	Discretion risk; owns and enforces the versioned Discretion Policy	May tighten exposure automatically; may never loosen it [B]

Agent	Role	Autonomy boundary
Guardian	Safety; veto over every other agent	No permanent action without a named human; never reads message content server-side
Aftercare	48 h private outcome capture → reward	Never shown to the counterparty; never a public rating

Two release-gate functions that are not agents: **Privacy Auditor** (continuous access, retention and model-leakage tests; can block a release) [B] and the **agent eval harness** (03 §9).

Growth agent (outside the member platform): Herald, the dynamic AI SEO agent (v1.2). It researches topics, drafts and optimises content for the public marketing site, manages internal linking, tracks rankings and drafts link-earning outreach. It has **no access to any member data**, and a human editor approves everything before it is published (03 §3.2, D-21).

8. Trust, verification and admission policy

8.1 Verification ladder

The tier definitions and what each one unlocks are canonical in 02 §3. In summary:

- **V0** = contact handle, onboarding only.
- **V1** = liveness.
- **V2** = highly effective age assurance. **Required for discovery, intents, threads, media, Envoy and meets** (E-01).
- **V3** = optional bank or digital ID.
- **VC** = couple co-signature: both partners independently at V2, each confirming from their own device.

VC is the most important trust feature in the product. It will cost 20–30% of couple sign-ups in year one. Pay it.

8.2 Screening reality and admission control

Source: A §19

In the UK you cannot lawfully buy sex-offender or criminal-record screening for platform users (ViSOR is not accessible; DBS checks are employer-only). TRYST builds five-layer admission control instead (03 §8):

Layer	What it does
L1 Identity binding	One human, one account, permanently: liveness challenge, ID match through the provider, device attestation, payment fingerprint, £9.99 refundable deposit
L2 Ban anchor	One-way identifiers that survive deletion. The declared exception to erasure.
L3 Voluntary attestation	Self-submitted Basic DBS shown as a dated fact; Clare's Law explainer (answer never asked for or stored); conduct declaration at V2
L4 Behavioural detection	Guardian M6a–M6e and M7b; most dangerous users have no record, but they do have a pattern
L5 Consortium	Industry hash sharing (e.g., Online Dating Association) of anchors plus reason codes only; gated on competition-law and data-protection opinions

"Declared unsafe" grounds (auditable): (1) self-declared intent to harm or coerce; (2) corroborated reports from two or more independent reporters; (3) a court order (SHPO, restraining or non-molestation order); (4) law-enforcement notification; (5) a model signal confirmed by human review; (6) a knowingly false conduct attestation.

Ladder: warn → shadow-limit → suspend → permanent ban + Ban Anchor → law-enforcement referral. Grounds 1, 3 and 4 go straight to permanent.

Due process: a named human makes every permanent decision; appeals go to a different reviewer within 10 working days; reporter-reputation scoring (retaliatory or coordinated reporting is itself removable — suspicious spouses and rejected counterparties are predictable abuse vectors); the transparency report publishes the overturn rate.

8.3 Safety floor

- **Safe Meet:** public venue first; check-in timer with silent escalation.

- **Share My Plan:** encrypted, unbranded, auto-deleting share with a trusted contact.
- **Panic:** escalation to 999 with location, separate from the discretion **Burn** control (E-21).
- **Emergency override:** safety beats discretion (P7).
- **Victim-support routing** after an incident.
- **24/7 T&S rota** from P3, with published SLAs.

9. Discretion promise

TRYST minimises platform exposure. It cannot guarantee secrecy from a device owner, bank, employer, network administrator, recipient screenshot, legal process or a determined third party, and it says so (P9).

The design starts from the eight adversaries in 03 §6.1:

- A1 partner holding the device — daily case
- A2 shared bank statement
- A3 social-graph recognition
- A4 mass breach — existential
- A5 extortion by a counterparty
- A6 civil discovery in divorce
- A7 insider
- A8 face-search de-anonymisation

Member-facing controls are in 04 §5. Retention defaults are canonical in 02 §7.

10. Business model and unit economics

10.1 Principle

Source: A §15.1

Charge for agency and discretion; make resolution the profitable outcome. Subscription plus success-weighted retention, never pay-per-message. The economic model and the loss function must agree.

10.2 Tiers (UK)

Tier	Price	Includes
Verified (free)	£0 + £9.99 refundable deposit (credited 100% as Keys)	Full profile; be discovered; reply to anyone free ; 3 outbound intents a week; Envoy handshake on mutual intent; full safety suite and core discretion suite
TRYST+	£24.99/mo · £14.99/mo billed annually	Unlimited intents, filters, incognito, travel mode, extended media, read-receipt control
ENVOY	£49.99/mo	The agent tier: Envoy <i>proposes</i> outreach and pre-qualification, screens risk, negotiates logistics. Every outbound intent needs a one-tap human approval (D-04)
DUO	£39.99/mo per couple	Linked VC profile, dual-consent controls, joint-veto model, third-screening Envoy
Keys	£9.99 / 60	Non-subscribers: intent = 5 Keys solo→solo, 15 Keys couple→solo; fully refunded on reply
GHOST add-on	£7.99/mo	Advanced discretion: multiple decoy skins, custom TTLs, extended ExclusionRing heuristics. Core decoy, duress PIN, Burn, masking and voucher billing stay free (D-05)

Source B's Black / Duo £34.99 / Signals / Vault+ / Concierge Black tiers are superseded (E-14). Concierge is deferred beyond P4.

10.3 Unit economics (illustrative; replace with beta cohorts)

Source: A §15.3

Metric	Value
Blended CAC per sign-up (paid + referral, UK)	£38
Sign-up → V2 verified	52%
Verified → payer (month 3)	11%
Blended ARPPU	£31.50/mo

Metric	Value
Payer gross churn	19%/mo (average life 5.3 months)
Gross revenue per payer (lifetime)	£167
less PSP (high-risk, 5.0%) / verification (2 × £1.10) / infra + inference	−£8 / −£2 / −£14
Contribution per payer	£143
CAC per payer at £38 per sign-up ($£38 \div 0.11 \div 0.52$)	£664 — paid acquisition alone does not work

Corrected viability thresholds (E-03): break-even blended CAC per sign-up = £143 × payer rate × 0.52.

Verified → payer	Break-even CAC per sign-up	CAC for 3× LTV:CAC
11%	£8.18	£2.73
16%	£11.90	£3.97
18%	£13.38	£4.46

Source A's "< £12" target is only viable at a payer rate of about 16% or more. Two levers decide the business; track both weekly from beta day one:

1. Raise verified → payer toward 18% (ENVOY is the instrument).
2. Drive blended CAC down through member referrals and organic channels.

Treat paid acquisition as a seeding cost, never a growth engine.

11. Liquidity and go-to-market

Source: A §16

Changed by the founder in v1.3 (D-35): public sign-up replaces Source A's 9-month invite-only period. The proven density-first sequence is kept: matching opens one city at a time, London first, and the balance waitlist stays.

1. **Density gate (G-City-1).** Sign-up is open from day one in all four launch cities, but **matching opens one city at a time, in launch order**: a city goes live at 2,000 verified members once the city before it is live, and stays live. Until then, members verify, set up their profile and discretion settings, and see how close their city is. An empty grid is never shown.
2. **Balance gate (G-City-2) — the demand-balance waitlist, gender-neutral (D-36).** Nobody waits because of who they are. Each member belongs to the groups that describe them and seeks the groups they asked to meet. A new member waits only if a group they want to meet is already saturated: more than 2.2 seekers per member of that group in their city (2.0 by P3). The same rule protects whoever is being overwhelmed: women in a straight pool, men in a gay pool, singles sought by couples. Waiting is first come, first served; members see their position and which group is busy, and are admitted automatically as that group grows. The per-city gender ratio stays a monitored KPI and is achieved through this rule, not by admitting by gender. Price and ranking never use any of it (D-06; 03 §3). Implemented and tested in backend/internal/launch (FR-090).
3. **Seed the scarce side first.** Outreach waves and founding-member status for S4 (and, subject to D-06, under-represented cohorts) with free ENVOY for 6 months.
4. **Public sign-up, no invitation needed** (founder, v1.3). Member referrals stay as a growth channel: 3 per member, refreshed monthly for members in good standing. Referral rewards never reveal the referrer to the invitee [B]. 4a. **Paid-acquisition brake** (v1.3, FR-091). Each week, per city, blended cost per verified sign-up is compared with the E-03 break-even line (£8.18 at the base case, recalculated from observed payer and V2 rates). Paid spend in the city **stops automatically** after two weeks above break-even, or immediately at 50% above it; restarting needs a named approver. Referrals and organic search continue. Acquisition cost therefore cannot run above break-even for more than one week (above 1.5× it, not at all) without a person deciding it should.
5. **Cities:** sign-up open in **London, Manchester, Birmingham and Brighton** from launch; matching opens **London → Manchester → Birmingham → Brighton**, and seeding spend is focused on the next city to open, not spread across all four. Dublin opens with the Irish launch in P3. Adults elsewhere in the UK can register interest, and a new city opens when its sign-ups approach the density gate.
6. **Channels:** assume Meta and Google reject the ad account. Use podcast host-reads in adjacent verticals, ENM and relationship creators, technology-led PR (agent handshake, breach-proof architecture), and an editorial presence. Never synthetic personas or staff posing as members [B].
7. **Positioning:** "Discretion is the product." Never "cheat better."
8. **Organic search as the growth engine (v1.2).** Paid social is closed and paid CAC is fatal (§10.3), so search is the main scalable low-CAC channel alongside referrals. The Herald agent runs it:

- **Dynamic content:** topic clusters on privacy, discretion, safety (Safe Meet, Clare's Law), ethical non-monogamy and couples, plus genuinely useful launch-city pages (no doorway pages). Pages are refreshed from search-performance data. Source B §11.2's "privacy-led content" channel is delivered here.
- **Hyperlinks:** automated internal-link graph across the topic clusters (hub and spoke), with no orphan pages and descriptive anchors.
- **Backlinks, earned and never bought:** digital PR on the technology story (agent handshake, breach-proof architecture, transparency reports); open-sourcing the audited crypto core; expert contributions; podcast and creator partnerships marked rel="sponsored" where paid; relationship and lifestyle publishers.
- **Top ranking:** target top-3 positions for the priority keyword set in each launch market (e.g., "discreet dating app UK", "couples looking for a third", "ethical non-monogamy app"), tracked weekly. Rankings can be targeted and measured, never guaranteed.
- **Never:** paid links, private blog networks, link exchanges, comment or forum spam, expired-domain or parasite-SEO abuse, cloaking, doorway pages, unreviewed mass AI content, or fake reviews. Each breaches search-engine spam policies and risks the domain, which is the channel the business depends on.

Markets: launch UK + IE · wave 2 NL, BE, DE, SE, DK, ES, PT · wave 3 CA, AU, NZ. The US is out of scope (E-16). Every market needs its own written legal opinion.

12. Roadmap, team, budget and phase gates

Source: A §17

Phase	Window	Scope	Exit gate
P0 Legal & foundation	M0-M2	SPV, counsel, DPIA, OSA assessments, PSP + backup MID, two age-assurance providers, threat model	G-P0: SPV incorporated · written counsel opinion · DPIA complete · OSA risk assessments drafted · high-risk PSP + backup MID approved · two assurance providers contracted · threat model signed off
P1 Core MVP	M2-M6	Public web beta: open sign-up in four UK cities, matching live in London first; Loops A, B, D; E2EE; V2 gating; VC co-sign; Guardian baseline (on-device); deletion and DSR machinery; web PWA + native shells. No Envoy — this is the control cohort	G-P1: matching live in London with ≥ 2,000 verified; sign-up open in all four launch cities · erasure < 60 s p99 · no-go list (§13.2) clear
P2 Differentiator	M6-M10	Envoy, Mirror v1, ExclusionRing PSI	G-P2: meets per intent sent up ≥ 40% vs the P1 control. If not met, stop and rethink
P3 Public UK + IE	M10-M15	Full public launch, Dublin added; DUO + GHOST; native apps approved on both stores; 24/7 T&S	G-P3: 25,000 verified · payer rate ≥ 11% · both stores approved
P4 EU wave 2	M15-M24	NL/BE/DE/SE/DK/ES/PT; per-market opinions; localised assurance and Guardian	G-P4: 150,000 verified · contribution-positive

Source B's phase-level exit criteria are folded into the gates above and into the phase backlogs in 03 §12 and 04 §10.

Team to P3 (17):

- **Product & engineering (12):** Head of Product · Principal Engineer, security-led (**hire first**) · 2 Backend (Go) · iOS · Android · Web · 2 ML/Agents · Data Engineer · Security Engineer · Design.
- **Trust, legal & growth (5):** Head of Trust & Safety (**hire in P0**) · 2 T&S ops · DPO/Counsel (fractional until P2) · Growth/Community.

24-month budget ≈ £5.9M:

Line	Amount
People	£3.4M
Legal, DPIA, counsel	£260k
Age assurance	£180k
Infra + inference	£420k
Security audits	£150k
Seeding & community	£480k

Line	Amount
PSP reserves & fees	£220k
Contingency (15%)	£780k

Plus insurance of £85-140k a year at P3 scale (\$15.3), which is not in Source A's total. **Plus biometric login checks** (v1.2): an estimated £0.10-0.60 per mobile MAU per month, pending the vendor quote (03 \$6.7).

13. KPIs, no-go conditions and risk register

13.1 North star and KPIs

Source: A §18

QCAM — Quality Connections per Active Month: confirmed meets with a positive aftercare score, per active member, per month. If a change raises engagement and lowers QCAM, it does not ship.

Metric	P1	P3
QCAM	0.18	0.45
Meets per intent sent	1.8%	4.5%
Mutual reply rate	22%	34%
Inbound Gini (per city, per segment)	< 0.60	< 0.55
Gender ratio (per city)	≤ 2.2:1	≤ 2.0:1
S4 30-day retention	45%	62%
Report rate per 1k threads	< 12	< 6
V2 completion	52%	68%
Chargeback ratio	< 0.65% always (above 0.9% the MID is lost)	
Erasure latency	< 60 s p99	

Secondary (Source B): qualified reciprocal conversations per 1,000 verified weekly active members; exposure incidents; location-inference success in red-team tests; notification leakage; calibration and hard-boundary violations (target 0).

Source B KPI dimensions (verbatim in 07) are kept as the secondary dashboard set. Match quality, discretion, authenticity, safety, AI, commercial and reliability are each reported weekly from P1. B's north star (qualified reciprocal conversations per 1,000 verified weekly active members) is the leading indicator for QCAM. v1.2 adds emergency-lock latency (FR-064; target < 5 s p95) and login factor failures (FR-057/058) to reliability.

Growth/SEO (v1.2): organic sign-ups and organic CAC; share of priority keywords ranking top 3 and top 10; quality referring domains earned; marketing-site Core Web Vitals pass rate (NFR-20); spam-policy incidents (target 0).

13.2 No-go conditions

Do not launch or scale if any of these hold. Each is backed by an automated or audited test (03 §9). [B §14, A]

ID	Condition
G-NG-1	Private media can be accessed outside an active reveal grant
G-NG-2	Deletion or erasure claims are not technically true (including the declared exceptions)
G-NG-3	Exact location can be triangulated in red-team testing
G-NG-4	Age assurance is not "highly effective" across all four Ofcom criteria, or any gated surface is reachable below V2
G-NG-5	A couple profile can represent an unverified person or be co-signed from one device
G-NG-6	High-severity reports lack 24/7 response within published SLAs
G-NG-7	Ranking violates any hard limit, exclusion zone or ExclusionRing pair
G-NG-8	Any third-party analytics, advertising or attribution SDK is present in a client
G-NG-9	Staff can read message content
G-NG-10	Any login or recovery path issues a session with fewer factors than 02 §3.3 requires, or allows SMS/email/staff-assisted sign-in

13.3 Risk register

Risk	Severity	Mitigation
Breach (existential)	Critical	E2EE, identity split, cryptographic erasure, short TTLs, zero standing access, quarterly pen test, annual red team on A4/A7, rehearsed 72 h playbook
App-store rejection or removal (Apple 1.1.4)	High	Sanitised store build; explicit media web-only; matched-only messaging; web-first revenue
Merchant account loss	High	Backup MID from day one; alert at 0.5% chargebacks; Ethoca/Verifi; one-click cancel; pre-renewal notice; voucher path
Ofcom enforcement	High	Two assurance providers; inbox-level gating; documented assessments; named accountable manager; quarterly compliance review
Liquidity failure in month 4	High	Density gate per city, opened in order (London first); balance waitlist; scarce side first; seeding focused on one city at a time (D-35)
Envoy does not move meets per intent	High	P1 control cohort; G-P2 at +40%; kill the thesis honestly
Agent seen as a bot / impersonation scandal	High	P4; labelled Briefs; publish the design
Harm incident attributed to the platform	High	T&S in P0; 24/7 from P3; published SLAs; Safe Meet; transparency reporting
CAC model unviable	High	\$10.3 thresholds; referral loop; organic search; weekly tracking per city
Reputational contagion to parent group	Medium	Ring-fenced SPV, brand, counsel, cap table and PR
Feeld or Match Group ships an agent layer	Medium	Speed to P2; the moat is discretion architecture plus the aftercare dataset
Search penalty or deindexing of the marketing domain	Medium	White-hat only (D-21); human editorial approval; no bought links; monthly link-profile audit; marketing site on its own host, separate from the app
Two-factor biometric login hurts conversion or excludes disabled members	Medium	Fast provider flow (NFR-16); 30-day sessions; accessible alternative (D-17); measure login drop-off weekly
Guardian false positives on consensual kink	High	Trajectory-not-vocabulary design; matched-pair ship gate (03 \$7)

14. Compliance envelope

This section gates the build. Everything here needs sign-off from UK counsel experienced in the OSA and Article 9; it says what to instruct counsel on, not what they will conclude. Statutory positions and commencement dates must be checked as at the date of use. [A \$14]

14.1 Online Safety Act 2023

TRYST is in scope as a user-to-user service. Ofcom treats dating services as posing a risk of significant harm to minors and published dating-specific guidance in May 2026.

Required before public launch:

- illegal-content and children's-access risk assessments;
- highly effective age assurance meeting all four criteria (accuracy, robustness, reliability, fairness), **applied to the inbox and media, not just sign-up**;
- clear terms, in-app reporting, rapid review and removal;
- controls mapped to Ofcom's named dating risks (romance fraud, intimate image abuse, sextortion, harassment, stalking, grooming);
- a named accountable senior manager.

Penalties reach the greater of £18M or 10% of global turnover, with possible criminal liability for senior managers.

14.2 UK GDPR / DPA 2018

- **Article 9 explicit consent** — granular, separate from the ToS, withdrawable, recorded in the consent ledger.
- **DPIA mandatory** — two triggers: large-scale special-category data and systematic evaluation.
- **Minimisation as architecture** — no legal name, exact DOB, ID images or coordinates.
- **Article 22** — human review path for suspensions.
- **DSR machinery from sprint one.**
- **Residency** — UK/EU only; no US sub-processor for Article 9 data without a TIA.

- **Article 6 lawful bases** (v1.2; counsel to confirm): contract, Art. 6(1)(b), for the core service; legitimate interests, 6(1)(f), for fraud and safety detection, with a documented LIA; legal obligation, 6(1)(c), for mandatory reporting and OSA duties. Each sits alongside the Article 9 condition for any special-category element (explicit consent, or DPA Sch. 1 for safeguarding and ban anchors).
- **Breach playbook** — 72 h, rehearsed twice a year.
- **Article 10 conviction data and the ban-anchor lawful basis** — DPA Sch. 1 Pt 2 para 18 assumed; counsel must confirm.

14.3 Distribution — two-surface strategy

- **Apple 1.1.4** rejects "hookup" apps that may include pornography or facilitate prostitution.
- **Google Play** requires minor-blocking for dating apps and restricts anonymous or random chat.

Native store builds are therefore **sanitised**: no nudity or explicit copy, matched-only messaging, 18+, "discreet dating" positioning. The **web PWA** carries the full product: explicit media behind V2, and all billing. Details: [04 §2](#).

14.4 Payments

- MCC 7273, high-risk; Visa Integrity Risk Program Tier 1 registration.
- **Stripe is the chosen payment processor (founder decision, v1.2; D-22)**: Stripe Checkout / Payment Element on the web, Stripe Billing for subscriptions, the Customer Portal for one-step cancellation, Radar for fraud screening, and Stripe Tax for VAT. **Condition**: Source A states that Stripe, Square and PayPal prohibit adult content, and TRYST's web surface carries explicit media behind V2. Before any build commitment, PO must obtain **Stripe's written approval of TRYST's exact business model** (affair-positioned dating, couples, explicit media on web). A dedicated high-risk acquirer (3–7%) is contracted in parallel as the backup MID from day one, and becomes primary if Stripe declines or later restricts the account.
- **Spousal chargebacks** are the category-specific risk. Mitigations: neutral but truthful descriptor with a customer-service URL, offline voucher/gift-code path, pre-renewal email, one-click cancel, Ethoca/Verifi alerts, backup MID, correct MCC.
- **DMCC Act 2024 subscription regime**: commencement has moved repeatedly; build to the full regime now.

14.5 Jurisdiction gating and structure

- Hard geo-block at the edge and at sign-up for any market that criminalises adultery or same-sex conduct, or prohibits the category.
- Travel into a blocked market auto-locks the app to the decoy.
- Ring-fenced SPV with its own controller registration, brand, counsel, cap table and PR, and no public link to the parent group.

15. Liability, claims and insurance

Source: A §19.8, §20

15.1 Claims discipline

Never say	Say
"Background checked", "Safe", "Vetted"	"Identity verified", "Age assured"
"No sex offenders on TRYST"	"We permanently remove and block re-registration where we find harmful conduct"
"Screened for domestic violence"	"We detect and act on coercive and abusive behaviour, and we'll show you how to use Clare's Law"
"Guaranteed real people"	"Every member completes liveness and ID verification"
Green "verified safe" tick	Factual, dated attestations only

Overclaiming is a misleading commercial practice under CPUT and the DMCC Act, and the first thing a claimant's solicitor will point to.

15.2 Liability architecture

- **CRA 2015 s.65** bars excluding liability for death or personal injury caused by negligence (including mental impairment). Under s.65(2), agreeing to a term does not mean accepting a risk.
- **s.62** strikes unfair terms in full.

The draft ToS therefore uses narrow, honestly carved limits:

- correct characterisation of the service;
- an explicit list of what is and isn't verified;
- no liability for consequences to members' relationships;
- a financial cap of the greater of 12 months' fees or £100;
- a load-bearing "what we do not exclude" clause.

Keep out: death/injury exclusions, "use at your own risk", mandatory arbitration or class waivers, unilateral no-notice amendment, any screening implication, and forfeiting Keys before an appeal concludes.

The real protection is §8, not the clauses: liability will turn on whether TRYST took reasonable care. The ToS draft is 05.

15.3 Insurance and structure

Cover needed:

- cyber and data liability, sized for an Article 9 breach;
- tech E&O / professional indemnity;
- public and product liability;
- D&O (OSA named-manager exposure);
- media liability.

The SPV provides structural containment. Indicative premium at P3 scale is **£85-140k a year**.

02 — TRYST Shared Contracts v1.1

Part of the **TRYST v1.1 baseline**. This document is the **single source of truth** for everything both backend (03) and frontend (04) depend on. Change it first (00 §4).

1. Conventions

Topic	Rule
Contract formats	Edge API: OpenAPI 3.1 (contracts/openapi/). Internal gRPC, the agent bus and events: Protobuf in the schema registry (contracts/proto/). Generated client and server types; no hand-written DTOs (D-07).
IDs	UUIDv7, opaque, never sequential or enumerable. Prefixed in logs only (usr_, prf_, thr_).
Time	RFC 3339 UTC on the wire. Availability windows carry an IANA tz.
Money	Integer minor units + ISO 4217 ({ "amount": 2499, "currency": "GBP" }).
Location	Geohash-5 (~5 km) or coarser on the wire and at rest. Exact coordinates are sent to the server only by the Panic flow (§5.8).
Age	age_band only (18-24, 25-29, 30-34, 35-39, 40-44, 45-49, 50-59, 60+). No DOB anywhere.
Free text	Everything a counterparty writes (bios, messages, agent turns) is untrusted input : schema-validated, never executed or interpreted as instructions.
Versioning	URL major version (/v1). Additive changes only within a major version. Enum additions are additive; clients must tolerate unknown values.

2. Glossary and enumerations

2.1 Glossary

Term	Definition
Member	A person with a TRYST account (one human, one account).
Entity	The matchable unit: a solo profile or a couple profile.
Couple profile	An entity linking two members, each verified to V2 and each co-signing from a distinct device (VC).
Intent	A costed outbound expression of interest from one entity to another.
Mutual intent	Both entities have sent an intent to each other. In P2+ it triggers a handshake.
Handshake	Envoy-to-Envoy negotiation over a fixed schema (§6).
Brief	The asymmetric, agent-generated summary each side receives after a handshake. Always labelled as machine-generated.
Thread	An E2EE conversation (MLS group) between the entities' members. Opens on double accept (P2+) or mutual intent (P1 control).
Keys	Prepaid credits. Spent on intents; refunded in full on reply.
Discretion Policy	A member's versioned privacy configuration, enforced on every read path (§6.4).
Reveal grant	A recipient-specific, expiring, revocable permission to view a media item or move up the reveal ladder.
ExclusionRing	A set of bidirectional, silent suppression pairs (PSI contacts, heuristics, manual).
Exclusion Identifier / Ban anchor	An irreversible one-way identifier kept after deletion for permanently removed members. The declared exception to erasure.
Burn	Discretion control: wipe local data, revoke device keys, sign out everywhere.
Panic	Safety control: escalate to emergency services with live location. Never the same control as Burn.
QCAM	Quality Connections per Active Month (north star; 01 §13).

2.2 Segments

segment ∈ {S1, S2, S3, S4, S5, S6} — definitions in 01 §5.1. A solo entity is S1, S2, S4 or S6; a couple entity is S3 or S5.

2.3 Intent shapes and modes

intent_shape (data)	Mode label (UI)	Entity pairs allowed
one_off	SPARK	solo↔solo
recurring	EMBER	solo↔solo
ongoing	EMBER	solo↔solo
exploratory	OPEN	any, subject to seeking
third	THIRD	couple(S3)↔solo(S4)
couple	QUAD	couple(S5)↔couple(S5)

2.4 Other enumerations

Enum	Values
structure	attached_undisclosed, attached_disclosed, open_relationship, unattached, couple_seeking_third, couple_seeking_couple, prefer_not_to_say
visibility	public (discoverable to eligible), matched (visible only to matches), incognito (discoverable only after own outbound intent)
veto_mode (couple)	either (either partner can decline), both (both must accept before an intent or decision counts)
handshake_decision	pending, accept, decline, amend
intent_state	sent, queued_inbound_cap, delivered, replied, declined, expired, refunded
desire_tag_level	yes, curious, no
consent_kind	art9_processing, media_view, meet, couple_link, retention, conduct_attestation, safety_training_use, biometric_login
reveal_scope	alias, blurred_gallery, photo_set, voice, video (the reveal ladder, in order)
report_category	minor, ncii, sextortion, coercion, threat_violence, stalking, romance_fraud, commercial_sex, harassment, impersonation, protected_trait_harassment, exposure_threat, other
enforcement_action	warn, shadow_limit, suspend, permanent, law_enforcement_referral
unsafe_ground	g1_self_declared, g2_corroborated_reports, g3_court_order, g4_law_enforcement, g5_model_confirmed, g6_false_attestation
guardian_reason_code	m6a_coercive_control, m6b_boundary_violation, m6c_off_platform_pressure, m6d_identity_location_pressure, m6e_image_coercion, m7_authenticity, m7b_ban_evasion, payment_solicitation
tier (commerce)	verified_free, tryst_plus, envoy, duo; add-on ghost

3. Verification tiers and gates

3.1 Tiers

Tier	Method	Retained	Source
V0	Email or phone OTP; device attestation	HMAC contact hash; device key	A §8.1
V1	Liveness selfie with randomised challenge, matched to profile photos	Boolean + match score. No biometric template.	A §8.1, §19.3
V2	Highly effective age assurance through a third-party provider (one of two, abstracted), including an ID-document match	is_adult boolean + provider token + ban-anchor candidates (§4.1). No ID image, no DOB.	A §8.1, §19.3 L1
V3	Optional open banking or digital ID	Boolean + token	A §8.1
VC	Couple co-signature: both partners at V2, each confirming from their own distinct device	Link record, both profile IDs, device hashes	A §8.1

3.2 Capability gates (enforced server-side; mirrored in UI)

Corrected in v1.1: Source A's tier table let V1 send intents and open threads, which contradicted ToS cl. 5.1, Source B FR-001 and the OSA inbox requirement. V2 is now the gate for all interaction (E-01).

Capability	V0	V1	V2	V3	Couple (VC)
Onboarding, Cartographer intake, discretion setup, waitlist	✓	✓	✓	✓	✓

Capability	V0	V1	V2	V3	Couple (VC)
See other members' profiles (slate)	—	—	✓	✓	✓
Send intents / receive deliveries	—	—	✓	✓	✓
Envoy handshake / Briefs	—	—	✓	✓	✓
Threads (E2EE)	—	—	✓	✓	✓
Media exchange; explicit media (web only)	—	—	✓	✓	✓
Meets, Safe Meet, Share My Plan	—	—	✓	✓	✓
Trust badge; set min_counterparty_tier = V3	—	—	—	✓	—
Publish couple profile	—	—	—	—	requires both partners at V2 + co-sign

Rules:

- No other member's content is shown before V2, and none during the verification flow itself.
- min_counterparty_tier (default V2, option V3) is a Stage 0 hard gate in both directions.
- Re-verification triggers: a Guardian minor signal, a change of device plus a profile-photo change, a credible report, or a 12-month expiry.

3.3 Login assurance — account holder only

v1.2, founder requirement (D-16). Only the account holder can sign in. The rules depend on the surface:

Surface	Factors required at every login	Both required?
Native app (iOS, Android) and PWA on a mobile device	B1 device-bound biometric passkey + B2 liveness face match to the account holder	Yes — two biometric factors
Desktop (web browser or installed PWA on a computer)	F1 passkey (platform authenticator with user verification, or FIDO2 security key) + F2 approval from the member's registered mobile device, which itself requires B1 (or a second registered FIDO2 security key)	Yes — two factors

Factor definitions:

- **B1 — Device biometric passkey.** A WebAuthn/FIDO2 credential bound to the device's secure element. userVerification = required and must be satisfied by **biometrics only** for login:
 - iOS: biometryCurrentSet;
 - Android: BIOMETRIC_STRONG, key invalidated on biometric enrolment change;
 - mobile PWA: platform authenticator. The biometric never leaves the device. **If anyone adds or changes a fingerprint or face on the device, the credential is invalidated** and the member must log in again with B1 re-registration + B2 (FR-059). This defeats the partner who enrolls their own fingerprint (A1).
- **B2 — Liveness face match.** A randomised challenge-response liveness check (ISO 30107-3 PAD Level 2 certified provider), matched 1:1 against the account holder's reference enrolled at V1. The reference is held by the provider under a deletion obligation, never by TRYST (D-18). Requires the explicit biometric_login consent.
- **F1 — Desktop passkey.** Windows Hello, Touch ID or a roaming FIDO2 key with PIN or biometric. Phishing-resistant, origin-bound.
- **F2 — Registered-device approval.** The desktop shows a QR code; the member scans it **in the TRYST app** on a registered mobile device and approves with B1. No push notification is sent, so nothing appears on a lock screen. Alternative: a second registered FIDO2 security key.

Never accepted as a login factor: passwords (TRYST has none), SMS or email OTP, magic links, security questions, or support-desk overrides. Contact OTP is used **only** at sign-up to prove ownership of the contact handle (V0).

Login vs unlock:

- **Login** creates a session on a device: first use on a device, after sign-out or Burn, after session expiry, after a biometric-enrolment change, or on a risk step-up.
- **Unlock** reopens the app within a live session: device biometric (or the duress PIN, which opens the decoy) per FR-027.

Session rule	Mobile (app / PWA)	Desktop
Maximum session	30 days	12 hours

Session rule	Mobile (app / PWA)	Desktop
Idle lock	2 min default (member-adjustable 1-15)	15 min → full re-login
Token binding	DPoP to the device key	DPoP to a non-extractable WebCrypto key

Step-up (repeat both factors for the surface, valid 5 min) is required to:

- add or revoke a device;
- broaden the Discretion Policy;
- change contact handle or payment method;
- co-sign or change a couple link or veto mode;
- export data;
- view the device list;
- start account recovery.

It is **never** required for Burn, Panic, report, block, subscription cancel or account deletion. Deletion requires B1 only, so it stays one step (ToS cl. 16.3, 25.1).

Before V1: a V0 member has no B2 reference, so they can only continue onboarding on the device they signed up on. Login on any other device is unavailable until V1 is complete.

Downgrade resistance: the server decides the surface from attestation (native) and from the registered authenticator's properties. It does not trust a client claim. The desktop path still requires a second, separately registered authenticator, so claiming "desktop" from a phone gains nothing.

Account recovery (lost or replaced device, FR-061):

1. B2 liveness match.
2. Fresh V2 age assurance with an ID match, whose document anchor must equal the account's stored anchor_doc (same document).
3. A 24-hour hold, cancellable from any still-registered device.
4. New passkey registration.

No staff member can bypass recovery.

4. Data model

Three isolated PostgreSQL 16 clusters with separate credentials and separate KMS domains. **No plaintext join between identity and profile.** The join key is `HMAC(identity_id, HSM_pepper_join)`, computed only inside the identity service. [A §9.2, §10]

4.1 IDENTITY-DB

```
CREATE TABLE identity (
  identity_id      UUID PRIMARY KEY,
  contact_hash     BYTEA UNIQUE NOT NULL,    -- HMAC(phone|email, HSM pepper_contact)
  contact_ct       BYTEA,                    -- encrypted under root key; for OTP/notices only
  age_assured      BOOLEAN NOT NULL DEFAULT FALSE,
  assurance_tier    TEXT CHECK (assurance_tier IN ('V0','V1','V2','V3')),
  assurance_token  TEXT,                    -- provider ref. NO ID image, NO DOB
  assurance_at     TIMESTAMPTZ,
  billing_ref      TEXT,                    -- PSP token. No PAN, ever
  jurisdiction     CHAR(2) NOT NULL,
  root_key_ref     TEXT NOT NULL,           -- KMS handle; destroy = erasure
  created_at       TIMESTAMPTZ DEFAULT now(),
  deleted_at       TIMESTAMPTZ
);

-- v1.1 addition (E-02): ban anchors can only be written at ban time if candidates exist.
-- Candidates are erased with the account unless a permanent ban copies them to BAN-DB.
CREATE TABLE anchor_candidate (
  identity_id      UUID PRIMARY KEY REFERENCES identity,
  anchor_doc       BYTEA NOT NULL,          -- HMAC(doc_number, HSM pepper_ban)
  anchor_device    BYTEA[] NOT NULL,        -- HMAC(attestation_id, pepper_ban)
  anchor_payment   BYTEA,                  -- HMAC(bsp_instrument_token, pepper_ban)
  created_at       TIMESTAMPTZ DEFAULT now()
);

CREATE TABLE device (
  device_id        UUID PRIMARY KEY,
  identity_id      UUID NOT NULL REFERENCES identity,
  device_pubkey    BYTEA NOT NULL,          -- DPoP + MLS credential binding
  device_hash      BYTEA NOT NULL,          -- HMAC(attestation_id, pepper_device)
  platform         TEXT CHECK (platform IN ('ios','android','web')),
  surface          TEXT CHECK (surface IN ('mobile_app','mobile_pwa','desktop')), -- v1.2
  revoked_at       TIMESTAMPTZ
);

-- v1.2 addition (D-16): login authenticators. No biometric data is stored by TRYST.
CREATE TABLE authenticator (
  credential_id    BYTEA PRIMARY KEY,      -- WebAuthn credential id
  identity_id      UUID NOT NULL REFERENCES identity,
  device_id        UUID REFERENCES device, -- NULL for roaming security keys
  public_key       BYTEA NOT NULL,
  aaguid           UUID,
  kind             TEXT CHECK (kind IN ('platform_biometric','platform_uv','roaming_key')),
  bio_enrolment_bound BOOLEAN NOT NULL,    -- invalidated on biometric enrolment change
  sign_count       BIGINT DEFAULT 0,
  created_at       TIMESTAMPTZ DEFAULT now(),
  revoked_at       TIMESTAMPTZ
);

CREATE TABLE login_reference (
  identity_id      UUID PRIMARY KEY REFERENCES identity,
  provider         TEXT NOT NULL,
  reference_token  TEXT NOT NULL,          -- provider ref only; deleted at provider on erasure
  enrolled_at     TIMESTAMPTZ NOT NULL
);

CREATE TABLE auth_event (
  event_id         UUID PRIMARY KEY, identity_id UUID, device_id UUID,
  kind             TEXT,                  -- login_ok|login_fail|step_up|recovery_*|bio_change
  factors TEXT[], surface TEXT, created_at TIMESTAMPTZ DEFAULT now()
);
```

4.2 BAN-DB (separate store, separate key domain; survives erasure)

```
CREATE TABLE ban_anchor (  
  anchor_id          UUID PRIMARY KEY,  
  anchor_doc         BYTEA, anchor_device BYTEA[], anchor_payment BYTEA,  
  anchor_biometric   BYTEA,                -- NULL unless counsel confirms lawful basis (D-08)  
  reason_code        TEXT NOT NULL,         -- unsafe_ground enum; NOT free text, NOT evidence  
  tier               TEXT CHECK (tier IN ('permanent','law_enforcement_referred')),  
  reviewed_by        TEXT NOT NULL,         -- staff id; always a human decision  
  created_at         TIMESTAMPTZ DEFAULT now(),  
  expires_at         TIMESTAMPTZ,          -- NULL for permanent; destroyed on successful appeal  
  consortium_share   BOOLEAN DEFAULT FALSE -- L5, gated on opinion (Q8)  
);
```

4.3 PROFILE-DB

```
CREATE TABLE profile (
  profile_id      UUID PRIMARY KEY,
  segment         TEXT CHECK (segment IN ('S1','S2','S3','S4','S5','S6')),
  entity_type     TEXT CHECK (entity_type IN ('solo','couple')),
  display_name    TEXT,
  pronouns        TEXT,
  gender_identity TEXT,
  seeking         TEXT[],
  modes           TEXT[],
  age_band        TEXT,
  city_geohash5   TEXT,
  bio_ct          BYTEA,
  visibility      TEXT CHECK (visibility IN ('public','matched','incognito')),
  status          TEXT DEFAULT 'active',
  version         INT NOT NULL DEFAULT 1,
  created_at      TIMESTAMPTZ DEFAULT now()
);

CREATE TABLE couple_invite (
  invite_id      UUID PRIMARY KEY,
  inviter_profile UUID NOT NULL REFERENCES profile,
  token_hash     BYTEA NOT NULL,
  inviter_device_hash BYTEA NOT NULL,
  expires_at     TIMESTAMPTZ NOT NULL,
  consumed_at    TIMESTAMPTZ
);

CREATE TABLE couple_link (
  couple_id      UUID PRIMARY KEY,
  couple_profile UUID NOT NULL REFERENCES profile,
  profile_a      UUID NOT NULL REFERENCES profile,
  profile_b      UUID NOT NULL REFERENCES profile,
  a_cosigned_at  TIMESTAMPTZ NOT NULL,
  b_cosigned_at  TIMESTAMPTZ NOT NULL,
  a_device_hash  BYTEA NOT NULL,
  b_device_hash  BYTEA NOT NULL,
  approved_version INT NOT NULL,
  veto_mode      TEXT DEFAULT 'either' CHECK (veto_mode IN ('either','both')),
  dissolved_at   TIMESTAMPTZ,
  CONSTRAINT distinct_devices CHECK (a_device_hash <> b_device_hash),
  CONSTRAINT distinct_members CHECK (profile_a <> profile_b)
);

CREATE TABLE intent_vector (
  profile_id      UUID PRIMARY KEY REFERENCES profile,
  intent_shape    TEXT[],
  structure       TEXT,
  desire_tags     JSONB,
  hard_limits     TEXT[] NOT NULL,
  discretion       JSONB,
  logistics        JSONB,
  min_counterparty_tier TEXT DEFAULT 'V2',
  embedding        VECTOR(128),
  private_fields  TEXT[],
  confirmed_at    TIMESTAMPTZ,
  updated_at      TIMESTAMPTZ
);

CREATE TABLE revealed_preference (
  profile_id      UUID PRIMARY KEY REFERENCES profile,
  embedding        VECTOR(256),
  alpha           REAL CHECK (alpha BETWEEN 0.15 AND 0.85),
  divergence       REAL, confidence REAL, n_signals INT,
  personalisation_enabled BOOLEAN DEFAULT TRUE,
  computed_at     TIMESTAMPTZ
);

CREATE TABLE discretion_policy (
  profile_id      UUID REFERENCES profile,
  version         INT,
  policy_ct       BYTEA NOT NULL,
  zones_geohash   TEXT[],
  created_at      TIMESTAMPTZ DEFAULT now(),
  PRIMARY KEY (profile_id, version)
);
```

```

CREATE TABLE media (
  media_id          UUID PRIMARY KEY,
  owner_profile     UUID REFERENCES profile,
  object_key        TEXT NOT NULL,          -- client-encrypted ciphertext
  tier               TEXT CHECK (tier IN ('profile_public','profile_private','thread')),
  explicit          BOOLEAN DEFAULT FALSE,   -- web surface only, V2 only
  moderation_state  TEXT DEFAULT 'pending',
  expires_at        TIMESTAMPTZ,            -- thread media: 7d default
  created_at        TIMESTAMPTZ DEFAULT now()
);

CREATE TABLE reveal_grant (
  grant_id          UUID PRIMARY KEY,
  owner_profile     UUID, recipient_profile UUID,
  scope             TEXT,                   -- reveal_scope enum
  media_id          UUID,                  -- NULL for voice/video scope
  owner_version     INT,
  wrapped_key_ref   TEXT,                   -- destroyed on revoke
  expires_at        TIMESTAMPTZ, revoked_at TIMESTAMPTZ,
  created_at        TIMESTAMPTZ DEFAULT now()
);

CREATE TABLE slate_item (
  slate_id          UUID, profile_id UUID, candidate_id UUID,
  score             REAL, model_version TEXT,
  explanation_codes TEXT[],                -- 2-4 non-sensitive codes
  exploration        BOOLEAN, paid_label BOOLEAN,
  created_at        TIMESTAMPTZ,
  PRIMARY KEY (slate_id, candidate_id)
);

CREATE TABLE intent (
  intent_id         UUID PRIMARY KEY,
  from_profile       UUID, to_profile UUID,
  keys_spent         INT NOT NULL DEFAULT 0,
  via_envoy_proposal UUID,                 -- ENV0Y tier; always human-approved (D-04)
  state              TEXT NOT NULL,         -- intent_state enum
  created_at         TIMESTAMPTZ, delivered_at TIMESTAMPTZ
);

CREATE TABLE handshake (
  handshake_id       UUID PRIMARY KEY,
  a_profile UUID, b_profile UUID,
  transcript_ct       BYTEA,                -- encrypted agent turns, 30d TTL
  brief_a_ct BYTEA, brief_b_ct BYTEA,
  compatibility REAL, meet_probability REAL,
  a_decision TEXT, b_decision TEXT,         -- handshake_decision
  opened_thread       UUID,                 -- non-null ONLY on double accept
  created_at         TIMESTAMPTZ
);

CREATE TABLE thread (
  thread_id          UUID PRIMARY KEY,
  handshake_id        UUID,                 -- NULL in P1 control cohort
  mls_group_id        BYTEA NOT NULL,
  participant_profiles UUID[] NOT NULL,
  retention_days       INT DEFAULT 30 CHECK (retention_days IN (7,30,90)),
  safety_state         TEXT DEFAULT 'normal', -- normal|frozen|closed
  burned_at           TIMESTAMPTZ,
  created_at          TIMESTAMPTZ
);

CREATE TABLE message_envelope (
  thread_id UUID, seq BIGINT, sender_device UUID,
  ciphertext BYTEA NOT NULL, expires_at TIMESTAMPTZ NOT NULL,
  PRIMARY KEY (thread_id, seq)
);

CREATE TABLE meet (
  meet_id UUID PRIMARY KEY, thread_id UUID,
  state TEXT,          -- proposed|confirmed|cancelled|occurred
  window_start TIMESTAMPTZ, checkin_due_at TIMESTAMPTZ,
  aftercare_due_at TIMESTAMPTZ -- +48h
);

CREATE TABLE exclusion_ring (
  profile_id UUID, excluded_profile_id UUID,

```

```

source TEXT,                                -- psi_contact|device_graph|manual|payment|geo|wifi
confidence REAL,
PRIMARY KEY (profile_id, excluded_profile_id)
); -- ALWAYS bidirectional, ALWAYS silent

CREATE TABLE consent_ledger (
  consent_id UUID PRIMARY KEY, profile_id UUID NOT NULL,
  kind TEXT NOT NULL,                        -- consent_kind enum
  granted BOOLEAN NOT NULL, scope JSONB,
  prev_hash BYTEA, entry_hash BYTEA NOT NULL, -- hash chain; append-only
  created_at TIMESTAMPTZ DEFAULT now()
);

CREATE TABLE aftercare (
  meet_id UUID PRIMARY KEY, reporter UUID, counterparty UUID,
  occurred BOOLEAN, as_represented SMALLINT, would_repeat BOOLEAN, safety_flag BOOLEAN,
  created_at TIMESTAMPTZ
); -- NEVER surfaced to counterparty. NEVER a public rating.

```

4.4 SAFETY-DB (T&S; access only through the trust console)

```

CREATE TABLE guardian_signal (              -- scores and reason codes only; no content
  signal_id UUID PRIMARY KEY, thread_id UUID, subject_profile UUID,
  model_id TEXT, model_version TEXT, score REAL, reason_code TEXT,
  source TEXT CHECK (source IN ('on_device','server_profile','server_graph')),
  created_at TIMESTAMPTZ
);
CREATE TABLE report_case (
  case_id UUID PRIMARY KEY, reporter_pseudo BYTEA, subject_profile UUID,
  category TEXT, severity SMALLINT, disclosed_content_ref TEXT, -- reporter-disclosed only
  training_consent BOOLEAN DEFAULT FALSE, -- D4 source; never a condition of action
  state TEXT, sla_due_at TIMESTAMPTZ, legal_hold BOOLEAN DEFAULT FALSE,
  created_at TIMESTAMPTZ
);
CREATE TABLE enforcement (
  enforcement_id UUID PRIMARY KEY, subject_profile UUID,
  action TEXT, ground TEXT, reason_code TEXT,
  decided_by TEXT, -- REQUIRED for permanent (human)
  appeal_state TEXT, appeal_reviewer TEXT, created_at TIMESTAMPTZ
);
CREATE TABLE staff_access_log (             -- immutable (WORM)
  entry_id UUID PRIMARY KEY, staff_id TEXT, approver_id TEXT,
  resource TEXT, reason_code TEXT, created_at TIMESTAMPTZ
);

```

4.5 Commerce and operations (PROFILE-DB schema commerce, ops)

```

CREATE TABLE commerce.entitlement (profile_id UUID, tier TEXT, addon_ghost BOOLEAN,
  status TEXT, source TEXT CHECK (source IN ('web','voucher','app_store','play')),
  renews_at TIMESTAMPTZ, renewal_notice_sent_at TIMESTAMPTZ, PRIMARY KEY (profile_id));
CREATE TABLE commerce.keys_ledger (entry_id UUID PRIMARY KEY, profile_id UUID,
  delta INT, reason TEXT, intent_id UUID, balance_after INT, created_at TIMESTAMPTZ);
CREATE TABLE commerce.deposit (profile_id UUID PRIMARY KEY, amount_minor INT,
  state TEXT, credited_keys INT, created_at TIMESTAMPTZ);
CREATE TABLE ops.deletion_job (job_id UUID PRIMARY KEY, subject_ref BYTEA,
  systems TEXT[], started_at TIMESTAMPTZ, completed_at TIMESTAMPTZ,
  completion_proof BYTEA); -- signed record of key destruction + purges

```

5. Edge API contract (v1)

Cross-cutting middleware, in order:

1. TLS 1.3
2. OAuth2 access token **DPoP-bound** to the device key
3. Jurisdiction gate
4. Tier gate (§3.2)
5. Article 9 consent check
6. Per-endpoint rate limit
7. Idempotency (Idempotency-Key header required on all POST/PUT/DELETE)

The **Min** column is the minimum verification tier. sys means server-to-server only.

5.1 Auth, devices, assurance

Method	Path	Min	Contract
POST	/v1/auth/start	—	Sign-up only. {contact} → OTP sent. Uniform response whether or not the account exists. Never a login factor (§3.3)
POST	/v1/auth/verify	—	Sign-up only. {contact, otp, jurisdiction?} → onboarding session (device binding via DPoP is a P1 hardening item, D-24). Refused with account_exists_sign_in if the account already has a passkey: an OTP can never add one. Ban-anchor check runs here for new identities (03 §8)
POST	/v1/auth/passkeys/register (begin / finish)	V0	Register B1 (mobile) or F1 (desktop) WebAuthn credential; step-up required after the first one
POST	/v1/auth/login/begin	—	{contact_hint?} → WebAuthn challenge + required_factors (decided by the server, §3.3). Uniform response
POST	/v1/auth/login/passkey	—	WebAuthn assertion (B1 or F1) → login_state = factor1_ok (no session yet)
POST	/v1/auth/login/liveness	—	Mobile: start B2 provider session → provider callback (sys) → session issued only on match
POST	/v1/auth/login/cross-device	—	Desktop: → QR token (60 s); GET .../{id} polls status → session issued only on approval
POST	/v1/auth/approvals/{qr_token}	V1	Called from the registered mobile app with a fresh B1 assertion; approves or denies the desktop login
POST	/v1/auth/step-up	V0	Same factor flow as login → step-up token (5 min) for §3.3 sensitive actions
POST	/v1/auth/enrol-device · /v1/auth/enrol-device/redeem	V0 (onboarding)	During onboarding only: one-time 10-minute link that lets the member's phone register its own passkey on the same account. Refused once the account is sealed (device add then needs step-up)
POST	/v1/auth/login/second-key/begin · /v1/auth/login/second-key	—	Desktop F2 alternative: assertion with a different registered security key
POST	/v1/auth/approvals/{qr_token}/begin	—	Assertion options restricted to the account's registered phone passkeys
GET	/v1/me	V0	Session summary (kind, surface, tier, passkey count). No contact or identity data
POST	/v1/auth/logout	V0	Ends this session
POST	/v1/auth/recovery/begin · /finish	—	B2 + fresh V2 with matching anchor_doc → 24 h hold → new passkey (FR-061)
POST	/v1/auth/refresh	V0	DPoP refresh
GET / DELETE	/v1/devices · /v1/devices/{id}	V0	List and revoke devices (step-up required , §3.3)
POST	/v1/assurance/session	V0	{kind: liveness age} → provider hand-off URL
POST	/v1/assurance/callback	sys	Signed provider token → is_adult boolean + anchor candidates. Raw document and biometric never stored
POST	/v1/attestations/conduct	V1	Conduct declaration (ToS cl. 3.5) → consent ledger. Required to complete V2

5.2 Account, consent, data rights

Method	Path	Min	Contract
GET / POST	/v1/consents	V0	Read the ledger; append grant or withdrawal (never mutate)
POST	/v1/account/pause	V0	Removed from discovery immediately
POST	/v1/emergency-lock	—	No session needed. {contact_hint, emergency_code} → freeze sessions, hide profile, disable media links (FR-064). Uniform response; rate-limited; the code is single-use and reissued after recovery
POST	/v1/account/vanish	V0	Hidden + threads closed; account kept
POST	/v1/account/export	V0	SAR/portability export (encrypted), < 30 days; target < 24 h

Method	Path	Min	Contract
DELETE	/v1/account	V0	Cryptographic erasure, < 60 s p99. Returns deletion_job_id and, later, completion_proof

5.3 Intake, profile, discretion

Method	Path	Min	Contract
POST	/v1/intake/turn	V0	Cartographer conversational turn
GET	/v1/intake/vector	V0	Proposed IntentVector (structured)
POST	/v1/intake/vector/confirm	V0	Member confirms or edits each field; stored with confirmed_at
PATCH	/v1/profile	V0	Pseudonymous fields only
POST / DELETE	/v1/profile/media · /{id}	V1	Client-encrypted blob + wrapped keys; public-tier photos go through server moderation (03 §7.5)
PUT	/v1/discretion/policy	V0	New version (§6.4). Rejects any change that broadens exposure unless confirm_broaden=true is sent from a fresh biometric/PIN session
PUT	/v1/discretion/zones	V0	Home/work exclusion geofences (geohash-6 cells)
POST	/v1/exclusion/psi/blind	V2	OPRF round 1
POST	/v1/exclusion/psi/commit	V2	OPRF round 2 → suppression written. Nothing returned about the intersection
POST	/v1/burn	V0	Renamed from /v1/panic (E-21) . Revoke device keys, sign out everywhere; the client wipes local data

5.4 Couples

Method	Path	Min	Contract
POST	/v1/couple/invite	V2	→ co-sign token (24 h)
POST	/v1/couple/cosign	V2	Must come from a device whose hash differs from the inviter's . Both partners must approve the exact profile version before publication
PUT	/v1/couple/{id}/veto-mode	VC	Change requires both partners' confirmation
DELETE	/v1/couple/{id}	V2	Either partner dissolves immediately; no reason shown to the other

5.5 Discovery and learning

Method	Path	Min	Contract
GET	/v1/slate	V2	Up to 18 cards incl. 15% exploration; each with explanation_codes (2-4) and paid_label. Re-authorised against the current Discretion Policy and block graph on every read
POST	/v1/intent	V2	{target, keys_spent} → 202 {intent_id, state}. Inbound-cap overflow returns state=sent to the sender and is queued (the cap is never revealed)
POST	/v1/pass · /v1/save	V2	Private
POST	/v1/signals:batch	V0	≤ 50 behaviour events (§8)
GET	/v1/preferences/insights	V2	Category-level view of what has been learned (never a raw vector)
POST	/v1/preferences/correction	V2	"More / less like this" and category corrections
POST	/v1/preferences/reset	V2	Neutralise RevealedPreferenceVector; α returns to 0.15; rebuild from the IntentVector
PUT	/v1/preferences/personalisation	V2	Turn personalisation on or off

5.6 Handshake, ENVOY tier, reveal

Method	Path	Min	Contract
—	(system) handshake created on mutual intent	—	Not client-callable in v1.1 (E-29)
GET	/v1/handshake/{id}/brief	V2	The caller's own asymmetric Brief only (§6.2)
POST	/v1/handshake/{id}/decision	V2	accept decline amend (+ amend fields). Thread opens only when both are accept (couples: per veto_mode)

Method	Path	Min	Contract
GET	/v1/envoy/proposals	V2 + envoy	Agent-proposed intents with rationale
POST	/v1/envoy/proposals/{id}/approve	V2 + envoy	One-tap human approval required for every outbound intent (D-04)
POST	/v1/reveal-grants	V2	{recipient, scope, media_id?, expires_at}
DELETE	/v1/reveal-grants/{id}	V2	Destroys the wrapped key; access ends at once (p95 < 5 s, including CDN)

5.7 Threads (E2EE; the server relays ciphertext only)

Method	Path	Min	Contract
POST / GET	/v1/threads/{id}/messages	V2	MLS ciphertext envelopes; expires_at set from retention_days
PUT	/v1/threads/{id}/retention	V2	7 / 30 / 90 days; the shortest elected by any participant wins
POST	/v1/threads/{id}/meet	V2	propose confirm cancel
POST	/v1/threads/{id}/checkin	V2	Safe Meet check-in; a missed check-in triggers silent escalation to the trusted contact
POST	/v1/threads/{id}/burn	V2	Revoke keys on both sides; thread closed
POST	/v1/aftercare/{meet_id}	V2	Private outcome (48 h after the meet)

5.8 Safety

Method	Path	Min	Contract
POST	/v1/report	V0	{subject, category, disclosed_content?, training_consent?} → SLA routing. The reporter is never revealed
POST	/v1/block	V0	Bidirectional and immediate; removes thread access
POST	/v1/guardian/signals	V2	On-device scores: {thread_id, model_id, version, score, reason_code}. Never content
POST	/v1/safety/panic	V0	Emergency. Live location + preserved evidence → emergency escalation. Overrides TTLs (P7)
POST	/v1/safety/share-plan	V2	Encrypted, unbranded, auto-deleting plan for a trusted contact
POST	/v1/appeals	V0	Appeal an enforcement action; a different reviewer decides within 10 working days
POST	/v1/attestations/dbs	V2	Self-submitted Basic DBS through the issuing route; shown as a dated fact (P3)

5.9 Commerce (web-first)

Method	Path	Min	Contract
GET	/v1/entitlements	V0	Tier, add-ons, Keys balance, renewal date
POST	/v1/deposit	V1	£9.99 refundable deposit → Keys credit
POST	/v1/keys/purchase	V1	Keys packs
POST	/v1/vouchers/redeem	V0	Offline voucher/gift code
POST	/v1/subscription	V1	Neutral, truthful descriptor enforced at the PSP; pre-contract and cooling-off notices
POST	/v1/billing/stripe/webhook	sys	Stripe events (signature-verified, idempotent) → entitlement and Keys-ledger updates (v1.2, D-22)
POST	/v1/subscription/cancel	V0	One step; no retention flow

5.10 Error model

RFC 9457 application/problem+json with a stable code.

- **Existence must not leak.** Anything not visible to the caller, for any reason (block, ExclusionRing, policy, jurisdiction of the target), returns the same 404 not_found.

HTTP	code	Meaning
401	auth_required, dpop_invalid, second_factor_required (+factor), step_up_required, authenticator_invalidated (biometric enrolment changed)	
409	account_exists_sign_in	Sign-up attempted for an account that already has a passkey
503	liveness_unavailable	B2 provider unavailable; login fails closed
403	tier_required (+required_tier), consent_required (+consent_kind), jurisdiction_blocked, entitlement_required	
404	not_found	Uniform for invisible resources
409	idempotency_conflict, version_conflict	
402	keys_insufficient	
422	validation_failed, broadening_requires_confirmation, same_device_cosign	
429	rate_limited (+retry_after)	

6. Agent and policy schemas

6.1 Envoy negotiation turn (internal bus; Protobuf mirror)

Six dimensions only: intent alignment, structure alignment, boundary intersection, logistics feasibility, discretion compatibility, verification parity. [A §5.4]

```
{
  "turn": 3,
  "from": "envoy:usr_7f3a",
  "to": "envoy:usr_9lbc",
  "assert": {
    "intent_shape": "one_off",
    "structure": "attached_undisclosed",
    "notice_required_hours": 48,
    "availability": [{"dow": [2,3], "start": "18:00", "end": "23:00", "tz": "Europe/London"}],
    "travel_radius_km": 25,
    "hard_limits": ["no_photos_exchanged", "no_home_visit", "no_overnight"],
    "verification_tier": "V2",
    "min_counterparty_tier": "V2"
  },
  "query": ["intent_shape", "notice_required_hours", "hosting_capability"],
  "private_fields": ["exclusion_zones"]
}
```

Rules:

- No free-text fields.
- private_fields are computable but **never** appear in any output.
- Counterparty turns are untrusted data.
- At most 6 turns per handshake.
- The transcript is encrypted with a 30-day TTL.

6.2 Brief (caller-specific, asymmetric)

```
{
  "label": "agent_generated",
  "compatibility": 0.81,
  "aligned": ["intent_shape", "availability", "travel_radius", "verification"],
  "friction": [{"field": "hosting", "detail_code": "neither_can_host",
    "suggested_resolution": "hotel_or_venue"}],
  "blocking": [],
  "recommended_next": "open_thread",
  "estimated_meet_probability": 0.34
}
```

detail_code values map to localised copy on the client. A Brief never contains anything the counterparty marked private, nor any hard limit the counterparty has not chosen to disclose.

6.3 IntentVector (wire view)

The fields follow intent_vector (§4.3). desire_tags uses yes|curious|no. hard_limits is drawn from a controlled vocabulary (contracts/vocab/hard_limits.yaml). Free-text goals entered during intake are converted by Cartographer into controlled values, and the member confirms them.

6.4 Discretion Policy (versioned)

```
{
  "version": 7,
  "visibility": "incognito",
  "distance_band_km": 10,
  "zones": {"home": ["gcpvj0"], "work": ["gcpuvx"]},
  "zone_behaviour": "hide_profile_while_inside",
  "reveal_rules": {"default_scope": "blurred_gallery", "require_mutual_handshake": true,
    "max_grant_hours": 72},
  "notifications": {"enabled": false, "text": "1 update"},
  "device_risk": "high",
  "min_counterparty_tier": "V2",
  "retention": {"messages_days": 30, "media_days": 7}
}
```

Monotonic safety: Curtain may tighten any field automatically. Broadening needs an explicit member action, a fresh local authentication and a new version. Every read path evaluates (requester, resource, reveal_grant, block graph, ExclusionRing, discretion version, safety_state, jurisdiction) — ABAC [B §9].

7. Retention and erasure

This table is canonical; 03 and 04 implement it, and the ToS (cl. 15) and Privacy Notice must match it.

Data	Default	Member options	Mechanism
Messages	30 d	7 / 90 d	Envelope expires_at; the client purges its copy too
Thread media	7 d	— (GHOST: shorter)	Per-object key destroyed at expiry
Profile media	Until removed	—	Per-object key; removal destroys it
Handshake transcripts	30 d	—	Encrypted; TTL
Behaviour events (raw)	180 d → aggregated features	Reset / opt-out	ClickHouse TTL
Aftercare	365 d, pseudonymous	—	
Logs	90 d, IP truncated at ingest	—	
Consent ledger	Life of account + statutory period	Export	Hash chain
Ban anchors	Permanent (or until a successful appeal)	Appeal	Declared exception (ToS cl. 15.3)
Evidence under legal hold	Duration of hold	—	Declared exception (ToS cl. 17.4); narrowly scoped
Panic-preserved location/evidence	As required by emergency services	—	Declared exception (ToS cl. 20; P7)
B2 login face reference (v1.2)	Life of account; held by the provider, not TRYST	—	Deleted at the provider on erasure; deletion receipt included in completion_proof
Auth events	90 d, IP truncated	—	TTL
Backups	35 d rolling, disclosed in the Privacy Notice	—	Stored under per-member keys, so an erased member's data in backups is undecryptable from the moment of erasure

Account deletion:

- 1. Destroy the root key in KMS.
- 2. Remove from discovery immediately.
- 3. Purge primaries.
- 4. Ciphertext in backups becomes undecryptable.
- 5. Issue a signed completion_proof.

Target < 60 s p99 (NFR-01).

8. Event taxonomy

Events go to Redpanda with schema-registry enforcement, are pseudonymous at rest, and are batched from clients through `/v1/signals:batch`. [A §6.1] **No message content, ever.**

Class	Events
Attention	card_impression, dwell_ms, photo_index_viewed, bio_scroll_depth, profile_revisit, expand_tag
Decision	intent_sent, pass, save, undo, hesitation_ms
Conversation	msg_sent, reply_latency_s, msg_len_bucket, typing_abandoned, question_ratio, thread_decay — computed on device from metadata
Handshake	brief_viewed, brief_accepted, friction_dismissed, amend_requested
Resolution	logistics_agreed, meet_scheduled, meet_confirmed, no_show, aftercare_score
Discretion	burn_invoked, session_aborted, geofence_hit, screenshot_detected, decoy_opened
Negative	report, block, unmatched, ghost_detected, refund_request

Envelope: {event_id, pseudo_subject, type, ts, context:{surface, app_version}, consent_basis, ttl}. The pseudo-subject rotates per training window (03 §4).

9. Couple and group semantics

- **Veto:** veto_mode=either means either partner can decline and a decline is final. both means both partners must accept before the couple's intent or handshake decision counts.
- **THIRD:** a couple(S3) ↔ solo(S4) handshake involves both partners' Envoy contexts. The solo's veto is always respected. The thread is an MLS group of 3 members.
- **QUAD:** couple ↔ couple; 4 MLS members. A thread opens only when both couples' veto rules are satisfied.
- **Leaving:** any participant can leave a group thread privately. The thread then closes for that member, and others see a neutral notice.
- **Dissolution:** immediately hides the couple profile, closes its threads and keeps both individual accounts.

10. Requirement catalogue

Owner: BE backend, FE frontend (native + web), Both. **Phase:** first phase in which the requirement must be met (01 §12). **Src:** A/B/C plus section.

10.1 Functional requirements

ID	Requirement	Owner	Phase	Src
FR-001	Discovery, intents, Envoy, threads, media and meets are locked below V2; nothing from other members is shown during verification	Both	P1	A §8.1, B FR-001, E-01
FR-002	Pseudonymous identity: no legal name, employer or exact address is ever required or shown	Both	P1	A §14.2, B FR-002
FR-003	Cartographer conversational intake produces an IntentVector; the member confirms every structured field; re-run quarterly or on drift	Both	P1	A §5.1, B
FR-004	A member can enable several modes, each with its own visibility and filters	Both	P1	B FR-003
FR-005	Versioned Discretion Policy enforced on every read path; cannot be broadened silently	BE	P1	B FR-004
FR-006	Couple profile requires VC: both at V2, distinct devices, 24 h invite, both approve the exact profile version	Both	P1	A §8.1, B FR-005
FR-007	Either partner can dissolve a couple immediately, with no reason shown to the other	Both	P1	Part B cl. 7.1
FR-008	Reveal ladder (alias → blurred → photo set → voice → video); each grant records recipient, scope, version and expiry, and is revocable	Both	P1	B FR-006
FR-009	Incognito members are discoverable only after their own outbound intent	BE	P1	B FR-007

ID	Requirement	Owner	Phase	Src
FR-010	ExclusionRing PSI: on-device blinding; the server never learns contacts; suppression is bidirectional and silent	Both	P2	A §9.3, B FR-008
FR-011	Home/work exclusion zones suppress both sides and hide the profile while inside them	Both	P1	A §5.5, §9.4
FR-012	Daily slate of up to 18 cards with a 15% exploration floor	Both	P1	A §4, §7
FR-013	Intent costing: free tier 3/week; Keys 5 (solo→solo) / 15 (couple→solo); refund on reply; inbound cap default 12/24 h (adjustable 4–40), overflow queued and re-scored	BE	P1	A §7.4, §15.2
FR-014	Hard limits, zones, structure compatibility, ExclusionRing, blocks, tier parity and jurisdiction are Stage 0 pre-retrieval gates; any violation is a P0 incident; enforced by property tests	BE	P1	A §6.2, §7
FR-015	Handshake on mutual intent: six-dimension Envoy negotiation; asymmetric Briefs; private fields never emitted; Briefs labelled	Both	P2	A §5.4
FR-016	A thread opens only on double human accept (P2+); in P1 the control cohort opens on mutual intent	Both	P1	A §4, §17
FR-017	E2EE messaging (MLS, RFC 9420); the server stores ciphertext only and cannot read content	Both	P1	A §9.2
FR-018	Thread media is client-encrypted with a per-object key wrapped per recipient; revocation destroys the wrapped key	Both	P1	A §9.2
FR-019	Thread burn revokes keys on both sides and closes the thread	Both	P1	A §11
FR-020	Meet propose/confirm/cancel; Safe Meet checklist (public venue first); check-in timer with silent escalation	Both	P1	A §19.7
FR-021	Share My Plan: encrypted, unbranded, auto-deleting share with a trusted contact	Both	P1	A §19.7
FR-022	Panic: escalation to 999 with live location and preserved evidence; a separate control from Burn	Both	P1	A §19.7, E-21
FR-023	Private Aftercare prompt 48 h after a meet; never visible to the counterparty	Both	P1	A §5.7
FR-024	Preference transparency: view learned categories, correct, disable personalisation, reset learning	Both	P2	B FR-010
FR-025	Each slate card shows 2–4 non-sensitive reasons and a correction control	Both	P2	B §6.2
FR-026	Notifications off by default; if enabled, content-free with member-chosen neutral wording (default "You have an update") — never a name, image or excerpt	FE	P1	A §9.4, B §4
FR-027	Decoy skin (neutral icon/name) and silent duress PIN opening a populated innocuous state (native)	FE	P1	A §9.4
FR-028	Burn: wipe local storage, revoke device keys, sign out everywhere; confirmable in < 2 s	Both	P1	A §9.4
FR-029	Quick exit hides content and clears the client view and cache in < 500 ms on all surfaces; onboarding states plainly that it cannot hide network or device history	FE	P1	B FR-011, B §4
FR-030	Screenshot protection (FLAG_SECURE) or detection with counterparty notice (iOS); obscured app switcher; no cloud backup; keys in Secure Enclave/StrongBox	FE	P1	A §9.4
FR-031	Entering a blocked jurisdiction auto-locks to the decoy and hides the profile	Both	P1	A §9.4, §14.5
FR-032	Report and block from every surface; the reporter is never revealed; reporters may disclose decrypted content and separately opt in to training use	Both	P1	A §21.3, Part B cl. 19
FR-033	Guardian M6a–M6e run on device for E2EE threads and send only score + reason code	Both	P1	A §21.6
FR-034	Escalation matrix routing and SLAs (03 §7.3)	BE	P1	A §8.4
FR-035	Enforcement ladder; permanent actions need a named human; appeal to a different reviewer within 10 working days	BE	P1	A §19.6
FR-036	Ban anchors are checked before account creation; a match is silently refused	BE	P1	A §19.3 L2
FR-037	No commercial sex: solicitation and payment patterns go to review; no payments between members	Both	P1	A §8.4, B FR-015

ID	Requirement	Owner	Phase	Src
FR-038	Account deletion = cryptographic erasure < 60 s; immediate discovery removal; signed completion proof	Both	P1	A §9.2, B FR-012
FR-039	DSRs: SAR export including the consent ledger, rectification, portability	Both	P1	A §14.2
FR-040	Hash-chained consent ledger; Article 9 consent separate from the ToS accept	BE	P1	A §8.2
FR-041	Web-first billing: TRYST+ and Keys (P1), DUO, ENVOY and GHOST (P3); neutral truthful descriptor; vouchers; one-step cancel; pre-renewal and cooling-off notices	Both	P1	A §14.4, §15
FR-042	£9.99 refundable verification deposit credited as Keys	Both	P1	A §15.2
FR-043	Jurisdiction gate at the edge and at sign-up	BE	P1	A §14.5
FR-044	Conduct attestation at V2, logged separately	Both	P1	A §19.3 L3
FR-045	Basic DBS self-submission shown as a dated fact; Clare's Law explainer before the first meet (outcome never requested or stored)	Both	P3	A §19.3 L3
FR-046	Pause and vanish	Both	P1	B §3
FR-047	Profile-photo protection: EXIF strip, face-blur default, AI-generation and reverse-image checks on every upload, adversarial perturbation on public-tier photos	Both	P1	A §9.1 A8, §19.4
FR-048	ENVOY tier: agent-proposed outreach, each intent approved by the human with one tap	Both	P3	A §15.2, D-04
FR-049	Couple veto_mode (either/both); THIRD and QUAD matching and MLS group threads	Both	P1	A §3, §10
FR-050	Exposure-fairness monitor (M9) adjusts the fairness penalty when inbound Gini > 0.55 per city/segment	BE	P1	A §7.4
FR-051	Native store builds are sanitised (no explicit media or copy); explicit media lives only on web behind V2	FE	P1	A §14.3
FR-052	Semi-annual transparency report: removals by ground, appeals, overturn rate, disclosures	BE	P3	A §19.6, Part B cl. 18.3
FR-053	Travel mode: set a planned location for discovery (premium)	Both	P3	A §15.2
FR-054	In-product safety notice at intake and before a first meet (04 §7)	FE	P1	A §20.3
FR-055	Envoy and other agents never send messages to a human as the member	BE	P2	A §5.4
FR-056	Passkey-only authentication: no passwords; SMS/email OTP, magic links and support overrides are never login factors (contact OTP only at sign-up)	Both	P1	D-16
FR-057	Native app and mobile PWA login requires two biometric factors: B1 device-bound biometric passkey and B2 liveness face match to the account holder (§3.3)	Both	P1	D-16
FR-058	Desktop login (browser or desktop PWA) requires two factors: F1 passkey and F2 approval from a registered mobile device via B1 (or a second FIDO2 key) (§3.3)	Both	P1	D-16
FR-059	A change to the device's biometric enrolment invalidates its passkey and forces a full login (B1 re-registration + B2)	Both	P1	D-16
FR-060	Step-up re-authentication for sensitive actions (§3.3); never for Burn, Panic, report, block, cancel or deletion	Both	P1	D-16
FR-061	Account recovery needs B2 + fresh V2 matching the stored document anchor + 24 h hold cancellable from registered devices; no staff bypass	Both	P1	D-16
FR-062	Session limits: mobile 30 d max with biometric unlock and 2 min idle lock; desktop 12 h max, 15 min idle → full login	Both	P1	D-16
FR-063	Brand: logo used exactly as supplied and palette tokens applied on every branded surface; logo never shown in decoy mode, on lock screens, in notifications, app-switcher snapshots, descriptors or Share My Plan	FE	P1	D-19
FR-064	Emergency lock from a separate path: with only the contact handle and a one-time emergency code issued at sign-up (no login), freeze all sessions, hide the profile and disable all media links; unlocking needs full recovery (FR-061). Uniform response; abuse only locks, never exposes	Both	P1	B §3, §4

ID	Requirement	Owner	Phase	Src
FR-065	Location veil: distances shown only as bands, computed from approximate cells whose offset rotates at least daily; never live or exact home/work coordinates	Both	P1	B §4
FR-066	Conversation suggestions (openers, boundary questions, respectful replies) are opt-in, generated on device, draft-only and never sent without the member's own action	FE	P3	B §5, D-20
FR-067	Timing and availability matching use only member-provided windows, never routines inferred from location or activity	BE	P2	B §5
FR-068	Behavioural signals decay over time (default half-life 90 days); members can reset all inferred preferences (FR-024)	BE	P2	B §6.2
FR-069	Encrypted media safety: the sending device accepts only allow-listed image/video formats, re-encodes them (stripping active content) before encryption, and shows a consent/NCII warning before any intimate image is sent; public-tier media is malware-scanned on the server	Both	P1	B §9, §10
FR-070	Role separation: finance and billing staff cannot access profile, intent or safety data; T&S staff cannot see billing identity; enforced in the ABAC policy engine	BE	P1	B §9
FR-071	Viewing reporter-disclosed intimate evidence needs dual approval and a reason code, and is logged immutably	BE	P1	B §9
FR-072	Community Rules and enforcement prohibit fetishisation of, and harassment based on, protected characteristics; reportable as protected_trait_harassment	Both	P1	B §10
FR-073	Public marketing site built for search: server-rendered/static pages, XML sitemaps, canonical tags, hreflang (en-GB, en-IE, later wave-2 locales), schema.org markup (Organization, WebSite, Article, FAQPage, BreadcrumbList), descriptive internal links, a quick-exit button and a neutral-favicon option	FE	P1	v1.2
FR-074	Herald SEO agent: keyword and topic research, content briefs and drafts, on-page optimisation, internal-link graph maintenance, content refresh from performance data, rank tracking, broken-link repair, and link-earning outreach drafts; every publication and outreach message is approved by a human editor	BE	P1	v1.2, D-21
FR-075	Herald is isolated from member data: separate cloud account and network, no credentials for any member store, inputs limited to public web data, Search Console and cookieless aggregate site analytics; no member content, profiles or stories (unless separately consented and anonymised) are ever published	BE	P1	v1.2
FR-076	No member surface is indexable: the app origin serves noindex, nofollow and robots disallow; there are no public profile URLs; shared links never reveal a member	Both	P1	v1.2
FR-077	White-hat link policy: paid links, PBNs, link exchanges, spam, cloaking, doorway pages, unreviewed scaled AI content and fake reviews are prohibited; any paid or sponsored link carries rel="sponsored"; the link profile is audited monthly	BE	P1	v1.2, D-21
FR-078	Payments run through a processor-agnostic PaymentProvider layer with Stripe as primary (Checkout, Billing, Customer Portal one-step cancel, Radar, Tax) and a high-risk acquirer as the live backup; Stripe receives only the opaque billing reference and payment data; the neutral descriptor never names TRYST	Both	P1	v1.2, D-22
FR-079	Every IntentVector change (member edit or Cartographer re-run) is kept as an append-only, member-viewable version history with audit records	BE	P1	B §13
FR-080	Every agent action passes the AI Governance gate before it runs: unknown agents or unlisted actions are denied; irreversible, member_safety, money_movement and production_change actions always need a registered named human; autonomy levels L0-L3 and daily budgets cap everything else (09 §5.4)	BE	P1	founder OS brief, D-25
FR-081	Agents are declared by manifest (owner, level, actions, data scopes, budget). A member-data scope is refused unless the agent runs inside the member platform on a self-hosted model; no manifest may contain an act-as-member action	BE	P1	founder OS brief, P4
FR-082	Agent decisions and admin actions are written to hash-chained, append-only logs anchored hourly to write-once storage; a kill switch per agent and for all agents takes effect immediately and is itself logged	BE	P1	founder OS brief
FR-083	Member Command Centre: private insights (own activity charts, security activity, "why this slate" from reason codes, what TRYST has learned with reset); no engagement counters or streaks; works with Leave, decoy and emergency lock	Both	P2	founder OS brief, P1, P5
FR-084	BitriPay adapter behind PaymentProvider: sandbox supported; live keys refused unless the market is enabled for TRYST and BitriPay with a recorded legal opinion; deterministic idempotency keys; only amount, currency and a neutral descriptor are sent; signed webhooks verified	BE	P1	founder OS brief, D-26
FR-085	Self-healing: SLO burn triggers incident creation and pre-approved runbooks (rollback, restart, scale out, halt canary) within budget; code or model changes reach production only through a reviewed pull request	BE	P1	founder OS brief, D-28

ID	Requirement	Owner	Phase	Src
FR-086	Admin Super Control Centre on an internal origin: redacted by default; member-level access only via a case-scoped just-in-time grant approved by a second person (max 8 h); break-glass needs two people; every view and action logged	Both	P1	founder OS brief
FR-087	Connector register: every third-party connector records purpose, module, data sent and received, lawful basis, DPA, region, health, kill switch and owner; no member data to CRM, messaging-app, enrichment or advertising platforms; live mode needs a named approver	BE	P1	founder OS brief, D-29
FR-088	Partner API (licensing): scoped tk_ keys shown once and stored hashed, ACU metering with idempotent usage records, signed outbound webhooks with retries and replay; no TRYST member data is reachable from any partner scope	BE	P4	founder OS brief, D-27
FR-089	Public sign-up in the launch cities (London, Manchester, Birmingham, Brighton) with no invitation required; matching opens one city at a time in that order, each at 2,000 verified members once the previous city is live, and never closes again; members see how many more are needed	Both	P1	Founder v1.3, D-35
FR-090	Demand-balance waitlist per city, gender-neutral: a member waits only if a group they seek would exceed the load cap (2.2 seekers per group member in P1, 2.0 from P3); the rule never reads the applicant's own characteristics; waiting is first come, first served per sought groups, shows position and the busy group, and releases automatically as supply grows; departures never reopen seeding; nothing here affects price or ranking	BE	P1	A \$16, D-35, D-36
FR-091	Paid-acquisition brake per city: blended cost per verified sign-up checked weekly against break-even (contribution per payer × payer rate × V2 rate); paid spend stops automatically after 2 weeks above it or 1 week above 1.5×; only a named approver can restart it	BE	P1	E-03, D-35

10.2 Non-functional requirements

ID	Requirement	Owner
NFR-01	Erasure < 60 s p99	BE
NFR-02	Retrieval < 40 ms p95; ranking < 120 ms p95 for 200 candidates; /v1/slate < 400 ms p95	BE
NFR-03	On-device Guardian inline < 80 ms p95 per thread turn	FE
NFR-04	Quick exit < 500 ms; Burn confirmable < 2 s; reveal revocation effective < 5 s p95	Both
NFR-05	Availability: core API 99.9%; message relay 99.95%	BE
NFR-06	UK/EU residency; primary AWS eu-west-2	BE
NFR-07	WCAG 2.2 AA (web), platform accessibility APIs (native)	FE
NFR-08	OWASP ASVS L3 (API), MASVS L2 + R (mobile); quarterly pen test; annual red team on A4/A7	Both
NFR-09	No third-party analytics, advertising or attribution SDKs in any client	FE
NFR-10	Zero standing production access; break-glass needs dual authorisation and writes an immutable log	BE
NFR-11	Logs kept 90 d, IP truncated at ingest	BE
NFR-12	Scale: P3 25k verified; P4 150k; design headroom 500k verified / 50k concurrent	BE
NFR-13	RPO ≤ 5 min, RTO ≤ 1 h; backups stay under per-member keys	BE
NFR-14	Localisation: EN-GB/EN-IE at P1; NL, DE, SV, DA, ES, PT at P4 (Guardian per-language evaluation)	Both
NFR-15	Platforms: iOS 17+, Android 10+ (StrongBox where present), evergreen browsers (last 2 versions)	FE
NFR-16	Login completion p95: mobile < 20 s including liveness; desktop < 30 s including cross-device approval	Both
NFR-17	B2 provider: ISO/IEC 30107-3 PAD Level 2 certified; face-match FAR ≤ 1:10,000 at FRR ≤ 3%; fairness across demographic groups reported by the provider and audited (D-17)	BE

ID	Requirement	Owner
NFR-18	Key management: automatic yearly rotation of KMS root keys and HSM peppers (dual-run); every KMS/HSM key use streamed to the SIEM with anomaly alerts	BE
NFR-19	Assurance: public bug bounty from the end of P1 (public beta); annual independent privacy audit; quarterly pen test and annual red team (NFR-08)	Both
NFR-20	Marketing site Core Web Vitals at p75: LCP < 2.5 s, INP < 200 ms, CLS < 0.1; WCAG 2.2 AA; no third-party tracking scripts (cookieless, self-hosted analytics only)	FE

03 — TRYST Backend Specification v1.1

Part of the TRYST v1.1 baseline. Implements the contracts in 02. Product intent is in 01.

1. Scope

This document covers server-side services, agents, ML, data stores, cryptography (server side), trust and safety tooling, commerce, compliance machinery and operations. Everything a client can see or call is defined in 02. This document does not redefine those contracts.

2. Service topology and stack

2.1 Stack

Source: A §12

Layer	Choice	Rationale
Edge API	Go, gRPC-gateway (REST/JSON out, gRPC in)	Concurrency, small attack surface, fast cold start
Agents	Python, FastAPI, typed state machines (LangGraph-style)	Deterministic orchestration; every agent turn can be replayed for evals
LLM, profile-touching	Self-hosted open-weight models, UK region, vLLM	Article 9 data must not leave the region or reach a third-party API. Non-negotiable
LLM, generic copy	Hosted frontier model, PII redacted at the boundary	Only where no personal data is involved
Primary stores	PostgreSQL 16 + pgvector — identity, profile, safety and ban clusters isolated	Identity split enforced in the infrastructure
Vector search	Qdrant	Filtered ANN after hard-gate pre-filtering
Events	Redpanda + Schema Registry (Protobuf)	Kafka API, lower ops burden
Analytics	ClickHouse	Behavioural aggregation; raw events purged at 180 d
Feature store	Feast	Train/serve parity (M3/M4)
Objects	S3-compatible; client-side encrypted	The server never holds plaintext media
Keys	AWS KMS + CloudHSM	Destroying the root key is the erasure primitive; peppers held in the HSM
E2EE	MLS (RFC 9420) via OpenMLS (delivery service + key packages)	Standardised; group threads for couples
Age assurance	Two providers behind one interface	Ofcom expects layered protection and due diligence; one provider is a single point of failure
Region	AWS eu-west-2 (London), UK/EU residency only	
IaC / CI / CD	Terraform, GitHub Actions, Argo CD	Auditable change control

2.2 Deployables

Source A requires each agent to be independently deployable. Core domain services start as a **modular monolith** with hard module boundaries, to fit a 2-person backend team, and are split when load or team size justifies it (D-09).

Deployable	Language	Owns	Store
edge-gateway	Go	AuthN (OAuth2 + DPoP), jurisdiction gate, tier gate, consent gate, rate limits, idempotency	Redis
identity-svc	Go	Identity, devices, authenticators and login/step-up/recovery flows (§6.7) , assurance adapters, anchor candidates, ban-anchor check, join-key HMAC, root keys	IDENTITY-DB, BAN-DB, KMS/HSM
core-svc (modular monolith)	Go	Modules: profile, couple, intent, discretion, discovery-api, reveal, threads (MLS delivery service), meet, consent, deletion, commerce	PROFILE-DB

Deployable	Language	Owns	Store
agent-cartographer, agent-mirror, agent-broker, agent-envoy, agent-curtain, agent-guardian, agent-aftercare	Python	One agent each (§3)	Own schemas / Feast / Qdrant
llm-gateway	Python	vLLM pool, prompt/policy registry, PII redaction, injection filters	—
media-svc	Go	Pre-signed upload/download, ciphertext lifecycle, public-tier moderation, perturbation	Object store
safety-svc	Go	Report intake, SLA routing, case management, enforcement, appeals, transparency data	SAFETY-DB
trust-console	Web (internal)	Moderator UI; redacted by default; break-glass flows	via safety-svc
jobs	Go	TTL purges, deletion orchestration, renewal notices, aftercare scheduling, Privacy Auditor probes	—

Bus. Agents communicate over typed Protobuf messages on Redpanda topics (agent.*), never with free text [A §5]. Synchronous calls between agents use gRPC with mTLS and SPIFFE identities.

2.3 Request path (slate)

edge-gateway (DPoP, tier $\geq$ V2, jurisdiction, consent) → core-svc/discovery-api → agent-broker:

- Stage 0 gates
- Qdrant retrieval
- M3/M4 ranking
- constraints

→ agent-curtain exposure check → re-authorisation against the current Discretion Policy and block graph → response. Cached slates are **re-authorised on read** [B §8.2].

3. Agent services

Every agent has a policy file (policies/<agent>.yaml), a tool allow-list, a memory scope, a 400-case golden eval set, and an injection red-team suite (§9). [A §5, §13]

Agent	Inputs	Writes	Tools	Hard constraints
Cartographer	Intake turns	Proposed IntentVector (unconfirmed)	llm-gateway (self-hosted), tag vocabulary	Structured output only; the member confirms each field (FR-003); re-runs quarterly or when Mirror reports drift
Mirror	Behaviour events, aftercare	RevealedPreferenceVector (sole writer), α , divergence	Feast, training pipeline	Message content excluded; personalisation off → α frozen at 0.15 and no updates; reset supported (FR-024)
Broker	Effective vectors, gates, fairness state	Slates, slate_item	Qdrant, M2/M3/M4/M9	Stage 0 gates are not scored and cannot be bypassed; paid labels only within the eligible set; no gender-based uplift
Envoy	Both IntentVectors (incl. private fields), Discretion Policies, tiers	Handshake transcript, two Briefs; ENVOY-tier proposals	Rules engine + small model for friction ranking; llm-gateway only to render Brief copy from codes	Six dimensions; $\leq$ 6 turns; no free text; private fields never emitted; never messages a human; cannot open a thread (FR-015, FR-055)
Curtain	Session context, geofence hits, device risk, discretion events	Policy tightening, suppressions, nudges	M8, zone matcher	May only tighten policy; broadening needs a member action (02 §6.4)
Guardian	On-device scores, profile/media uploads, reports, graph	Signals, freezes, holds, case escalation	M6a-e, M7, M7b, hash lists	Veto over all agents; no permanent action without a human; never reads message content server-side (§7)

Agent	Inputs	Writes	Tools	Hard constraints
Aftercare	Meet events	Aftercare records → reward	Scheduler	Never visible to the counterparty; never aggregated into a public rating

3.1 Envoy negotiation algorithm

1. **Trigger:** on mutual intent (P2+), core-svc/intent publishes agent.handshake.requested.
2. **Load both sides** inside Envoy's memory scope. Private fields stay inside it.
3. **Score each dimension deterministically:**
4. *intent*: shape compatibility matrix.
5. *structure*: segment pair, veto mode, partner participation.
6. *boundaries*: set intersection of hard limits and desire tags; any no × yes collision is **blocking**.
7. *logistics*: overlap of availability windows (tz-aware), intersection of travel radii, notice, hosting. Windows come **only** from what the member entered, never from inferred routines (FR-067).
8. *discretion*: zone intersection (must be empty), reveal pace, device risk.
9. *verification*: both sides' min_counterparty_tier satisfied.
10. **Up to 6 turns** in schema form to settle open queries (02 §6.1). Unknown fields → friction.
11. **Emit asymmetric Briefs** (02 §6.2) built from codes; llm-gateway renders only localised copy for codes, with no counterparty free text.
12. **Encrypt** the transcript and both Briefs (30 d TTL). Publish agent.handshake.completed.
13. **Guardian** may veto at any step (blocking: ["safety_hold"], no detail disclosed).

ENVOY tier: Envoy generates envoy_proposal items (candidate + rationale codes). An intent exists **only** after POST /v1/envoy/proposals/{id}/approve (D-04).

3.2 Herald — dynamic SEO and growth agent (v1.2)

Herald is a growth agent that sits **outside the member platform** (FR-074, FR-075, FR-077; D-21).

Aspect	Specification
Deployment	Separate cloud account and VPC; no network route or credentials to IDENTITY-, PROFILE-, SAFETY- or BAN-DB, the event bus or Feast. Python typed state machines, like the other agents, with their own policy file and eval set
Inputs	Public web and SERP data (licensed SEO data provider), Google Search Console / Bing Webmaster, cookieless aggregate site analytics, the editorial style guide and claims-discipline rules (01 §15.1)
Tools	Keyword and topic clustering; content-brief generation; draft writing through the hosted frontier LLM (no personal data involved, so allowed per §2.1); on-page optimiser (titles, meta, headings, schema.org); internal-link graph builder; rank tracker; broken-link and redirect checker; outreach-draft generator; link-profile auditor
Dynamic loop	Weekly: pull rankings and Search Console data → find decaying or near-miss pages (positions 4–15) → propose refreshes, internal-link changes and new briefs → editor approves → publish through the CMS → re-measure. Monthly: link-profile audit and disavow review
Hyperlinks	Maintains the internal-link graph: hub and spoke per topic cluster, no orphans, ≤ 3 clicks from home, descriptive anchors, automatic repair of broken internal links
Backlinks	Finds link-earning opportunities (journalists covering privacy or security, relationship publishers, podcasts, unlinked brand mentions) and drafts pitches. A human sends every message. It never buys links or automates outreach volume
Guardrails	Claims discipline and "never cheat better" checks on every draft; a factual-claim checklist with sources; no fabricated statistics, testimonials or reviews; AI-assisted content is human-edited and adds original value (search-engine helpful-content and spam policies); it never targets or mentions real individuals
Outputs	Draft pages and refresh proposals in the CMS (state awaiting_editor); rank and visibility dashboards; an outreach queue. Nothing is published without approval

4. Behaviour learning (Mirror)

4.1 Vectors

Source: A §6.2

```

IntentVector          stated    dim 128    source Cartographer    updated quarterly / on confirm
RevealedPreferenceVector  revealed dim 256    source Mirror          nightly batch + hourly delta
DivergenceScore = 1 - cosine(project(IntentVector), RevealedPreferenceVector)
EffectiveVector =  $\alpha \cdot \text{Revealed} + (1 - \alpha) \cdot \text{project(Intent)}$ 
 $\alpha = \sigma(w \cdot [\text{interaction\_count}, \text{days\_active}, \text{divergence\_stability}, \text{aftercare\_count}])$ ,  $\alpha \in [0.15, 0.85]$ 

```

- Divergence is used as a signal, not corrected away.
- The 0.15 floor on stated preference is a **product rule** (P5) and is enforced as a constraint, not left to the optimiser.

4.2 Privacy constraints

- **Inputs:** only events in 02 §8. **No message content.** Conversation features are computed on device from metadata.
- **Pseudonyms:** pseudo-subject IDs rotate every training window. Deleted members are dropped from serving immediately and from the next training run.
- **Raw events:** kept 180 d, then aggregated features only. Cohort reporting needs $k \geq 50$.
- **Decay:** behavioural signals are time-weighted (default half-life 90 days) so old behaviour fades (FR-068).
- **Emergency lock** (FR-064) is handled by identity-svc: it verifies the emergency-code HMAC, revokes all sessions and authenticators' sessions, sets the profile to vanished and destroys all active reveal-grant wrapped keys.
- **Excluded from training:** protected or sensitive characteristics are never targets or inferred labels. Orientation, kinks and structure are used only as filters or features the member chose.

4.3 Model inventory

Source: A §6.3, §19.3, §21

ID	Model	Type	Serving	Target
M1	Intent classifier	Multi-label transformer, 9 classes	Batch + on edit	Macro F1 > 0.82
M2	Two-tower retrieval	Dual encoder, 256-d, ANN	< 40 ms p95	Recall@200 > 0.91
M3	Cross-encoder ranker	GBDT + neural blend	< 120 ms p95 / 200 candidates	Mutual-reply AUC > 0.78
M4	Meet propensity	Survival (time-to-meet)	Nightly	C-index > 0.70
M5	No-show / flake risk	GBDT on on-device conversation features	Pre-meet	Precision@20% > 0.55
M6a-e	Guardian behavioural family	§7	On device (+ server graph)	§7.6
M7	Authenticity (bot/catfish)	Ensemble: image, text, device, graph	Sign-up + continuous	Catch rate > 0.93
M7b	Ban-evasion graph	Style similarity, device/network adjacency	Server	Human review before any anchor extension
M8	Discretion risk	Contextual bandit	Per session	Reduction in Burn invocations
M9	Exposure fairness monitor	Gini + constraint solver	Hourly	Inbound Gini < 0.55

4.4 Reward

Source: A §6.4

```

R = 0.10·mutual_reply + 0.15·sustained_thread(≥6 turns, both sides) + 0.20·logistics_agreed
+ 0.35·meet_confirmed + 0.25·aftercare_positive
- 0.40·report - 0.20·block - 0.15·no_show - 0.10·ghost_after_intent
- 0.05·session_time_beyond_8min

```

Time in app is a **cost**. A model that improves AUC but lowers ΔR or the meet rate is rejected (§9).

4.5 Cold start and exploration

- **Turns 1-20:** content-based from the IntentVector; $\alpha = 0.15$.
- **Exploration:** a permanent 15% of every slate, chosen by Thompson sampling.
- **New-profile boost:** 72 h decaying multiplier, capped by the inbound limiter.

- **Couples:** a blended vector from both partners plus a learned `joint_veto_model`. The true preference is the **intersection** of the two partners', not the union.

5. Matching and ranking (Broker)

5.1 Pipeline

Source: A §7

```

STAGE 0  HARD GATES (non-scored, non-bypassable, bidirectional)
| verification tier ≥ V2 and ≥ each side's min_counterparty_tier
| hard-limit intersection: zero collisions
| structure compatibility matrix (02 §2.3) and couple VC validity
| ExclusionRing (PSI + heuristics)
| geofence exclusion zones (both sides)
| block graph / report holds / safety_state
| jurisdiction gate
| visibility (incognito: only if caller already received their intent)
STAGE 1  RETRIEVAL (M2)           ~50k → 200      < 40 ms
STAGE 2  RANKING (M3 + M4)        200 → 40        < 120 ms
STAGE 3  CONSTRAINTS              exposure fairness, inbound caps, diversity, paid labels
STAGE 4  SLATE                    ≤ 18 cards incl. 15% exploration, 2–4 explanation codes each

```

5.2 Launch baseline (P1, before M2/M3 are trained)

Source B's transparent weighted score is the **P1 ranker** and later the fallback:

```

score = 0.24·reciprocal_attraction + 0.20·intent_fit + 0.15·discretion_fit
       + 0.12·communication_fit + 0.10·availability_logistics + 0.08·reliability
       + 0.06·novelty + 0.05·model_confidence - risk_penalties

```

5.3 Multi-person scoring (THIRD/QUAD)

Edge scores are computed for every required pairing. The entity score is a **max-min** (soft-min with $\tau \approx 0.1$ for differentiability) over required edges, so one person's strong interest cannot mask another participant's poor fit [B §6.2]. Each couple's veto rules (02 §9) and the solo's veto apply on top.

5.4 Exposure fairness

1. **Inbound cap:** 12 intents per 24 h by default (member-adjustable 4–40). Overflow is queued and re-scored, never delivered in a burst, and the cap is never revealed to the sender.
2. **Gini constraint:** M9 runs hourly per city and segment. If Gini > 0.55, the ranker's fairness penalty rises until it recovers. Fairness is a constraint, not the objective.
3. **Asymmetric outreach cost:** couple→solo costs 3× the Keys (15 vs 5), fully refunded on reply.

S4 inbound above the cap is a release-blocking regression (§9).

5.5 Hard-limit inviolability

Stage 0 is implemented as a pure function with **property-based tests**: for random entities and constraint sets, no output may violate a gate. These tests run on every PR touching broker/ or contracts/. A failing case is a P0 incident, even in staging [A §6.2].

6. Security, cryptography and discretion (server side)

6.1 Threat model

Source: A §9.1

ID	Adversary	Likelihood	Primary server-side controls
A1	Partner with physical access to the device	Very high	Device revocation, Burn endpoint, session TTLs, masked push payloads (no content in APNs/FCM)
A2	Partner reading a shared bank statement	Very high	Neutral truthful descriptor, voucher path, 7-day pre-renewal notice

ID	Adversary	Likelihood	Primary server-side controls
A3	Recognition through the shared social graph	High	ExclusionRing PSI + heuristics, progressive reveal, perturbation, geofence
A4	Mass data breach	Medium / existential	E2EE, client-side media encryption, identity/profile split, cryptographic erasure, short TTLs
A5	Targeted extortion by a counterparty	Medium	Watermarked media, screenshot signals, M6e freeze, evidence preservation
A6	Civil discovery in divorce proceedings	Medium	Minimal retention, erasure, published legal-process policy, transparency report
A7	Insider	Medium	Zero standing access, dual-authorised break-glass, WORM audit, no content access by design
A8	Face-search de-anonymisation	Medium, rising	Adversarial perturbation on public-tier photos; private photos never served unencrypted

6.2 Cryptographic design

Source: A §9.2

- **Messaging:** MLS (RFC 9420). The server is the Delivery Service and Authentication Service: it stores key packages and ciphertext only. It **cannot** read content.
- **Media:** the client generates a per-media AES-256-GCM key; ciphertext goes to the object store; the key is wrapped per recipient and delivered in band. **Revoke = destroy the wrapped key.**
- **Identity split:** IDENTITY-DB and PROFILE-DB use separate credentials and KMS domains. The join key is $\text{HMAC}(\text{identity_id}, \text{pepper_join})$, computed in identity-svc from an HSM-held pepper. Neither dump alone identifies anyone.
- **Peppers** (contact, join, device, ban) are HSM-resident and non-exportable, with yearly rotation and dual-run.
- **Field encryption:** sensitive PROFILE-DB columns (bio_ct, policy_ct, zones, transcripts, Briefs) use envelope encryption under the **member's data key**, which is derived from the member's root key in KMS.

6.3 Erasure (< 60 s p99)

DELETE /v1/account runs deletion_job:

1. Mark the profile deleted and remove it from Qdrant and slates **synchronously**.
2. Revoke all devices and tokens.
3. **Destroy the root key** (KMS schedule-deletion with the minimum window, plus immediate disable; the HSM-wrapped data key is zeroised).
4. Delete rows and object keys (best effort; they are already undecryptable).
5. Drop features from Feast online and mark the pseudo-subject for training exclusion.
6. If no ban applies, delete anchor_candidate. Revoke all authenticators and request deletion of the provider-held B2 login reference; the provider's deletion receipt is attached to the proof (it may arrive after the 60 s on-platform erasure).
7. Write a signed completion_proof (key-destruction receipt + purge summary).

Exceptions (02 §7): ban anchors, legal hold, Panic preservation. Each is written to the job record.

6.4 ExclusionRing PSI (server half)

Source: A §9.3

- OPRF over ristretto255 (RFC 9497). The server key k lives in the HSM.
- Round 1: the server evaluates blinded points.
- Round 2: the client uploads unblinded PRF outputs. The server intersects them with the member-identifier PRF set (each member's own contact hash, PRF-evaluated at enrolment) and writes **bidirectional** exclusion_ring rows.
- The server **returns nothing** about the intersection, and stores no uploaded PRF outputs beyond the intersection computation.
- Rate-limit enrolment (contacts $\leq 5,000$ per call; ≤ 4 calls per day) to block membership-oracle attacks.
- **Heuristic suppression** (suppress only, never notify): shared payment fingerprint, sustained shared home geohash, device-graph adjacency, shared Wi-Fi BSSID history (only where the client opts in).

6.5 Authorisation (ABAC)

Every read evaluates (requester, resource, reveal_grant, block graph, ExclusionRing, discretion version, safety_state, jurisdiction, tier) in a central policy engine (OPA/Cedar) shared by all core-svc modules [B §9]. Denials

are indistinguishable from not-found (02 §5.10). Policy tests are generated per read path; the Privacy Auditor runs them continuously against staging and production canaries.

6.6 Engineering non-negotiables

Source: A §12

- No third-party analytics, advertising or attribution SDK (server or client); NFR-09.
- Zero standing production data access; break-glass needs dual authorisation and writes `staff_access_log` (WORM); NFR-10.
- Staff cannot read message content; this is an architectural property.
- Quarterly external pen test; annual red team scoped to A4/A7. OWASP ASVS L3.
- 72-hour breach playbook, rehearsed twice a year, with specific runbooks for outing, sextortion, intimate-media leak, insider abuse, account takeover and mass breach [B §9].

6.7 Authentication — account holder only (v1.2)

Implements 02 §3.3 (FR-056–FR-062, D-16).

- **WebAuthn relying party** in `identity-svc.userVerification=required`. Attestation is verified for native authenticators; the AAGUID allow-list is maintained.
 - B1 credentials must be platform authenticators flagged `bio_enrolment_bound`. The native apps send an App Attest / Play Integrity statement that the key was created with biometric-only access control and invalidation on enrolment change.
 - On the mobile PWA the browser cannot prove biometric-only verification. The server therefore requires B2 on **every** PWA login and on every step-up, while native B1 attestation lets step-ups inside an existing session use B1 alone, except for device add and recovery.
- **Login state machine** (Redis, 5 min TTL): `begin` → `factor1_ok` → (`liveness_pending` | `approval_pending`) → `session`. **No session token exists until both factors pass**. Partial states are bound to the DPoP key that started them.
- **Surface decision** is made on the server: native attestation → `mobile_app`; a registered platform authenticator on a mobile UA with no attestation → `mobile_pwa` (B1 + B2); anything else → `desktop` (F1 + F2). An unrecognised authenticator falls back to the strictest path available.
- **B2 provider adapter**: shares its interface with the two age-assurance providers (02 §3.1). Uses a randomised challenge and a 1:1 match against `login_reference`. Only `{match: bool, score, pad_pass: bool}` comes back to TRYST; no image or template. **Fail closed**: a provider outage blocks new logins, not existing sessions. An in-session unlock never needs B2.
- **Cross-device approval (F2)**: the QR code carries a single-use 60 s token bound to the desktop's DPoP key and coarse location. The approving app shows the desktop's browser, approximate city and time, and needs a fresh B1 assertion. No push notification is sent.
- **Risk controls**: rate limits per contact hint, device and IP; lockout escalation (5 failures → 15 min; 10 → recovery only). Every success and failure writes `auth_event`. A new-device login sends a masked in-app notice to the other registered devices. Impossible-travel and new-AAGUID signals force recovery-grade checks.
- **Recovery**: B2 + fresh V2 whose HMAC `anchor_doc` must equal the stored candidate (02 §4.1), then a 24 h hold. The hold is cancellable from any registered device and is never bypassable by staff (no break-glass path exists for authentication).
- **Lawful basis**: B2 is biometric processing for unique identification (Art. 9). It uses explicit consent `biometric_login`, recorded in the ledger. The accessible alternative (D-17) keeps the consent freely given (ToS Q14).
- **Cost**: about 1–2 B2 checks per mobile MAU per month (30 d sessions plus new devices and step-ups). At an assumed £0.10–0.30 per check this is £0.10–0.60 per MAU per month; it is added to the infra line in 01 §12. Replace with the vendor quote in P0.

6.8 Source B §9–§10 control coverage (v1.2)

Every control in the founder's Security and privacy doctrine (§9) and Trust, safety and misuse controls (§10) (07) maps to an implementation.

Source B control	Implementation
Article 6 + Article 9 + DPIA + consent records + transfers	01 §14.2; consent ledger (FR-040); UK/EU residency (NFR-06)
Minimisation	FR-002; age band, not DOB; geohash-5; no contact upload (PSI only)

Source B control	Implementation
Separation (identity, profile, media, analytics, billing)	Isolated clusters and peppered join (§6.2); media-svc and ciphertext-only store; ClickHouse pseudonymous; billing_ref only in IDENTITY-DB
Encryption, key rotation, access telemetry	§6.2; NFR-18
Authorisation (ABAC)	§6.5; FR-005; FR-070 role separation
Private media	Signed short-lived access + recipient-bound wrapped keys (FR-018); no public CDN; watermarking; revocation < 5 s (NFR-04); FR-069 format allow-list and re-encode on device, server malware scan for public tier
Payments	Tokenised PSP, neutral truthful descriptor (§10); FR-070 finance staff cannot browse profiles
Admin	Zero standing access, dual-approved break-glass, WORM log (NFR-10); FR-071 dual approval for intimate evidence
Deletion	§6.3; backups 35 d disclosed (02 §7); legal holds narrowly scoped (report_case.legal_hold)
Incident readiness	§6.6 runbooks: outing, sextortion, intimate-media leak, insider abuse, account takeover, mass breach
Assurance	ASVS L3 / MASVS L2+R (NFR-08); threat model (P0); NFR-19 bug bounty + annual privacy audit
Fake profiles / bots	§8 L1, M7, M7b, fake-profile stack, human review
Blackmail / sextortion	M6e freeze, evidence preservation, victim resources, Action Fraud path (§7.3)
Stalking / doxxing	Coarse location + rotating cells (FR-065); visibility (FR-009); no free-text member search exists , and slate size is capped; M6d; targeted network blocks as an enforcement action only (VPNs are scored, never blanket-blocked); M7b repeat-account detection
Non-consensual media	Hash matching where lawful (D-10, StopNCII); FR-069 upload warning; reveal/revoke controls; 1 h takedown; sanctions
Coercion in threesomes	Independent accounts and VC (FR-006); per-person approval and veto; private withdrawal from group threads (02 §9); either partner dissolves; no "couple owner"
Underage access	V2 gate (FR-001), re-check triggers (02 §3.2), 15 min escalation and mandatory reporting
Commercial sex / trafficking	No member-to-member payments; solicitation detection; specialist review; jurisdictional reporting (FR-037)
Bias / predatory targeting	No advertising at all; exposure audits (M9, fairness audits); FR-072 protected-trait harassment rule
Retaliatory reports	Reporter-reputation scoring; evidence review; appeal; no permanent ban from one uncorroborated report (§8)

7. Guardian — safety models (M6a-M6e)

7.1 Central problem

Source: A §21.1

Consensual power exchange looks identical to coercion at the token level. Guardian therefore classifies the **trajectory** of a conversation, not its content:

```

consent_state(t) = f(explicit_affirmations, explicit_refusals, limit_declarations,
                    enthusiasm_markers, reciprocity, initiative_balance)
RISK fires on:  refusal(t) ∧ escalation(t+k) · limit_stated(t) ∧ limit_approached(t+k)
               disengagement(t) ∧ persistence(t+k) · sustained asymmetric initiative
RISK never fires on: intensity, vocabulary, kink taxonomy, power-exchange framing, explicitness

```

Any model that separates consensual from coercive only through topic vocabulary is rejected at eval.

7.2 Models and thresholds

Source: A §19.3 L4, §21.2

ID	Target	Threshold action
M6a	Coercive control (isolation, monitoring, guilt leverage, secrecy framing imposed on the counterparty, identity erosion)	> 0.75 → queue; > 0.92 → thread freeze pending review
M6b	Boundary violation (escalation after a refusal, repeated declined request, approaching a hard limit, recasting a "no" as a "maybe")	> 0.70 → send-block + warning; 2nd → strike; 3rd → suspend

ID	Target	Threshold action
M6c	Off-platform pressure	> 0.65 → caution banner to the counterparty; feeds M5 and M6a
M6d	Identity/location pressure	Auto-redact in transit (client) + banner + strike on repeat
M6e	Image coercion and sextortion	Any hit > 0.60 → immediate freeze , human within 1 h. Tuned for recall

Serving:

- M6a–e run **on device** for E2EE threads, inline < 80 ms p95. Only a score + reason code + model version is sent to /v1/guardian/signals.
- Server-side: cross-account correlation on reason codes (the same script used against many people), M7/M7b, and profile/media moderation.
- Scores are calibrated (isotonic) so moderators can triage by probability. Thresholds are set per model, locale and segment, and reviewed monthly against the overturn rate.

7.3 Escalation matrix

Source: A §8.4

Signal	Automatic	Human SLA	External
Suspected minor	Immediate full freeze of both accounts; content preserved under legal hold	15 min	Mandatory report; law-enforcement liaison
Intimate image abuse / NCII	Media quarantine, hash to StopNCII, sender frozen	1 h	Takedown log; victim-support routing
Sextortion / coercion	Thread freeze, evidence preserved, victim safety banner	1 h	Action Fraud referral path
Threat of violence / stalking	Account suspended, mutual geo-suppression, device ban	1 h	Police referral on request
Romance fraud pattern	Payment-link stripping, shadow-limit, cohort sweep	4 h	Bank/PSP intelligence sharing
Commercial sexual services	Immediate removal	4 h	—
Harassment (M6 > 0.9)	Send-block, warning, strike	24 h	—

Evidence. Server-side evidence exists only when a member reports and **chooses** to disclose decrypted content, or when Panic preserves it. Reporter identity is stripped before evidence reaches labelling.

7.4 Training data

Source: A §21.3–21.4

Permitted sources only:

- D1: expert-authored synthetic corpus — 12k threads, **matched consensual/coercive pairs**, written with DV and kink-education specialists.
- D2: published taxonomies, used as the label schema.
- D3: public toxicity transfer, **not** as a source of coercion labels.
- D4: reported content with separate, revocable training consent (from P3).
- D5: active learning on the moderator queue (overturns are the most informative labels).

Never: scraped platform conversations without a lawful basis, purchased corpora of uncertain provenance, or survivors' disclosures collected for another purpose.

Annotation:

- Triple annotation on the coercion classes; Krippendorff's $\alpha > 0.70$ to accept a batch.
- Senior adjudication.
- Annotator welfare budgeted: rotation caps, clinical supervision, exposure limits on M6e material.

Bootstrapping:

Phase	Data	Expected performance
P1	Rules + D3	recall ~0.55 / precision ~0.30
P2	+ D1, D2	~0.82 / 0.61
P3	+ D4, D5	~0.91 / 0.74
P4	Per-language evaluation	

7.5 Profile and media moderation (server side)

Public-tier profile photos and bios are **not** E2EE and are moderated on the server:

- EXIF strip
- AI-generation detector and reverse-image search on **every** upload
- face consistency across the photo set
- C2PA provenance where present
- nudity classification (explicit content is web-only and V2-only)
- CSAM hash matching (IWF/NCMEC lists) before storage
- adversarial perturbation for A8

Private-tier media is client-encrypted; the sending client runs on-device checks before encrypting, including hash matching against licensed lists where licensing permits (D-10).

7.6 Evaluation gates and limits

Source: A §21.5–21.9

Ship gates:

Gate	Requirement
Matched-pair separation (hard gate)	AUC > 0.85
M6b	Recall > 0.90 at FPR < 0.03
M6e	Recall > 0.95, precision ≥ 0.35
Dialect FPR parity	Within 1.5× across groups
Counterfactual gender swap	Mean score delta < 0.05
Appeal overturn	< 8%
Moderator agreement	> 75%

Fairness slices: dialect/register (AAVE, MLE, regional), non-native English, gender, orientation/structure across S1-S6, age band. Audited quarterly.

Adversarial resistance:

- Quarterly red-team retainer.
- Rolling windows across messages and sessions.
- Cross-account correlation.
- Weekly drift checks. A sudden drop in positives means an evasion technique has been found, not that behaviour improved.

Guardian never:

- takes a permanent action without a named human;
- scores protected or inferred characteristics;
- produces a public or counterparty-visible score;
- reads message content on the server;
- flags on vocabulary alone.

8. Admission control (L1-L5)

Layer	Server responsibilities
L1 Identity binding	Liveness with a randomised challenge (provider); ID-document match through the provider; device attestation verification (App Attest / Play Integrity), emulator/root detection; payment-instrument fingerprint (token); £9.99 deposit. VPN/datacentre IPs are scored, never blocked

Layer	Server responsibilities
L2 Ban anchor	At V2, identity-svc stores anchor_candidate HMACs (E-02). On a permanent ban (human decision, reason_code), the candidates are copied to BAN-DB. On auth/verify (new identity), attestation and payment checks, and V2 completion, every anchor is checked; a match → silent refusal with a generic error that never names the anchor. Biometric anchor only if counsel confirms a lawful basis (D-08). A successful appeal destroys the anchors
L3 Attestation	Conduct declaration stored in the consent ledger; Basic DBS verified through the issuing route (no PDF uploads) and stored as {dated_fact}; Clare's Law content is static, with no field for the outcome
L4 Behavioural	Guardian M6a-e (on device), M7b ban-evasion graph → shadow-limit → human review → anchor extension
L5 Consortium	Interface built behind a feature flag: anchors + reason code only. Disabled until competition-law and data-protection opinions (ToS Q8)

Fake-profile stack [A §19.4]:

Stage	Checks
Sign-up	Liveness · ID match · GAN/deepfake detection on selfie and video · attestation · deposit
Media	C2PA · reverse-image · EXIF anomaly · face consistency · AI-generation classifier on every upload
Network	Score, never block
Graph	Sign-up bursts · shared device or payment · referral-chain anomalies
Behaviour	Typing cadence · templating · copy-paste · latency uniformity · slate-interaction entropy

Detectors are refreshed quarterly; a static detector is a known future failure.

Enforcement pipeline (FR-035): warn → shadow-limit → suspend → permanent + anchor → law-enforcement referral. Grounds g1, g3 and g4 go straight to permanent. enforcement.decided_by is required for permanent (DB constraint). Appeals are assigned to a different reviewer (enforced in the console) within 10 working days. Reporter-reputation scoring down-weights coordinated or retaliatory reports; one uncorroborated report never leads to a permanent ban [B §10].

9. MLOps and release gates

Model ship gates:

- Offline replay on held-out cohorts measuring **ΔR**, not AUC alone.
- Shadow → 5% → 25% → 100% rollout, with automatic rollback on report-rate or S4-churn regression.
- Model card + DPIA annex for each model.

Agent gates:

- 400-case golden set per agent: Envoy negotiation fidelity, Guardian adversarial recall, Cartographer extraction accuracy.
- Prompt-injection corpus.
- Run on every PR touching a prompt, policy file or agent code.

Audits:

- Quarterly fairness audit per protected characteristic and segment.
- Monthly counterfactual audit for popularity collapse (the 15% exploration floor is the structural defence).

Privacy Auditor [B §5]: continuous probes. It **blocks a release** if any of the following holds:

Probe	Blocks release if...	Mapped to
Media access	Media is reachable without an active grant	G-NG-1
Erasure	Erasure proof is missing or slow	G-NG-2
Location	Location can be triangulated (distance-band oracle tests)	G-NG-3
Tier	Any gated endpoint is reachable below V2	G-NG-4
Couple co-sign	A same-device co-sign is accepted	G-NG-5
Hard gates	Any hard-gate property test fails	G-NG-7
Client SDKs	A third-party SDK appears in a client SBOM	G-NG-8
Staff access	Any staff role can reach message plaintext	G-NG-9

10. Commerce

- **PSP (v1.2): Stripe** is the primary processor, subject to Stripe's written approval in P0 (D-22). A dedicated high-risk acquirer (MCC 7273, Visa IRP Tier 1) is integrated as the **backup MID from day one**. Card data is tokenised and never touches TRYST.
- **Payment abstraction:** commerce talks to a PaymentProvider interface (create_checkout, create_subscription, cancel, refund, webhook_verify, descriptor). The Stripe adapter is implemented first and the high-risk acquirer adapter second. Entitlements live in commerce.entitlement and never depend on a single processor, so switching is a configuration change with token migration where the processors support it.
- **Stripe usage:** Checkout Sessions / Payment Element (web only; native follows D-11); Billing (Products/Prices for TRYST+, ENVOY, DUO, GHOST; one-off Prices for Keys packs and the £9.99 deposit); Customer Portal configured for one-step cancel with no retention flow; Radar rules for card testing and burner patterns (feeding L1); Stripe Tax for UK/IE VAT; refunds through the API to the original method within 14 days (ToS 16.7).
- **Statement descriptor:** neutral, truthful prefix with a customer-service URL, set per account and per charge suffix; it never includes "TRYST" (ToS 16.2).
- **Webhooks:** POST /v1/billing/stripe/webhook (sys), with signature verification, idempotent processing of event IDs, and entitlement updates only from verified events. Stripe metadata carries **only** the opaque billing_ref: no profile ID, segment, intent or anything that reveals the nature of the service beyond the product name.
- **Data minimisation at Stripe:** the Stripe Customer is created with the billing email chosen by the member (it may differ from the contact handle) and no name unless the card network requires it. Stripe is a processor for payment data only and receives no Article 9 data.
- **Descriptor:** neutral and truthful, with a customer-service URL; enforced per MID.
- **Vouchers:** offline codes redeemed through /v1/vouchers/redeem, single-use, with rate limits and fraud scoring.
- **Subscriptions** (TRYST+, ENVOY, DUO, GHOST add-on): state machine with pre-contract information, cooling-off notice (CCR 2013), pre-renewal email (≥ 7 days), one-step cancel and end-of-contract notices. Built to the full DMCC regime now.
- **Keys ledger:** double-entry, idempotent. Intent spend is held and **refunded on reply**. Keys don't expire while the account is open. They are refunded on removal unless the ground is in ToS cl. 17.3, and are never forfeited before an appeal ends.
- **Chargebacks:** Ethoca/Verifi alerts $\rightarrow$ auto-refund rules. Alert at 0.5%; hard stop review at 0.65%.
- **Payment fingerprint** feeds L1 and the ExclusionRing heuristics (suppress only).
- **App-store purchases:** web-first. Where a store build must offer in-app purchase, entitlements are unified in commerce.entitlement (D-11).

11. Compliance engineering and operations

- **Consent ledger:** append-only, hash-chained (entry_hash = H(prev_hash || record)); verified nightly; exportable on SAR.
- **DSR service:** SAR export, rectification, erasure (§6.3) and portability, all built in sprint one [A §14.2].
- **OSA:** report SLAs are measured and published. Transparency data is produced from SAFETY-DB (FR-052). A named accountable manager owns the on-call escalation.
- **Article 22:** human review path for suspensions; decisions record decided_by.
- **Jurisdiction gate:** GeoIP at the edge + sign-up declaration + travel detection from coarse client location. The block list is versioned config, and each enabled market references its legal opinion ID.
- **Legal process:** published policy. Warrants or court orders are required. The member is notified where lawful. Only what exists is disclosed (ciphertext and minimal metadata).
- **Observability:** OpenTelemetry. Logs carry no content and truncate IPs at ingest (NFR-11). SIEM, plus anomaly alerts on staff access.
- **SLOs:** NFR-01, -02, -05, -13.

12. Backend backlog by phase

Phase	Epics (backend)	Definition of done
P0 (M0-M2)	Threat model; architecture review; assurance-provider adapters ($\times 2$) spike; PSP + backup MID integration spike; IaC landing zone (eu-west-2), KMS/HSM, isolated clusters; contracts repo (OpenAPI + Protobuf) generated from 02	Threat model signed off; G-P0 legal items done by others

Phase	Epics (backend)	Definition of done
P1 (M2-M6)	edge-gateway (DPoP, gates); identity-svc (sign-up OTP, devices, passkeys + two-factor login, step-up, recovery , assurance, anchors, join HMAC); core-svc modules: profile, couple (VC), intent (Keys, caps), discretion (policy, zones), reveal, threads (MLS DS), meet, consent, deletion, commerce (TRYST+, Keys, deposit, vouchers); Cartographer; Broker P1 (Stage 0 + weighted baseline); Curtain v1; Guardian server (signals intake, escalation, M7, media moderation); Aftercare; safety-svc + trust console; Privacy Auditor probes; DSR	FR-001-009, 011-014, 016-023, 026-044, 046-047, 049-051, 054, 056-062, 064-065, 069-072 met (backend side); NFR-01/02/05/06/10/11 met; no-go list clear; control cohort running (no Envoy)
P2 (M6-M10)	Envoy (agent + Briefs); Mirror v1 (M1-M4, α , reset/insights); ExclusionRing PSI; Guardian P2 models (D1/D2); FR-024/025	G-P2: meets per intent sent $\geq +40\%$ vs control
P3 (M10-M15)	DUO, ENVOY, GHOST billing; ENVOY proposals; travel mode; DBS attestation; transparency report; 24/7 T&S tooling; scale to 25k verified	G-P3
P4 (M15-M24)	Multi-market jurisdiction config; localisation; per-language Guardian; consortium interface (if cleared); scale to 150k	G-P4

04 — TRYST Frontend Specification v1.1

Part of the TRYST v1.1 baseline. Consumes the contracts in 02. Product intent is in 01.

1. Scope

This document covers the native iOS and Android apps, the web PWA, the shared client crypto core, on-device models, UX flows, discretion and safety UI, copy rules, accessibility and store compliance. It does not redefine API shapes, enums or retention; those live in 02.

2. Two-surface strategy

Apple Guideline 1.1.4 and Google Play's age-restricted and anonymous-chat policies make two surfaces a **requirement** [A §14.3]:

	Native store builds (iOS, Android)	Web PWA
Role	Discreet companion: discovery, Briefs, threads, meets, safety, discretion	The full product and the primary revenue surface
Content	Sanitised. No nudity, no explicit imagery or copy in the app, listing or screenshots	Explicit media allowed behind V2 only , never shown during verification
Messaging	Matched only — never random or anonymous chat	Same
Billing	Web-first; in-app purchase only if a store requires it (D-11)	All billing, vouchers, deposit
Positioning	"Discreet dating"; 18+; Play Console minor-blocking on; Apple age questionnaire done	Full brand
Discretion strength	Strongest: decoy skin, duress PIN, Secure Enclave/StrongBox, FLAG_SECURE, backup exclusion	Reduced (browser limits): quick exit, neutral title/favicon option, session lock, private-browsing guidance

Explicit media sent from the web is delivered to native clients as a placeholder ("View on web"). It is never rendered in the store build (FR-051).

Store basics that cause most rejections: in-app account deletion, a web deletion link, privacy manifests, a working reviewer demo account (sanitised seeded data), and a Data Safety form that exactly matches the Privacy Notice.

2.1 Public marketing site and SEO surface (v1.2)

- **Separate host** from the app (e.g., `www.` marketing vs app. product) so search indexing, caching and analytics never touch member surfaces (FR-076).
- **Next.js SSG/ISR** with a headless CMS fed by Herald drafts after editor approval (FR-073, FR-074). XML sitemaps, canonical tags, hreflang, schema.org, clean URLs and breadcrumb navigation. City pages exist only for launched or waitlisted cities and carry real local content (waitlist status, safety resources), never doorway pages.
- **Performance and accessibility:** Core Web Vitals and WCAG targets in NFR-20. Brand tokens and the logo per §7.2.
- **Discretion on a public site:** quick exit, a neutral-favicon option, no third-party trackers or pixels, and cookieless analytics. "Join" leads to the app origin, which is `noindex`.

3. Client stacks

Source: A §12

Surface	Stack	Notes
iOS	Swift / SwiftUI, iOS 17+	Keys in Secure Enclave; NSFileProtectionComplete; files excluded from backup; alternate app icons for the decoy
Android	Kotlin / Jetpack Compose, Android 10+	Keys in StrongBox where present (TEE fallback); FLAG_SECURE; allowBackup=false and data-extraction rules; activity-alias for the decoy

Surface	Stack	Notes
Web	Next.js PWA (installable), TypeScript	No third-party scripts; strict CSP; non-extractable WebCrypto keys in IndexedDB; service worker caches no personal data
Crypto core	Rust (OpenMLS, AES-GCM, OPRF client) compiled to Swift (UniFFI), Kotlin (JNI) and WASM	One audited implementation across all three surfaces
On-device ML	Core ML (iOS), TFLite/NNAPI (Android), ONNX Runtime Web (WASM/WebGPU)	Guardian M6a-e, nudity pre-send check, conversation-feature extraction
API client	Generated from contracts/openapi	DPoP proof per request, signed with the device key

React Native is not used. Source A rejects cross-platform shells because Secure Enclave/StrongBox, FLAG_SECURE, screenshot detection, decoy skin and backup exclusion are unreliable through them (E-07).

Client non-negotiables:

- **NFR-09** — no third-party analytics, advertising, attribution or crash SDKs that send data off-platform. Crash reporting is self-hosted with scrubbed payloads.
- **No social login. No contact upload** — contacts are only ever blinded on device for PSI.

4. Screens and flows (by loop)

4.0 Sign-in — account holder only (v1.2)

Implements 02 §3.3 (FR-056–FR-062).

Native app and mobile PWA — two biometric factors:

1. **B1:** a system biometric prompt (Face ID / Touch ID / BiometricPrompt BIOMETRIC_STRONG) unlocks the device-bound passkey. Device-passcode fallback is **disabled** for login.
2. **B2:** an in-app liveness capture with a randomised challenge (turn your head, read a short phrase) from the provider SDK (native) or web flow (PWA). The camera view uses brand colours with the night background. No logo is shown if a decoy skin is active.
3. The session is issued only after both succeed. Failure copy never says which factor failed.

Desktop browser / desktop PWA — two factors:

1. **F1:** the browser passkey prompt (Windows Hello / Touch ID / security key).
2. **F2:** a QR code appears on screen. The member opens TRYST on their phone, goes to **Approve sign-in**, scans it, and confirms with B1. The phone shows the desktop's browser, approximate city and time. Nothing is pushed to the phone's lock screen. Alternative: tap a second registered security key.

Other screens:

- **Biometric change detected:** "Your phone's fingerprints or face settings changed, so we've signed you out to protect your account." Then a full B1 re-registration + B2 login.
- **Unlock** (existing session): biometric or PIN per FR-027. The duress PIN opens the decoy.
- **Step-up sheet** for sensitive actions (02 §3.3). **Burn, Panic, report, block, cancel and delete never show a step-up.**
- **Recovery:** liveness → age/ID check with the same document → a 24 h countdown that any registered device can cancel → new passkey. The copy explains that support cannot bypass this.
- **Consent:** a separate biometric_login explicit-consent screen before V1, alongside the Article 9 consent, with a link to the accessible alternative (D-17).

4.1 Loop A — Quiet entry and intake

1. **Neutral landing / first launch.** On native, the member picks a decoy skin and neutral name at first launch.
2. **Contact (V0).** Email or phone OTP. Device key generated in the secure element; attestation sent.
3. **Discretion setup** (before any personal data): app PIN/biometric, duress PIN, notification mode (default **off**), Burn gesture, home/work zones (map picker → geohash-6 cells), and an honest **"What discretion can't cover"** card (P9).
4. **Liveness (V1)** through the provider SDK/web flow.
5. **Age assurance (V2)** with the ID match, then the **conduct declaration** (separate screen, logged) and the **Article 9 explicit consent** (separate from the ToS accept).
6. **Cartographer interview** (~7 min, conversational). Each proposed structured field is shown as an editable chip; the member confirms (FR-003). Desire tags offer three levels: yes / curious / no. Hard limits come from a controlled list with plain-language descriptions.

7. **ExclusionRing enrolment** (P2+). Opt-in. Explains that contacts are blinded on the device and that nobody learns who matched. Runs the OPRF client.
8. **Profile**: pseudonym, pronouns, free-text gender, age band (from assurance), photos. EXIF is stripped on device and face-blur is on by default; reveal ladder defaults come from the Discretion Policy.
9. **Safety notice** (FR-054) → first slate, or the waitlist if the city is gated.

Couples (VC). Partner A creates the couple and shows a QR or link (24 h). Partner B must complete V2 **on their own device**. Both review and approve the exact profile version. Veto mode is chosen jointly. The UI explains that either partner can dissolve the couple at any time without giving a reason.

4.2 Loop B — Discovery

- **Slate view**: up to 18 cards; no infinite scroll, no swipe mechanic.
- **Each card shows**: pseudonym, age band, distance band, modes, verification tier, and 2–4 reason chips with a "not quite" correction control (FR-025). Paid placement is labelled visibly.
- **Actions**: Send Intent (shows the Keys cost and the "refunded if they reply" promise), Pass, Save.
- Intent sending is **never** blocked by the counterparty's inbound cap, and the UI never reveals it.
- **My Signals** (P2): categories the system has learned, with correct, reset and personalisation on/off (FR-024).

4.3 Loop C — Handshake

- **Brief screen**: a persistent "**Prepared by Envoy (automated)**" label; aligned items; friction items with suggested resolutions; blocking items; estimated meet probability shown as a band (low / medium / high), never a precise number to the member.
- **Decisions**: Accept / Decline / Amend (amend chooses alternative values from the schema; no free text).
- **Couples**: the decision UI reflects the veto mode (e.g., "Waiting for your partner").
- **ENVOY tier** (P3): a proposal inbox where every proposed intent needs **one-tap approval**. Nothing is sent automatically (D-04).

4.4 Loop D — Thread and resolution

- **Thread (MLS)**: text, view-once/timed media, voice/video calls (reveal-ladder gated). Retention selector (7/30/90 days; the shortest wins). Burn-thread action.
- **Reveal ladder control**: grant or revoke per recipient and scope. A revoked item immediately shows as unavailable on the recipient's device.
- **Meet planner**: propose/confirm. Safe Meet checklist (public venue first). Check-in timer. **Share My Plan** (FR-021). Clare's Law explainer before the first meet (P3).
- **Aftercare** (48 h): a two-tap private prompt (happened? as represented? would repeat? + optional safety flag).

4.5 Account

Pause, Vanish, Delete (shows the erasure promise **and** the declared exceptions in plain words), Export, Devices, Consents, Subscription (one-step cancel), Vouchers, Appeals.

5. Client discretion controls

Control	Native	Web	Req
Decoy skin — generic utility (calculator / notes) with a neutral icon and name; the real app needs a distinct PIN or biometric	✓	Neutral title/favicon option only	FR-027
Duress PIN — opens a populated, innocuous decoy state; silent	✓	—	FR-027
Quick exit — one action hides content and clears the view in < 500 ms (web: replace history and navigate to a neutral page)	✓	✓	FR-029
Emergency lock — if a phone is lost or taken, the member visits a neutral web page and enters their contact handle + emergency code (shown once at sign-up, to store offline) to freeze everything (FR-064)	—	✓	FR-064
Location veil — distance bands only, from rotating cells (FR-065)	✓	✓	FR-065
Burn — long-press, confirm in < 2 s: wipe local storage, revoke device keys (POST /v1/burn), sign out everywhere. Visually and gesturally distinct from Panic	✓	✓	FR-028, E-21
Notification masking — off by default; if on, the payload is content-free ("1 update") and the push contains no name, preview or image	✓	✓ (web push)	FR-026
Screenshot protection — Android FLAG_SECURE; iOS capture detection that blurs content and notifies the counterparty; obscured app-switcher snapshot	✓	Best-effort blur on visibility change	FR-030

Control	Native	Web	Req
No cloud backup — excluded from iCloud/Google backup; keys non-exportable	✓	Keys non-extractable	FR-030
Zone awareness — coarse client location checked against zones on device; "You're inside your work zone — hiding profile for 2 h"	✓	Optional	FR-011
Jurisdiction auto-lock — entering a blocked country switches to the decoy and hides the profile	✓	Edge-enforced	FR-031
Watermarking — recipient-specific visible/invisible watermark on revealed media; no-download viewer	✓	✓	A §9.1 A5

Location never leaves the device at better than geohash-5, except in the Panic flow.

6. Client cryptography and on-device intelligence

- **Device identity:** a key pair generated in the secure element, used for DPoP proofs and as the MLS credential. Rotated on re-install; revoked by Burn.
- **Messaging:** OpenMLS. One MLS group per thread (2, 3 or 4 members). Key packages are uploaded to the delivery service.
- **Media:** a per-media AES-256-GCM key; ciphertext uploaded through a pre-signed URL; the key is wrapped per recipient (MLS exporter secret) and sent in band. **Revocation** asks the server to destroy the wrapped key, and the recipient client purges any cached plaintext.
- **PSI client:** normalise contacts (E.164 / lowercased email) → hash → blind with a random r → server evaluates → unblind → upload the PRF set. Raw contacts never leave the device; nothing is shown to the member about matches.
- **Guardian on device (FR-033):** M6a-e run on each thread turn over a rolling window (< 80 ms p95). Only {model_id, version, score, reason_code} is posted. Local UI responses:
 - M6d auto-redacts addresses and phone numbers in outgoing text and asks the sender to confirm;
 - M6c shows a caution banner to the counterparty;
 - M6e locks the thread pending review.
- **Pre-send media check:** on-device nudity classifier (explicit media only to V2 web, with the counterparty's media-view consent), plus licensed hash matching where permitted (D-10).
- **Conversation features** for M5 and Mirror (length bucket, latency, question ratio) are computed on device. **No content leaves the device.**
- **Reporting:** the reporter may choose to attach decrypted messages or media. This is an explicit, separately confirmed disclosure, with a separate opt-in for safety-model training (never a condition of action).

7. Safety UX

Safety notice (FR-054; shown at intake and before a first meet) [A §20.3]:

Before you meet anyone. We verify that members are real adults. We do not — and cannot — check anyone's criminal record, their honesty, or whether they are who they say they are once you are off this app. No UK platform can. Meet in public the first time. Tell someone where you are going; Share My Plan does it privately. Trust your instincts and leave if something feels wrong. If anyone pressures you, threatens you, or pushes past a boundary you set, report them — we act on it, and we do it fast.

- **Panic** (FR-022): a prominent, distinct control in threads and the meet screen. It confirms with a hold, sends live location + preserved evidence through POST /v1/safety/panic, and offers a one-tap call to 999. It works even when notifications are masked or the app is in decoy mode (reachable through a gesture from the decoy).
- **Share My Plan:** pick a trusted contact (stored locally only). It sends a neutral, unbranded, auto-deleting message with venue and time window through the device's share sheet or an encrypted link. The wording is chosen for someone who cannot tell their partner where they are.
- **Check-in timer:** a missed check-in triggers silent escalation to the trusted contact.
- **Report / block:** one tap on every profile, Brief, thread and media item. The reporter is never revealed.

7.1 Copy rules (claims discipline, P6)

Source: A §19.8

Never	Use
"Background checked", "Safe", "Vetted", green "safe" ticks	"Identity verified", "Age assured", dated facts ("Submitted a Basic DBS dated 14 Mar 2026")
"No sex offenders"	"We permanently remove and block re-registration where we find harmful conduct"
"Guaranteed real"	"Every member completes liveness and ID verification"
"Undetectable", "No trace"	"We minimise what we store, and here's what we can't control"
"Cheat better" or any affair-shaming or affair-glamorising copy in the store build	"Discreet dating"; "Discretion is the product" (web)

7.2 Brand application (v1.2)

Logo. Use `assets/brand/tryst-logo.png` **exactly as supplied**: no redraws, recolours, crops or effects. Place it on night, oxblood-deep or oxblood only. Never show it in decoy mode, on lock screens, in notifications, app-switcher snapshots or Share My Plan (FR-063; 01 §3.1). For small sizes (favicon, ≤ 48 px) use the PNG scaled down. If it is illegible at that size, the web favicon defaults to the neutral option (FR-027) rather than a modified logo.

Colour. Import `assets/brand/tokens.css` (web) or generate native colour assets from `tokens.json`. Defaults:

Element	Token
Background	night
Cards	oxblood-deep
Body text	cream
Headings and accents	gold
Primary button	oxblood with champagne text
Borders	oxblood-light
Disabled	bronze

Type: serif display with tight tracking for headings, matching the wordmark; a clean sans for body text. **Never** pink, hearts or silhouettes [A §2]. The decoy skins use their own neutral system palettes, never brand tokens.

8. Accessibility, localisation, performance

- **Accessibility:** WCAG 2.2 AA on web (brand text pairs pass AA, 01 §3.2); VoiceOver/TalkBack labels on native; Dynamic Type; reduced motion. The decoy and quick-exit controls must be operable by assistive technologies **without announcing "TRYST"** (NFR-07). Members who cannot use face or fingerprint biometrics get the accessible sign-in alternative (D-17).
- **Localisation:** EN-GB/EN-IE at P1; NL, DE, SV, DA, ES, PT at P4 (NFR-14). Brief and Guardian copy is code-based for localisation.
- **Performance:** cold start < 2 s on mid-tier devices; slate first paint < 1 s after the API response; quick exit < 500 ms; Burn confirmable < 2 s (NFR-04).
- **Platforms:** iOS 17+, Android 10+, evergreen browsers (NFR-15).

9. Client testing and release

- Contract tests against generated mocks of 02. A UI tier-gate test checks that no screen shows other members' content below V2.
- Security: MASVS L2 + R checks, jailbreak/root behaviour, screenshot and backup tests, SBOM scan with zero third-party telemetry SDKs (G-NG-8).
- On-device model benchmarks per device class (NFR-03).
- **Store review pack:** sanitised build flag, reviewer demo account, privacy manifest, Data Safety form diffed against the Privacy Notice in CI.

10. Frontend backlog by phase

Phase	Epics (frontend)	Definition of done
P0 (M0–M2)	Design system (oxblood/bone); discretion UX prototypes (decoy, duress, quick exit, Burn vs Panic) usability-tested; Rust crypto core spike on three targets; store pre-review consultation	Prototypes validated; crypto core builds on iOS, Android and WASM

Phase	Epics (frontend)	Definition of done
P1 (M2-M6)	Web PWA (full product) + native apps (TestFlight / internal track): onboarding V0-V2, conduct and Article 9 consent, Cartographer UI, profile/media with client encryption, couple co-sign, slate, intents/Keys, threads (MLS), reveal ladder, meet/Safe Meet/Share My Plan/Panic, Aftercare, discretion suite, on-device Guardian baseline, report/block, account/DSR, web billing (TRYST+, Keys, deposit, vouchers)	FR-001-009, 011-013, 016-023, 026-033, 037-039, 041-042, 044, 046-047, 049, 051, 054, 056-065, 069, 072 met (client side); NFR-03/04/07/09/15 met
P2 (M6-M10)	Brief screens + decisions; My Signals; ExclusionRing enrolment (PSI client); reason chips	FR-010, 015, 024, 025
P3 (M10-M15)	DUO, ENVOY (proposal inbox), GHOST UI; travel mode; DBS attestation and Clare's Law explainer; store submissions approved on both platforms	G-P3 store approval; FR-045, 048, 053
P4 (M15-M24)	Localisation (6 languages); market-specific copy and assurance flows	NFR-14

05 — TRYST Terms of Service — Draft 0.9.3

Part of the TRYST v1.1 baseline.

Draft for counsel review. Not legally reviewed. Do not publish or rely on in this form. Prepared by: Product. For review by: UK consumer, Online Safety Act and data-protection counsel. Governing law assumed: England & Wales.

Changes from draft 0.9 (Source A, Part B) are marked **[v1.1]**. They align the terms with specification v1.1 (06 §3: E-02, E-09, E-10, E-11) and add counsel questions Q11-Q13. Clause numbers are unchanged; one clause (3.5) is added.

Note to counsel. This draft is written by the product team to be reviewed, not approved. Clauses carrying the highest legal risk are marked **REVIEW** with a specific question. Three areas need particular attention:

- **clause 15** — retention of exclusion identifiers after account deletion: the lawful-basis question;
- **clause 16** — the subscription regime, whose commencement date is unsettled;
- **clauses 21-23** — the liability structure, drafted deliberately narrow so that it survives s.62 rather than broad so that it gets struck.

1. These terms, and who we are

1.1 TRYST is operated by [SPV NAME] LIMITED, a company registered in England and Wales (No. [—]), registered office [—]. In these terms, "we", "us" and "our" mean that company. "You" means the person using TRYST.

1.2 These terms are a contract between you and us. They apply whenever you use the TRYST website, apps or any related service.

1.3 They are written to be readable. Where a clause matters to you, we have added a plain-English box explaining it. Those boxes are part of these terms and can be relied on.

1.4 Our Privacy Notice, Community Rules and Safety Policy form part of these terms.

2. Definitions

2.1 **Member** — a person with a TRYST account. **Couple Profile** — a profile representing two Members who have each independently confirmed the link. **Keys** — the prepaid credits described in clause 16. **Intent** — an outbound expression of interest. **Brief** — the agent-generated summary described in clause 7. **Exclusion Identifier** — the irreversible identifier described in clause 15.

3. Who can use TRYST

3.1 You must be 18 or over. We verify this using age-assurance methods that meet the "highly effective" standard under the Online Safety Act 2023. If we cannot verify that you are an adult, you cannot use TRYST, and there is no appeal against that.

3.2 You must have the legal capacity to enter this contract, and you must not be prohibited from using a service of this kind by any court order, licence condition or notification requirement.

3.3 TRYST is available only in the countries listed at [URL]. We block access elsewhere, including in countries where using a service of this kind could expose you to criminal liability. You must not attempt to circumvent that block.

3.4 You may hold one account. If you previously held an account that we permanently closed under clause 17, you may not open another.

3.5 **[v1.1]** When you verify your age, we ask you to declare that you have no unspent convictions for sexual or violent offences, and that you are not subject to a Sexual Harm Prevention Order, a restraining order or a non-molestation order. We record that declaration separately. A declaration you knew to be false is grounds for permanent removal under clause 17.3.

4. What TRYST is

4.1 TRYST is an introduction and discovery platform. We provide software that lets adults find one another, communicate, and decide for themselves whether to meet.

4.2 We are not a matchmaking agency, an introduction agency acting on your behalf, an escort or companionship service, or a party to anything you arrange with another Member. We do not act as your agent and we owe you no duty to find you anyone.

4.3 We do not select, endorse, recommend, approve or vouch for any Member. Our software ranks and suggests profiles; that is a computational output, not a recommendation from us that a person is suitable, honest or safe.

IN PLAIN ENGLISH — We introduce people. What happens after that is between you and them. Nothing we show you is us telling you that someone is a good idea.

5. What we verify — and what we do not

5.1 We do verify, before a Member can send an Intent or exchange media: that they are an adult; that they are a real, living person, by liveness check; that their photographs are of them; and, for a Couple Profile, that both partners have independently confirmed the link from their own separate devices.

5.2 We do not verify, and cannot verify:

- your relationship or marital status;
- whether a partner of yours knows about, or consents to, your use of TRYST;
- your criminal record, or whether you are subject to any order or notification requirement;
- your employment, income, health or personal circumstances;
- the truth of anything you or any other Member says.

5.3 There is no lawful route by which a UK platform can check a Member against the sex offenders register or obtain a criminal-record check on them. We do not claim to do so, and you should treat any platform that claims otherwise with caution.

5.4 Membership of TRYST is not, and must never be read as, a statement by us that a Member is safe, honest, available, unattached, or suitable for you.

IN PLAIN ENGLISH — We check you are a real adult who looks like your photos. We do not check your marriage, your record, or your honesty. No UK app can check a criminal record on you — and we will never pretend we have.

REVIEW — Clause 5 is drafted to anchor the whole liability structure and to keep us clear of the CPUT Regulations and the DMCC Act consumer regime on misleading practices. Please confirm the wording is strong enough to defeat an argument that our verification badges, ranking or marketing create an implied representation of safety.

6. Your account

6.1 Everything you tell us about yourself must be accurate, and your photographs must be current and of you. Profiles that misrepresent who you are will be removed.

6.2 Keep your access credentials, PIN and device secure. Tell us immediately if you think someone else has access to your account.

6.5 **[v1.2]** Only you may sign in to your account. Signing in needs two checks. On a phone (app or website), these are your phone's fingerprint or face unlock **and** a live face check matched to you. On a computer, they are a passkey **and** approval from your registered phone. We never let anyone sign in by text message or email code, and our staff cannot sign in for you or bypass these checks. If you cannot use face or fingerprint checks, we offer an alternative (see our Accessibility page at [URL]).

IN PLAIN ENGLISH — Even someone holding your unlocked phone cannot get into your account. If anyone adds their fingerprint or face to your phone, we sign you out.

6.3 You are responsible for everything done through your account, unless someone gained access through our failure to take reasonable care.

6.4 You may not sell, transfer, rent or share your account.

7. Couple Profiles, and our agent features

7.1 A Couple Profile requires both partners to hold their own verified accounts and each to confirm the link from their own device. We will not create a Couple Profile on one person's say-so. Either partner may dissolve the link at any time, and we will act on that immediately without telling the other partner why.

7.2 Some features use automated agents. Before you and another Member speak, our agents may exchange structured information about intentions, boundaries, availability and verification, and produce a Brief for each of you.

7.3 Our agents never pose as you, and never speak to another Member as if they were you. Every Brief is labelled as automatically generated. A Brief is information to help you decide; it is not advice, and it is not a judgement about whether someone is safe. **[v1.1]** Where you use a paid agent feature that suggests people to contact, nothing is sent until you approve it.

7.4 A conversation only opens when both of you have seen your Brief and accepted.

IN PLAIN ENGLISH — Our software does the awkward pre-screening so you don't waste weeks. It never pretends to be you, and it never tells you someone is safe.

8. Community Rules — what you must not do

You must not, on TRYST or in any contact that begins on TRYST:

- be under 18, or attempt to contact anyone you believe or suspect is under 18;
- share, request, or threaten to share intimate images of anyone without their consent;
- pressure anyone for images, contact, money, or sexual activity, or continue after someone has said no;
- threaten, harass, stalk, intimidate, or attempt to coerce or control another person;
- misrepresent your identity, age, or appearance, or use another person's images;
- offer or seek sexual services in exchange for money or anything of value;
- attempt to identify, locate, contact, or expose a Member outside TRYST without their agreement, or disclose that someone is a Member to anyone;
- solicit money, investment, or financial information;
- record, screenshot or copy another Member's private content or messages;
- use TRYST for any commercial, promotional or research purpose;
- scrape, reverse-engineer, or automate access to TRYST;
- fetishise or harass anyone because of their race, ethnicity, religion, disability, age, sex, gender identity or reassignment, sexual orientation, or any other protected characteristic **[v1.2]**;
- evade a suspension or removal.

8.1 Exposing a Member's use of TRYST to a third party is one of the most serious things you can do on this platform, and we treat it as grounds for immediate permanent removal.

9. Your conduct and your decisions

9.1 You decide who to speak to, what to tell them, whether to meet, where, when, and what happens when you do. Those are your decisions.

9.2 You are responsible for your own conduct. Another Member is responsible for theirs. We are not responsible for what any Member says, does, or fails to do — subject always to clause 22.

9.3 You accept that other Members may misrepresent themselves despite our checks, and you agree to take reasonable care of your own safety, including using the safety features in clause 11.

10. Consequences outside TRYST

10.1 We are not responsible for the effect of your use of TRYST on your relationships, marriage, family, employment, finances or reputation. Those consequences are yours.

10.2 No person who is not a party to this contract acquires any right under it. The Contracts (Rights of Third Parties) Act 1999 does not apply.

REVIEW — Clause 10.2 is intended to head off a claim by a Member's spouse or partner. Please advise whether it is sufficient, and whether any free-standing duty could arise in tort notwithstanding it.

11. Safety features — and their limits

11.1 We provide safety features including Safe Meet guidance, Share My Plan (which privately tells a contact of your choosing where you are going), a check-in timer, in-app reporting, blocking, and a Panic control that contacts emergency services.

11.2 These features are aids. They depend on your device, your network, and third parties we do not control. We cannot guarantee that any of them will work at the moment you need it, and you must not rely on any of them as your only safeguard. If you are in danger, call 999.

11.3 Using a safety feature does not make a meeting safe, and choosing not to use one does not make what happens to you your fault.

REVIEW — Clause 11.2 is a statement about technical limits, not an exclusion of liability for negligent design or operation of the features themselves — which we accept we cannot exclude under CRA 2015 s.65 where personal injury results. Please confirm the drafting does not stray across that line, since a term that did would be void and would jeopardise clause 21.

12. Your content

12.1 You keep ownership of everything you upload. You give us a limited, non-exclusive licence to host, store, encrypt, transmit and display your content solely to operate TRYST for you. That licence ends when your content is deleted, save as set out in clause 15.

12.2 Your private messages are end-to-end encrypted. We relay them; we cannot read them. This means we cannot recover them for you, and it also means we cannot moderate their content directly — we rely on your reports and on signals available to us outside the message content.

12.3 Media you send is encrypted on your device. You can revoke access at any time, which makes it permanently unreadable to the recipient. We cannot stop someone photographing their own screen with another device, and you should assume anything you send may be kept.

IN PLAIN ENGLISH — We genuinely cannot read your messages. That protects you, and it also means we can only act on what you report. Assume any photo you send could be kept, whatever our software does.

13. Intimate images

13.1 Sharing or threatening to share an intimate image of someone without their consent is a criminal offence. On TRYST it results in immediate permanent removal, preservation of evidence, and referral to the police.

13.2 If an intimate image of you has been shared without your consent, report it to us. We will act within the time stated in our Safety Policy, support you in submitting a hash to StopNCII where you wish, and preserve evidence for any investigation.

14. Your personal data

14.1 Information about your sex life and sexual orientation is special category data under the UK GDPR. We process it only on your explicit consent, which you give separately from accepting these terms, and which you can withdraw at any time.

14.2 Our Privacy Notice explains what we collect, why, how long we keep it, and your rights. It is not a summary — it is the operative document, and it prevails over anything in these terms about personal data.

14.3 We do not sell your data, we do not show advertising, and we do not embed third-party tracking in our apps.

15. Keeping and deleting your information

15.1 You can delete your account at any time. When you do, we destroy the encryption key that protects your data, which makes it permanently unreadable. This happens in under a minute and cannot be reversed. We cannot restore a deleted account.

15.2 Messages are deleted after 30 days by default, or after 7 or 90 days if you choose. Media **[v1.1] that you send in a conversation** is deleted after 7 days by default. **[v1.1] Your profile photos stay until you remove them or delete your account.**

15.3 There is one exception, and we want you to read it. If we permanently remove you because your conduct endangered another person, we keep a small set of Exclusion Identifiers — irreversible one-way codes derived from your identity document, device and payment method. We cannot reverse them into your personal details. We keep them so that you cannot open another account, and we keep them after your account is deleted. We do this to protect other people. We keep no evidence, no messages and no profile alongside them — only the codes, a category of reason, and the date. **[v1.1]** To make this possible, we create these codes when you verify and keep them protected with the rest of your account. If you delete your account and have not been removed under clause 17.3, they are destroyed with everything else.

15.4 If you are removed on that basis you may appeal under clause 18, and if your appeal succeeds we destroy the Exclusion Identifiers.

IN PLAIN ENGLISH — Delete your account and it is genuinely gone. The one thing that survives is a code that stops a removed predator coming straight back under a new name. If that is you and we got it wrong, appeal and we delete it.

REVIEW — This is the highest-risk clause in the document. We need: (a) the specific DPA 2018 Sch. 1 condition relied on — we have assumed Pt. 2 para. 18 (safeguarding of individuals at risk) — confirmed in writing; (b) a view on whether retention after erasure is defensible against an Art. 17 request; (c) confirmation that the disclosure is prominent enough, given that the rest of clause 15 promises complete erasure. We have deliberately refused to bury this. Please tell us if it needs to be more prominent still. **[v1.1]** See also Q13 on creating the codes for every verified member.

16. Paying for TRYST

16.1 Basic membership is free. Replying to any Member is always free. Paid tiers and Keys are described at [URL], with prices shown before you pay.

16.2 **Billing discretion.** Your card or bank statement will show a neutral descriptor that does not name TRYST. You can also pay using prepaid vouchers bought elsewhere. We do this because your privacy matters, not to obscure what you are paying for — the descriptor includes a customer service address.

16.3 **Cancelling.** You can cancel a subscription at any time, from inside the app or website, in one step. We will not ask you to call us, email us, or answer questions to cancel. Cancellation takes effect at the end of your current paid period.

16.4 **Your 14-day right to cancel.** Under the Consumer Contracts Regulations 2013 you normally have 14 days to cancel a distance contract. If you ask us to start the service immediately, you may lose that right for digital content, or owe us a proportionate amount for services already supplied. We tell you this clearly at checkout and ask you to confirm it separately.

16.5 **Reminders.** We email you before any renewal and before any trial converts to a paid subscription. The reminder says what you will be charged, when, and how to stop it.

16.6 **Keys.** Keys spent on an Intent are refunded in full if the Member replies. Unused Keys do not expire while your account is open. If we permanently remove you, unused Keys are refunded unless removal was on a ground in clause 17.3.

16.7 **Refunds.** Where we owe you a refund we pay it to the original payment method within 14 days.

REVIEW — The DMCC Act 2024 subscription regime is not yet in force, and its commencement date has moved repeatedly. We have drafted to comply with the full regime now, on the view that pre-contract information, reminder notices, cooling-off notices, end-of-contract notices and one-step cancellation are things we want anyway, and that retrofitting a payments flow is expensive. Please confirm the current commencement position

and whether clause 16 already meets it, particularly on the form and prominence of the notices. Note CMA fining powers of up to 10% of global turnover.

17. Suspension and removal

17.1 We may warn you, limit your account, suspend it, or remove you permanently if you break clause 8, these terms, or the law, or if we reasonably believe you present a risk to another person.

17.2 We will tell you what ground we acted on, unless telling you would put someone at risk or prejudice an investigation. A permanent removal is always decided by a named person at TRYST, never by software alone.

17.3 Immediate permanent removal, with Exclusion Identifiers retained under clause 15.3, applies where:

- you are under 18, or sought contact with someone you believed to be under 18;
- you shared or threatened to share intimate images without consent;
- you coerced, threatened or attempted to control another person;
- you exposed a Member's use of TRYST;
- you are subject to a relevant court order; or
- you knowingly gave a false declaration under clause **[v1.1] 3.5**.

17.4 We may preserve evidence relating to a removal beyond the periods in clause 15, where we need it for an investigation, an appeal, or a legal claim.

17.5 Reporting another Member falsely, or organising others to report someone, is itself grounds for removal.

18. Appeals

18.1 You can appeal any suspension or removal by [route]. A different person from the one who made the decision will review it, and we will respond within 10 working days.

18.2 If we were wrong, we restore your account, refund anything you lost, and destroy any Exclusion Identifiers.

18.3 We publish, twice a year, how many accounts we removed, on what grounds, how many appealed, and how many appeals succeeded.

19. Reporting, and what we do about it

19.1 Report anything from inside any conversation or profile. Reporting takes one tap and you never have to explain yourself to the person you are reporting.

19.2 Our response times for each category of report are published in our Safety Policy. We do not hide them.

19.3 We will never tell the person you reported that it was you.

20. Emergencies come before privacy

20.1 Our privacy and deletion commitments are strong. But if there is a risk to someone's life or of serious harm, or where the law requires it, we will preserve and disclose whatever is necessary to the emergency services, the police, or a court.

20.2 We publish what we disclose, in aggregate, in our transparency report.

IN PLAIN ENGLISH — If someone's life is at risk, safety wins over discretion. Every time.

21. Our liability to you

21.1 We will supply TRYST with reasonable care and skill, as the Consumer Rights Act 2015 requires. Nothing in these terms limits that obligation.

21.2 We are not liable to you for the acts or omissions of another Member, or for anything that happens between you and another person, except where our own failure to take reasonable care caused it.

21.3 Where we are liable to you for financial loss, our total liability is limited to the greater of (a) the amount you paid us in the 12 months before the event, or (b) £100.

21.4 We are not liable for loss of profit, loss of opportunity, loss of reputation, or any indirect or consequential loss.

21.5 We are not liable for failures caused by events outside our reasonable control, or by your device, network or operating system.

22. What we do not exclude

22.1 Nothing in these terms limits or excludes our liability for:

- death or personal injury caused by our negligence, including any impairment of mental health;
- fraud or fraudulent misrepresentation;
- any breach of the obligations the Consumer Rights Act 2015 says we cannot exclude;
- our duties under the Online Safety Act 2023 and data protection law;
- anything else that cannot lawfully be limited or excluded.

22.2 Your statutory rights as a consumer are not affected by these terms. Agreeing to these terms does not mean you have accepted any risk.

REVIEW — Clause 22 is deliberately included and deliberately broad. We understand it removes any argument that we tried to exclude what CRA 2015 s.65 bars, and that s.65(2) means a "use at your own risk" framing would buy us nothing anyway. Our working assumption is that a narrow, honestly carved limitation survives the s.62 fairness test where an all-encompassing one would be struck out in full, taking clause 21 with it. Please confirm — and please push back if you think clause 21.3's cap is too low to be reasonable given what Members pay.

23. Your liability to us

23.1 If you break these terms or the law and that causes us financial loss — including a claim by another Member, a fine, or the cost of responding to an investigation — you agree to cover that loss.

23.2 This covers financial loss only. It does not apply where the loss was caused by our own failure to take reasonable care, and it is limited to what is reasonably foreseeable.

24. Changing these terms

24.1 We may change these terms. If a change materially affects you, we will tell you at least 30 days before it takes effect and explain what has changed and why.

24.2 If you do not accept a material change, you can close your account before it takes effect, and we will refund any unused prepaid amount in full.

24.3 We will not change these terms retrospectively.

25. Ending this agreement

25.1 You can close your account at any time, from inside the app, in one step.

25.2 Clauses 9, 10, 12, 13, 15, 17.4, 21, 22 and 23 survive the end of this agreement.

26. Complaints

26.1 Complain to [email]. We acknowledge within 2 working days and respond substantively within 10.

26.2 If you are not satisfied, you may refer the matter to [ADR provider], or contact Citizens Advice. Nothing here stops you going to court.

26.3 Our Online Safety Act complaints route, and the name of the senior manager accountable for it, are at [URL].

27. Law and courts

27.1 These terms are governed by the law of England and Wales.

27.2 You can bring proceedings in the courts of England and Wales. If you live in Scotland or Northern Ireland, you can also bring proceedings there. If you live elsewhere in a country where consumer law gives you the right to sue locally, that right is not affected.

27.3 There is no arbitration clause and no class-action waiver in these terms.

REVIEW — We have deliberately not included mandatory arbitration or a class-action waiver. Our view is that both are vulnerable under CRA 2015 s.62 and Sch. 2 in a consumer contract, and that including them invites CMA attention for limited benefit. Please confirm or overrule.

28. Contacting us

28.1 General: [email]. Safety and urgent reports: [route], monitored 24 hours.

28.2 Data protection: our Data Protection Officer at [email]. You can also complain to the Information Commissioner's Office.

28.3 Online Safety Act accountable senior manager: [name/role], [email].

29. Open questions for counsel

#	Question	Clause
Q1	Confirm the DPA 2018 Sch. 1 condition for retaining Exclusion Identifiers after erasure, and its resilience against an Art. 17 request.	15.3
Q2	Does any part of clause 11 stray into excluding liability for negligent design of a safety feature, which would be void under s.65 and could taint clause 21?	11.2
Q3	Is the £100 / 12-month cap reasonable, or does the low free-tier ARPU make it vulnerable under s.62?	21.3
Q4	Current commencement position for the DMCC subscription regime, and whether clause 16 meets it in full.	16
Q5	Could a Member's spouse establish a free-standing duty of care in tort notwithstanding clause 10.2?	10.2
Q6	Do our verification badges, ranking outputs or marketing create an implied representation of safety that clause 5 fails to displace?	5
Q7	Does clause 12.2 (we cannot read E2EE messages) weaken our OSA position, and how should the risk assessment address content moderation without content access? (See spec D-10.)	12.2
Q8	Is a cross-platform exclusion consortium (spec [v1.1] Source A §19.3 L5 ; v1.1 03 §8) viable under competition law and joint controllership?	—
Q9	Confirm the geo-block list, and whether operating in any listed market creates exposure for directors.	3.3
Q10	Review the Community Rules and Safety Policy as incorporated documents for consistency with these terms.	1.4
Q11 [v1.1]	Is scarce-side seeding that offers free paid features to women-identifying members lawful under the Equality Act 2010 (services), or must it be segment-based? (Spec D-06.)	16
Q12 [v1.1]	Does gating all discovery, Intents, conversations and media at age assurance (V2), with no other members' content visible before then, satisfy Ofcom's highly-effective-age-assurance expectations for dating services, including the May 2026 dating guidance?	3.1, 5.1
Q14 [v1.2]	Biometric login (clause 6.5) processes biometric data to identify the account holder. Is explicit consent valid (freely given) when login requires it, given the accessible alternative? Is a provider-held face reference compatible with clause 15's erasure promise if the provider's deletion completes after our 60-second erasure?	6.5, 15
Q15 [v1.2]	Does requiring biometrics to sign in, with the alternative path, meet the Equality Act 2010 reasonable-adjustment duty for disabled members?	6.5
Q13 [v1.1]	Is creating Exclusion Identifier codes for every verified member at verification (destroyed on deletion unless the member is removed) proportionate, and does it need its own Art. 9/10 and Sch. 1 analysis separate from post-removal retention? Should a biometric code ever be used (spec D-08)?	15.3

06 — Decisions, errata and change log (v1.1)

Part of the TRYST v1.1 baseline.

1. How v1.1 was reconciled

- **Precedence:** Source A (consolidated Spec + ToS, 4 Oct 2026) > Source B (Developer Specification, 17 Sep 2026) > Source C (PRD draft v0.1). See 00 §2.
- **Where A is silent,** B's requirements are kept: the Discretion Policy object, reveal ladder, preference transparency, quick-exit timing, Privacy Auditor, no-go conditions, the weighted P1 ranker and mode names.
- **From C:** non-functional requirements, the couple dissolution and data-export APIs, and the competitor rows. C's US expansion, its pricing and its unit economics are dropped.
- **Where A contradicts itself or contains an arithmetic error,** v1.1 fixes it and records it in §3.

2. Decision log

Status: *Resolved* = decided by a source document. *Default* = decided in v1.1 to remove ambiguity, and reversible by the owner. *Open* = needs an owner decision by the gate shown; build against the stated default until then.

ID	Decision	Status	Resolution / default	Owner · deadline
D-01	Are unattached singles allowed?	Open (v1.2)	The founder scope statement (01 §1) says TRYST is designed for adults already in a relationship, while Source A §3 includes solos open to couples (S4) and the curious (S6), who may be unattached. Default: the product and marketing target partnered adults and couples; unattached solos are <i>accepted</i> as S4 (needed for THIRD liquidity) and S6, but not targeted. Partner awareness is self-declared and never verified or displayed.	Founder · before P1 public beta
D-02	Chat encryption vs server-side scanning	Resolved (A §9.2, §21.6)	Full E2EE (MLS). Guardian runs on device and sends score + reason code only. Content reaches the server only through member-disclosed reports or Panic. Source B FR-009's "media scanning in real time" is met on device.	—
D-03	Financial base case	Resolved (A §15.3)	A's illustrative model is the base case (CAC £38, 11% payer rate, 19% churn). B's targets (< 7% churn, > 3× LTV:CAC) are superseded. Thresholds corrected in E-03.	—
D-04	What does the ENVOY tier's "runs outreach" mean?	Default	Envoy <i>proposes</i> intents with rationale; every outbound intent needs a one-tap human approval. Envoy never sends anything to a human as the member (P4; ToS cl. 7.3).	Head of Product · before P3
D-05	GHOST add-on charges for decoy skin, duress PIN and voucher billing, which conflicts with B's "never monetise safety" and C's "discretion is free"	Open	Default: core decoy, duress PIN, Burn, masking, zones and voucher billing are free . GHOST sells multiple decoy skins, custom TTLs and extended ExclusionRing heuristics at £7.99. Revenue impact is to be modelled.	CEO · before P3 pricing
D-06	"Free ENVOY for 6 months for women-identifying users" vs gender-neutral pricing (B) and the Equality Act 2010	Open (counsel)	Default: scarce-side seeding by segment (S4) and city-balance need, not gender. The gender-based variant ships only with a written counsel opinion.	DPO/Counsel · before P1 public beta
D-07	Contract formats	Default	OpenAPI 3.1 at the edge; Protobuf for gRPC, the agent bus and events (schema registry); generated types on both sides.	Principal Eng · P0
D-08	Biometric ban anchor	Open (counsel)	Default off . Anchors use document, device and payment HMACs only. V1 retains no biometric template (A §8.1).	DPO/Counsel · before P1
D-09	Microservices vs modular monolith for core domain services	Default	Agents are separate deployables (as A requires). Core domain services form a modular monolith (core-svc) with enforced module boundaries, given a 2-person backend team; split on evidence. Identity, ban, media and safety are always separate.	Principal Eng · P0

ID	Decision	Status	Resolution / default	Owner · deadline
D-10	Hash matching of private (E2EE) media	Open (counsel + vendor)	Default: on-device pre-send nudity classifier always; on-device perceptual hash matching against licensed lists where licensing permits. Otherwise rely on reports, StopNCII and the NCII escalation path. Must be addressed in the OSA risk assessment (ToS Q7).	Head of T&S · P0
D-11	In-app purchase in store builds	Default	Web-first. Store builds use reader/link-out entitlements where rules allow; where IAP is required, offer TRYST+ only through IAP, with entitlements unified server-side.	Head of Product · P3
D-12	Delivery timeline	Resolved (A §17)	A's P0-P4 over 24 months, with phase gates, replaces B's 68-week sequential plan.	—
D-13	Launch markets	Resolved (A §14.5)	UK + IE, then EU wave 2, then CA/AU/NZ. US out of scope; never Africa or MENA.	—
D-14	Mode naming	Default	B's SPARK/EMBER/THIRD/QUAD/OPEN as UI labels; A's intent shapes as data values (02 §2.3).	Design · P1
D-15	Concierge tier (B, C)	Default	Deferred beyond P4. A's tiers are canonical.	CEO · P4 review
D-16	Who can sign in, and with what	Resolved (founder, v1.2)	Only the account holder. App and mobile PWA: two biometric factors (device-bound biometric passkey + liveness face match). Desktop: two factors (passkey + registered-phone approval or a second security key). No passwords, SMS or email codes, or staff overrides (02 §3.3).	—
D-17	Members who cannot use face or fingerprint biometrics (disability, device)	Open (counsel)	Default alternative: two registered FIDO2 security keys with PIN, plus a provider video liveness check with assisted human review in place of B2. Required so biometric consent is freely given and to meet the Equality Act (ToS Q14, Q15).	DPO/Counsel · before P1 beta
D-18	Where the B2 face reference is held	Open (counsel + vendor)	Default: held by the liveness provider under contract (TRYST stores only a token); deleted at the provider on erasure. Alternative rejected by default: TRYST-held encrypted template (contradicts A §8.1 "no biometric template").	DPO/Counsel · P0
D-19	Brand identity	Resolved (founder, v1.2)	Founder-supplied logo used exactly as supplied; palette sampled from it (oxblood and gold on night, cream replacing "bone"). Discretion rules still keep the logo off decoy, lock-screen and notification surfaces (01 §3).	—
D-20	Source B's Conversation Agent (suggested openers and replies) is not among Source A's seven agents	Default	Ship in P3 as an opt-in, on-device, draft-only assistant inside the member's own client. It never sends anything and never talks to another member as the member (FR-066; P4).	Head of Product · P3
D-21	SEO approach for "dynamic AI agent SEO, hyperlinks, backlinks and top ranking" (founder request, v1.2)	Default	Build Herald as a dynamic AI SEO agent with human editorial approval, isolated from member data. Links are earned through digital PR, partnerships and open-source work, never bought or spammed. Top-3 rankings are targets tracked weekly, not guarantees. Black-hat tactics are prohibited because a penalty would remove the business's main low-CAC channel (FR-073-FR-077).	Growth lead · P1
D-22	Payment processor	Open — conditional (founder chose Stripe, v1.2)	Stripe is primary, conditional on Stripe's written approval of TRYST's business model in P0, because Source A §14.4 states Stripe prohibits adult content and the web surface carries explicit media. A high-risk acquirer is integrated as the backup MID from day one and becomes primary if Stripe declines. Stripe's restricted-businesses policy could not be checked from the build environment. Fallbacks if Stripe approves dating but not explicit media: (a) keep explicit media off the Stripe-billed surface, or (b) bill everything through the high-risk acquirer.	CEO + Head of Payments · G-P0

ID	Decision	Status	Resolution / default	Owner · deadline
D-23	Source B's final directive ("partnered adults, affairs and threesomes are explicit product categories — not hidden behind generic dating language") vs Source A §14.3 (store builds positioned as "discreet dating" to survive Apple 1.1.4)	Default	Explicit categories (SPARK/EMBER/THIRD/QUAD/OPEN, affairs, threesomes) on the web PWA, marketing site and Herald SEO content . Only app-store listings and in-store-build copy use neutral "discreet dating" wording, with no explicit imagery. Positioning never says "cheat better".	Head of Product · before store submission (P3)
D-24	Identity service iteration 1 scope	Default	Implemented and tested: contact OTP sign-up (dev returns the code), passkey registration (user verification required, discoverable), mobile login B1 + B2 (pluggable liveness provider; fails closed; dev provider refused in prod), desktop login F1 + F2 (registered-phone approval or a second security key), session limits, logout, Burn. Next: DPoP-bound tokens (bearer over TLS until then), PostgreSQL IDENTITY-DB store (in-memory now), native biometric-only attestation, real liveness provider, step-up and recovery endpoints.	Principal Eng · P1
D-25	The founder OS brief asks for a platform that runs "without requiring human intervention"	Default (founder brief, v1.3)	Adopted with bounded autonomy : autonomy levels L0-L3 let agents observe, recommend, make reversible changes and run pre-approved runbooks on their own; irreversible, member-safety, money-movement and production-change actions always need a named human, and no agent may act as a member. Required by UK GDPR Art 22 and principle P4; enforced in code (aigov, FR-080)	Principal Eng + DPO · P1
D-26	BitriPay as a core payment gateway	Default (founder brief, v1.3)	BitriPay is integrated as a third PaymentProvider adapter (built, sandbox-tested). Live use is refused in code unless a market is enabled for both TRYST and BitriPay, with a legal opinion and BitriPay's written acceptance, because its documented rails serve markets TRYST excludes (D-13). Near-term use: partner commissions and licensee flows via connected accounts	CEO + Head of Payments · P0 check, P4 use
D-27	AI credit / ACU system	Default	ACU meters partner and licensee API usage only. Members never buy ACU: Keys already serve that role, and a second credit currency would breach P8 ("unreadable credit mechanics"). The founder's earlier instruction to ignore the ACU sections of the pasted KODA/BitriPay docs is respected: TRYST's ACU is defined here, not copied	CEO · P4 pricing
D-28	"Auto-Repair Agent automatically patches issues and redeploys components"	Default	Auto-Repair runs pre-approved runbooks (rollback, restart, scale out, halt canary, rotate a leaked key) on its own within a budget. It may open a fix pull request; that PR follows normal review and CI, and promotion to production is a human decision	Principal Eng · P1
D-29	Connectors that would expose members (WhatsApp, member data in CRM, data enrichment, ad platforms, live location)	Default	Declined for member data: each would reveal a pseudonymous member to a third party or on a shared device (FR-002, P3). Operator- and B2B-only uses are kept (CRM for licensee pipeline, enrichment for B2B leads)	DPO · P1
D-30	Dynamic pricing and data-intelligence revenue	Default	Pricing experiments are market-level, published and identical for every member in a market; no personalised or surge pricing. No revenue from member data (P8); the alternative is free published research using differential privacy, and licensing of technology trained on licensed or synthetic data	CEO + DPO · P3
D-31	Executive agents (CEO, COO, CFO, CTO, CMO, CRO)	Default	Analyst copilots for the named executive, reading aggregate data only, with no actions beyond recommendations. Accountability stays with people	CEO · P1

ID	Decision	Status	Resolution / default	Owner · deadline
D-32	KYC, KYB and AML agents	Default	Members get age assurance and verification, not identity KYC (pseudonymity, FR-002). TRYST holds no customer funds, so AML duties sit with regulated processors; TRYST runs KYB and sanctions/PEP screening on licensees and commission partners	DPO/Counsel · P4
D-33	Should TRYST's name or brand mention AI?	Resolved (founder, v1.3: "Do not add AI on TRYST name or brand")	Never. No "AI", "A.I." or "artificial intelligence" in the name, tagline, logo, store listings, marketing, Herald content, UI or notifications (01 §3.3). Internal engineering docs keep technical terms. Enforced by a build test. Doc 09 is renamed "Operating System (internal engineering)"	—
D-34	Are people in open relationships in scope?	Resolved (founder, v1.3)	Yes, explicitly welcome as S2 (solo) or S3/S5 (couples), shown on the marketing site and in the OPEN mode. Complements D-01	—
D-35	Launch model: Source A's 9-month invite-only period and London-first city sequence	Resolved (founder, v1.3: open public sign-up, keep the gender-balance waitlist; "if this is the best and proven strategy, go for it")	Public sign-up from launch in London, Manchester, Birmingham and Brighton, with matching opening one city at a time in that order (London first), each at 2,000 verified members once the previous city is live. This is the proven density-first pattern for two-sided products (campus-by-campus and city-by-city launches, public waitlists); opening four cities to matching at once was considered and rejected as less proven because it splits density and seeding spend. Kept: a balance waitlist (now gender-neutral, D-36); referrals as a growth channel; the P1 control cohort without Envoy (London). Enforced in backend/internal/launch (FR-089, FR-090). Risks — closed in v1.3: (1) acquisition cost without the invite loop is capped by the automatic paid-acquisition brake (FR-091); (2) the Equality Act exposure of a gender-based waitlist is removed by replacing it with the gender-neutral demand-balance rule (D-36, FR-090)	Founder · closed
D-36	How to keep cities balanced without admitting people by gender	Resolved (founder, v1.3: "close these risks")	The waitlist no longer reads the applicant's gender or any other characteristic. It paces admission by the load on the groups the applicant asked to meet (seekers per group member, cap 2.2), applied identically to everyone; a person waits only when the people they want to meet are already saturated. This removes direct discrimination (which, for sex, cannot be justified under the Equality Act 2010 s.13). Any uneven effect between groups is a proportionate means to legitimate aims (members' safety from harassment volume; OSA duties), the standard s.19 justification. Confirmation is part of the written UK legal opinion that is already a G-P0 launch condition, not a separate gate. The gender ratio KPI stays monitored. D-06 is unaffected	Founder · closed (confirmed in the G-P0 opinion)

3. Errata fixed in v1.1

ID	Issue (source)	Fix in v1.1	Where
E-01	A §8.1 lets V1 (liveness only) send intents and open threads, contradicting ToS cl. 5.1 ("adult verified before Intent"), B FR-001 and the OSA expectation that age assurance covers the inbox.	V2 required for discovery, intents, Envoy, threads, media and meets. V0/V1 = onboarding only. min_counterparty_tier default raised to V2.	02 §3.2; FR-001
E-02	A §19.3 L2 computes ban anchors from the ID document number and payment token at ban time , but A §8.1/§10 retain neither, so anchors could never be written.	anchor_candidate HMACs are computed at V2 and held in IDENTITY-DB under the member's keys. They are erased with the account unless a permanent ban copies them to BAN-DB. Disclosed in ToS cl. 15.3.	02 §4.1; 03 §8
E-03	A §15.3 says viability requires "blended CAC < £12". At the stated 11% payer rate and 52% V2 rate, break-even blended CAC per sign-up is £8.18 (£143 × 0.11 × 0.52). £12 breaks even only at ≥ 16% payer rate.	Thresholds table added: £8.18 at 11%, £11.90 at 16%, £13.38 at 18%; 3× LTV:CAC needs £2.73-£4.46.	01 §10.3
E-07	B specifies React Native for mobile; A requires native Swift/Kotlin for security.	Native (A). React Native not used. Shared Rust crypto core instead.	04 §3
E-09	ToS Q8 cites "spec §19.5" for the consortium; it is §19.3 L5.	Corrected in ToS draft 0.9.1.	05 cl. 29
E-10	ToS cl. 17.3 relies on "a false declaration under clause 3.2", but cl. 3.2 contains no declaration. A §19.3 L3 defines a conduct attestation at V2.	New ToS cl. 3.5 (conduct declaration); cl. 17.3 now refers to it.	05 cl. 3.5, 17.3

ID	Issue (source)	Fix in v1.1	Where
E-11	A §9.2 media retention of "7 d default" was ambiguous between profile and thread media.	7 d applies to thread media; profile media persists until removed.	02 §7
E-12	A §8.1 says V1 retains "no biometric template", but A §19.3 L2 lists anchor_biometric.	Biometric anchor off by default (D-08).	02 §4.2
E-13	Sources A and B both call themselves "v1.0 / 17 Sep 2026" (A's cover says 4 Oct 2026, its footer says 17 Sep).	Renamed Source A / Source B with file dates.	00 §2
E-14	B's tiers (Black £24.99, Duo £34.99, Signals, Vault+, Concierge Black) conflict with A's.	A's tiers are canonical; Signals $\approx$ Keys; Vault+ folded into TRYST+ (extended media); Concierge deferred (D-15).	01 §10.2
E-16	C planned a US launch, which A excludes.	US out of scope. C's US legal notes archived for reference only.	01 §11
E-21	A §11 POST /v1/panic performs a discretion wipe, while A §19.7 defines Panic as emergency escalation and says it must be "distinct from Burn".	Wipe endpoint renamed POST /v1/burn; emergency is POST /v1/safety/panic. Distinct UI controls.	02 §5.3, §5.8; 04 §5, §7
E-23	A §10 couple_link requires both co-sign timestamps to be NOT NULL but has no pending state for the invite flow.	couple_invite table added; couple_link written only on the second co-sign.	02 §4.3
E-26	Slate size given as "12-18" (A §4) and "18" (A §11).	"Up to 18" cards including a 15% exploration floor.	02 §5.5
E-29	A §11 lists POST /v1/handshake as client-callable, but handshakes must start only on mutual intent.	System-initiated only; clients read Briefs and post decisions.	02 §5.6
E-30	A's §17 budget (£5.9M) excludes the £85-140k/yr insurance from A §20.5.	Called out as an addition in 01 §12.	01 §12
E-32	Ashley Madison membership cited as 91M+ (B §1) and 80M+ (A §1).	Both kept as attributed, unaudited claims; external materials must use a current, dated figure.	01 §2
E-33	Source B §11.1 targets verification + safety cost of £1.20-£2.20 per MAU alongside gross margin > 72%. At B's own conversion (7-10%) and ARPPU (£27-£34), revenue per MAU is only about £1.89-£3.40 before fees, so that cost alone consumes most of the revenue and makes 72% unreachable.	B's economics remain superseded by Source A's model (D-03); the inconsistency is recorded so B's table is not reused externally.	01 §10.3; 07
E-31	B places THIRD/QUAD in delivery Phase 4; A treats S3/S5 as first-class and ships VC in P1.	THIRD/QUAD matching in P1 (A); DUO billing in P3.	FR-049

(Gaps in the E-numbering are intentional; IDs are never reused.)

4. Change log

Version	Date	Change	Author
1.0 (Source B)	17 Sep 2026	Developer Specification	Founder
0.1 (Source C)	4 Oct 2026	PRD & business plan draft	Claude Code (draft)
1.0 (Source A)	4 Oct 2026	Consolidated Spec + ToS draft 0.9	Office of the Group CEO
1.1	4 Oct 2026	Merged A + B + C; split into Product / Shared Contracts / Backend / Frontend / ToS; decisions D-01-D-15; errata E-01-E-31; baseline and change control established	Claude Code, for founder review
1.2	4 Oct 2026	Founder logo added exactly as supplied (assets/brand/); brand palette and tokens from the logo; account-holder-only login: two biometric factors on app/mobile PWA, two factors on desktop (D-16-D-18, FR-056-FR-063, NFR-16/17, ToS 0.9.2 cl. 6.5, Q14-Q15); source register (00 §5); D-23 market-clarity vs store wording; Stripe as primary processor, conditional on approval (D-22, FR-078); intent change history (FR-079); E-33; Herald dynamic SEO agent and marketing site (FR-073-FR-077, NFR-20, D-21); full Source B added verbatim as appendix 07 with traceability; Source B §9-10 coverage (03 §6.8), Article 6 bases, FR-069-FR-072, NFR-18-19, ToS 0.9.3; FR-064-FR-068, D-20; founder product-scope statement, executive product decision, one-sentence pitch and market thesis added verbatim (01 §1-§2), D-01 reopened, E-32; spec folder renamed docs/spec	Claude Code, for founder review

Version	Date	Change	Author
1.3	4 Oct 2026	Operating system for TRYST (new doc 09, internal engineering) from the founder's Master founder OS brief: brief coverage register, market-gap review, user ecosystem, command centres for every user type, agent workforce and core agents, self-managing layer, BitriPay API door, connector register, architecture, ops and partner schemas, ops and partner APIs, monetisation, security, Admin Super Control Centre, roadmap; D-25-D-32, FR-080-FR-088. Implemented and tested: AI Governance gate with workforce registry, kill switch and hash-chained audit (backend/internal/aigov); BitriPay payment adapter, sandbox-only (backend/internal/payments/bitripay.go) ; no "AI" in name or brand (D-33, 01 §3.3, build test); open relationships explicitly welcome (D-34); public sign-up in four UK cities, matching opening city by city from London, with a balance waitlist (D-35, FR-089, FR-090, backend/internal/launch) ; launch risks closed: gender-neutral demand-balance waitlist (D-36, FR-090) and automatic paid-acquisition brake (FR-091)	Claude Code, for founder review

07 — Source B: Founder Developer Specification (verbatim) and traceability

Part of the TRYST specification baseline.

This appendix reproduces **Source B** — *TRYST AI Discreet Relationship Platform — Developer Specification, v1.0, 17 September 2026* — **word for word**, extracted directly from `source/TRYST_Developer_Specification_2026-09-17.docx`. The founder re-confirmed it on 4 October 2026.

- **Headings, tables and callouts** keep the original wording. Only the formatting has been converted to Markdown.
- **Precedence:** where Source B conflicts with Source A (the consolidated Spec + ToS), 00 §2 decides, and 06 records the outcome.
- **Traceability:** §1 below shows where each Source B section is implemented and what v1.2 added to cover its gaps.

1. Traceability — where each Source B section lives

Source B section	Implemented in	Notes
Product scope	01 §1 (verbatim)	Unattached solos: D-01
Executive product decision, pitch	01 §1.1 (verbatim + mapping)	
1. Market thesis	01 §2 (verbatim)	Member counts: E-32
2. Users and modes	01 §5.2, 02 §2.3, 02 §9	Mode names are the UI labels (D-14)
3. Experience architecture	04 §4; loops in 01 §6	Stage 9 "emergency lock" → FR-064 (new in v1.2)
4. Discretion product system	04 §5, 02 §6.4	New in v1.2: FR-064 emergency lock, FR-065 rotating location cells; FR-026 wording and FR-029 disclosure updated
Privacy reality	01 §4.2 P9, 01 §9	
5. AI agentic operating system	01 §7, 03 §3	Agent mapping in §2 below; new FR-066 (conversation drafts) and FR-067 (member-provided windows only)
6.1 Learned dimensions	03 §4	New FR-068 behaviour decay
6.2 Ranking model	03 §5.2 (P1 baseline), 03 §5.3	Optimisation target: Source A's QCAM (meets) is the north star; B's "qualified reciprocal conversations" is a secondary KPI (01 §13.1)
7. Functional requirements (FR-001-015)	02 §10.1	B's FR numbers were folded into the v1.1+ catalogue (each row cites "B")
8. Architecture and data, 8.1 entities, 8.2 API	02 §4, 02 §5, 03 §2	Stack: Source A's native apps replace React Native (E-07)
9. Security and privacy doctrine	03 §6, 01 §14.2; full coverage table 03 §6.8	New in v1.2: Article 6 bases, FR-069-071, NFR-18-19, backup disclosure
10. Trust, safety and misuse controls	03 §7, 03 §8; coverage table 03 §6.8	New in v1.2: FR-072 (protected-trait harassment), FR-069 upload warning
11. Business model, 11.1, 11.2	01 §10, 01 §11	Tiers and economics superseded by Source A (E-14, D-03); the revenue ethic is kept as P8
12. Delivery plan, 12.1 team	01 §12	Source A's P0-P4 and 17-person team replace B's plan (D-12)
13. MVP backlog and acceptance	03 §12, 04 §10	
14. KPIs, governance, no-go	01 §13 (G-NG-1 to G-NG-7 come from B)	
15. Source register	§3 below (verbatim)	Revalidate before launch
Final directive	01 §4, 01 §12, no-go list 01 §13.2	Market clarity vs store wording: D-23. Readiness = gates G-P0-G-P4 and G-NG-1-10

1.1 Requirement-ID crosswalk (Source B FR → catalogue FR)

Caution: Source B numbers its own FR-001–FR-015, and those numbers do **not** mean the same as the catalogue's (02 §10.1). Inside the verbatim text in §3, an FR number is Source B's. Everywhere else, FR numbers are the catalogue's.

Source B ID	Source B requirement	Catalogue requirement(s)
B FR-001	18+ entry	FR-001 (V2 gate)
B FR-002	Pseudonymous identity	FR-002
B FR-003	Intent modes	FR-004
B FR-004	Discretion policy	FR-005
B FR-005	Couple representation	FR-006
B FR-006	Reveal ladder	FR-008
B FR-007	Private discovery	FR-009
B FR-008	Contact suppression	FR-010
B FR-009	Chat safety	FR-032 (report, block), FR-033 (real-time on-device scanning, D-02), FR-034 (escalation), FR-013 (rate and inbound limits)
B FR-010	Behaviour controls	FR-024
B FR-011	Quick exit	FR-029
B FR-012	Deletion	FR-038
B FR-013	Neutral notifications	FR-026
B FR-014	Reports/appeals	FR-034, FR-035
B FR-015	No commercial sex	FR-037

1.2 Architecture, entity and endpoint crosswalk (Source B §8 → contracts)

Layers. Every Source B layer is kept. The one exception is "React Native iOS/Android", which is replaced by native Swift/Kotlin apps with a shared Rust crypto core (E-07, 04 §3).

- "OIDC/passkeys" → OAuth2 + DPoP with passkey-only, two-factor login (02 §3.3).
- "Alias service" → pseudonymous `display_name` in PROFILE-DB, joined to identity only through the HSM pepper.
- "Relationship/block graph projection" → the Stage 0 gates (03 §5.1).

Source B entity	Contract entity (02 §4)
User	identity (IDENTITY-DB) + profile.status
AliasProfile	profile (display_name, bio_ct, visibility, version) + media
IntentProfile	intent_vector
DiscretionPolicy	discretion_policy (versioned)
CoupleUnit	couple_link + couple_invite
Introduction	slate_item
RevealGrant	reveal_grant
Room	thread (MLS group)
BehaviourEvent	event envelope (02 §8)
ReportCase	report_case + enforcement
DeletionJob	ops.deletion_job

Source B endpoint	Contract endpoint (02 §5)
POST /v1/identity/verify	POST /v1/assurance/session (+ /v1/assurance/callback)
PUT /v1/discretion-policy	PUT /v1/discretion/policy
POST /v1/intents	POST /v1/intake/vector/confirm (member-confirmed intent profile); outbound interest is POST /v1/intent
POST /v1/couple-units	POST /v1/couple/invite + POST /v1/couple/cosign

Source B endpoint	Contract endpoint (02 §5)
GET /v1/discovery	GET /v1/slate
POST /v1/introductions/{id}/decisions	POST /v1/intent · /v1/pass · /v1/save · /v1/preferences/correction
POST /v1/reveal-grants / DELETE .../{id}	Same paths
POST /v1/reports	POST /v1/report
POST /v1/preferences/reset	Same path
DELETE /v1/users/me	DELETE /v1/account

2. Agent mapping (Source B's ten → v1.2's seven)

Source B agent	v1.2 home	Boundary carried over
Chemistry Agent	Mirror (learning) + Broker (ranking)	Ranks; never fabricates interest or sends messages (FR-055)
Intent Agent	Cartographer	Member confirms every structured field (FR-003)
Discretion Agent	Curtain	May tighten, never loosen (02 §6.4)
Timing Agent	Envoy (logistics dimension)	Member-provided windows only, never inferred routines (FR-067)
Conversation Agent	On-device drafting assistant (D-20)	Draft-only, opt-in, never auto-sent (FR-066)
Authenticity Agent	Guardian (M7, M7b) + admission control L1	Friction/hold automatic; permanent ban only by a human (FR-035)
Risk Agent	Guardian (M6a-e, escalation matrix)	Escalates, restricts, preserves evidence (FR-034)
Group Agent	Broker (max-min scoring) + Envoy + couple semantics	Cannot add anyone or reveal anyone without consent (FR-006, FR-049)
Learning Agent	Mirror	Approved features only; no message content (03 §4.2)
Privacy Auditor	Release gate (not an agent)	Can block a release (03 §9)

3. Source B — full text (verbatim)

CONFIDENTIAL PRODUCT BLUEPRINT

TRYST

Private chemistry. Intelligent discretion.

AI-AGENTIC PLATFORM FOR AFFAIRS, EXTRA-RELATIONSHIP CONNECTIONS AND THREESOMES

Developer-ready product, AI, privacy, commercial and delivery specification Version 1.0 | 17 September 2026

PRODUCT SCOPE

TRYST is designed specifically for adults who are already in a relationship and want a discreet one-off affair, a continuing extramarital connection, another couple, or a threesome. Existing-partner permission is not a condition of membership. Every person using TRYST must nevertheless be an independently consenting adult; the platform prohibits coercion, stalking, blackmail, impersonation, minors, non-consensual imagery and commercial sexual services.

Executive product decision

Build TRYST as a privacy-first, affair-specific connection platform rather than a generic dating app. Its strategic advantage is an AI relationship-intelligence layer that learns attraction, discretion requirements, availability, risk tolerance and interaction quality while deliberately collecting less identifying data than incumbents.

Decision	Specification	Why it wins
Name	TRYST	One sharp word meaning a private romantic meeting; instantly communicates the category.

Decision	Specification	Why it wins
Audience	Partnered adults of any gender/orientation; couples seeking a third or another couple	Clear, high-intent niche rather than mass-market dating.
Promise	Private chemistry. Intelligent discretion.	Combines emotional benefit with the functional moat.
Core modes	One-off, Ongoing, Couple+, Threesome, Couple-to-Couple	Intent is explicit before matching.
Moat	Discretion Graph + Behavioural Compatibility Engine	Learns fit without exposing identity or partner data.
Revenue	Subscription-led with credits and premium privacy	Predictable recurring revenue plus high-margin intent transactions.
Launch	Web/PWA first, then native apps subject to store review	Reduces store dependency and permits rapid privacy iteration.

ONE-SENTENCE PITCH

TRYST is the intelligent private network where attached adults and couples find precisely matched affairs, extra-relationship connections and threesome partners without exposing more of their identity than necessary.

1. Market thesis

The category is already validated. Ashley Madison markets private, anonymous connections and claims more than 91 million members. Gleeden uses a credit economy for extramarital connections; Victoria Milan markets a rapid-exit “panic button.” The opportunity is not to prove demand but to rebuild the category around modern privacy engineering, verified authenticity, gender-neutral pricing, transparent AI and higher-quality matching.

Incumbent pattern	Observed weakness	TRYST response
Large anonymous inventory	Fake profiles, low trust and expensive dead-end messaging	Liveness, authenticity confidence and conversation-quality scoring.
Credit-heavy messaging	Revenue can conflict with real connection outcomes	Subscription provides core value; credits enhance but do not manufacture access.
Basic privacy theatre	A panic button cannot protect server-side data	Pseudonymity, data minimisation, isolated vaults, short retention and discreet billing.
Pairwise matching	Poor fit for couples and threesomes	Native multi-person matching and room consent.
Swipe engagement	High noise and fatigue	Small daily set of high-confidence, intent-aligned introductions.

2. Users and modes

Mode	User need	Eligibility / system rule
SPARK	Discreet one-off encounter	Attached status accepted; match on availability, distance and chemistry.
EMBER	Ongoing affair or recurring connection	Match on cadence, emotional intensity, communication windows and expectations.
THIRD	A couple seeks one additional person	Both couple members must separately verify the shared couple profile; the third is treated as an equal user.
QUAD	Two couples seek one another	Four independent accounts; shared match only after each side’s configured quorum.
OPEN	Flexible extra-relationship exploration	User selects allowed configurations and visibility rules.

TRYST does not ask a user to prove permission from a spouse or external partner. It only requires consent from the people who are actually participating through TRYST. A couple profile cannot represent the second member without that member’s own account and confirmation.

3. Experience architecture

Stage	User experience	Platform behaviour
1. Quiet entry	Neutral landing page, email alias or private number, discreet notifications	No contact upload; no social-login requirement; risk-screen session.
2. Adult/authenticity check	Age assurance and liveness; legal name stays private	Store verification token/result separately from profile.

Stage	User experience	Platform behaviour
3. Discretion setup	Choose icon, notification wording, app lock, visibility and location fuzzing	Generate a Discretion Policy applied across all services.
4. Intent map	Select SPARK, EMBER, THIRD, QUAD or OPEN and define boundaries	Build explicit intent vector and exclusions.
5. Discovery	Receive a limited ranked set with compatibility explanations	Candidate retrieval excludes blocks, contacts hash list if opted in, risk conflicts and impossible schedules.
6. Reveal ladder	Alias → blurred gallery → selected photos → voice/video	Each reveal is separately approved and revocable.
7. Private room	Chat, disappearing media, availability coordination	Encrypted content, screenshot friction, abuse/scam detection.
8. Outcome learning	Private “more/less like this” and quality feedback	Update preference model; feedback is not shown to the other user.
9. Exit	Pause, vanish, delete or emergency lock	Immediate discovery removal; deletion and retention workflow starts.

4. Discretion product system

Control	Requirement
Neutral shell	Optional generic icon/name on device where technically possible; neutral web favicon and page title option.
Notification privacy	Default notification says “You have an update”; never includes names, images or message excerpts.
App lock	Independent PIN/biometric lock; session re-authentication after configurable inactivity.
Quick exit	Immediate neutral screen plus client cache clearing; clearly state this cannot hide network/device history.
Location veil	Use distance bands and rotating approximate cells; never reveal live or exact home/work coordinates.
Invisible mode	Only people the member explicitly likes can see them, subject to mutual eligibility.
Recognition shield	Optional on-device contact hashing to suppress known contacts; raw address book never uploaded.
Media vault	Face blur, selective reveal, recipient watermark, expiry, revocation and no public URLs.
Discreet billing	Neutral merchant descriptor where lawful and accurate; transparent receipt settings; no deceptive descriptor.
Emergency lock	Freeze logins and sessions, hide profile and disable media links from a separate recovery path.
True deletion	Free account deletion; documented purge and backup expiry; no false “full delete” claim.

PRIVACY REALITY

TRYST must never promise that an affair is undetectable. It can minimise platform exposure, but cannot guarantee secrecy from a device owner, bank, employer, network administrator, recipient screenshot, legal process or determined third party.

5. AI agentic operating system

Agent	Role	Autonomy boundary
Chemistry Agent	Learns attraction patterns from explicit choices and reciprocal outcomes	Ranks; never fabricates interest or sends messages.
Intent Agent	Transforms free-text goals into SPARK/EMBER/THIRD/QUAD/Open vectors	User confirms every structured field.
Discretion Agent	Creates and enforces each member’s privacy policy	May reduce exposure automatically; cannot weaken settings.
Timing Agent	Matches compatible communication and meeting windows	Uses user-provided windows, not inferred home routines.
Conversation Agent	Suggests openers, boundary questions and respectful replies	Draft-only; no impersonation or autonomous chat.
Authenticity Agent	Detects bots, scams, stolen images and identity inconsistency	Can add friction/hold; permanent ban follows policy/review.
Risk Agent	Detects harassment, blackmail, coercion, trafficking and underage signals	Escalates, restricts and preserves evidence under defined rules.

Agent	Role	Autonomy boundary
Group Agent	Solves multi-person compatibility and quorum for THIRD/QUAD	Cannot add anyone or reveal one user to another without product consent.
Learning Agent	Updates behavioural embeddings and experiment assignments	Uses approved features only; sensitive raw content excluded by default.
Privacy Auditor	Continuously tests access, retention and model leakage	Can block a release when thresholds fail.

6. Behaviour-learning design

6.1 Learned dimensions

Dimension	Signals	Treatment
Attraction	Likes, passes, private preference corrections	Embeddings; never disclose inferred sensitive traits.
Intent fit	Mode, desired duration, emotional depth, exclusivity expectation	Hard/soft constraints controlled by user.
Discretion fit	Visibility, response windows, reveal pace, distance bands	Mismatch strongly penalised.
Communication fit	Message cadence, length preference, voice/video comfort	Aggregate features; raw content short-lived and excluded from training by default.
Reliability	Mutual attendance feedback, cancellations, ghosting pattern	Private score with anti-retaliation safeguards.
Respect/safety	Blocks, reports, boundary violations, pressure patterns	Safety model; reports never become visible reputation badges.
Group symmetry	Each participant's attraction and boundaries	Maximise minimum compatibility, not average score alone.

6.2 Ranking model

TRYST MATCH SCORE

Eligibility filter first. Then: 24% reciprocal attraction + 20% intent fit + 15% discretion fit + 12% communication fit + 10% availability/logistics + 8% reliability + 6% novelty + 5% model confidence – risk penalties. THIRD/QUAD adds a max-min fairness constraint so one person's strong interest cannot overwhelm another participant's low fit.

- Hard exclusions: age uncertainty, block edges, missing participating-couple confirmation, prohibited purpose, incompatible mode, critical safety risk.
- Behaviour decays over time; the user can reset all inferred preferences.
- Every introduction displays two to four non-sensitive reasons and offers a correction control.
- No ranking uplift based on gender. Paid visibility applies only inside an already eligible set and is visibly labelled.
- Optimisation target: qualified reciprocal conversations and user-rated fit—not time spent, message volume or credit burn.

7. Functional requirements

ID	Requirement	Acceptance criterion
FR-001	18+ entry	Discovery remains locked until age-confidence threshold and liveness pass.
FR-002	Pseudonymous identity	Public profile never requires legal name, employer or exact address.
FR-003	Intent modes	A member may enable multiple modes with separate visibility and filters.
FR-004	Discretion policy	Every API returning profile data evaluates the requester against current policy.
FR-005	Couple representation	THIRD/QUAD shared profile publishes only after all represented members approve exact version.
FR-006	Reveal ladder	Each media/identity disclosure records grant, recipient, version and expiry.

ID	Requirement	Acceptance criterion
FR-007	Private discovery	Invisible members are discoverable only following their outbound interest.
FR-008	Contact suppression	Optional hashing occurs on device; server stores salted blinded tokens only.
FR-009	Chat safety	Rate limits, block, report, media scanning and blackmail/threat escalation operate in real time.
FR-010	Behaviour controls	Member can view categories, correct assumptions, disable personalisation and reset learning.
FR-011	Quick exit	One action hides content and clears client view in under 500 ms on supported clients.
FR-012	Deletion	Profile disappears immediately; purge jobs meet published retention SLA and generate audit proof.
FR-013	Neutral notifications	Lock-screen content contains no relationship, profile or message detail.
FR-014	Reports/appeals	24/7 urgent queue; reasoned enforcement and appeal workflow.
FR-015	No commercial sex	Solicitation/payment-for-contact patterns route to review and enforcement.

8. Architecture and data

Layer	Developer specification
Clients	Next.js PWA/web; React Native iOS/Android after policy clearance; local encrypted vault; biometric/PIN lock.
Edge	WAF, bot management, API gateway, rate limiting, device attestation and privacy-safe abuse fingerprints.
Identity plane	OIDC/passkeys, verification adapter, alias service and separated identity vault.
Core services	Profile, Intent, Discretion Policy, Discovery, Match, Group, Conversation, Media, Safety, Billing, Deletion.
Data plane	PostgreSQL transactions; relationship/block graph projection; Redis; Kafka-compatible events; encrypted object storage.
AI plane	Feature registry/store, vector retrieval, ranker, LLM gateway, model registry, evaluation and experiment platform.
Operations	Trust console with redacted defaults; SIEM; audit service; privacy/deletion orchestrator; incident command.

8.1 Core entities

Entity	Essential fields
User	id, status, jurisdiction, verification_state, deletion_state
AliasProfile	user_id, display_alias, bio, approved_media, visibility, version
IntentProfile	user_id, modes[], boundaries, cadence, emotional_depth, availability
DiscretionPolicy	user_id, visibility, distance_band, reveal_rules, notification_policy, version
CoupleUnit	id, member_ids, approval_version, active_state
Introduction	subjects, candidate, score, model_version, explanation_codes, paid_label
RevealGrant	resource_id, owner, recipient, scope, expires_at, revoked_at
Room	id, participant_ids, mode, membership_version, safety_state
BehaviourEvent	pseudonymous_subject, type, context, consent_basis, ttl
ReportCase	reporter, subject, taxonomy, severity, evidence_refs, decision, appeal
DeletionJob	user_id, systems[], due_at, completion_proof

8.2 API surface

Method / endpoint	Purpose / rule
POST /v1/identity/verify	Begin age/liveness flow; returns status token, not document.
PUT /v1/discretion-policy	Versioned update; cannot silently broaden exposure.
POST /v1/intents	Create user-confirmed mode-specific intent.
POST /v1/couple-units	Invite represented participant; no spouse approval requirement unless spouse is represented in-platform.

Method / endpoint	Purpose / rule
GET /v1/discovery	Policy-filtered ranked introductions; cache result re-authorised on read.
POST /v1/introductions/{id}/decisions	Private like/pass/correction; idempotent.
POST /v1/reveal-grants	Recipient-specific, expiring access.
DELETE /v1/reveal-grants/{id}	Immediate revocation and CDN/key invalidation.
POST /v1/reports	Create protected report and evidence references.
POST /v1/preferences/reset	Delete/neutralise learned vectors and rebuild from explicit choices.
DELETE /v1/users/me	Immediate disable and deletion workflow initiation.

9. Security and privacy doctrine

TRYST processes exceptionally sensitive information: sex life, sexual orientation, relationship status, private messages, intimate media and location. Under UK GDPR, sex-life and sexual-orientation data is special category data. The controller needs an Article 6 lawful basis, an Article 9 condition, a DPIA, explicit-consent records where relied upon, strict retention, rights handling and international-transfer controls.

Control family	Mandatory implementation
Minimisation	No real name on profile; no mandatory contact upload; no home/work address; age result instead of DOB where possible.
Separation	Identity vault, profile store, media vault, analytics and billing linked through scoped pseudonymous identifiers.
Encryption	TLS 1.3; envelope encryption; per-object media keys; HSM/KMS; key rotation and access telemetry.
Authorisation	ABAC using requester, resource, reveal grant, block graph, discretion version, safety state and jurisdiction.
Private media	Short-lived signed access, recipient binding, no public CDN, malware scanning, watermark option, rapid key revocation.
Payments	Tokenised PSP; store minimum references; truthful neutral descriptor; finance staff cannot browse profiles.
Admin	Just-in-time access, dual approval for intimate content, reason codes, immutable logs and anomaly detection.
Deletion	Free, simple, verified purge across primary stores; backup expiry disclosed; legal holds narrowly scoped.
Incident readiness	Specific runbooks for outing, sextortion, intimate-media leak, insider abuse, account takeover and mass breach.
Assurance	OWASP ASVS/MASVS, threat modelling, independent penetration tests, bug bounty and annual privacy audit.

ASHLEY MADISON LESSON

The FTC reported that the 2015 Ashley Madison breach exposed information associated with 36 million users and alleged security and deletion failures. TRYST must treat privacy claims as verifiable product requirements: collect less, separate identity, delete truthfully and assume attackers will target the platform specifically because of its members' vulnerability to exposure.

10. Trust, safety and misuse controls

Risk	Controls
Fake profiles / bots	Liveness, device risk, behaviour graph, photo provenance checks, friction and human review.
Blackmail / sextortion	Payment-demand and threat detection, rapid room freeze, evidence preservation, victim resources and law-enforcement policy.
Stalking / doxxing	Coarse location, visibility controls, search throttling, network-level blocks and repeat-account detection.
Non-consensual media	Hash matching where lawful, upload warnings, recipient controls, rapid takedown and account sanctions.
Coercion in threesomes	Independent accounts, individual approval, private withdrawal, no single "couple owner."
Underage access	Age assurance, re-check triggers, content/report triage and legally required reporting.
Commercial sex / trafficking	No payments between members, solicitation detection, specialist review and jurisdictional reporting.
Bias / predatory targeting	No sensitive-trait ad targeting; exposure audits; prohibit fetishisation/harassment based on protected traits.

Risk	Controls
Retaliatory reports	Reporter credibility signals, evidence review, appeal and no automatic permanent ban from one uncorroborated report.

11. Business model

Offer	Indicative UK price	Included value
TRYST Free	£0	Verified profile, intent setup, limited introductions, mutual matches, block/report and deletion.
TRYST Black	£24.99/month or £179/year	Unlimited curated discovery, invisible mode, advanced filters, travel, preference insights and priority support.
TRYST Duo	£34.99/month or £249/year per couple	THIRD/QUAD tools, joint profile, multi-person matching and shared room coordination.
Private Signals	£6.99 / £14.99 / £29.99 packs	Labelled high-intent signals or temporary visibility; never bypass filters or consent.
Vault+	£7.99/month	Extended encrypted media vault, granular expiry and multi-device recovery.
Concierge Black	£299-£799/month	Human-curated introductions for verified members; no escorting or outcome guarantee.

REVENUE ETHIC

Do not monetise fake profiles, unreadable credit mechanics, gender discrimination, safety, deletion or access to one's data. TRYST earns more when members trust the platform enough to stay—not when it manufactures frustration.

11.1 Illustrative unit economics

Metric	12-month target	Scale target
Free-to-paid conversion	7-10%	12-16%
Blended ARPPU	£27-£34/month	£35-£45/month
Monthly paid churn	<7%	<4.5%
Gross margin	>72%	>82%
CAC payback	<8 months	<5 months
LTV:CAC	>3x	>4.5x
Verification + safety cost / MAU	£1.20-£2.20	£0.55-£1.10
Chargeback rate	<0.7%	<0.4%

11.2 Go-to-market

1. Launch a private web beta in London, then Birmingham and Manchester only after local liquidity thresholds are met.
2. Acquire through privacy-led content, discreet affiliates, relationship/lifestyle publishers and invitation codes—not explicit advertising that triggers platform rejection.
3. Seed verified, balanced cohorts; never use synthetic personas or employees posing as members.
4. Lead PR with modern privacy engineering and authenticity rather than moral provocation.
5. Create referral rewards that reveal no referrer relationship to invitees and expire automatically.
6. Expand by city only when compatible inventory, moderation coverage and safety SLAs are proven.

12. Delivery plan

Phase	Duration	Deliverables	Exit gate
0. Discovery	6 weeks	Research, legal/DPIA, naming clearance, threat model, prototypes, store review	No fatal legal/platform blocker; user value validated.
1. Privacy core	12 weeks	Identity vault, profile, intent, discretion engine, deletion, admin audit	Independent architecture/security review.
2. Closed beta	12 weeks	Discovery, chat, reveal ladder, billing, reports and moderation	High-severity response SLA and zero critical access failures.

Phase	Duration	Deliverables	Exit gate
3. AI learning	12 weeks	Retrieval/ranking, explanations, reset, experiments and model governance	Quality lift without safety/privacy regression.
4. Group modes	10 weeks	Duo profiles, THIRD/QUAD matching, quorum rooms	Independent-participant consent tests pass.
5. City launch	16 weeks	Growth, SRE, 24/7 operations, native app decision	Retention, paid conversion and liquidity gates.

12.1 Initial team

Function	Capacity
Leadership	Product Development Director, CTO/Engineering Director, Trust & Safety Director
Product/Design	Senior PM, UX researcher, 2 product designers
Engineering	Tech lead, 4 backend, 2 mobile/web, platform/SRE, QA automation, security engineer
AI/Data	ML lead, 2 ML engineers, data engineer/analyst
Trust/Operations	Policy lead, investigator lead, moderation coverage, customer operations
Legal/Privacy	DPO, product counsel, app-store counsel, external penetration/privacy specialists
Commercial	Growth lead, lifecycle marketer, partnerships/affiliate manager

13. MVP backlog and acceptance

Epic	P0 delivery	Definition of done
Quiet onboarding	Alias, private contact, age/liveness, device lock	No public/legal identity leakage; unverified users excluded.
Intent engine	Five modes, hard boundaries, availability and discretion	User confirms structured output; change history audited.
Discretion engine	Visibility, neutral notifications, location veil, invisible mode	Every read path policy-tested; stale caches cannot expose.
Discovery	Candidate retrieval, ranker baseline, explanations	Hard exclusions are deterministic; paid status cannot override.
Reveal ladder	Blurred/private media, grants, expiry and revoke	Revocation invalidates active access quickly and is measured.
Chat	Encrypted transport/storage, disappearing options, block/report	Blocked users lose all room access; evidence flow is isolated.
Safety console	Cases, severity, actions, appeals and audit	Urgent cases page on-call; admin access is least privilege.
Billing	Free, Black, Duo and signals	Clear prices/refunds; neutral truthful descriptor; safety remains free.
Deletion	Pause, vanish and full delete	Published SLA; primary-store purge verified; backups age out.
AI controls	Preference view/correct/reset/opt-out	Reset removes inferred vectors; model card and rollback exist.

14. KPIs, governance and no-go gates

Dimension	Metric
North star	Qualified reciprocal conversations per 1,000 verified weekly active members.
Match quality	Mutual interest, reply rate, private fit rating, repeat-connection rate, unwanted-contact rate.
Discretion	Exposure incidents, location inference success in red-team tests, notification leakage, deletion completion.
Authenticity	Verified ratio, bot prevalence, scam report rate, successful account-takeover rate.
Safety	Severe incidents per 10,000 rooms, response time, repeat offender, appeal overturn and victim-support SLA.
AI	Calibration, hard-boundary violation, explanation usefulness, exposure fairness, drift and privacy leakage.
Commercial	Conversion, ARPPU, churn, CAC payback, contribution margin, refunds and chargebacks.
Reliability	Crash-free sessions, p95 latency, message delivery, RPO/RTO and emergency-lock latency.

NO-GO CONDITIONS

Do not launch or scale if private media can be accessed outside an active grant, deletion claims are not technically true, exact location can be triangulated, age assurance is ineffective, couple profiles can represent unverified people, high-severity reports lack 24/7 response, or ranking violates a hard boundary.

15. Source register

Accessed 17 September 2026. Policies, prices and legal requirements must be revalidated before launch.

Source	URL
Ashley Madison official positioning	https://www.ashleymadison.com/
Gleeden official site	https://en.gleeden.com/
Victoria Milan official privacy features	https://www.victoriamilan.com/
FTC Ashley Madison settlement and breach findings	https://www.ftc.gov/news-events/news/press-releases/2016/12/operators-ashleymadisoncom-settle-ftc-state-charges-resulting-2015-data-breach-exposed-36-million
ICO special category data guidance	https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/lawful-basis/special-category-data/what-is-special-category-data/
UK GDPR Article 9	https://www.legislation.gov.uk/eur/2016/679/article/9
Apple App Review Guidelines	https://developer.apple.com/app-store/review/guidelines/
Google Play inappropriate content policy	https://support.google.com/googleplay/android-developer/answer/9878810
NIST AI Risk Management Framework	https://www.nist.gov/itl/ai-risk-management-framework
OWASP Mobile Application Security	https://mas.owasp.org/

Final directive

BUILD TRYST

Build the platform the brief demands: an adult-only, gender-inclusive, privacy-first system for discreet affairs, continuing extra-relationship connections and threesomes. The investable difference is not secrecy theatre. It is the combination of high-intent modes, authentic users, a Discretion Policy enforced at every data access, multi-person matching, transparent behavioural learning and security designed for the consequences of exposure.

Non-negotiable	Launch standard
Market clarity	Partnered adults, affairs and threesomes are explicit product categories—not hidden behind generic dating language.
Discretion	Every exposure path is controlled by policy, tested adversarially and described without false guarantees.
Intelligence	AI learns attraction, intent, timing and communication fit while users can inspect, correct, reset or disable learning.
Authenticity	Age, liveness and anti-fraud controls make the network credible without publishing legal identity.
Economics	Subscriptions fund core value; optional signals add margin without manufacturing frustration or overriding consent.
Readiness	No public launch until privacy, deletion, media access, safety response and city liquidity gates are proven.

08 — Source A: traceability (consolidated Spec + ToS, 4 Oct 2026)

Part of the TRYST specification baseline.

Source A is the top-precedence source (00 §2). Its full text is in [source/TRYST_Spec_and_ToS_Consolidated_2026-10-04.pdf](#). This table shows where each section lives in the specification and in the code scaffold, and what v1.2 changed.

Source A section	Specification	Code scaffold	v1.2 changes
1. Thesis & market gap; operating stance	01 §2, 01 §4.1	—	—
2. The Name; brand codes	01 §3	assets/brand/, web/	"Oxblood and bone" → logo palette, oxblood + gold, cream for bone (D-19)
3. Segments & jobs (S1-S6)	01 §5.1	agents/.../broker/entities.py	Unattached solos: D-01
4. Product loops	01 §6	—	—
5. Agentic architecture (7 agents)	01 §7, 03 §3	agents/.../envoy/ (schema-strict, no free text)	Herald SEO agent outside the member platform (D-21)
6. Behaviour learning	03 §4	agents/.../mirror/vectors.py ($\alpha \in [0.15, 0.85]$, decay)	Decay FR-068
7. Matching & ranking; exposure fairness	03 §5	agents/.../broker/gates.py (property-tested), slate.py; backend/internal/keys (5/15 Keys, refund on reply, inbound cap)	—
8. Trust, verification, consent	02 §3, 03 §7.3	backend/internal/gate (V2 gate), backend/internal/couple (VC co-sign)	V2 required for intents and threads , not V1 (E-01)
9. Discretion architecture; threat model; crypto; ExclusionRing	03 §6, 04 §5	backend/internal/erasure (destroy key → undecryptable), crypto-core (per-media key wrap and revoke)	Two-factor login (D-16); emergency lock FR-064
10. Data model	02 §4	backend/migrations/ (generated from 02 §4)	Anchor candidates, authenticators, couple invite, discretion policy, reveal grants, threads, safety tables
11. API surface	02 §5	contracts/openapi/tryst.v1.yaml	/v1/panic → /v1/burn + /v1/safety/panic (E-21); handshake system-only (E-29); login, emergency-lock and Stripe webhook endpoints
12. Stack & infrastructure	03 §2, 04 §3	Go, Python, Rust, Next.js as specified; native apps documented in mobile/	Modular monolith for core services (D-09)
13. MLOps	03 §9	Property and invariant tests in CI	—
14. Compliance	01 §14	backend/internal/gate (jurisdiction gate)	Stripe as primary processor, conditional (D-22); Article 6 bases
15. Business model	01 §10	backend/internal/keys, backend/internal/payments	CAC threshold corrected (E-03); GHOST scope (D-05)
16. Liquidity & GTM	01 §11	agents/.../broker/slate.py (caps)	Gender-based seeding: D-06; organic search (Herald)
17. Roadmap, team, budget	01 §12	—	Insurance and biometric-check costs added
18. KPIs & risk register	01 §13	—	G-NG-8 to G-NG-10; new risks
19. Admission control	01 §8, 03 §8	backend/internal/admission (anchors, silent refusal, human-only bans, appeal lifts)	Anchor candidates at V2 (E-02); biometric anchor off (D-08)
20. Liability & disclaimer architecture	01 §15	web/ safety copy	—

Source A section	Specification	Code scaffold	v1.2 changes
21. Guardian classifier spec	03 §7	Reason codes in the OpenAPI contract; on-device only	—
Part B — Terms of Service draft 0.9	05 (draft 0.9.3)	—	cl. 3.5, 6.5, 15.2–15.3, 17.3, protected-trait rule; Q11–Q15

09 — TRYST Operating System (internal engineering) v1.3

Part of the **TRYST v1.3 baseline**. Written in response to the founder's *Master AI Operating System Architect* brief (4 Oct 2026), applied to TRYST. Decisions **D-25-D-32**; requirements **FR-080-FR-088**.

How to read this document. The brief asks for an autonomous, self-improving AI infrastructure OS with a full agent workforce, a BitriPay payment door, a connector ecosystem and a commercial engine. **This document keeps every item in the brief.** Each item is either built as asked, adapted so it fits TRYST's product principles (**01 §4**), or recorded with a stated reason it does not apply to a members-only discreet platform. The full register is in **§0**. Nothing in the brief has been dropped silently.

The one rule that shapes everything here: *the AI runs the machine; people make the decisions that cannot be undone*. Agents observe, predict, recommend, draft and carry out reversible or pre-approved actions on their own. A permanent action against a member, money movement, a production code change and any contact made on a member's behalf always need a named human. This is enforced in code by the AI Governance gate (backend/internal/aigov, tested), not by policy text (**§5.4**).

0. Brief coverage register

Status: **Built** = in the codebase and tested. **Spec** = fully specified here and scheduled. **Adapted** = delivered in a changed form, with the reason. **Not for members** = applies to operators or partners only, never to member data. **Declined** = conflicts with a v2.0-locked principle; recorded with the reason and a stated alternative.

Brief item	Status	Where / why
Forensic review: modules, workflows, journeys, API dependencies, data flows, monetisation, automation	Spec	§2.3 platform weakness review (technical, commercial, scalability, security, operational, AI gaps), each with a fix
"Do not remove any existing functionality"	Built	No v1.2 requirement removed; FR-001-FR-079 unchanged
Proven patterns only, from the named companies	Spec	Every pattern below cites the company or standard it comes from (§9.1)
Learn, adapt, predict, optimise, automate, secure, scale, self-heal	Spec	§5.6 self-managing layer; §9 architecture; Mirror learning loop (03 §4)
"...without requiring human intervention"	Adapted	Autonomy levels L0-L3 (§5.2). Reversible and runbook actions run without a person; four action classes always need one. UK GDPR Art 22 and principle P4 make full autonomy over members unlawful and unsafe (D-25)
AI Command Centre for every user	Spec	§4: one per user type, including a private Member Command Centre
Chief of Staff, Analyst, Research, Automation, Growth, Security, Knowledge agents in each centre	Adapted	§4.1 maps all seven for members (Growth = "my chances", not acquisition); full set for operators (§4.2-§4.9)
Executive agents (CEO, COO, CFO, CTO, CMO, CRO)	Adapted	Analyst copilots for the named executive; no authority of their own (D-31)
Product, Engineering, Quality, Cybersecurity, Revenue, Customer, Compliance agents	Spec	§5.3 workforce register; §5.5 core agents
Self-managing layer: System Health, Bug Detection, Auto-Repair, Infra Optimisation, Release Management, AI Governance	Built (governance) · Spec (rest)	§5.6. Auto-Repair = runbooks + rollback; code fixes go through a reviewed pull request (D-28)
Cybersecurity Command Centre (zero trust, identity, threat, fraud, anti-hacking, data protection)	Built (identity) · Spec	§13; login factors already built (02 §3.3)
Data intelligence layer (lake, warehouse, vector DB, knowledge graph, streaming, analytics, predictive, behavioural, recommendation, decision engines)	Spec	§9.6. All ten built on the existing stack; behavioural data follows Mirror's privacy rules

Brief item	Status	Where / why
BitriPay Integration Gateway (merchant portal, keys, webhooks, sandbox, production, QR, wallet, card, bank, mobile money, SDKs, docs, testing, fees, revenue share, settlement)	Built (adapter, sandbox) · Spec	\$7. Live use is gated by market (D-26)
Connector ecosystem (payments, banking, identity, comms, CRM, cloud, AI, productivity, storage)	Spec	\$8.2 connector register, one row per provider, with data sent and received
WhatsApp, CRM for member data, data enrichment, maps tracking, logistics	Declined for member data	\$8.2 states why for each (D-29); operator-only uses kept
Revenue engine (subscription, transaction, marketplace, AI usage, API, enterprise, licensing, white label)	Spec / Adapted	\$12. Member revenue is the v1.2 tier set; B2B lines are licensing of the discretion stack, never member data
Dynamic pricing engine	Adapted	Experiments on list price per market, disclosed and identical for every member in a market; no personalised pricing (D-30)
AI revenue optimisation, CLV, churn prevention, upsell, cross-sell engines	Adapted	\$12.5. Optimise for QCAM-weighted retention; no dark patterns; one-step cancel stays (ToS cl. 16)
AI credit / ACU system	Adapted	ACU meters B2B and partner API usage only. Members never buy ACU; Keys already serve that role (D-27)
Data intelligence revenue	Declined for member data	P8 forbids data sale. Alternative: licensing of models and privacy technology, and free published aggregate research (\$12.4)
PRD, TRD, system, infra, AI, agent, security, DB, API, event, DevOps, monitoring, DR, BC, data governance, compliance, scalability, commercial architecture	Spec	\$9-\$13; index in \$17
Complete ERD, API specs, user journeys, agent workflows, DB schemas, security controls, deployment, testing, production-readiness review	Spec	\$10 (ERD + DDL), \$11 (API + webhooks), \$3.2 journeys, \$5.5 workflows, \$13 controls, \$9.9 deployment, \$9.10 testing, \$15.6 PRR
17 required sections	Spec	\$1-\$17 below

1. Executive product vision

1.1 What the OS is

The TRYST operating system is the operating layer that runs TRYST: **a governed workforce of specialised AI agents, the data and event fabric they work on, and the command centres through which members, operators, partners and regulators see and steer them.** It turns TRYST from an app with some models into infrastructure that runs its own operations. It matches, protects, bills, supports, secures, monitors and repairs itself, and puts people in charge of every decision that cannot be undone.

It is built for one product first. Discreet, verified, consent-first connection for partnered adults and couples is the hardest trust problem in consumer software. A system that can run that safely can be licensed to others (§12.3).

1.2 What problem it solves

For	Problem today	What the OS does
Members (S1–S6)	Volume without outcome; exposure risk; harassment; fakes; endless chat that never becomes a meeting	Agents pre-qualify, negotiate and protect; the member decides. Measured by QCAM, not time in app
Couples (S3, S5)	Coordinating two people's consent with a third or another couple collapses	Envoy negotiates over a fixed schema; joint-veto modes; dual-consent controls
TRYST operations	A 17-person team cannot run 24/7 trust and safety, payments, SRE and growth by hand at 150k verified members	Agents carry the volume; people handle the exceptions (§5.6, §14)
Regulators and auditors	Opaque recommendation systems; unverifiable safety claims	Every agent decision is logged in a hash-chained audit trail; transparency reporting is generated from it (§4.10)
Partners and licensees	Discretion, verification and consent tooling is expensive to build	The same governed stack, licensed with no member data attached (§12.3)

1.3 Why it is different

1. **Resolution over engagement (P1).** Every agent's reward is tied to QCAM, the confirmed and positively rated meets per active member per month. Competitors optimise for time in app.
2. **Agents negotiate; humans decide (P4).** The double opt-in agent handshake (Loop C) is the product. No competitor runs agent-to-agent pre-negotiation for one-to-one, threesome or couple-to-couple matching.
3. **Governance is code.** The always-human classes, the kill switch, autonomy budgets and the tamper-evident audit log are in aigov and covered by tests.
4. **Breach-assumed privacy (P3).** E2EE threads, the identity/profile split, cryptographic erasure in under 60 s and self-hosted models for anything touching special-category data.
5. **Payment-processor independence.** One PaymentProvider interface across Stripe, the high-risk acquirer and BitriPay. Entitlements never depend on one processor.

1.4 Why the market needs it and why it can win

Category leaders built on volume and on paid messaging, and they are held back by fake-profile and breach history (01 \$2). The demand is proven (attributed, unaudited: Ashley Madison 80M+/91M+, E-32). The unmet need is **trust that is engineered, not claimed**. TRYST can win because the OS makes the trustworthy behaviour the cheapest to run:

- automation carries trust and safety and support volume, so gross margin survives verification costs;
- Envoy raises meets per intent (target +40% at G-P2), and that outcome is what members pay for;
- the governed agent stack becomes a licensable B2B asset with no member data in it.

2. Market gap deep review

2.1 Competitors

Platform type	Does well	Fails to solve	Where users are underserved	Where businesses lose money	Missing automation	How the OS fills it
Affair sites (Ashley Madison, Gleeden, Victoria Milan)	Category awareness; explicit positioning	Fake and inactive profiles; breach legacy; pay-per-message mechanics	Women and couples face harassment volume; no real couple model	Chargebacks, refunds and trust loss from fake profiles; ads banned	No pre-qualification; manual moderation	V2 gate (E-01), Envoy handshake, Guardian on device, Keys refunded on reply, no fakes (P8)
Mainstream dating (Tinder, Hinge, Bumble)	Scale; polished UX; ranking ML	Will not serve the partnered; engagement-maximising loops	Couples, thirds and the discreet are pushed off-platform	Paywall churn; moderation cost per MAU	Agents limited to profile tips	Purpose-built segments S1-S6 and modes; agents tied to QCAM
Kink and couples platforms (Feeld and similar)	Inclusive structures; couple profiles	Weak verification; little discretion engineering	S4 harassment; discretion relies on member habits	Trust and safety cost; store risk	No negotiation layer; no exposure-risk model	ExclusionRing PSI, Curtain, S4 inbound caps, VC co-sign
Swinger directories and forums	Community; events	Unverified; open directories; scraping	Discretion; safety	Little monetisation beyond membership	None	Verified, private, agent-mediated
General AI agent platforms (OpenAI Agents, Microsoft Copilot Studio, Salesforce Agentforce, ServiceNow)	Orchestration tooling; enterprise governance	Not built for special-category data or member-to-member consent	Not applicable to members	High per-seat cost for SMEs	Domain policy is left to the customer	TRYST packages domain policy (consent, discretion, verification) with orchestration

2.2 The four structural gaps (from 01 §2.2, extended)

Gap	Evidence	OS answer	Measured by
Outcome gap: volume without meets	Reply rates, ghosting, message volume as a KPI	Broker + Envoy + Aftercare reward loop	QCAM; meets per intent
Trust gap: fakes, bots, sextortion	Category history; OSA duties	V2 gate, admission L1-L5, Guardian M6a-e, Fraud agent	Report rate per 1k threads; fake-profile prevalence
Discretion gap: exposure by collision, notification, device	Members' primary anxiety (S1)	Curtain, ExclusionRing, decoy/duress/Burn, emergency lock	Exposure incidents; leak tests
Operations gap: a small team cannot run 24/7 trust, payments and SRE	Team of 17 to P3 (01 §12)	\$5.6 self-managing layer, \$5.5 core agents, \$14 control centre	Tickets per 1k MAU handled by agents; MTTR; on-call pages per week

2.3 Platform weakness review (v1.2 baseline → fix)

The brief asks for a forensic review of TRYST itself. Findings against the v1.2 codebase and spec:

Area	Weakness found	Fix (this document)	Phase
Technical	Identity store is in memory; DPoP not yet bound (D-24)	PostgreSQL IDENTITY-DB and DPoP are P1 items; OS adds no dependency on them before P1	P1
Technical	No central registry of which agent may do what	aigov registry and gate (built, FR-080)	P1
Technical	Agent actions are not uniformly audited	Hash-chained audit log for every agent decision (built, FR-082)	P1
Commercial	One processor family; Stripe approval uncertain (D-22)	PaymentProvider with three adapters: Stripe, high-risk acquirer, BitriPay (built, sandbox)	P1
Commercial	Revenue limited to B2C	B2B licensing of the discretion and verification stack, metered in ACU (§12.3)	P4+
Scalability	Modular monolith (D-09) with no stated split triggers	Split triggers in §9.8	P2+
Security	No automated containment for agent misbehaviour	Kill switch per agent and global (built)	P1
Security	Connectors could leak member data to third parties	Connector register with a data map and a no-member-data rule per provider (FR-087)	P1
Operational	On-call relies on people reading dashboards	System Health → Auto-Repair runbooks with SLO-burn rollback (FR-085)	P1
Operational	No single admin view across safety, payments, agents and health	Admin Super Control Centre (§14, FR-086)	P1
AI capability	No member-facing explanation surface for agent decisions	Member Command Centre "why" panels from reason codes (FR-083)	P2
AI capability	Agent evals exist per agent but not for cross-agent workflows	Workflow-level evals and replay in the release gate (§9.10)	P2

3. Complete user ecosystem

3.1 User types

Every user type in the brief is listed. A type that does not exist in TRYST is marked as such and kept for the licensing model.

#	User type (brief)	In TRYST	Examples	Authenticates with	Sees member data?
U1	Customers	Members (S1-S6, solo and couple)	Attached solo; couple seeking a third	Two biometric factors (mobile) / two factors (desktop) (02 §3.3)	Their own, and what counterparties reveal
U2	Businesses	Licensees (P4+)	A dating platform licensing the verification and discretion stack	OIDC SSO + hardware key; API keys	Never TRYST's members

#	User type (brief)	In TRYST	Examples	Authenticates with	Sees member data?
U3	Admins	Platform admins	Head of T&S, Principal Engineer, DPO	Hardware security key + device posture + just-in-time role	Redacted by default; break-glass only
U4	Operators	T&S moderators, support agents, SRE on-call, payments ops	Moderator on the trust console	As admins, with role scope	Case-scoped and redacted
U5	Partners	Assurance, payment and content partners	Age-assurance providers; voucher resellers; digital-PR partners	Mutual TLS + signed webhooks	Only what the contract and data map allow
U6	Agents	AI agents (§5)	Envoy, Guardian, Auto-Repair	SPIFFE workload identity + manifest	Only their manifest's scopes
U7	Developers	TRYST engineers; licensee developers	Backend engineer; licensee integrator	SSO + hardware key; sandbox keys	Never in production without break-glass
U8	Merchants	Not in TRYST's member product. Kept for BitriPay connected accounts in the licensing model (§7.4)	Voucher reseller receiving a commission	BitriPay account + TRYST partner portal	None
U9	Service providers	Processors and sub-processors	Cloud, liveness, CDN, email	Contract + DPA	Per the data map (§8.2)
U10	Regulators	Ofcom, ICO, DPC (IE), card schemes, auditors	Ofcom information request	Read-only regulator portal; time-boxed	Aggregates and specific lawful disclosures only
U11	Third-party API partners	Connector providers	Stripe, BitriPay, Persona	Their own auth; our keys in the vault	Per the data map

3.2 User journeys

The member journeys are Loops A-E (01 §6) and the screens in 04 §4. The operating system adds these journeys:

Journey	Steps	Agents	Human checkpoint
J1 Member sees why	Opens Command Centre → "Why this slate?" → reason codes rendered as plain sentences → corrects a preference	Chief of Staff (member), Mirror, Broker	Member edits; P5 floor holds
J2 Member safety brief	Before a meet: Safe Meet checklist → venue sanity (public, open) → check-in timer → aftercare at 48 h	Guardian, Aftercare	Member sets the check-in; Panic is always available
J3 Operator case	Report → Guardian triage score → case queue with redacted evidence → moderator decision → appeal path	Guardian, Compliance	Moderator decides any removal (always-human)
J4 Incident self-heal	SLO burn alert → System Health opens incident → Auto-Repair runs rollback runbook → on-call is told, not woken, if the rollback resolves it	System Health, Auto-Repair, Release Manager	Human reviews the post-incident report; any code fix is a reviewed PR
J5 Payment exception	Webhook says dispute → Payment agent assembles evidence → Fraud agent scores → payments ops approves refund or contest	Payment, Fraud	Refund is money movement (always-human), unless covered by the pre-approved auto-refund rule (§5.5, Payment)
J6 Licensee onboarding	Partner applies → KYB → sandbox keys → integration tests → contract → live keys	Onboarding (B2B), Compliance, API Integration	Live keys need a named approver
J7 Regulator request	Request logged → Compliance agent assembles the lawful minimum from audit and transparency data → DPO approves → disclosure	Compliance	DPO approves every disclosure
J8 Content publish	Herald drafts → editor approves → publish → rank tracking	Herald (Marketing)	Editor approves every publish

4. AI-agent command centres

Each command centre is a role-scoped surface over the same agent fabric. **The data a centre sees is decided by the caller's role and the agents' manifests, never by the UI.** All actions go through the governance gate.

4.1 Member Command Centre (U1)

Design constraints that override the brief's defaults: private by default; nothing is shown that would help someone holding the member's phone (P3); no engagement counters, streaks or "you're missing out" nudges (P1); the Leave control and decoy behaviour work on every screen; agents never act as the member (P4). Built on the existing /v1/me and /v1/me/security routes; web route /account/insights (FR-083).

Brief agent	Member version	What it sees	Actions it can perform	What it may decide	Automations the member controls
Personal AI Chief of Staff	Concierge brief: one calm daily summary: new Briefs waiting, a meet to confirm, renewal date	Own Briefs, intents, meets, entitlement	Draft a reply (on device, D-20), set a check-in timer, remind	Nothing; it recommends	Summary time; quiet hours; off switch
AI Analyst	My patterns: charts of own activity (intents sent, mutual rate, meets, aftercare trend), security activity	Own aggregates only	Explain a chart	—	Personalisation on/off (FR-024)
AI Research Agent	Explainers: verification, discretion and safety guides; local laws for travel mode	Public content	—	—	—
AI Automation Agent	ENVOY tier proposals: Envoy pre-qualifies and drafts outreach; each needs one tap	Envoy proposals with rationale codes	Approve / discard	Never sends (D-04)	Cadence; filters
AI Growth Agent	My chances: which settings narrow the slate (radius, tier floor, modes) and the expected effect	Own settings; aggregate liquidity bands	Suggest a setting change	—	Accept or ignore
AI Security Agent	Security activity: sign-ins, factors, new devices, locks (built: /v1/me/security); session risk from Curtain	Own auth events (90 d)	Sign out everywhere; emergency lock; revoke device	Curtain may only tighten (02 §6.4)	Lock behaviour; decoy
AI Knowledge Agent	My memory: what TRYST has learned about me (stated vs revealed, α weight) with reset	Own IntentVector, RevealedPreferenceVector summary, version history (FR-079)	Correct, reset, export (DSR)	—	Reset learning

4.2 Couple Command Centre (U1, VC)

Shared panel, visible to both partners, plus each partner's private panel. It shows the joint-veto state, pending thirds or couples, dual-consent requests and Envoy screening for a third. **Neither partner sees the other's private panel, private fields or security activity.** Actions are a co-sign (both partners, two devices; same_device_cosign is refused) and dissolve (either partner, immediate, 02 §9).

4.3 Operator Command Centre: Trust and Safety (U4)

Element	Specification
Sees	Case queue ordered by Guardian severity and SLA; redacted evidence; member-disclosed content only; prior enforcement; appeal history
Agents	Guardian (triage, holds), Compliance (SLA and statement of reasons), Support (drafts), Risk (ban-anchor matches)
Actions	Hold (reversible, agent may do); warn, restrict, remove, ban, report to authorities (always-human); request more evidence
Automations	SLA routing; duplicate merge; known-hash auto-hold pending review

Element	Specification
Decides or recommends	Guardian recommends with a confidence and reason code; the moderator decides; a second moderator reviews permanent bans (03 §7.3)

4.4 Operator Command Centre: Support (U4)

Sees ticket, entitlement and Keys-ledger state for the ticket's member only (case-scoped). The Support agent drafts and the person sends. Goodwill Keys and refunds are money movement (always-human), except for refunds covered by a published automatic rule (Keys refunded on reply; 14-day statutory refunds, ToS 16.7), which run as ledger operations.

4.5 Operator Command Centre: SRE and engineering (U4, U7)

Sees SLOs, error budgets, traces, deploys, cost and the agent decision log. Agents: System Health, Bug Detection, Auto-Repair, Release Manager, Infra Optimiser, Vulnerability. Automations: canary analysis, automatic rollback on SLO burn, scale-out, dependency PRs. A **production change** (promote, patch, rightsizing apply) is always-human.

4.6 Operator Command Centre: Payments and finance (U4)

Sees processor health per adapter, authorisation rates, chargeback ratio (alert 0.5%, hard stop 0.65%), disputes, settlement and VAT. Agents: Payment, Fraud, Revenue Analyst, Pricing. Actions: switch the primary processor (always-human), approve refunds and dispute responses, publish a price experiment (always-human, §12.2).

4.7 Executive Command Centre (U3)

The brief's CEO, COO, CFO, CTO, CMO and CRO agents are **analyst copilots** for the named executive (D-31). They read warehouse aggregates (never member-level data) and produce: the weekly QCAM and KPI pack (01 §13.1), phase-gate readiness (G-P0 to G-P4), no-go monitor (01 §13.2), unit economics against E-03 thresholds, budget burn and risk register deltas. They hold no actions beyond `recommend`.

4.8 Partner and Licensee Command Centre (U2, U5, U8)

API keys (sandbox and live), webhooks with delivery log and replay, usage in ACU, invoices, SLA status, integration test results, data-processing terms. Connected BitriPay accounts (for commission-earning partners such as voucher resellers) show settlement statements from BitriPay directly. **No member data is ever visible here.**

4.9 Developer Command Centre (U7)

OpenAPI reference (contracts/openapi/tryst.v1.yaml), sandbox keys, webhook tester, the BitriPay sandbox magic numbers (§7.6), SDK downloads, changelog, status. Same surface for internal engineers and licensee developers, with scopes.

4.10 Regulator Portal (U10)

Read-only, time-boxed accounts. Shows the transparency report (OSA), the aggregate enforcement statistics, the agent register (which agents exist, what they may do, autonomy level and owner) and the audit-chain verification result. Specific disclosures are made only through J7 with DPO approval.

4.11 Agent Command Centre (U6)

Agents also have a "centre": the registry entry (manifest), budget use today, kill state, recent decisions and eval scores. It is rendered inside the Admin Super Control Centre (§14).

5. Core AI agents

5.1 Agent model

Every agent is a separately deployable service (03 §2.2) declared by a **manifest** in the agent registry. The manifest is the contract; anything not in it is denied.

Manifest field	Meaning	Enforced by
id, owner	Agent name and accountable human team	Registry refuses a manifest without an owner
level	Autonomy level L0-L3 (§5.2)	Gate

Manifest field	Meaning	Enforced by
actions	Action name → effect class (read, recommend, reversible, runbook, irreversible, member_safety, money_movement, production_change)	Gate; unlisted action → deny
data_scopes	Stores and datasets it may read	Registry: a member-data scope (IDENTITY-, PROFILE-, SAFETY-, BAN-DB, event bus, feature store, vector store) is refused unless the agent runs inside the member platform and on a self-hosted model
daily_budget	Maximum autonomous reversible or runbook actions per UTC day	Gate; above budget → needs a human
model_version, policy_file, eval_suite	What it runs and how it is tested	Release gate (03 §9)
act_as_member	Not a valid class	Registry rejects any manifest that contains it (P4)

Agents talk over typed Protobuf messages on the bus (agent.* topics), never free text (A §5), and call each other with gRPC + mTLS under SPIFFE identities.

5.2 Autonomy levels

This is how the brief's "without requiring human intervention" is delivered safely. Pattern: the levels of driving automation (SAE J3016, as used by Tesla and others), Google SRE's automation ladder, and Microsoft and Anthropic's published human-in-the-loop guidance for agents.

Level	Name	Runs without a person	Typical agents
L0	Observe	Read, report	Monitors in shadow mode
L1	Recommend	Read, recommend, draft	Cartographer, Envoy, Support, Chief of Staff, executive copilots, Infra Optimiser
L2	Reversible	+ reversible changes inside a daily budget	Mirror, Broker, Curtain, Guardian (holds), Aftercare, Fraud (holds), System Health, Herald (link repair)
L3	Runbook	+ pre-approved runbooks (rollback, restart, scale out, halt canary)	Auto-Repair, Release Manager

Always-human classes (no level unlocks them): irreversible (delete, publish, send to an outside party), member_safety (freeze, restrict, remove, ban), money_movement (refund, payout, price change, credit), production_change (code, model or prompt change reaching production). **Never-allowed class:** act_as_member.

5.3 Workforce register

All agents named in the brief, mapped to TRYST. "Plane" = **M** member platform (self-hosted models, member data allowed by manifest) or **O** operations plane (no member data). The first 23 rows are loaded from aigov.Workforce() and pinned by tests; the rest are specified for the phases shown.

Brief category	Brief agent	TRYST agent	Plane	Level	Phase	Notes
Matching (01 §7)	—	Cartographer	M	L1	P1	Member confirms every field
Matching	—	Mirror	M	L2	P2	Sole writer of RevealedPreferenceVector
Matching	—	Broker	M	L2	P1	Stage 0 gates cannot be bypassed
Matching	—	Envoy	M	L1	P2	Never sends; advisory Briefs
Matching	—	Curtain	M	L2	P1	Tighten only
Matching	—	Guardian	M	L2	P1	Holds only; removals are human
Matching	—	Aftercare	M	L2	P1	Private outcome only
Core (brief §5)	Onboarding Agent	Onboarding	M	L2	P1	\$5.5
Core	Risk Agent	Risk	M	L2	P1	\$5.5
Core	Fraud Detection Agent	Fraud	M	L2	P1	\$5.5
Core / Customer	Customer Support Agent	Support	M	L1	P1	\$5.5

Brief category	Brief agent	TRYST agent	Plane	Level	Phase	Notes
Command centre	Personal AI Chief of Staff	Chief of Staff	M	L1	P2	\$4.1; on device where possible
Core / Compliance	Compliance Agent, GDPR Agent	Compliance	M	L2	P1	\$5.5
Core	Payment Agent	Payment	M	L2	P1	\$5.5
Core / Marketing	Marketing Agent (Herald)	Herald	O	L2	P1	Isolated (FR-075)
Revenue / Core	Revenue Agent, CFO Agent	Revenue Analyst	O	L1	P1	Aggregates only
Self-managing	System Health Agent	System Health	O	L2	P1	\$5.6
Self-managing / Quality	Bug Detection Agent	Bug Detection	O	L2	P1	\$5.6
Self-managing	Auto-Repair Agent	Auto-Repair	O	L3	P1	\$5.6
Self-managing	Infrastructure Optimisation Agent	Infra Optimiser	O	L1	P2	\$5.6
Self-managing	Release Management Agent	Release Manager	O	L3	P1	\$5.6
Cybersecurity	Threat Hunter Agent	Threat Hunter	O	L2	P1	\$13
Cybersecurity	Vulnerability Agent	Vulnerability	O	L2	P1	\$13
Self-managing	AI Governance Agent	The gate itself + Governance reviewer	O	—	P1	Deterministic code, not a model (\$5.4)
Cybersecurity	SOC Agent	SOC triage	O	L2	P1	Alert dedupe and enrichment; containment runbooks
Cybersecurity	Identity Agent	Identity risk	M	L2	P1	Login risk, step-up triggers; never bypasses two factors
Cybersecurity	Fraud Agent	= Fraud	M	L2	P1	—
Core	Pricing Agent / Revenue: Pricing	Pricing	O	L1	P3	Market-level experiments only (\$12.2)
Revenue	Monetisation Agent	Monetisation	O	L1	P3	Proposes packaging; never dark patterns
Revenue	Sales Agent	B2B Sales	O	L1	P4	Licensee pipeline; drafts only
Customer	Success Agent	Licensee Success	O	L1	P4	B2B only
Customer	Retention Agent	Retention	M	L1	P3	\$12.5; no retention flow on cancel
Compliance	AML Agent	AML (payments)	O	L1	P3	Monitors processor alerts; AML duties sit with the regulated PSP (D-32)
Compliance	KYC Agent	= Onboarding (age assurance) + KYB for licensees	M / O	L2	P1 / P4	—
Compliance	Regulatory Agent	Regulatory	O	L1	P1	OSA codes, ICO guidance, store policies; change alerts
Core	Data Intelligence Agent	Data Intelligence	O	L1	P2	Warehouse aggregates; DP noise on member-derived metrics
Core	Operations Agent / COO Agent	Operations	O	L1	P1	Queue health, SLA, staffing forecast
Core	API Integration Agent	API Integration	O	L2	P1	Connector health, key rotation reminders, contract tests
Core	Workflow Automation Agent	Orchestrator	O / M	L2	P1	Runs declared workflows (\$5.7)
Core	Predictive Growth Agent	Predictive Growth	O	L1	P2	Liquidity forecasts per city and segment
Core	Admin Control Agent	Admin Control	O	L1	P1	\$14; summarises and proposes, never grants access
Executive	CEO, COO, CFO, CTO, CMO, CRO Agents	Executive copilots	O	L1	P1	\$4.7; D-31

Brief category	Brief agent	TRYST agent	Plane	Level	Phase	Notes
Product	Product Architect, UX, Journey, Feature Agents	Product copilots	O	L1	P2	Funnel and journey analysis on aggregates; design-review checklists (04 §7)
Engineering	Frontend, Backend, Infrastructure, API, Database Agents	Engineering copilots	O	L1	P1	Coding assistants that open PRs; every PR is human-reviewed and passes CI
Quality	QA, Testing, Performance Agents	Quality copilots	O	L2	P1	Generate tests and load profiles; run them in CI; cannot merge

5.4 AI governance gate (implemented)

Code: backend/internal/aigov (Go, standard library only). Tests: aigov_test.go, 9 tests, including a mutation check that removing member_safety from the always-human set fails the suite.

Behaviour	Implementation	Test
Unknown agent, unknown action → deny	decideLocked	TestLevelCapsAutonomy
Always-human classes → needs_human unless a registered named human approved	alwaysHuman set; approver allow-list	TestAlwaysHumanClassesNeedAPerson
An agent cannot approve its own action	Approver must be in the human allow-list	TestUnknownApproverIsDenied
Autonomy level caps classes	minLevel per class	TestLevelCapsAutonomy
Daily budget, then a person	Per-agent per-UTC-day counter	TestBudgetThenHuman
Kill switch per agent and global; global overrides approvals	Kill, KillAll, Resume (each recorded with who did it)	TestKillSwitch
Member-data scopes only for self-hosted member-plane agents; act_as_member rejected	Register	TestManifestRules, TestOpsAgentsHoldNoMemberScopes
Hash-chained, append-only audit log; any edit or deletion detected	Entry.digest, VerifyChain	TestAuditChain

Production wiring (P1): the gate runs as a sidecar library in each agent and as a central agent-registry service. The audit chain is written to an append-only table (§10, ops.agent_decision) and anchored hourly to an object-lock (WORM) bucket, so even a database administrator cannot rewrite history unnoticed. Pattern: Certificate Transparency and AWS QLDB-style ledgers.

5.5 Core agents (the brief's fifteen)

Each entry gives purpose, inputs, outputs, permissions, triggers, workflow, escalation, APIs used and business value. "Gate" means the action passes through §5.4.

Onboarding Agent

Purpose	Get a new member from sign-up to V2 with the fewest drop-offs, without ever weakening verification
Inputs	Onboarding state, provider responses (age assurance, liveness), device type, jurisdiction
Outputs	Next-step routing, plain-language explanations, retry guidance, provider fallback choice
Permissions	L2; identity_db (state only), profile_db (onboarding fields); actions route_verification_step (reversible), explain_step (recommend)
Triggers	auth.verified, assurance.failed, onboarding idle > 10 min
Workflow	1 read state → 2 choose next step (passkey, B2 enrolment, age assurance, Cartographer) → 3 if provider A fails or is down, route to provider B (03 §2.1) → 4 explain without revealing why a check failed (anti-gaming)
Escalation	Two provider failures → manual review queue; suspected minor → Guardian, onboarding stopped (always-human)
APIs	/v1/auth/*, /v1/assurance/*, provider adapters
Value	V2 completion 52% → 68% target (01 §13.1); fewer support tickets

Compliance Agent

Purpose	Keep TRYST provably compliant: DSRs, retention, OSA duties, consent records
Inputs	DSR requests, retention policies (02 §7), consent ledger, enforcement records, regulatory feed

Outputs	DSR exports, retention violations, statements of reasons, transparency report drafts, regulator packs
Permissions	L2; run_dsr_export (reversible), check_retention (read), erase_account (irreversible → human, except member-initiated erasure, which is the member's own action)
Triggers	dsr.requested, nightly retention sweep, enforcement.decided, regulator request
Workflow	DSR: verify requester via step-up → assemble export across stores → member downloads (30-day statutory clock tracked). Retention: diff actual vs policy → open a ticket per violation → Privacy Auditor blocks release if critical
Escalation	Any regulator request, any breach indicator → DPO within 1 h; 72 h ICO clock starts on awareness
APIs	/v1/me/export, /v1/me DELETE, ops.deletion_job, safety-svc
Value	Avoids fines; makes the transparency report a by-product

Risk Agent

Purpose	Score sign-up and account risk (ban evasion, fake profiles, coordinated abuse)
Inputs	Admission L1-L5 signals (03 §8), device and payment fingerprints (HMAC only), graph features
Outputs	Risk score with reason codes; step-up requirement; ban-anchor match proposals
Permissions	L2; score_signup (recommend), require_step_up (reversible), ban_anchor_match (member_safety → human)
Triggers	auth.verified, assurance.passed, graph anomaly
Workflow	Score → if medium, require step-up → if a ban-anchor match, open a T&S case with the evidence
Escalation	Every anchor match → moderator; false-positive appeals go to a second moderator
APIs	identity-svc admission, safety-svc
Value	Keeps fakes out (the category's main failure)

Revenue Agent

Purpose	Explain revenue movement and forecast it
Inputs	Warehouse aggregates: entitlements, Keys ledger, processor settlements, churn
Outputs	Daily revenue pack, cohort LTV, forecast with intervals, anomaly alerts
Permissions	L1; warehouse_aggregates only
Triggers	Daily 06:00 UTC; anomaly > 3σ
Workflow	Aggregate → compare with forecast → explain drivers → post to the finance centre
Escalation	Chargeback ratio ≥ 0.5% → payments lead; revenue drop > 15% day on day → CFO
APIs	Warehouse SQL (read-only role)
Value	Finance runs on facts, not spreadsheets

Pricing Agent

Purpose	Design and analyse market-level price experiments
Inputs	Conversion by market and tier, willingness-to-pay surveys, competitor price list
Outputs	Experiment proposals with power analysis; results with confidence intervals
Permissions	L1; recommend only; change_price is money_movement → human
Triggers	Quarterly review; G-P3 pricing decisions (D-05)
Workflow	Propose → pricing committee approves → Stripe Prices created by a person → measure → report
Escalation	Any proposal that personalises price per member is rejected by policy (§12.2)
APIs	Warehouse; Stripe Prices (human-operated)
Value	Better price points without eroding trust

Customer Support Agent

Purpose	Answer members quickly and discreetly
Inputs	Ticket text (member-provided), the member's entitlement and Keys ledger (case-scoped), help centre
Outputs	Draft replies; ticket classification; suggested refunds under published rules
Permissions	L1; draft_reply (recommend); issue_goodwill_keys (money_movement → human)

Triggers	ticket.created
Workflow	Classify → retrieve help-centre answer (RAG over public docs) → draft → agent or person sends. Safety keywords → Guardian queue immediately
Escalation	Safety, legal or payment disputes → human queue with SLA
APIs	Support desk connector (§8.2), /v1/entitlements (case-scoped)
Value	Lower cost per ticket; faster first response

Marketing Agent (Herald)

Specified in 03 §3.2. Purpose: organic growth through earned search visibility. Inputs: public web, Search Console, aggregate cookieless analytics. Outputs: drafts and refresh proposals. Permissions: L2 for internal-link repair; publish_page is irreversible → editor. Escalation: claims-discipline flags → editor. Value: low-CAC channel (D-21).

Data Intelligence Agent

Purpose	Answer business questions over the warehouse with correct, privacy-safe SQL
Inputs	Natural-language question from an operator; semantic layer (metric definitions)
Outputs	Query, result, chart, caveats
Permissions	L1; read-only warehouse role over aggregate marts; member-level tables not granted
Triggers	Operator request; scheduled packs
Workflow	Map question to metric definitions → generate SQL → run with row-count floor ($k \geq 20$) → add DP noise to member-derived counts → render
Escalation	A question that needs member-level data → refused, with the route to a DPO-approved analysis
APIs	ClickHouse (read-only), semantic layer
Value	Self-serve analytics without exposing members

Operations Agent

Purpose: keep queues inside SLA. Inputs: queue depths, SLA timers, staffing rota. Outputs: staffing forecasts, rebalancing proposals, SLA breach alerts. Permissions: L1. Triggers: every 5 min. Escalation: forecast breach of OSA-relevant SLAs → Head of T&S. Value: 24/7 coverage with a small team.

Fraud Detection Agent

Purpose	Stop card testing, stolen cards, voucher abuse and Keys fraud
Inputs	Processor risk signals (Stripe Radar, acquirer scores, BitriPay status), payment fingerprint HMAC, velocity, device risk
Outputs	Transaction score, holds, block-list proposals
Permissions	L2; score_payment (recommend), hold_payment (reversible), refund (money_movement → human)
Triggers	payment.attempted, voucher.redeem_attempted, chargeback alert (Ethoca/Verifi)
Workflow	Score → hold above threshold → auto-refund on a confirmed fraud alert where the published rule allows (03 §10) → feed L1 admission
Escalation	Chargeback ratio $\geq 0.5\%$ → payments lead; 0.65% → hard-stop review
APIs	PaymentProvider webhooks, voucher service
Value	Keeps the MID alive (above 0.9% it is lost, 01 §13.1)

Payment Agent

Purpose	Keep money flowing correctly across processors
Inputs	Webhooks from Stripe, the acquirer and BitriPay; ledger; processor health
Outputs	Reconciliation reports, processor routing proposals, dispute evidence packs
Permissions	L2; reconcile_ledger and route_processor (reversible: failover to the backup MID), change_price (money_movement → human)
Triggers	Webhook received; nightly reconciliation; authorisation-rate drop

Workflow	Verify signature → dedupe by event id → update entitlement and ledger → reconcile daily against settlement statements → if the primary's authorisation rate drops below its floor for 15 min, fail new checkouts over to the backup (reversible) and page payments ops
Escalation	Ledger imbalance → stop entitlement writes, page on-call; dispute deadline < 48 h → payments ops
APIs	/v1/billing/*/webhook, PaymentProvider
Value	Revenue continuity; dispute win rate

API Integration Agent

Purpose: keep every connector healthy. Inputs: connector contract tests, error rates, key ages. Outputs: health board, rotation reminders, deprecation alerts. Permissions: L2 (open tickets, disable a failing optional connector). Escalation: a required connector failing (age assurance, payments) → on-call. Value: fewer outages caused by third parties.

Workflow Automation Agent (Orchestrator)

Runs declared workflows (§5.7) step by step, with retries, compensation and the gate at every action. It never invents a workflow; workflows are code-reviewed definitions. Value: reliable multi-step operations (DSR, erasure, dispute, incident).

Predictive Growth Agent

Purpose: forecast liquidity per city, segment and mode (S4 scarcity, gender ratio $\leq 2.2:1$). Inputs: aggregates. Outputs: waitlist release forecasts, city-opening and seeding recommendations (01 §11). Permissions: L1. Escalation: a forecast breach of a no-go condition → CEO. Value: avoids launching cities that cannot work.

Admin Control Agent

Purpose: help admins run the Super Control Centre (§14). It summarises alerts, proposes actions and drafts runbook steps. It **never grants access**; just-in-time access is approved by a second admin. Permissions: L1.

5.6 Self-managing platform layer

Agent	Monitors / does	Autonomous (inside budget)	Always a person	Pattern
System Health	Uptime, latency, errors, saturation (golden signals); SLO burn rates (multi-window, multi-burn-rate)	Open incident, page on-call, attach traces	—	Google SRE workbook
Bug Detection	Error-tracking clusters, regression detection on deploy, flaky-test detection, crash-free rate	File issue with repro and suspected commit (git bisect in CI)	—	Meta SapFix/Sapienz-style triage; Sentry clustering
Auto-Repair	Runs pre-approved runbooks: rollback the last release, restart a crash-looping pod, scale out, drain a bad node, rotate a leaked key	L3 runbooks (budget 50/day)	Deploying a code patch (production_change). It may <i>open</i> a fix PR, which follows normal review and CI (D-28)	Kubernetes self-healing; Argo Rollouts; Netflix Spinnaker automated canary analysis
Infrastructure Optimisation	Compute, storage, bandwidth, cloud spend; GPU utilisation of the vLLM pool	Proposals only	Applying rightsizing	AWS Compute Optimizer; FinOps Foundation
Release Management	Canary analysis, progressive delivery (1% → 10% → 50% → 100%), feature flags, freeze windows	Halt a canary; roll back	Promoting a release	Google, Facebook progressive rollouts
AI Governance	The gate (§5.4), model and prompt registry, eval results, drift monitors, red-team results	Deny; kill an agent on policy breach (e.g. a guardrail violation rate above threshold)	Resuming a killed agent; changing a manifest	NIST AI RMF; ISO/IEC 42001; EU AI Act logging duties

Self-healing sequence (J4): SLO burn > 14.4x over 1 h → System Health opens INC → Auto-Repair checks whether a deploy happened in the last 2 h → if yes, rolls it back (runbook) → if the burn stops within 10 min, closes with a report for review; if not, pages on-call with everything gathered. Every step is a gate decision in the audit log.

5.7 Orchestration and agent memory

Concern	Design	Pattern
Orchestration	Typed state machines per agent (LangGraph-style, 03 §2.1); cross-agent workflows on a durable workflow engine (Temporal) with compensation steps	Uber Cadence/Temporal; Stripe's idempotent job runners
Workflow catalogue (P1)	onboarding, handshake (Loop C), meet_resolution (Loop D), dsr_export, erasure (< 60 s), report_to_decision, dispute, incident_selfheal, processor_failover, herald_publish	—
Memory: short-term	Per-run state in the workflow engine; expires with the run	—
Memory: member long-term	Only the structured vectors (IntentVector, RevealedPreferenceVector) and the Discretion Policy. No free-text memory of members , no conversation memory (E2EE content is never available server-side)	P3
Memory: operational	Knowledge graph of services, runbooks, incidents and decisions; vector index over runbooks, ADRs and the spec for retrieval	Palantir-style ontology; RAG
Tool use	Tools are gRPC methods listed in the manifest; arguments validated against Protobuf schemas	OpenAI and Anthropic tool-use with schema validation
Prompt and model registry	Versioned prompts and policies (policies/<agent>.yaml); a change is a production_change	MLflow-style registry
Injection defence	Untrusted text (profiles, tickets, web pages) is never placed in an instruction position; outputs are schema-validated; the injection red-team suite runs per release (03 §9)	OWASP LLM Top 10

6. Full platform modules

Every module the brief lists, mapped to TRYST. Existing member modules are in 04 §4 and 03 §2.2; this table adds the OS surfaces.

Module	Users	What it does	Key screens / endpoints	Agents	Phase
Dashboards	All	Role-scoped overviews (§4)	/account/insights (member); /ops/* (operators)	Chief of Staff, executive copilots	P1-P2
User accounts	Members	Sign-up, two-factor sign-in, devices, security activity, Burn, emergency lock, export, delete	/join, /sign-in, /account; /v1/me, /v1/me/security (built)	Onboarding, Identity risk	P1 (built in part)
Admin panels	Admins	Super Control Centre (§14)	/ops/control	Admin Control	P1
Merchant portals	Partners (U8)	BitriPay connected-account view, commission statements	/partners/merchant	Payment	P4
Partner portals	U2, U5	Keys, webhooks, usage, invoices, contracts	/partners	API Integration, Licensee Success	P4
API settings pages	U2, U7	Create, scope and rotate keys; IP allow-list; webhook endpoints	/developers/keys, /developers/webhooks	API Integration	P4
Payment settings	Members, payments ops	Member: method, billing email, cancel (one step). Ops: processor routing, descriptors	/account/billing; /ops/payments	Payment	P1
Notification system	All	Masked push (no names or content, 04 §5), email for statutory notices, in-app inbox	Notification engine (§9.4)	Curtain decides masking	P1
Reporting	Operators, regulators	Transparency report, KPI pack, finance pack	/ops/reports; regulator portal	Compliance, Revenue	P1
Analytics	Operators	Semantic-layer metrics, funnels, cohorts, experiment readouts	/ops/analytics	Data Intelligence	P2
Billing	Members	Invoices, receipts with neutral descriptor	/account/billing	Payment	P1

Module	Users	What it does	Key screens / endpoints	Agents	Phase
Subscription	Members	TRYST+, ENVOY, DUO, GHOST; renewal notices; one-step cancel	/v1/subscription*	Payment, Retention	P1/P3
Wallet	Members	Keys balance is TRYST's wallet: a closed-loop ledger, not stored value. No cash-out, no transfers between members	/v1/entitlements	Payment	P1
Commission	Partners	Voucher-reseller and licensee referral commission, paid by split (BitriPay application_fee) or invoice	/partners/commission	Payment	P4
Transaction	Ops	Ledger explorer: every entitlement and Keys movement with its processor event	/ops/payments/ledger	Payment, Fraud	P1
Audit trail	Admins, regulators	Agent decisions (hash chain), admin actions, break-glass, data access	/ops/audit	AI Governance	P1
Security	Members, SOC	Member: security activity, devices. SOC: alerts, incidents, containment	/account, /ops/security	Identity risk, SOC, Threat Hunter	P1
Agent registry	Admins	Manifests, levels, budgets, kill switches, eval scores	/ops/agents	AI Governance	P1
Connector registry	Admins	Each connector's data map, DPA status, health, kill switch	/ops/connectors	API Integration	P1

7. BitriPay payment gateway API door

7.1 Position in TRYST

BitriPay is integrated as a **core payment gateway layer behind the processor-agnostic PaymentProvider interface** (FR-078), next to Stripe and the high-risk acquirer (D-26, FR-084).

- **Built now:** the adapter (backend/internal/payments/bitripay.go), tested against a BitriPay-shaped sandbox server.
- **Live use is gated by market.** BitriPay's documented rails are mobile-money operators in markets TRYST will never enter (01 §4.3; D-13). Live keys are therefore refused in code unless a market is enabled for both TRYST and BitriPay, with a recorded legal opinion and BitriPay's written acceptance of the business model (as for Stripe, D-22).
- **Where BitriPay is commercially useful to TRYST without member exposure:** (a) settling **partner commissions** through connected accounts (voucher resellers, licensees) in any market BitriPay serves; (b) the **licensing model** (§12.3), where a licensee in a BitriPay market runs its own product and is the merchant of record.

7.2 API door: what TRYST implements against BitriPay

From the BitriPay developer documentation supplied by the founder. Base URL <https://api.bitripay.com/v1>; test and live keys share it.

Capability (brief)	BitriPay surface	TRYST implementation	Status
API key generation	POST /api_keys (sk_, rk_, pk_ with scopes); POST /api_keys/{id}/rotate (live needs step-up)	TRYST holds an rk_restricted key with only the scopes it needs (below), in the vault; rotation every 90 days by the API Integration agent's reminder, done by a person	Spec
Sandbox mode	sk_test_...; magic MSISDNs (§7.6)	Adapter accepts test keys; CI contract tests use a BitriPay-shaped server	Built
Live mode	sk_live_... / rk_live_...	Refused unless LiveMarkets[market] is set	Built (gate)
Webhooks	POST /webhook_endpoints (whsec_once); BitriPay-Signature HMAC + platform Ed25519 key from GET /keys; at-least-once, dedupe by event id; replay	VerifyWebhook (HMAC, 5-min tolerance), shared Dedup; Ed25519 second check in P0 once the header is confirmed	Built (HMAC)

Capability (brief)	BitriPay surface	TRYST implementation	Status
Merchant onboarding	POST /accounts (business_name, type, email, country, application_fee_bps) → claim link POST /accounts/{id}/account_links	Partner portal creates the partner's connected account; the partner claims it and owns it	Spec (P4)
QR payment	POST /qr_codes (static BitriQR, EMVCo + signed extension); qr_payload on intents	Not used for members (QR payment at a venue would link a member to a place). Used for partner invoices only	Spec
Payment links	POST /payment_links	Licensee and partner invoices	Spec
Wallet payments	GET /wallets, POST /transfers	Not used for members (no stored value at TRYST). Partners can hold commission in their own BitriPay wallet	Spec
Settlement	/settlement_profiles (T0/T1/T2/weekly/manual), /settlement_cycles/{id}/statement (json, csv, pdf; numbered, hashed)	Payment agent reconciles statements against the ledger nightly	Spec
Refunds	POST /refunds (atomic reservation), GET /payment_intents/{id}/refundable	Refund (idempotent key per refund)	Built
Disputes	GET /disputes, POST /disputes/{id}/respond	Payment agent assembles evidence; payments ops submits	Spec
Commission split	application_fee_minor on intents; BitriPay-Account header for connected accounts; /payment_intents/{id}/splits	Licensee and reseller splits	Spec (P4)
Transaction monitoring	GET /events, /payment_intents/{id}/timeline, /payment_resolution, GET /status (degraded flags)	Payment agent watches status and fails over (reversible)	Spec
Developer documentation	API reference + OpenAPI 3.1	Linked from the Developer Centre; our adapter's contract tests pin what we use	Spec
Plugin-ready integration logic	Hosted checkout, embedded widget (pk_), WooCommerce and Shopify plugins, SDKs	TRYST uses hosted checkout only (no third-party script on TRYST pages, NFR-20 and CSP)	Spec

Scopes TRYST requests (least privilege): payment_intents:write, payment_intents:read, checkout_sessions:write, refunds:write, refunds:read, subscriptions:read, subscriptions:write, disputes:read, disputes:write, settlements:read, balance:read, webhooks:manage, events:read; for P4 partner payouts, accounts:read, accounts:write. **Never requested:** payouts:approve (four-eyes stays with people), credit:read, remittances:*, ai:run.

7.3 Merchant and partner onboarding flow

1. Partner applies in the partner portal → KYB by TRYST's KYB connector (§8.2) for licensees; resellers get a lighter check.
2. TRYST calls POST /accounts with the partner's business details and the agreed application_fee_bps → acct_....
3. TRYST sends the claim link (POST /accounts/{id}/account_links, valid 7 days). The partner owns the account; TRYST appears under the partner's Team as a developer, with no access to settlement or payouts.
4. Sandbox integration tests pass (§7.6) → a person approves live mode in the connector registry (always-human).
5. The partner can remove TRYST at any time; POST /accounts/{id}/detach ends TRYST's access, and the money and history stay the partner's.

7.4 Commission split

For a partner sale billed through BitriPay, the intent carries application_fee_minor and the BitriPay-Account: acct_... header. The partner is the merchant of record; TRYST's fee is a transparent split shown on both statements. Refund behaviour on splits follows the intent's split-refund policy (pro_rata or merchant_absorbs, chosen per contract).

7.5 Payment state machine (as consumed by TRYST)

requires_payment_method → processing → succeeded → settled, with branches ambiguous (manual review at BitriPay; TRYST grants **no** entitlement until resolved), failed (back to requires_payment_method), canceled. Entitlements are written only from verified payment_intent.succeeded events; settlement events update the ledger only. An ambiguous_hold older than 24 h opens a payments-ops task.

7.6 Sandbox

BitriPay's magic MSISDNs drive the real attempt machine with test keys: +243000000501 succeed; +243000000404 fail (invalid_msisdn); +243000000408 ambiguous → manual review; +243000000500 timeout then succeed; +243000000503 provider unavailable (retryable); any number ending 0000 declined. TRYST's P0 contract suite runs all six against BitriPay's sandbox and asserts the entitlement outcome for each (no entitlement on ambiguous, failed or declined).

7.7 Adapter implementation (built)

Property	Code	Test
Bearer auth; Idempotency-Key on every money-moving POST, derived deterministically from the logical operation, so a retry can never charge twice	post, idempotencyKey	TestBitriPayCheckoutSandbox
Only amount_minor, currency, description are sent: no member identifier, nothing that reveals the service beyond a neutral descriptor	CreateCheckout	TestBitriPayCheckoutSandbox (asserts the exact field set)
Descriptor must be neutral and never name TRYST (ToS 16.2)	ValidDescriptor	TestBitriPayRejectsPublishableKeyAndBadDescriptor
Publishable keys refused server-side; live keys refused outside enabled markets; refused calls never reach the network	check	TestBitriPayLiveRefusedOutsideEnabledMarkets
Refunds and one-step cancel	Refund, CancelSubscription	TestBitriPayRefundCancelAndErrors
BitriPay error codes surfaced (scope_denied etc.)	APIError	same
Webhook HMAC with replay tolerance; empty secret fails closed	VerifyWebhook	TestBitriPayWebhook

Two wire details are not in the supplied documentation and are marked **ASSUMED** in the code: the refund request field names and the BitriPay-Signature header layout. They are a P0 check against BitriPay's API reference (00 §5).

8. Third-party API connectors

8.1 Connector rules (FR-087)

1. **Every connector is registered** with: purpose, the module it connects to, the exact data sent and received, lawful basis, DPA and transfer mechanism, region, health check, kill switch and owner.
2. **No special-category or member-identifying data goes to a connector unless the connector is the processor for exactly that purpose** (e.g. the age-assurance provider). Nothing reaches CRM, marketing, enrichment, ad or messaging-app platforms.
3. **Two providers for anything on the critical path** (age assurance, payments, email for statutory notices).
4. **Adapters, not SDK sprawl.** Each category sits behind one internal interface (as PaymentProvider does), so switching providers is configuration.
5. **AI providers:** hosted frontier models only for generic copy with PII redacted at the boundary; anything touching member data runs on self-hosted open-weight models in the UK (03 §2.1).

8.2 Connector register

Category	Why needed	Connects to	Sends / receives	Provider options (primary · alternative)	TRYST rule
Payments	Subscriptions, Keys, deposit, refunds	commerce via PaymentProvider	Sends amount, currency, neutral descriptor, opaque billing ref; receives payment events	Stripe (conditional, D-22) · high-risk acquirer (backup MID) · BitriPay (gated, D-26) · Adyen · Checkout.com · PayPal	Adyen and Checkout.com are evaluated as the high-risk backup in P0; PayPal is unlikely to accept the category (verify)
Banking-as-a-Service	Not needed for members (TRYST holds no stored value)	Treasury (ops)	Company accounts only	Business bank with API (e.g. Modulr, ClearBank partners)	Operator-only

Category	Why needed	Connects to	Sends / receives	Provider options (primary · alternative)	TRYST rule
Open banking	Optional pay-by-bank in UK/EU, lower fees, no card data	commerce	Payment initiation; no account data retained	TrueLayer · Yapily · Stripe (pay by bank)	P3 evaluation; descriptor rules apply
KYC / age assurance	Highly effective age assurance (OSA); V2	identity-svc	Sends selfie/document to provider; receives pass/fail + token	Two providers: Yoti · Persona · Veriff · Sumsu · Onfido	No document image retained by TRYST (A §8.1)
KYB	Licensees and commission partners	Partner onboarding	Business registration data	Sumsu KYB · Persona · Companies House API	B2B only
AML screening	Sanctions and PEP screening of businesses (licensees, partners). Member AML duties sit with the regulated PSP	Partner onboarding	Business and director names	ComplyAdvantage · Sumsu	D-32
Fraud prevention	Card testing, ATO, bots	Payment, identity, edge	Device signals, payment risk	Stripe Radar · Ethoca/Verifi alerts · Cloudflare Bot Management · Fingerprint (device, hashed)	Device IDs hashed; no cross-site tracking
Identity verification (liveness B2)	Login factor B2 (02 §3.3)	identity-svc	Liveness session; match result	iProov · FaceTec · Persona	PAD Level 2 (NFR-17); D-18
Email	Statutory notices (renewal, cooling-off), sign-in alerts	Notification engine	Billing email (member-chosen), neutral subject lines	Postmark · Amazon SES · SendGrid · Brevo	No tracking pixels; neutral sender
SMS	Sign-up OTP only ; never a login factor (D-16)	identity-svc	Phone number, code	Twilio Verify · MessageBird	Neutral sender ID
WhatsApp	—	—	—	—	Declined for members (D-29): exposes the contact on a shared device and to a third party. Operator alerts may use Slack/PagerDuty instead
Push notifications	Masked alerts	Notification engine	Opaque payload; content fetched in app	APNs · FCM	No names or content (04 §5)
Maps	Venue sanity (public venue), travel mode, zone drawing	Client, Envoy logistics	Coarse geohash only; no live tracking	Mapbox (self-hosted tiles where possible) · Google Maps Platform	Exact locations never leave the device (FR-067)
Logistics	Not applicable (no physical goods)	—	—	—	Recorded as not applicable
Accounting	Revenue recognition, VAT returns	Finance (ops)	Aggregated revenue journals; no member data	Xero · NetSuite	Aggregates only
Tax	VAT/OSS calculation	commerce	Country + amount	Stripe Tax · Avalara	—
CRM	B2B pipeline only (licensees, partners, press)	Sales copilot	Business contacts	HubSpot · Salesforce	No member data, ever (D-29)

Category	Why needed	Connects to	Sends / receives	Provider options (primary · alternative)	TRYST rule
Analytics	Product and marketing-site analytics	Warehouse; marketing site	Cookieless aggregate; pseudonymous product events	Self-hosted: ClickHouse + Plausible/Umami	No third-party tracking scripts (NFR-20)
AI model providers	Generic copy (Herald, help centre); coding copilots	llm-gateway	Redacted prompts	Anthropic · OpenAI · Google Gemini / Vertex AI · Cohere · Mistral; self-hosted open-weight (vLLM) for member data	Zero-retention terms required for hosted providers
Cloud	Compute, storage, KMS/HSM	All	—	AWS eu-west-2 (primary) · Azure / GCP for DR evaluation	UK/EU residency
Cloud storage	Encrypted media, backups, WORM audit anchors	media-svc, ops	Ciphertext only	S3 (object lock) · Cloudflare R2 (egress-free CDN origin for marketing assets)	Client-side encryption for member media
Authentication	Members: passkeys (built, go-webauthn). Staff: SSO	identity-svc; ops	—	Members: in-house WebAuthn. Staff: Okta / Google Workspace SSO + hardware keys	No third-party IdP for members
Productivity	Staff collaboration	Ops	No member data	Google Workspace · Microsoft 365	DLP rules block member identifiers
Document generation	Statements of reasons, DSR exports, invoices, transparency report	Compliance, Payment	Generated server-side	In-house (HTML → PDF with headless Chromium, as the spec PDF)	Not sent to third parties
E-signature	Partner, licensee and vendor contracts	Partner onboarding	Business signatories	DocuSign · Dropbox Sign	B2B only
Customer support	Ticketing	Support agent	Ticket text, opaque member ref	Zendesk · Help Scout · Intercom (without its tracking widget)	Case-scoped; neutral email domain
Data enrichment	—	—	—	—	Declined for members (D-29): enrichment would link a pseudonymous member to a real identity, breaking FR-002. B2B lead enrichment allowed in the CRM
Currency exchange	Display prices in EUR, SEK, DKK in P4	commerce	Rates only	Processor-native pricing (Stripe multi-currency) · ECB reference rates	Prices set per market, not converted live
Subscription billing	Plans, renewals, notices	commerce	Via the payment adapters	Stripe Billing · BitriPay /plans and /subscriptions (gated)	DMCC/CCR notices (03 §10)
Observability	Metrics, logs, traces	All services	No member content; pseudonymous IDs	Self-hosted Grafana stack (Prometheus, Loki, Tempo) + OpenTelemetry	Logs scrubbed at source
Security	WAF, DDoS, secrets, scanning	Edge, CI	—	Cloudflare · AWS Shield · HashiCorp Vault / AWS Secrets Manager · Snyk / GitHub Advanced Security	—
Incident and on-call	Paging	System Health	Incident metadata	PagerDuty · Opsgenie	—

9. Production-grade architecture

9.1 Proven patterns used, and where they come from

Pattern	Proven at	Used for
Idempotency keys on every mutating request	Stripe	Edge API (02 \$5), all payment adapters
Signed webhooks with timestamp tolerance and replay	Stripe, GitHub	Inbound processor webhooks; outbound partner webhooks (\$11.4)
Zero trust / BeyondCorp: identity- and device-aware access, no trusted network	Google	Staff access; service-to-service mTLS (SPIFFE)
Event streaming with a schema registry	LinkedIn (Kafka), Uber	Redpanda + Protobuf registry (03 \$2.1)
Durable workflows with compensation	Uber (Cadence → Temporal), Netflix (Conductor)	\$5.7 workflow catalogue
Progressive delivery with automated canary analysis	Netflix (Kayenta), Google	Release Manager (\$5.6)
SLOs and error budgets; multi-window burn-rate alerts	Google SRE	System Health
Feature store for train/serve parity	Uber (Michelangelo), Airbnb (Zipline)	Feast (03 \$2.1)
Two-tower retrieval + learned ranking with constraints	YouTube, Pinterest, Airbnb search	Broker (03 \$5)
Ontology / knowledge graph over operations	Palantir	Operational memory (\$5.7)
Risk engine with explainable reason codes	Stripe Radar, Goldman Sachs Aladdin risk factors	Fraud, Risk
Cell-based architecture to cap blast radius	AWS, Slack	\$9.8 (P4)
Lakehouse with governed semantic layer	Databricks, Snowflake	\$9.6
Edge WAF, bot management, DDoS absorption	Cloudflare	\$13
Endpoint and identity threat detection	CrowdStrike	SOC and Threat Hunter design
Workflow-centric operations consoles	ServiceNow	Operator command centres (\$4)
Agent tool use with schema validation and human approval	OpenAI, Anthropic, Microsoft	\$5

9.2 Layered architecture

Clients – iOS / Android (native, Rust crypto core) · Web PWA (Next.js) · Marketing site · Ops console · Partner portal
TLS 1.3, passkeys, DPoP-bound tokens
Edge – Cloudflare (WAF, DDoS, bot) → edge-gateway (Go): authN, jurisdiction, tier, consent, rate limit, idempotency
gRPC + mTLS (SPIFFE)
Domain – identity-svc · core-svc (profile, couple, intent, discretion, discovery, reveal, threads/MLS, meet, consent, deletion, commerce)
media-svc · safety-svc · jobs · webhook-engine · notification-engine · partner-api (P4)
typed Protobuf events (Redpanda + schema registry)
Agents – member plane: Cartographer, Mirror, Broker, Envoy, Curtain, Guardian, Aftercare, Onboarding, Risk, Fraud, Support, Payment, Compliance, Chief of Staff
ops plane: Herald*, System Health, Bug Detection, Auto-Repair, Release Manager, Infra Optimiser, SOC, Threat Hunter, Vulnerability, analysts
every action → aigov gate → audit chain (*Herald: separate cloud account)
AI runtime – llm-gateway (vLLM pool, self-hosted, UK) · hosted LLM proxy (redacted, generic copy only) · model + prompt registry · eval harness
Data – IDENTITY-DB · PROFILE-DB · SAFETY-DB · BAN-DB (isolated PostgreSQL clusters) · Qdrant · Feast · ClickHouse · S3 (ciphertext, object lock) · KMS/HSM
Platform – Kubernetes (EKS) · Argo CD · Terraform · Vault · OpenTelemetry → Grafana stack · PagerDuty · Temporal

9.3 Component specification

Component	Specification
Frontend	Native Swift / Kotlin with a shared Rust crypto core (E-07); Next.js 16 PWA (built: join, enrol, sign-in, approve, account); ops console and partner portal as separate Next.js apps on separate origins with their own CSP
Backend	Go services (03 \$2.2); modular monolith core-svc with split triggers (\$9.8)

Component	Specification
Database	PostgreSQL 16 clusters isolated per domain; pgvector; migrations generated from 02 §4 (scripts/gen_migrations.py, drift-checked in CI)
Authentication	Members: account-holder-only passkeys + second factor (built). Staff: SSO + FIDO2 hardware key + device posture. Services: SPIFFE mTLS. Partners: scoped API keys + mTLS for webhooks
Role-based access control	RBAC for staff roles, plus ABAC (03 §6.5) on case, purpose and data class; just-in-time elevation with a second approver; all access logged
AI orchestration layer	Agent services + Temporal workflows + aigov gate (§5)
Agent memory layer	Structured member vectors only; operational knowledge graph and runbook vector index (§5.7)
Vector database	Qdrant for candidate retrieval (filtered ANN after Stage 0 gates); a separate Qdrant collection in the ops plane for runbook and spec retrieval
Event-driven workflows	Redpanda topics member.*, agent.*, payment.*, safety.*, ops.*; Protobuf schemas with compatibility checks; outbox pattern from PostgreSQL to avoid dual writes
API gateway	edge-gateway (built skeleton): middleware order in 02 §5; per-key and per-subject rate limits (token bucket in Redis)
Webhook engine	Inbound: verify signature → dedupe → enqueue → process idempotently. Outbound (partners, P4): signed TRYST-Signature (§11.4), exponential back-off for 72 h, delivery log, replay, endpoint auto-disable after 5 days of failure
Payment gateway layer	PaymentProvider with Stripe, acquirer and BitriPay adapters; the Payment agent's failover (§5.5)
Notification engine	Channel router (push, email, in-app); Curtain-controlled masking; quiet hours; statutory notices bypass quiet hours but stay neutral
Audit log	Two chains: agent decisions (aigov) and admin actions; both anchored hourly to S3 object lock
Admin control layer	Super Control Centre (§14) over the ops API; break-glass with two-person approval
Security layer	§13
Observability	OpenTelemetry traces, metrics and logs; member IDs replaced by rotating pseudonyms; content never logged
Monitoring	SLOs per user-facing journey (§9.5); synthetic probes for sign-in, slate, checkout; Privacy Auditor probes (03 §9)
Error handling	RFC 9457 problem details with stable codes (02 §5.10); uniform 404 not_found for anything invisible; retries only on idempotent calls; circuit breakers per connector; fail closed for auth, payments and safety
Scalability logic	§9.8

9.4 Event architecture

Topic family	Producers	Consumers	Retention	Notes
member.* (02 §8 taxonomy)	Clients via /v1/signals:batch, core-svc	Mirror, Broker, analytics	Raw 180 d, then purged	Pseudonymous; no content
agent.*	Agents	Agents, audit sink	30 d	Typed turns (02 §6.1)
payment.*	Webhook engine	Payment, Fraud, ledger	7 y (financial records), pseudonymised	—
safety.*	Guardian, safety-svc	T&S console, Compliance	Per SAFETY-DB policy	Access only via the trust console
ops.*	System Health, CI, deploys	Auto-Repair, Release Manager, Bug Detection	90 d	No member data
aigov.decision	Gate	Audit sink, Admin Control	7 y	Hash chain

9.5 SLOs

Journey	SLI	SLO
Sign-in (both factors)	Successful completions / attempts with valid factors	99.9% / 30 d; p95 < 3 s excluding the member's biometric time
Slate	/v1/slate success and latency	99.9%; p95 < 400 ms (NFR-02)
Thread delivery	MLS message relayed	99.95%

Journey	SLI	SLO
Checkout	Checkout creation success	99.9% across processors
Erasure	Completed < 60 s	99% (NFR-01)
Emergency lock	Effective < 5 s	95% (FR-064)
Agent gate	Decision latency	p99 < 5 ms in-process

9.6 AI data intelligence layer

Component (brief)	TRYST implementation	Privacy rule
Data lake	S3 (eu-west-2), Iceberg tables, partitioned by day; raw events land here pseudonymised	Raw behavioural events purged at 180 d
Data warehouse	ClickHouse marts (dbt-modelled); semantic layer of metric definitions (QCAM, mutual rate, etc.)	Aggregate marts only for analysts; k ≥ 20 row floor
Vector database	Qdrant (member plane) and a separate ops-plane collection	Member vectors never leave the member plane
Knowledge graph	Ops: services, runbooks, incidents, decisions (Neo4j or PostgreSQL + Apache AGE). Member plane: the block and exclusion graph used by Broker Stage 0	The member graph is used for gates and abuse detection only
Event streaming platform	Redpanda + schema registry	\$9.4
Real-time analytics engine	ClickHouse materialised views over Redpanda for ops dashboards (queue depth, authorisation rate, SLOs)	No member-level real-time views outside T&S
Predictive intelligence engine	Liquidity forecasting (city × segment × mode), churn propensity (aggregate cohorts), capacity forecasting	Forecasts on aggregates; churn scores never shown to support staff as a label on a member
Behavioural intelligence engine	Mirror (03 §4): RevealedPreferenceVector, divergence, α weighting with the 15% stated-preference floor	Message content never used; personalisation can be switched off (FR-024)
Recommendation engine	Broker: Stage 0 gates → retrieval → M3/M4 ranking → fairness constraints (S4 caps, Gini)	Hard limits inviolable (P2)
Decision intelligence engine	Guardian triage, Risk and Fraud scoring, Payment failover rules, all with reason codes and gate decisions	Every decision explainable and audited

Data governance architecture: a data catalogue with owner, classification (public, internal, personal, special-category, safety), lawful basis, retention and permitted purposes per dataset; column-level tags drive access policies; lineage from dbt; the Privacy Auditor tests that the catalogue matches reality each release.

9.7 Disaster recovery and business continuity

Tier	Systems	RPO	RTO	Mechanism
T0	identity-svc, edge, IDENTITY-DB, BAN-DB, KMS/HSM	1 min	30 min	Multi-AZ synchronous replicas; cross-region (eu-west-1) async replica; HSM cluster across AZs; quarterly restore drill
T1	core-svc, PROFILE-DB, threads, payments	5 min	1 h	As above; Redpanda tiered storage
T2	Agents, analytics	24 h	8 h	Rebuild from IaC and the lake

Erasure must survive restore: deletion jobs are replayed against any restored backup before it serves traffic (crypto-shredded keys stay destroyed, 03 §6.3). **Business continuity:** backup MID live from day one; two age-assurance providers; staff runbooks for a full agent kill (the platform keeps running with agents at L0, falling back to the deterministic P1 baseline ranker, 03 §5.2).

9.8 Scalability

Trigger	Action
A core-svc module > 40% of CPU or a separate team owns it	Split it into its own service behind the same gRPC contract
PROFILE-DB write p95 > 20 ms or > 2 TB	Partition by region cell, then shard by profile_id hash
> 500k verified members or a second region (P4)	Cell-based architecture: a cell = full stack for a market group (UK+IE, EU-North, EU-South); global services limited to identity routing and ban anchors

Trigger	Action
vLLM pool GPU utilisation > 70% p95	Infra Optimiser proposes scale-out; a person applies it
Qdrant p95 retrieval > 40 ms	Add replicas; tighten Stage 0 pre-filters

9.9 Deployment strategy

Trunk-based development; every merge runs CI (spec check, contracts lint, Go, Python, Rust, web, Playwright e2e, as in .github/workflows/ci.yml). Images are signed (Sigstore cosign) with SLSA provenance; Argo CD syncs from Git; progressive delivery 1% → 10% → 50% → 100% with automated canary analysis; feature flags for every user-visible change; freeze windows around launches. Agents ship with their eval suites as a release gate (03 §9). Infrastructure is Terraform with policy checks (OPA) in CI.


9.10 Testing strategy

Layer	What	Where today
Unit	Each package	Go (go test ./..., including aigov and payments), Python (pytest + Hypothesis), Rust (cargo test), web (Vitest)
Property	Hard-limit inviolability, ranking invariants	agents/tests (Hypothesis)
Contract	OpenAPI lint; Protobuf compatibility; processor adapters against recorded sandbox behaviour	Redocly in CI; bitripay_test.go
End to end	Passkey sign-up and two-factor sign-in with virtual authenticators	Playwright (web/e2e/passkey.spec.ts)
Agent evals	400-case golden set per agent; injection red team; workflow-level replay	03 §9 (P1)
Governance	Always-human classes, kill switch, audit-chain tamper detection; mutation-checked	aigov_test.go
Privacy	Privacy Auditor probes: access, retention, model leakage	jobs (P1)
Load	k6 profiles per journey at 3× forecast peak	P1
Chaos	AZ loss, connector outage (age assurance, payments), agent kill-all	Game days quarterly from P2
Security	SAST, DAST, dependency and container scanning; quarterly pen test; annual red team; bug bounty (NFR-19)	CI + vendors

10. Database schema

The member data model is owned by 02 §4 and is not repeated here. The operating system adds the **ops** and **partner** schemas below, in PROFILE-DB's cluster but with separate roles, so no member-plane role can read partner data and no ops role can read member tables.

10.1 ERD (OS additions)



10.2 DDL

```
CREATE SCHEMA ops; CREATE SCHEMA partner;

CREATE TABLE ops.agent (
  agent_id      TEXT PRIMARY KEY,                -- e.g. 'guardian'
  owner_team    TEXT NOT NULL,
  level         SMALLINT NOT NULL CHECK (level BETWEEN 0 AND 3),
  plane         TEXT NOT NULL CHECK (plane IN ('member','ops')),
  self_hosted   BOOLEAN NOT NULL,
  daily_budget  INT NOT NULL CHECK (daily_budget >= 0),
  data_scopes   TEXT[] NOT NULL,
  model_version TEXT, policy_version TEXT,
  killed        BOOLEAN NOT NULL DEFAULT false,
  updated_at    TIMESTAMPTZ NOT NULL DEFAULT now(),
  CHECK (plane = 'member' OR NOT data_scopes &&
ARRAY['identity_db','profile_db','safety_db','ban_db','event_bus','feature_store','vector_store']),
  CHECK (plane = 'ops' OR self_hosted)
);
CREATE TABLE ops.agent_action (
  agent_id TEXT REFERENCES ops.agent ON DELETE CASCADE,
  action   TEXT NOT NULL,
  class    TEXT NOT NULL CHECK (class IN ('read','recommend','reversible','runbook','irreversible',
                                          'member_safety','money_movement','production_change')),
  PRIMARY KEY (agent_id, action)
);
CREATE TABLE ops.staff_user (
  staff_id UUID PRIMARY KEY, display_name TEXT NOT NULL, roles TEXT[] NOT NULL,
  can_approve BOOLEAN NOT NULL DEFAULT false, hw_key_enrolled BOOLEAN NOT NULL,
  disabled_at TIMESTAMPTZ
);
CREATE TABLE ops.agent_decision (
  seq          BIGINT PRIMARY KEY,                -- append-only; hash chain (aigov)
  at           TIMESTAMPTZ NOT NULL,
  kind         TEXT NOT NULL,                    -- register | decide:allow | decide:needs_human |
decide:deny | kill | kill_all | resume
  agent_id     TEXT NOT NULL,
  action       TEXT,
  approver     UUID REFERENCES ops.staff_user,
  subject      TEXT,                            -- pseudonymous reference only
  detail       TEXT,
  prev_hash    TEXT NOT NULL,
  hash         TEXT NOT NULL UNIQUE
);
CREATE INDEX agent_decision_agent_at ON ops.agent_decision (agent_id, at DESC);
CREATE INDEX agent_decision_pending ON ops.agent_decision (at) WHERE kind = 'decide:needs_human';
REVOKE UPDATE, DELETE, TRUNCATE ON ops.agent_decision FROM PUBLIC;

CREATE TABLE ops.admin_action (
  seq          BIGINT PRIMARY KEY, at TIMESTAMPTZ NOT NULL, staff_id UUID NOT NULL REFERENCES ops.staff_user,
  action       TEXT NOT NULL, target TEXT, justification TEXT NOT NULL, second_approver UUID REFERENCES ops.staff_user,
  break_glass  BOOLEAN NOT NULL DEFAULT false, prev_hash TEXT NOT NULL, hash TEXT NOT NULL UNIQUE,
  CHECK (NOT break_glass OR (second_approver IS NOT NULL AND second_approver <> staff_id))
);
CREATE TABLE ops.access_grant (
  grant_id UUID PRIMARY KEY, staff_id UUID NOT NULL REFERENCES ops.staff_user, role TEXT NOT NULL,
  case_ref TEXT, approved_by UUID NOT NULL REFERENCES ops.staff_user,
  starts_at TIMESTAMPTZ NOT NULL, expires_at TIMESTAMPTZ NOT NULL,
  CHECK (approved_by <> staff_id), CHECK (expires_at <= starts_at + interval '8 hours')
);
CREATE TABLE ops.connector (
  connector_id TEXT PRIMARY KEY,                -- 'stripe', 'bitripay', 'yoti', ...
  category TEXT NOT NULL, purpose TEXT NOT NULL, module TEXT NOT NULL,
  data_sent TEXT[] NOT NULL, data_received TEXT[] NOT NULL,
  member_data BOOLEAN NOT NULL, lawful_basis TEXT, dpa_signed_at DATE, region TEXT NOT NULL,
  live_enabled BOOLEAN NOT NULL DEFAULT false, live_markets TEXT[] NOT NULL DEFAULT '{}',
  live_approved_by UUID REFERENCES ops.staff_user, owner_team TEXT NOT NULL,
  CHECK (NOT member_data OR dpa_signed_at IS NOT NULL),
  CHECK (NOT live_enabled OR live_approved_by IS NOT NULL)
);
CREATE TABLE ops.connector_secret_ref (
  connector_id TEXT REFERENCES ops.connector, environment TEXT CHECK (environment IN ('sandbox','live')),
  vault_path TEXT NOT NULL, rotated_at TIMESTAMPTZ NOT NULL, PRIMARY KEY (connector_id, environment)
);
CREATE TABLE ops.payment_event (
  provider TEXT NOT NULL, event_id TEXT NOT NULL, type TEXT NOT NULL, billing_ref TEXT,
  received_at TIMESTAMPTZ NOT NULL, processed_at TIMESTAMPTZ, payload_hash TEXT NOT NULL,
  -- every processor webhook, verified
```

```

PRIMARY KEY (provider, event_id)          -- idempotency
);

CREATE TABLE partner.partner (
  partner_id UUID PRIMARY KEY, kind TEXT NOT NULL CHECK (kind IN ('licensee','reseller','assurance','content')),
  legal_name TEXT NOT NULL, country TEXT NOT NULL, kyb_status TEXT NOT NULL DEFAULT 'none',
  bitripay_account TEXT, application_fee_bps INT CHECK (application_fee_bps BETWEEN 0 AND 3000),
  created_at TIMESTAMPTZ NOT NULL DEFAULT now()
);
CREATE TABLE partner.api_key (
  key_id UUID PRIMARY KEY, partner_id UUID NOT NULL REFERENCES partner.partner,
  prefix TEXT NOT NULL CHECK (prefix IN ('tk_test','tk_live')), secret_hash BYTEA NOT NULL,  -- Argon2id; secret
shown once
  scopes TEXT[] NOT NULL, ip_allowlist CIDR[], created_at TIMESTAMPTZ NOT NULL, revoked_at TIMESTAMPTZ
);
CREATE TABLE partner.webhook_endpoint (
  endpoint_id UUID PRIMARY KEY, partner_id UUID NOT NULL REFERENCES partner.partner,
  url TEXT NOT NULL CHECK (url LIKE 'https://%'), secret_ref TEXT NOT NULL, event_types TEXT[] NOT NULL,
  disabled_at TIMESTAMPTZ
);
CREATE TABLE partner.webhook_delivery (
  delivery_id UUID PRIMARY KEY, endpoint_id UUID NOT NULL REFERENCES partner.webhook_endpoint,
  event_id UUID NOT NULL, attempt INT NOT NULL, status_code INT, next_attempt_at TIMESTAMPTZ,
  delivered_at TIMESTAMPTZ
);
CREATE INDEX webhook_delivery_due ON partner.webhook_delivery (next_attempt_at) WHERE delivered_at IS NULL;
CREATE TABLE partner.usage_record (
  -- ACU metering (B2B only, D-27)
  usage_id UUID PRIMARY KEY, partner_id UUID NOT NULL REFERENCES partner.partner,
  operation TEXT NOT NULL, acu NUMERIC(12,3) NOT NULL CHECK (acu >= 0),
  idempotency_key TEXT NOT NULL, at TIMESTAMPTZ NOT NULL, UNIQUE (partner_id, idempotency_key)
);
CREATE TABLE partner.acu_balance (
  partner_id UUID PRIMARY KEY REFERENCES partner.partner, prepaid_acu NUMERIC(14,3) NOT NULL CHECK (prepaid_acu >=
0),
  updated_at TIMESTAMPTZ NOT NULL
);

```

Permissions. Roles member_plane_rw (02 \$4 schemas only), ops_rw (ops), partner_rw (partner), analyst_ro (aggregate marts in ClickHouse, not PostgreSQL), auditor_ro (ops.agent_decision, ops.admin_action). No role spans member and partner schemas. **Audit requirements:** append-only tables have UPDATE/DELETE revoked; hourly hash anchoring to S3 object lock; every access_grant and break-glass event is itself an admin_action.

11. API specification

11.1 Surfaces

Surface	Base	Auth	Who
Member API (existing)	/v1 (02 \$5; contracts/openapi/tryst.v1.yaml)	Passkey session; DPOP from P1	Members
Ops API	/ops/v1 (internal origin, not internet-routable)	Staff SSO + hardware key + JIT role	Operators, admins
Partner API (P4)	/partner/v1	Authorization: Bearer tk_live_... + optional mTLS; scopes	Licensees, resellers
Inbound webhooks	/v1/billing/{provider}/webhook	Provider signature	Stripe, acquirer, BitriPay
Outbound webhooks	Partner URLs	TRYST-Signature	Partners

11.2 Ops API endpoints (OS)

Method	Path	Permission	Contract
GET	/ops/v1/agents	agents:read	Registry: manifests, level, budget used today, killed
POST	/ops/v1/agents/{id}/kill	agents:kill	{reason} → kill; audit entry. id = * kills all
POST	/ops/v1/agents/{id}/resume	agents:resume + second approver	Resume a killed agent
GET	/ops/v1/decisions?agent=&outcome=needs_human	decisions:read	Pending and past gate decisions

Method	Path	Permission	Contract
POST	/ops/v1/decisions/{seq}/approve	decisions:approve (named human with can_approve)	Approves a needs_human request; the agent retries with approver set
GET	/ops/v1/audit/verify	audit:read	Runs VerifyChain; returns {ok, entries, last_anchor}
GET	/ops/v1/connectors	connectors:read	Register with health
POST	/ops/v1/connectors/{id}/live	connectors:live + second approver	Enable live in listed markets
POST	/ops/v1/payments/failover	payments:route	Switch new checkouts to the backup processor
GET	/ops/v1/slo	ops:read	SLO status and error budgets

Example: approve a pending removal

```
POST /ops/v1/decisions/18421/approve
Authorization: Bearer <staff session>
Idempotency-Key: 6flc...
Content-Type: application/json
```

```
{"justification": "Case C-2207: verified sextortion evidence; second review by moderator M-114"}
```

```
{"seq": 18422, "kind": "decide:allow", "agent_id": "guardian", "action": "remove_member",
  "approver": "b6b1...", "reason": "member_safety approved by named human", "hash": "9c4e..."}
```

11.3 Partner API endpoints (P4)

Method	Path	Scope	Contract	ACU
POST	/partner/v1/consent/handshakes	handshake:write	Run a schema-only agent handshake between two of the licensee's own users (the Envoy engine as a service)	2.0
POST	/partner/v1/discretion/exclusion-check	psi:write	ExclusionRing-style PSI round (client library does the blinding)	0.5
POST	/partner/v1/safety/classify	safety:write	On-premises Guardian model licence check-in (models run in the licensee's environment)	0 (licence)
GET	/partner/v1/usage	usage:read	ACU used by day and operation	—
POST	/partner/v1/webhook_endpoints	webhooks:manage	Register an endpoint; secret returned once	—

Example: handshake request

```
POST /partner/v1/consent/handshakes
{"a": {"intent_shape": "one_off", "hard_limits": ["no_photos"], "windows": [{"2026-11-02T18:00Z", "2026-11-02T23:00Z"}]},
  "b": {"intent_shape": "one_off", "hard_limits": [], "windows": [{"2026-11-02T20:00Z", "2026-11-03T01:00Z"}]}}
→ 200 {"handshake_id": "hs_...", "outcome": "compatible", "blocking": [], "friction": ["reveal_pace"],
  "briefs": {"a": {...codes...}, "b": {...codes...}}, "acu": 2.0}
```

11.4 Webhooks

Inbound (processor → TRYST): verify signature (Stripe Stripe-Signature; BitriPay BitriPay-Signature; acquirer per contract) → 5-minute tolerance → dedupe on (provider, event_id) in ops.payment_event → process → 2xx. Unverified → 400 with no detail.

Outbound (TRYST → partner): header TRYST-Signature: t=<unix>,v1=<hex HMAC-SHA256(secret, "<t>.<body>")>; events handshake.completed, usage.threshold_reached, invoice.issued, key.rotation_due; at-least-once; partners dedupe on event_id; retries at 1 m, 5 m, 30 m, 2 h, 6 h, then every 12 h for 72 h; replay via the partner portal.

11.5 Rate limits

Surface	Limit
Member API	Per endpoint (02 §5); sign-in begin 10/min per IP-range bucket and 5/min per contact hint
Ops API	600/min per staff session; kill and approve endpoints 30/min
Partner API	Plan-based: Sandbox 60/min; Growth 600/min; Enterprise contractual; burst 2× for 10 s

Surface	Limit
Webhook intake	2,000/min per provider; above that, 429 and the provider retries

11.6 Error codes (additions to 02 §5.10)

HTTP	code	Meaning
403	scope_denied	Key lacks the scope
403	approval_required	Action is always-human; a pending decision was created (decision_seq returned)
409	agent_killed	Agent is stopped by the kill switch
402	acu_insufficient	Partner prepaid ACU exhausted
403	live_not_enabled	Connector or key not enabled for live in this market
400	signature_invalid	Webhook signature failed (no further detail)

12. Monetisation model

12.1 Principles that bind the model

Charge for agency and discretion; make resolution the profitable outcome (01 §10.1). P8 stays in force: no advertising, no data sale, no affiliates, no fake profiles, no unreadable credit mechanics, no engagement paywalls, no charge for deletion or data access, replying always free. Every revenue line below is checked against it.

12.2 Revenue lines (brief → TRYST)

Brief line	TRYST line	Price	Status	Phase
Subscription plans	Verified (free + £9.99 refundable deposit), TRYST+ £24.99/mo (£14.99 annual), ENVOY £49.99/mo, DUO £39.99/mo per couple	01 §10.2	Canonical	P1/P3
Add-ons	GHOST £7.99/mo (D-05)	01 §10.2	Canonical	P3
Transaction fees	Keys packs £9.99/60 (intent costs refunded on reply)	01 §10.2	Canonical	P1
AI credit / ACU system	Partner API metering only (D-27). Members never buy ACU	Illustrative: £0.04 per ACU prepaid; volume bands	Adapted	P4
API usage fees	Partner API plans: Sandbox free; Growth £1,500/mo incl. 50k ACU; Enterprise contract	Illustrative	Spec	P4
Commission model	Reseller commission on voucher sales (paid out, a cost to TRYST); licensee referral split via BitriPay application_fee or invoice	Contract	Spec	P3/P4
Merchant fees	Not applicable to members (no merchants on the member platform). For licensees on BitriPay rails, BitriPay's fees apply to them; TRYST takes an application fee where it provides the integration	Contract	Adapted	P4
Partner fees	Assurance and content partners are suppliers, not fee payers	—	Not applicable	—
Premium automation fees	The ENVOY tier is the premium automation product for members	£49.99/mo	Canonical	P2/P3
Data intelligence revenue	Declined for member data (P8; D-30). Alternative: free published aggregate research with differential privacy as a PR asset (Herald)	—	Declined	—
White-label licensing	Licensing the governed stack (handshake engine, PSI exclusion, on-device safety models, governance gate) to other platforms; no member data and no shared member pool	Illustrative: £60–250k/yr per licensee + ACU	Spec	P4+
Enterprise plans	Enterprise licensing with SLA, on-premises models, dedicated support	Contract	Spec	P4+
BitriPay gateway revenue	Application fees on licensee and reseller flows routed through BitriPay connected accounts	bps per contract	Spec (gated)	P4

Brief line	TRYST line	Price	Status	Phase
Marketplace revenue	No marketplace of people or services (commercial sexual services are prohibited, 01 §4.1). A vetted integrations marketplace for licensees (e.g. assurance providers) with referral fees is possible	Contract	Adapted	P4+
Dynamic pricing engine	Market-level list prices tested by experiment; the same price for every member in a market; published; no personalised or surge pricing	—	Adapted (D-30)	P3

12.3 B2B licensing, white label and API revenue

What is licensed: the Envoy handshake engine (schema-only negotiation, asymmetric Briefs), the ExclusionRing PSI protocol and client library, the on-device Guardian models, Curtain's discretion policy engine and the aigov governance gate. **What is never licensed or shared:** TRYST members, their data, vectors or any model trained on them without a separate, lawful basis (models for licensing are trained on licensed or synthetic data). Licensees run the stack in their own environment; TRYST's Partner API meters calls in ACU.

12.4 Unit economics impact

The OS changes 01 §10.3 in three ways (to be measured in beta, not assumed):

1. **Cost per MAU down:** agent-handled support and T&S triage lower the variable ops cost, the main threat to gross margin identified in E-33.
2. **Payer rate up:** ENVOY's value depends on G-P2 (meets per intent +40%); the OS's job is to make that measurable and explainable.
3. **New revenue not dependent on member growth:** B2B licensing from P4, with near-zero marginal member risk.

12.5 Revenue optimisation, CLV, churn, upsell and cross-sell engines

Engine	What it does	What it may never do
AI revenue optimisation	Finds revenue leaks: failed renewals (dunning through the processor), processor authorisation drops, refund causes	Raise price for a member individually; obscure costs
Customer lifetime value	Cohort CLV by acquisition channel and segment, for CAC decisions (E-03 thresholds)	Label individual members by value in any staff-facing view
Paid-acquisition brake (FR-091)	Stops paid seeding in a city automatically when cost per verified sign-up is above break-even for 2 weeks (or 1.5× in one week)	Restart without a named approver
Churn prevention	Detects cohort-level churn drivers; fixes product causes (slate quality, liquidity)	Retention flows on cancel (one step, ToS cl. 16); guilt copy; hiding cancel
Upsell	Shows the relevant tier when a member hits a real limit (e.g. the 3 weekly intents)	Engagement paywalls (P8); upsell during a safety flow
Cross-sell	GHOST to members who use discretion features; DUO to linked couples	Selling safety features (core discretion and safety are free, D-05)

13. Security, compliance and risk

13.1 Cybersecurity Command Centre

Brief element	TRYST control	Status
Zero trust: never trust, always verify	Every request authenticated and authorised; mTLS + SPIFFE between services; staff via SSO + hardware key + device posture; no VPN trust	Spec (P1)
MFA	Members: two factors always (02 §3.3). Staff: FIDO2 hardware key mandatory	Built (members)
Biometric authentication	B1 device-bound biometric passkey + B2 liveness match on mobile	Built (B2 provider pluggable; fails closed)
Device fingerprinting	Device binding by passkey and DPoP key (not cross-site fingerprinting); hashed device risk signals for admission L1	Spec
Risk-based authentication	Identity risk agent can add step-up; it can never remove a factor	Spec
Real-time monitoring	SIEM over OpenTelemetry + audit chains; SOC agent triage	Spec

Brief element	TRYST control	Status
Behavioural analytics	Login anomaly (impossible travel, new device + new network), admin behaviour baselines	Spec
AI anomaly detection	Unsupervised models on ops telemetry and admin actions (not on member content)	Spec
Transaction scoring	Fraud agent + Radar + acquirer scores	Spec
Behaviour scoring	Admission L1-L5; Guardian graph signals	Spec
Device intelligence	Platform attestation (App Attest, Play Integrity) for native builds	Spec (P3)

13.2 Anti-hacking framework

Threat	Control
DDoS	Cloudflare and AWS Shield; per-route rate limits; autoscaling with caps
SQL injection	Parameterised queries only (Go database/sql placeholders, pglast-validated migrations); least-privilege DB roles; SAST rule
XSS	React escaping; strict CSP with no inline scripts on app origins (web/next.config.ts); no third-party scripts; Trusted Types in P1
CSRF	No cookie-authenticated mutating endpoints without SameSite=Strict and Origin checks; tokens are DPoP-bound bearer (not ambient)
Session hijacking	Sessions bound to the device (DPoP); short lifetimes; sign-out everywhere; emergency lock (built)
Account takeover	No passwords, no SMS or email login codes; passkey + second factor; no staff overrides (D-16)
Credential stuffing	Nothing to stuff: no passwords exist
API abuse	Scoped keys, idempotency, rate limits, anomaly detection on partner usage
Bot attacks	Bot management at the edge; V2 verification before any interaction (E-01)
Prompt injection and model abuse	§5.7 injection defence; output schemas; red team per release
Supply chain	Signed images, SLSA provenance, dependency pinning and scanning, reviewed PRs only (including agent-authored ones)

13.3 Data protection

Brief element	Control
Encryption at rest	AES-256 via KMS for every store; per-member keys for profile data; crypto-shredding for erasure (03 §6.3)
Encryption in transit	TLS 1.3 at the edge; mTLS inside
Encryption in use	E2EE threads (MLS); client-side media encryption (Rust core, built); on-device Guardian; Nitro Enclaves evaluated for the PSI server (P2)
Tokenisation	Card data tokenised by processors (TRYST is SAQ-A scope); contact handles stored as HMAC with an HSM pepper

13.4 Compliance

Area	How the OS meets it
UK GDPR / DPA 2018 (Art 9 special category)	Explicit consent, DPIA, self-hosted models for member data, minimisation, DSR automation, < 60 s erasure
UK GDPR Art 22 (automated decisions)	Always-human classes; explanations from reason codes; appeal path (FR-035)
Online Safety Act 2023	Highly effective age assurance; risk assessments; T&S SLAs; transparency reporting generated from audit data
EU AI Act (P4 markets)	Agent register, logging, human oversight and transparency labels ("machine-generated" on Briefs, P4) are already in place
KYC / KYB	Members: age assurance, not identity KYC (pseudonymous, FR-002). Partners and licensees: KYB
AML	TRYST holds no customer funds; AML obligations sit with regulated processors; TRYST screens business partners (D-32)
PCI DSS	SAQ-A: hosted checkout only; no card data touches TRYST
Audit logs	Hash-chained agent and admin logs; WORM anchoring

Area	How the OS meets it
Secure API keys	Shown once; Argon2id hashed; scopes; IP allow-lists; rotation; leaked-key scanning (GitHub secret scanning)
Webhook signing	HMAC with timestamp tolerance in and out; dedupe
Transaction monitoring	Payment and Fraud agents; chargeback thresholds
Admin action tracking	ops.admin_action, two-person break-glass

13.5 AI risk register (additions to 01 §13.3)

Risk	Likelihood	Impact	Mitigation	Owner
An agent takes a harmful action at scale	Medium	Severe	Gate, budgets, kill switch, canary rollout of agent versions	AI Governance
Model leaks special-category data	Low	Severe	Self-hosted models; Privacy Auditor leakage tests; no free-text member memory	DPO
Prompt injection via profile or ticket text	High	Medium	Injection defence and red team (§5.7)	Security
Auto-repair loop worsens an outage	Medium	High	Runbooks only; budget 50/day; rollback-first; pages a person if not resolved in 10 min	SRE
Over-reliance: staff rubber-stamp approvals	Medium	High	Approval sampling and second review for permanent bans; approval-rate monitoring per approver	Head of T&S
Connector sends member data to a third party	Low	Severe	Connector register checks; egress allow-list; DLP	Security

14. Admin Super Control Centre

Internal origin, staff SSO + hardware key, redacted by default, every view and action logged (FR-086).

Area	Visibility	Actions (and who may take them)
Users	Counts and cohorts by default; a specific member only through a case (JIT grant)	Hold (Guardian/moderator); permanent actions by moderator + second review
Businesses	Licensees and partners with KYB status, contracts, usage	Suspend keys (admin); terminate (two admins)
Transactions	Ledger explorer, payment events by processor	Refund (payments ops); failover (payments ops)
APIs	Traffic, errors, latency by route and key; deprecations	Throttle a key (admin); revoke (admin)
Payments	Authorisation rate, chargebacks, disputes, settlement by processor	Approve dispute responses; switch primary processor (two admins)
Agents	Registry, levels, budgets, decision stream, pending approvals, eval scores	Approve pending decisions (named approvers); kill (any on-call); resume (two people)
Automations	Workflow runs, failures, compensations	Retry, cancel (ops)
Disputes	Member appeals (FR-035) and payment disputes	Decide appeals (moderator not involved in the original decision)
Compliance	DSR clock, retention violations, OSA SLAs, regulator requests	Approve disclosures (DPO)
Revenue	Daily pack, forecasts, experiment results	Approve price experiments (pricing committee)
System health	SLOs, error budgets, incidents, deploys	Freeze deploys; approve promotion
Logs	Structured logs (scrubbed), traces	Read (engineers, JIT in production)
Alerts	Unified alert stream with dedupe and ownership	Acknowledge, assign, escalate
Platform performance	Web vitals, app crash-free rate, latency by journey	—
Audit	Chain verification status; last anchor; break-glass events	Export for auditors (DPO)

Break-glass: a member-level view outside a case needs a justification and a second approver, expires in 8 hours (DDL CHECK), and is reported in the monthly access review.

15. Developer build roadmap

The OS phases align with TRYST's phases and gates (01 §12). The brief's five phases map as: **MVP = P1, Beta = P1 public beta → P2, Commercial launch = P3, Enterprise version = P4, Global scale version = post-P4.**

15.1 MVP (P1, M2-M6)

Modules	Accounts and two-factor sign-in (built in part), Loops A, B, D, E2EE threads, V2 gating, VC co-sign, Guardian baseline on device, DSR and erasure, payments (Stripe + backup, BitriPay sandbox), agent registry and gate, audit chain, Admin Super Control Centre v1, connector register
APIs	Member API (02 §5), Ops API (§11.2), inbound webhooks
User flows	J1 (basic), J3, J4, J5, J7, J8
AI agents	Cartographer, Broker (P1 baseline ranker), Curtain, Guardian, Aftercare, Onboarding, Risk, Fraud, Support, Compliance, Payment, Herald, System Health, Bug Detection, Auto-Repair, Release Manager, Threat Hunter, Vulnerability, SOC triage, executive copilots
Technical milestones	PostgreSQL IDENTITY-DB; DPoP; real liveness provider; gate as a central service with WORM anchoring; SLOs live; erasure < 60 s p99
Commercial objectives	G-P1: matching live in London with 2,000 verified; sign-up open in all four launch cities; no-go list clear

15.2 Beta (P2, M6-M10)

Modules	Envoy, Mirror v1, ExclusionRing PSI, Member Command Centre (insights, "why" panels, my memory), analytics
AI agents	+ Envoy, Mirror, Chief of Staff (member), Data Intelligence, Predictive Growth, Infra Optimiser, product copilots
Technical milestones	Workflow-level evals; chaos game days; Temporal in production
Commercial objectives	G-P2: meets per intent +40% vs the P1 control. If not met, stop and rethink

15.3 Commercial launch (P3, M10-M15)

Modules	DUO and GHOST billing, native apps in both stores, 24/7 T&S, Couple Command Centre, pricing experiments
AI agents	+ Pricing, Monetisation, Retention, AML (payments monitoring)
Technical milestones	App attestation; open-banking evaluation; second processor live-tested
Commercial objectives	G-P3: 25,000 verified; payer rate ≥ 11%; both stores approved

15.4 Enterprise version (P4, M15-M24)

Modules	Partner API, partner and developer portals, ACU metering, KYB, BitriPay connected accounts for partner commissions, regulator portal, EU localisation
AI agents	+ B2B Sales, Licensee Success, KYB onboarding
Technical milestones	Cell architecture design; EU AI Act conformity file; SOC 2 Type II and ISO 27001 / 27701
Commercial objectives	G-P4: 150,000 verified; contribution-positive; first two licensees signed

15.5 Global scale version (post-P4)

Modules	Multi-cell deployment (UK+IE, EU-North, EU-South, then CA/AU/NZ per 01 §14.5), localised assurance and Guardian, licensee self-serve onboarding
AI agents	Per-cell agent instances with a global governance plane
Technical milestones	Cell isolation tests; regional key domains; follow-the-sun T&S

Commercial objectives	Licensing revenue $\geq$ 20% of total; per-market contribution margin positive within 12 months of launch
-----------------------	---

15.6 Production readiness review (every phase gate)

Area	Must be true
Reliability	SLOs defined and met for 30 days in staging/beta; rollback tested; DR restore drill passed this quarter
Security	Pen test closed (no high or critical open); threat model updated; secrets rotated; break-glass tested
Privacy	DPIA updated; Privacy Auditor green; erasure and DSR timings within NFRs
AI	Every agent registered with owner, level, budget, eval pass and red-team pass; kill-all drill done
Payments	Webhook replay and dedupe tested per processor; chargeback alerting live; backup MID transaction tested
Compliance	OSA risk assessment current; transparency report pipeline tested; store policy check (P3+)
Operations	On-call rota; runbooks for top 20 alerts; support macros reviewed against claims discipline (01 §15.1)

16. Competitive advantage

Dimension	How the OS wins	Why it is hard to copy
More powerful	Agent-to-agent negotiation turns intent into meetings; agents carry safety, support and operations volume	Needs the consent schema, the asymmetric Brief design and member trust built together
More profitable	Revenue tied to outcomes (ENVOY, QCAM); lower ops cost per MAU; B2B licensing from P4	Competitors' revenue depends on message volume and paywalls they cannot drop
More scalable	Event-driven services, cell architecture, agents that scale with load rather than headcount	—
More intelligent	Learns from outcomes (aftercare reward), not clicks; stated-preference floor keeps members in control	Outcome data only exists if the product produces meetings
More trustworthy	Two-factor account-holder-only sign-in, E2EE, crypto-erasure, governance as tested code, no data sale	Incumbents carry breach history and ad-funded or volume-funded models
More compliant	OSA, GDPR Art 9 and 22, and EU AI Act controls designed in	Retrofitting into an engagement-optimised stack is slow and expensive
More commercially competitive	Processor independence (Stripe, acquirer, BitriPay) protects revenue; licensing opens a second business	Category players rely on a single high-risk processor relationship

17. Final output format and index

This document is the developer-ready OS specification for TRYST. The brief's required technical outputs map as follows:

Required output	Where
Product Requirements Document	01 Product + §1–§4, §6, §12 here
Technical Requirements Document	02 (FR/NFR) + FR-080–FR-088
System, infrastructure, AI, agent architecture	§5, §9.1–§9.3; 03
Security architecture; complete security controls	§13; 03 §6
Database architecture; complete ERD and schemas	§10; 02 §4; generated migrations
API architecture; complete API specifications	§11; 02 §5; contracts/openapi/tryst.v1.yaml
Event architecture	§9.4; 02 §8
DevOps and deployment architecture	§9.9
Monitoring architecture	§9.3 (observability, monitoring), §9.5 SLOs
Disaster recovery and business continuity	§9.7

Required output	Where
Data governance architecture	§9.6
Compliance architecture	§13.4; 01 §14
Scalability architecture	§9.8
Commercial architecture	§12
Complete user journeys	§3.2; 01 §6; 04 §4
Complete agent workflows	§5.5, §5.6, §5.7
Complete testing strategy	§9.10
Production readiness review	§15.6
Code that implements this document today	backend/internal/aigov (governance gate, workforce registry, audit chain), backend/internal/payments/bitripay.go (BitriPay adapter), plus the v1.2 identity, payments and agent code